

Lecture 4

- **Instruction Set Architecture (II)**

Procedures

■ `int f1 (int i, int j, int k, int g)`

`{ ::::`

`add $3, $4, 1;`

`return 1;`

`}`

callee

`int f2 (int s1, int s2)`

`{`

`:::::`

`add $3, $4, $3`

`i = f1 (3,4,5, 6);`

`add $2, $3, $3`

`::::`

`}`

caller

- How to pass parameters & results?
- How to preserve caller register values?
- How to alter control? (I.e. go to callee, return from callee)

Procedures

■ How to pass parameters & results

- ❑ \$a0-\$a3: four argument registers. What if # of parameters is larger than 4? – push to the **stack**
- ❑ \$v0-\$v1: two value registers in which to return values

■ How to preserve caller register values?

- ❑ Use **stack**
- ❑ Caller saved register
- ❑ Callee saved register

■ How to switch control?

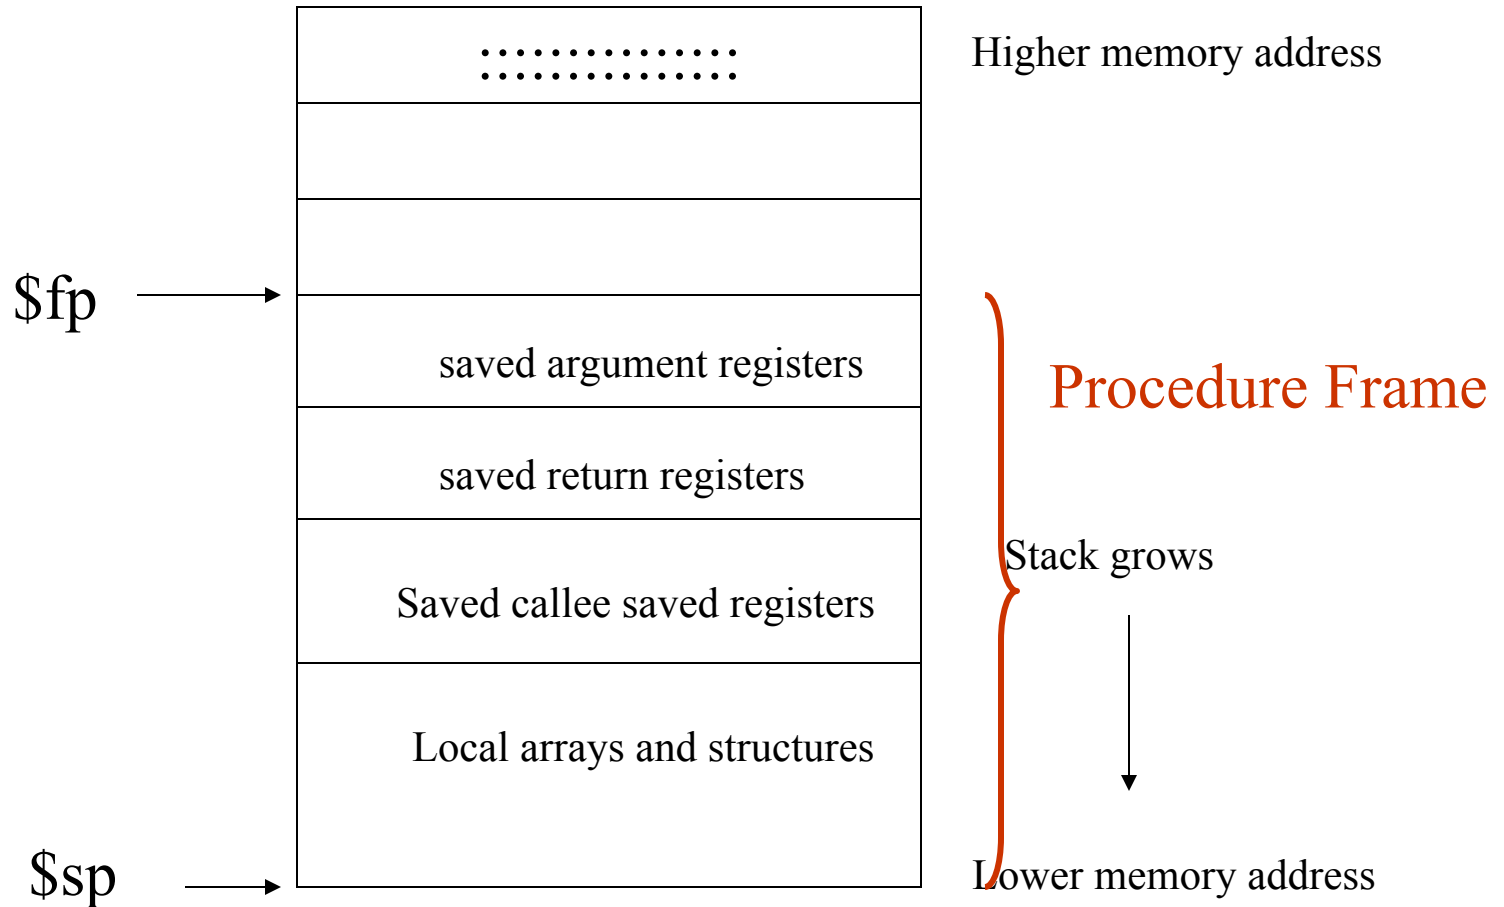
- ❑ How to go to the callee
 - **Jal procedure_address** (jump and link)
 - ❑ Store the the return address (PC +4) at \$ra
 - ❑ set PC = procedure_addres
- ❑ How to return from the callee
 - ❑ Callee exectues **jr \$ra**

```
int f1 (int i, int j, int k, int g)
{
    ::::
    add $3, $4, 1;
    return 1;
}
```

```
int f2 (int s1, int s2)
{
    :::::
    add $3, $4, $3
    i = f1 (3,4,5, 6);
    add $2, $3, $3
    ::::
}
```

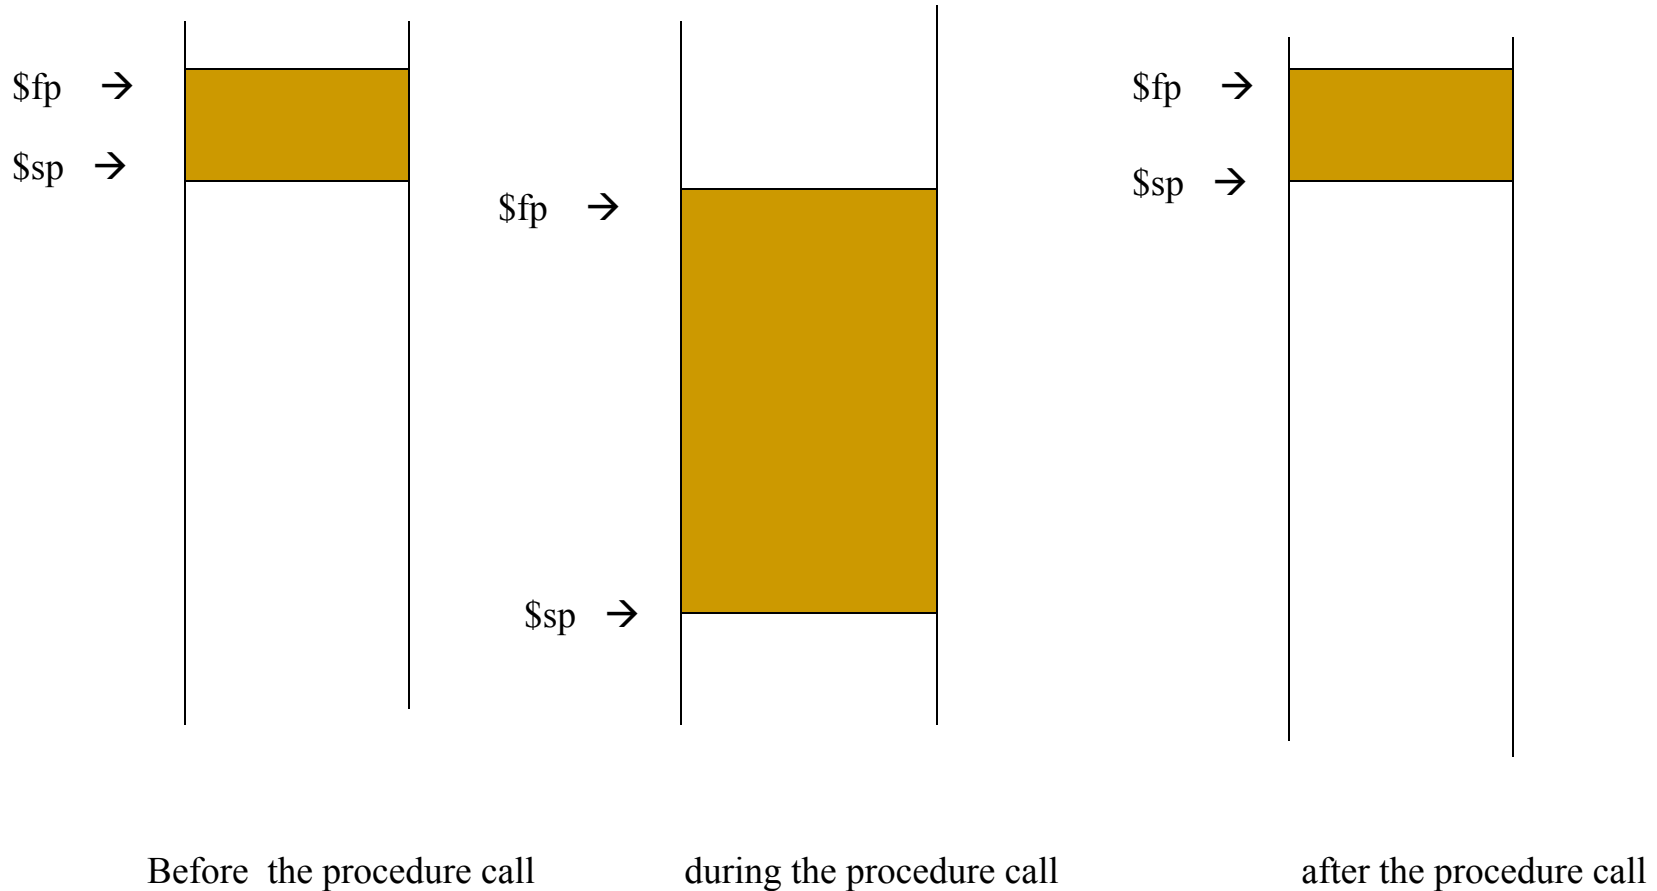
PC →

Procedure Call Stack (Frame)



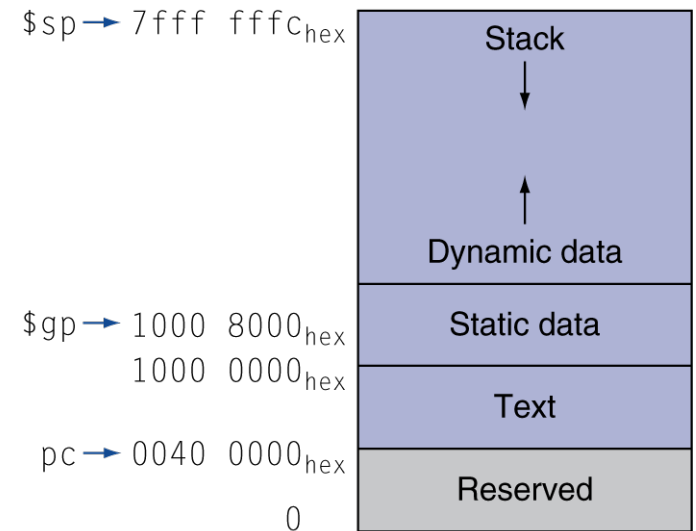
- Frame pointer points to the first word of the procedure frame

Procedure Call Stack (Frame)



Memory Layout

- Text: program code
- Static data: global variables
 - e.g., static variables in C, constant arrays and strings
 - \$gp initialized to address allowing \pm offsets into this segment
- Dynamic data: heap
 - E.g., malloc in C, new in Java
- Stack: automatic storage



Procedure Calling Convention

■ Calling Procedure

- ❑ Step-1: pass the argument
- ❑ Step-2: save caller-saved registers
- ❑ Step-3: Execute a **jal** instruction

```
foo1 ()  
{  
    .....  
    i = i+1; t3  
    x=foo(4);  
    i = x + i  
}
```

```
.....  
li    $a0, 4          # passing argument  
sw    $t3, 4($sp)     # save $t3  
jal   foo  
lw    $t3, 4($sp)     # restore $t3  
add   $t3, $v0, $t3
```

```
.....
```

Procedure Calling Convention

Called Procedure

- ❑ Step-1: establish stack frame
 - `subi $sp, $sp, <frame-size>`
- ❑ Step-2: saved callee saved registers
 - `$ra, $fp, $s0-$s7`
- ❑ Step-3: establish frame pointer
 - `Add $fp, $sp, <frame-size>-4`
- On return from a call
 - ❑ Step-1: put returned values in register `$v0, $v1`.
 - ❑ Step-2: restore callee-saved registers
 - ❑ Step-3: pop the stack
 - ❑ Step-4: return: `jr $ra`

```
subi $sp, $sp, 32
sw   $ra, 20($sp)
sw   $fp, 16($sp)
addi $fp, $sp, 28
....
....
....

addi $v0, $zero, 1
lw   $fp, 16($sp)
lw   $ra, 20($sp)
addi $sp, $sp, 32
jr   $ra
```


Preserved Registers

0	zero	constant 0
1	at	reserved for assembler
2	v0	expression evaluation &
3	v1	function results
4	a0	arguments
5	a1	
6	a2	
7	a3	
8	t0	temporary: caller saves
	...	
15	t7	

16	s0	callee saves
----	----	--------------

...

23	s7	
----	----	--

24	t8	temporary (cont'd)
----	----	--------------------

25	t9	
----	----	--

26	k0	reserved for OS kernel
----	----	------------------------

27	k1	
----	----	--

28	gp	Pointer to global area
----	----	------------------------

29	sp	Stack pointer
----	----	---------------

30	fp	frame pointer
----	----	---------------

31	ra	Return Address (HW)
----	----	---------------------

Nested Procedures

fact:

```
addi    $sp, $sp, -8
sw      $ra, 4($sp)  # save $ra
slti    $t0, $a0, 1  # n < 1?
beq     $t0, $zero, L1
addi    $v0, $zero, 1 # return 1
addi    $sp, $sp, 8  # fix up the stack pointer & return
jr      $ra

L1: sw   $a0, 0($sp)  # save argument $a0
addi    $a0, $a0, -1  # n = n-1
jal     fact          # jal(n-1)
lw      $a0, 0($sp)  # restore argument $a0
mul     $v0, $a0, $v0 # return n x fact(n-1)
lw      $ra, 4($sp)  # restore $ra
addi    $sp, $sp, 8  # restore stack pointer
jr      $ra          # return to the caller
```

```
int fact (int n)
{
    if (n < 1) return 1;
    else return (n x fact(n-1));
}
```

Representation of Characters

■ ASCII (American Standard Code for Information Interchange)

- Uses 8 bits to represent a character
- MIPS provides instructions to move bytes:
 lb \$t0, 0(\$sp) #Read byte from source
 sb \$t0, 0(\$gp) #Write byte to destination

■ Unicode (Universal Encoding)

- Uses 16 bits to represent a character
- Used in Java
- MIPS provides instructions to move 16 bits:
 lh \$t0, 0(\$sp) #Read halfword from source
 sh \$t0, 0(\$gp) #Write halfword to destination

Array vs. Pointer

```
Clear1(int array[ ], int size)
{
    int i;
    for (i=0, i< size; i+= 1)
        array[i] = 0;
}
```



```

      move    $t0, $zero    # i =0
Loop1: sll     $t1, $t0, 2    # i * 4
      add     $t2, $a0, $t1  # t2 = address of array[i]
      sw      $zero, 0($t2)  # array [i] = 0
      addi    $t0, $t0, 1    # i = i +1
      slt     $t3, $t0, $a1  # compare i and size
      bne     $t3, $zero, loop1
```

Array vs. Pointer (cont.)

```
Clear 2(int *array, int size)
{
    int *p,
    for (p = &array[0]; p < &array[size]; p = p+1)
        *p = 0;
}
```



move	\$t0, \$a0	# p = &array[0]
sll	\$t1, \$a1, 2	# t1 = size x 4
add	\$t2, \$a0, \$t1	# t2 = &array[size]
Loop2: sw	\$zero, 0(\$t0)	# memory[p] = 0
addi	\$t0, \$t0, 4	# p = p + 4
slt	\$t3, \$t0, \$t2	# compare p, & array[size]
bne	\$t3, \$zero, Loop2	

Array vs. Pointer (cont.)

Array

```
        move    $t0, $zero    # i = 0
Loop1 : sll     $t1, $t0, 2    # i * 4
        add     $t2, $a0, $t1 # t2 = address of array[i]
        sw      $zero, 0($t2) # array[i] = 0
        addi    $t0, $t0, 1    # i = i + 1
        slt     $t3, $t0, $a1  # compare i and size
        bne     $t3, $zero, loop1
```

of Instruction per iteration
= 6

Pointer

```
        move    $t0, $a0      # p = &array[0]
        sll     $t1, $a1, 2    # t1 = size x 4
        add     $t2, $a0, $t1  # t2 = &array[size]
Loop2: sw      $zero, 0($t0)    # memory[p] = 0
        addi    $t0, $t0, 4    # p = p + 4
        slt     $t3, $t0, $t2  # compare p, & array[size]
        bne     $t3, $zero, Loop2
```

of Instruction per iteration
= 4

Parallelism and Instructions: Synchronization

- Parallel tasks must **synchronize** to avoid **data race**, where the results of the program can change depending on how events happen to occur.
- Lock/unlock: ensure only one task entering the critical section

P(1)

Acquire Lock;

If Lock = 0

 enter critical section;

Release Lock;

P(2)

Acquire Lock;

If Lock = 0

 enter critical section;

Release Lock;

Parallelism and Instructions: Synchronization

■ **Atomic SWAP**: **atomically** interchange a value in a register for a value in memory; nothing else can interpose itself between the read and the write to the memory location

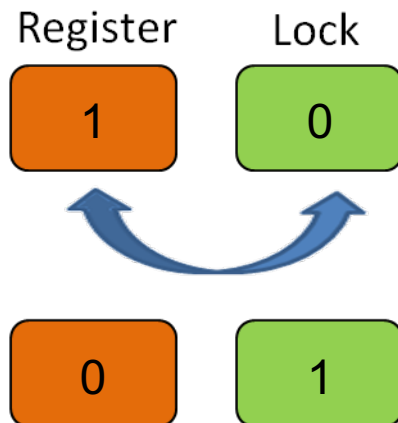
P(1)

`$S4 = 1;`

`Swap ($S4, Lock);`

If Lock = 0

enter critical section;



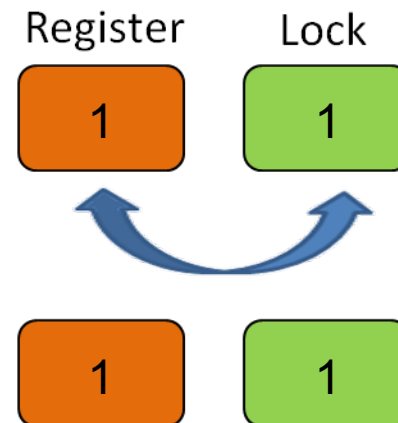
P(2)

`$S4 = 1;`

`Swap ($S4, Lock);`

If Lock = 0

enter critical section;



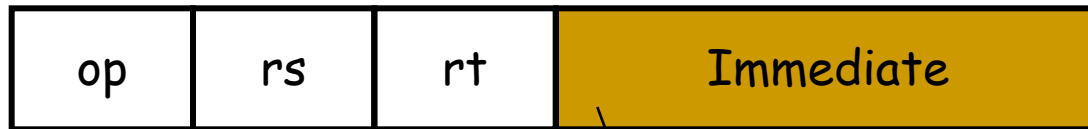
Synchronization in MIPS

- Load linked: `ll rt, offset(rs)`
- Store conditional: `sc rt, offset(rs)`
 - Succeeds if location not changed since the `ll`
 - Returns 1 in `rt`
 - Fails if location is changed
 - Returns 0 in `rt`
- Example: atomic swap (to test/set lock variable)

```
try: add $t0,$zero,$s4 ;copy exchange value
      ll $t1,0($s1)    ;load linked
      sc $t0,0($s1)    ;store conditional
      beq $t0,$zero,try ;branch store fails
      add $s4,$zero,$t1 ;put load value in $s4
```

MIPS Addressing Mode (1)

- Immediate addressing

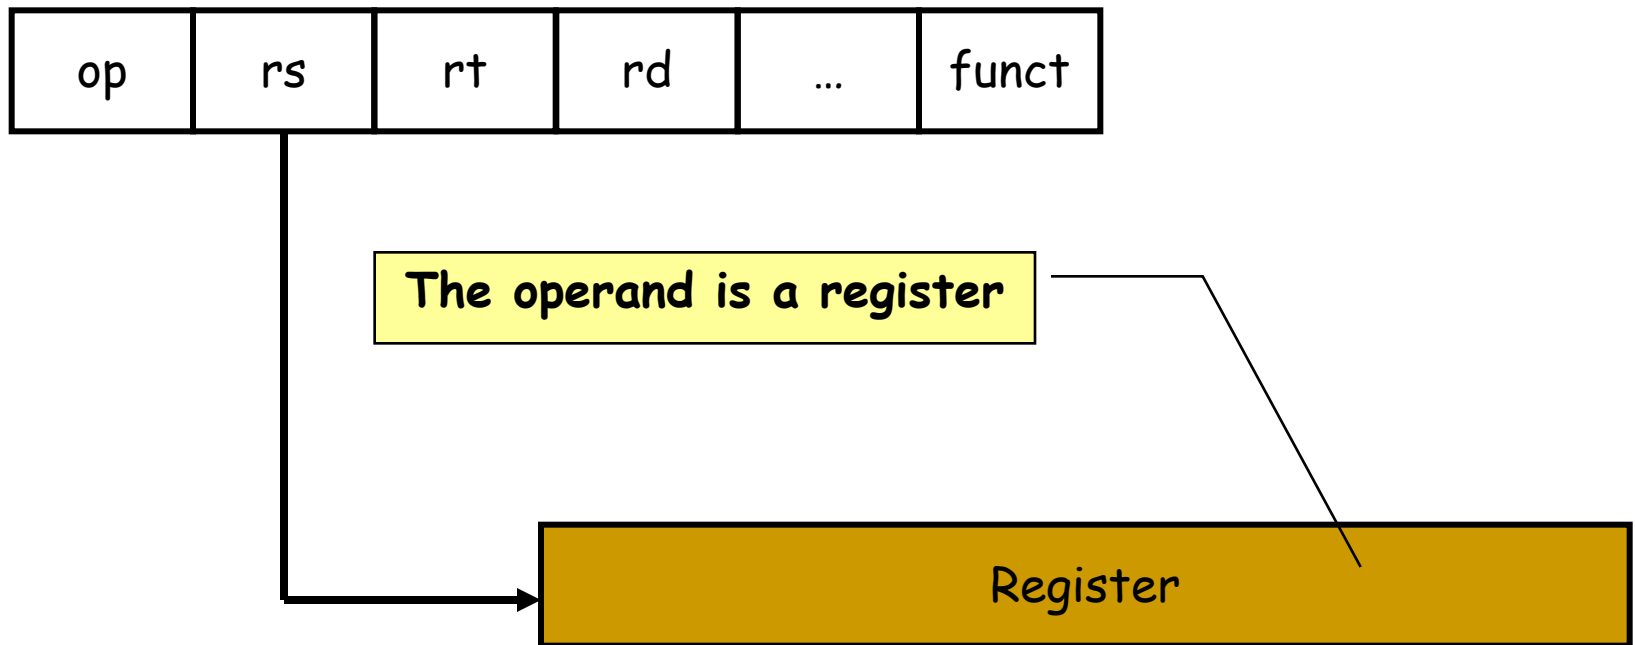


Example: `addi $2, $3, 4`

The operand is a constant within the instruction itself

MIPS Addressing Mode (2)

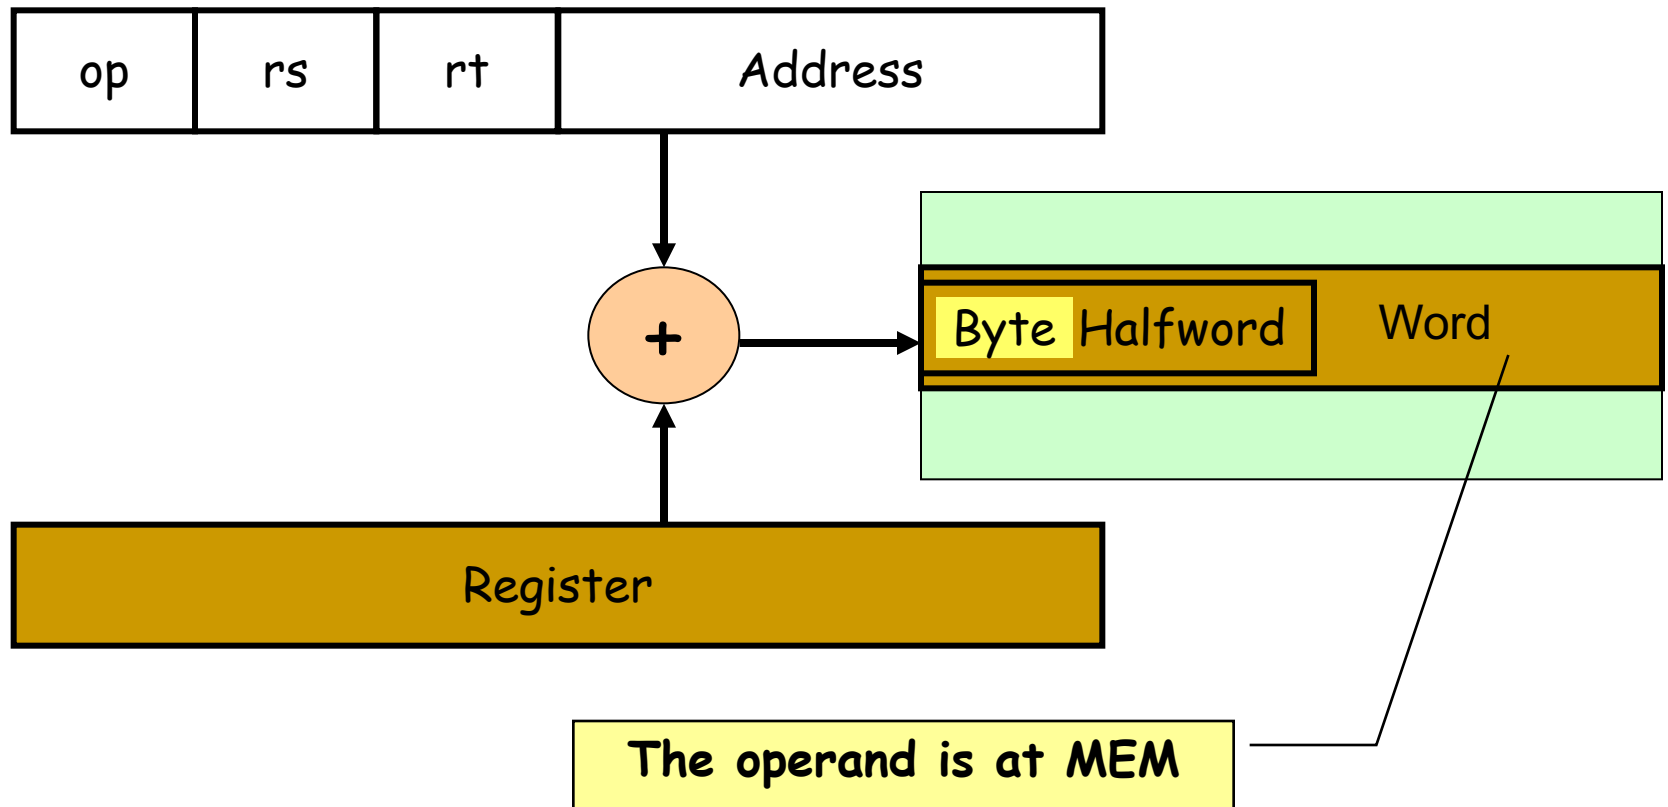
- Register addressing



Example : add \$r1, \$r2, \$r3

MIPS Addressing Mode (3)

- Base addressing



Example : lw \$2, 100(\$3)

How to Get the Base Address in the Base Register

Method 1.

```
.data          # define prog. data section
xyz: .word 1    # some data here
...           # possibly some other data
.text         # define the program code
...           # lines of code
lw $5,xyz      # loads contents of xyz in r5
```

- the assembler generates an instruction of the form:
`lw $5, offset($gp)` # gp is register 28, the global pointer

Note : `.data`, `.word`, `.text` are assembler directives

How to Get the Base Address in the Base Register (cont.)

Method 2.

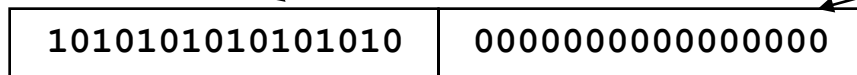
la \$6,xyz #r6 contains the address of xyz
lw \$5,0(\$6) #r5 contains the contents of xyz

Method 3. If the address is a constant

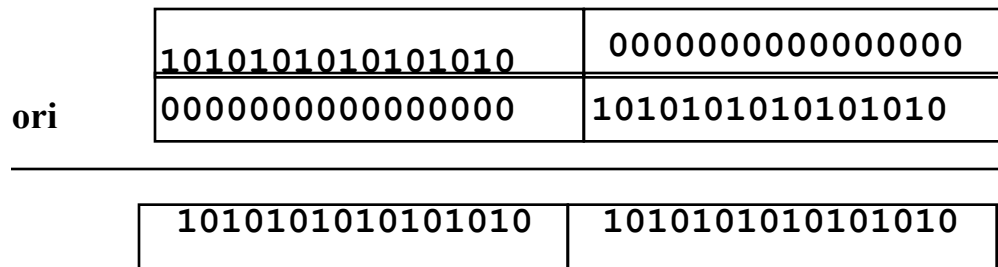
li, *if less than $\pm 32K$*
lui and ori, otherwise.

How about larger constants

- We'd like to be able to load a 32 bit constant into a register
- Must use two instructions, new "load upper immediate" instruction
- Example: load `10101010101010101010101010101010` into register `$t0`
- `lui $t0, 1010101010101010`

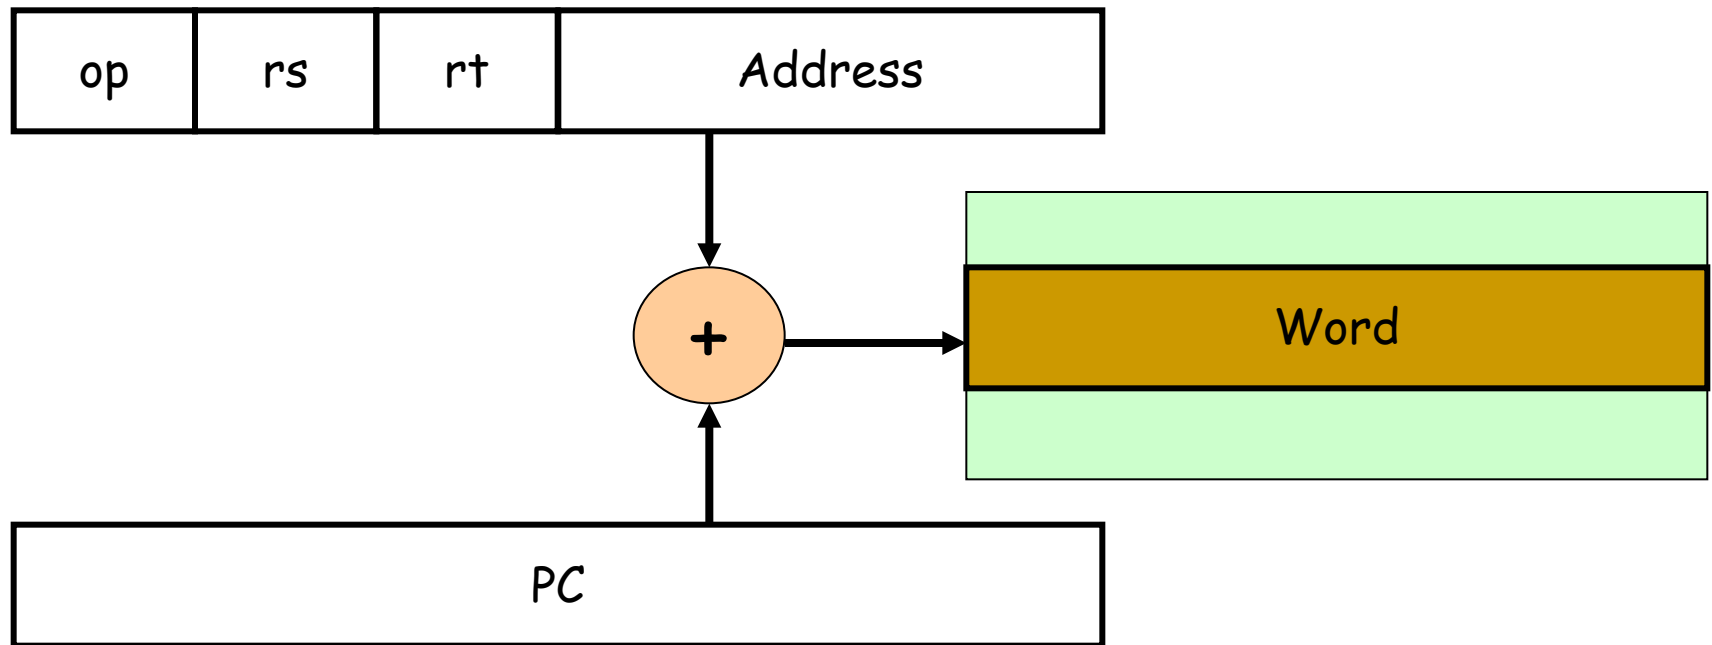


- Then must get the lower order bits right, i.e.,
`ori $t0, $t0, 1010101010101010`



MIPS Addressing Mode (4)

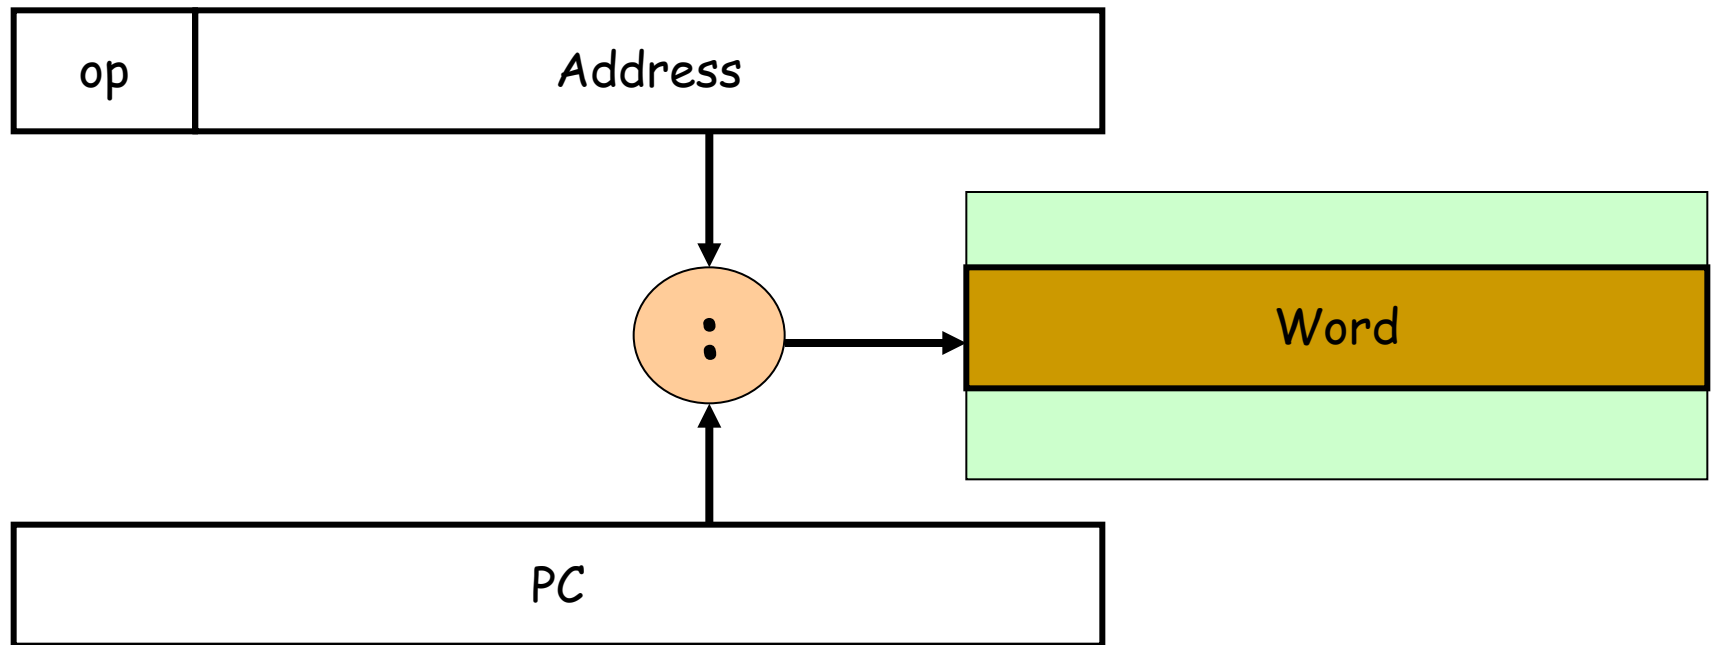
- PC-relative addressing



Example : beq \$2, \$3, 100

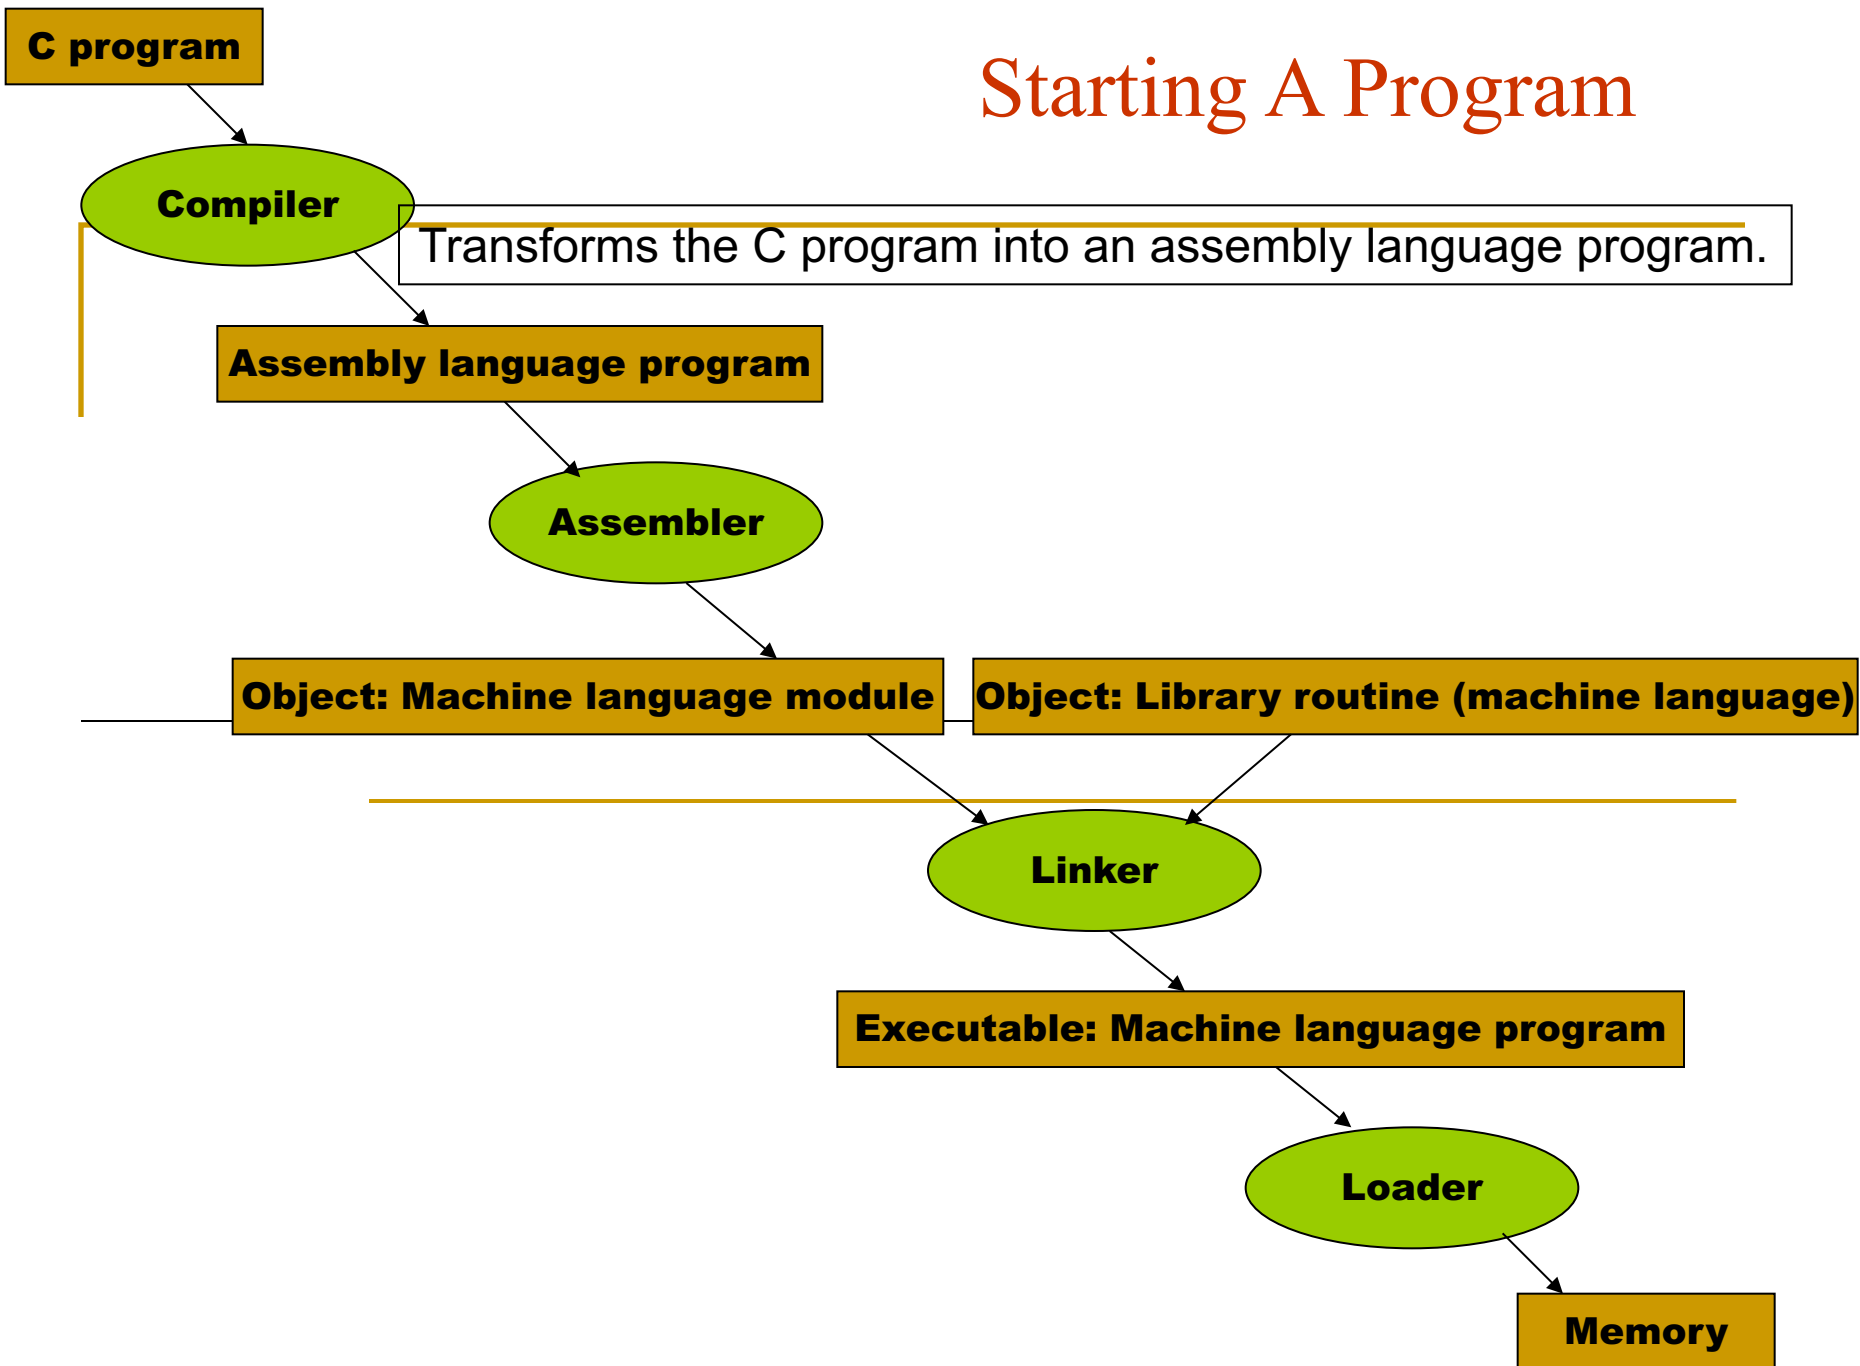
MIPS Addressing Mode (5)

- Pseudodirect addressing



Example : j 100

Starting A Program



Assembler

■ Assembler

- ❑ The assembler turns the assembly language program into an object file.
- ❑ Symbol table: A table that matches names of labels to the addresses of the memory words that instruction occupy.

Object file header			
	Name	Procedure A	
	Text size	100hex	
	Data size	20hex	
Text segment	Address	Instruction	
	0	lw \$a0, 0(\$gp)	
	4	jal 0	
	
Data segment	0	(X)	
	
Relocation information	Address	Instruction Type	Dependency
	0	lw	X
	4	jal	B
Symbol table	Label	Address	
	X	—	
	B	—	

lw \$a0, x
jal B

Assembler (cont.)

- Pseudoinstruction: a common variation of assembly language instructions often treated as if it were an instruction in its own right.
 - `move $t0, $t1` -> `add $t0, $zero, $t1`
 - `blt` -> `slt` & `bne`

Linker (Link editor)

- Linker takes all the independently assembled machine language programs and “stitches” them together to produce an executable file that can be run on a computer.
- There are three steps for the linker:
 1. Place code and data modules symbolically in memory.
 2. Determine the addresses of data and instruction labels.
 3. Patch both the internal and external references.

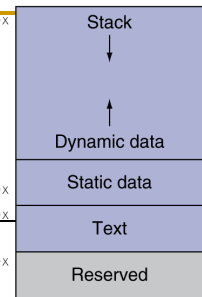
Object file header			
	Name	Procedure A	
	Text size	100hex	
	Data size	20hex	
Text segment	Address	Instruction	
	0	lw \$a0, 0(\$gp)	
	4	jal 0	
	
Data segment	Address	Instruction	
	0	(X)	
	
Relocation information	Address	Instruction Type	Dependency
	0	lw	X
	4	jal	B
Symbol table	Label	Address	
	X	—	
	B	—	

Object file header			
	Name	Procedure B	
	Text size	200hex	
	Data size	30hex	
Text segment	Address	Instruction	
	0	sw \$a1, 0(\$gp)	
	4	jal 0	
	
Data segment	Address	Instruction	
	0	(Y)	
	
Relocation information	Address	Instruction Type	Dependency
	0	sw	Y
	4	jal	A
Symbol table	Label	Address	
	Y	—	
	A	—	

\$sp → 7fff fffc_{hex}

\$gp → 1000 8000_{hex}
1000 0000_{hex}

pc → 0040 0000_{hex}
0



Executable file header		
	Text size	300hex
	Data size	50hex
Text segment	Address	Instruction
A	0040 0000 _{hex}	lw \$a0, 8000 _{hex} (\$gp)
	0040 0004 _{hex}	jal 40 0100 _{hex}
B
	0040 0100 _{hex}	sw \$a1, 8020 _{hex} (\$gp)
	0040 0104 _{hex}	jal 40 0000 _{hex}

Data segment	Address	Instruction
	1000 0000 _{hex}	(X)

	1000 0020 _{hex}	(Y)

sw \$a0, y
jal A

Loader

- Read the executables file header to determine the size of the text and data segments
- Creates an address space large enough for the text and data
- Copies the instructions and data from the executable file into memory
- Copies the parameters (if any) to the main program onto the stack
- Initializes the machine registers and sets the stack pointer the first free location
- Jump to a start-up routine

```
main();
```

```
_start_up:
```

```
lw  a0, offset($sp)  ## load arguments
```

```
jal  main;
```

```
exit
```


Dynamically Linked Libraries (DLL)

- Disadvantages with traditional statically linked library
 - ❑ Library updates
 - ❑ Loading the whole library even if all of the library is not used
- Dynamically linked library
 - ❑ The libraries are not linked and loaded until the program is run.
 - ❑ Lazy procedure linkage
 - Each routine is linked only after it is called.

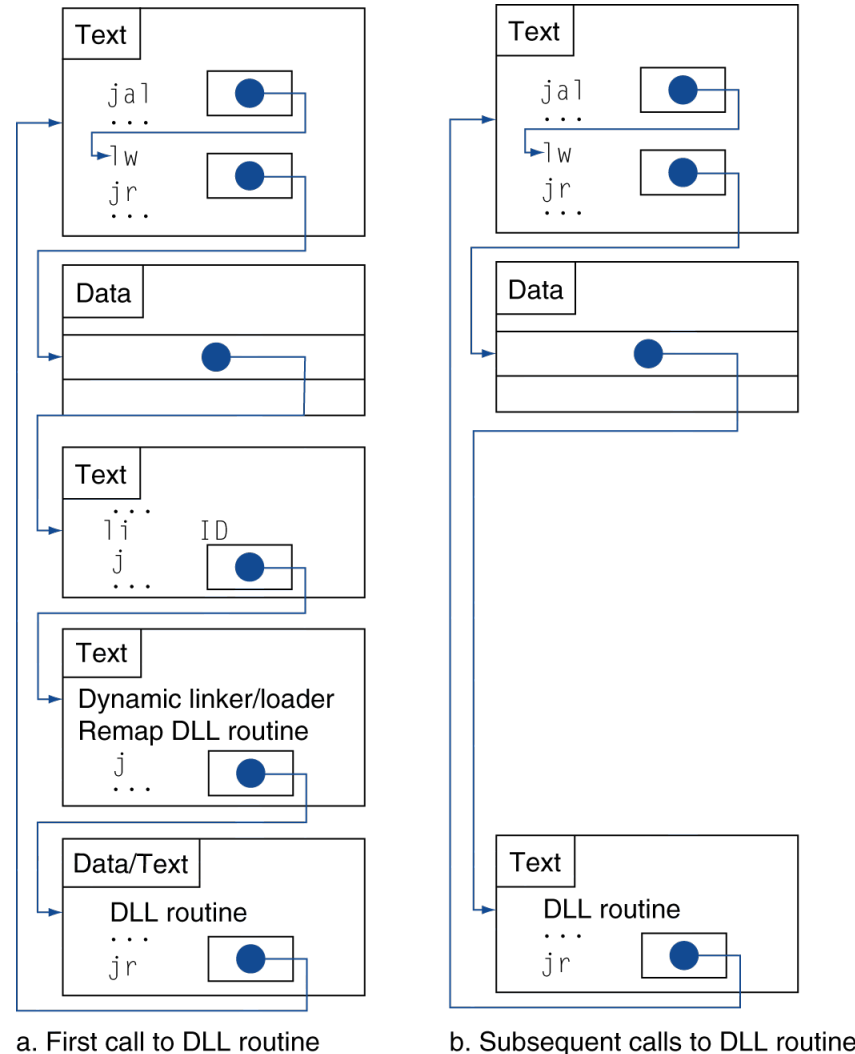
Lazy Linkage

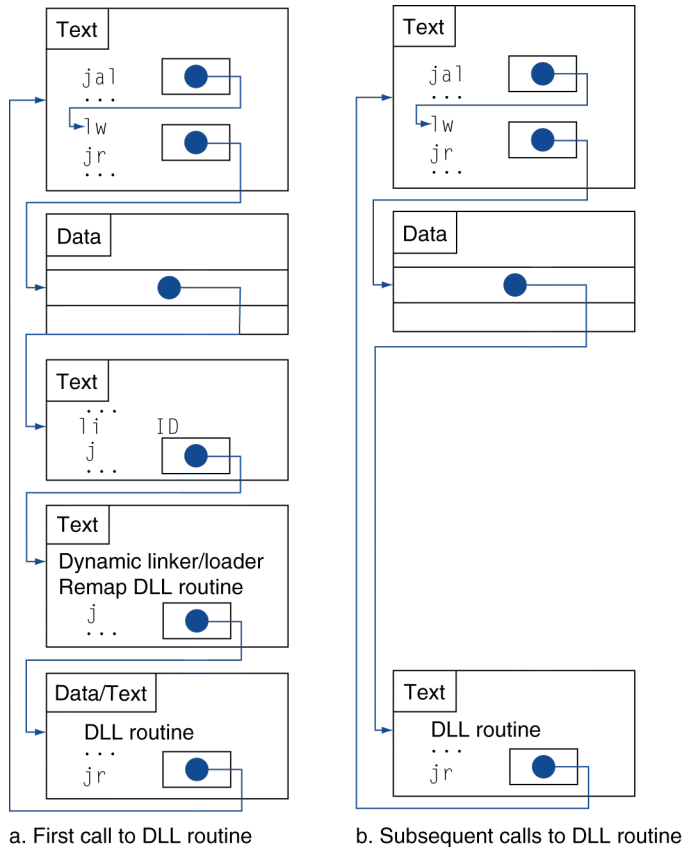
Indirection table

Stub: Loads routine ID,
Jump to linker/loader

Linker/loader code

Dynamically
mapped code



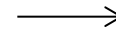


Text

```
jal printf();
printf ()
{ lw $r1, printf_addr
  jr  $r1}
```

Data

printf_addr .word L1



Data

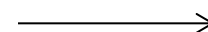
printf_addr 0x400000

Text

```
L1: li ID
     j DLL;
```

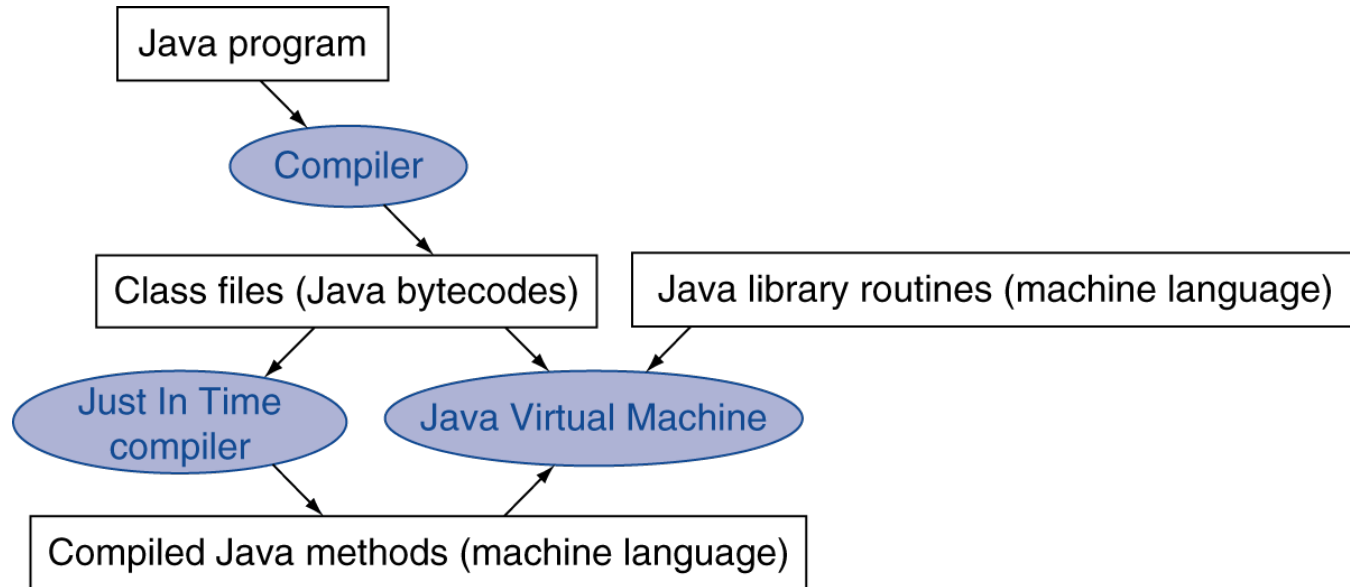
Text

```
printf()
.....
jr $ra
```



Assume this is loaded into
0x400000

Starting a Java Program



Java bytecode:

Instruction from an instruction set designed to interpret Java programs.

Java Virtual Machine

(JVM): The program that interprets Java bytecodes

Just In Time Compiler

(JIT): A compiler that operates at runtime, translating the interpreted code segments into the native code of the compiler.

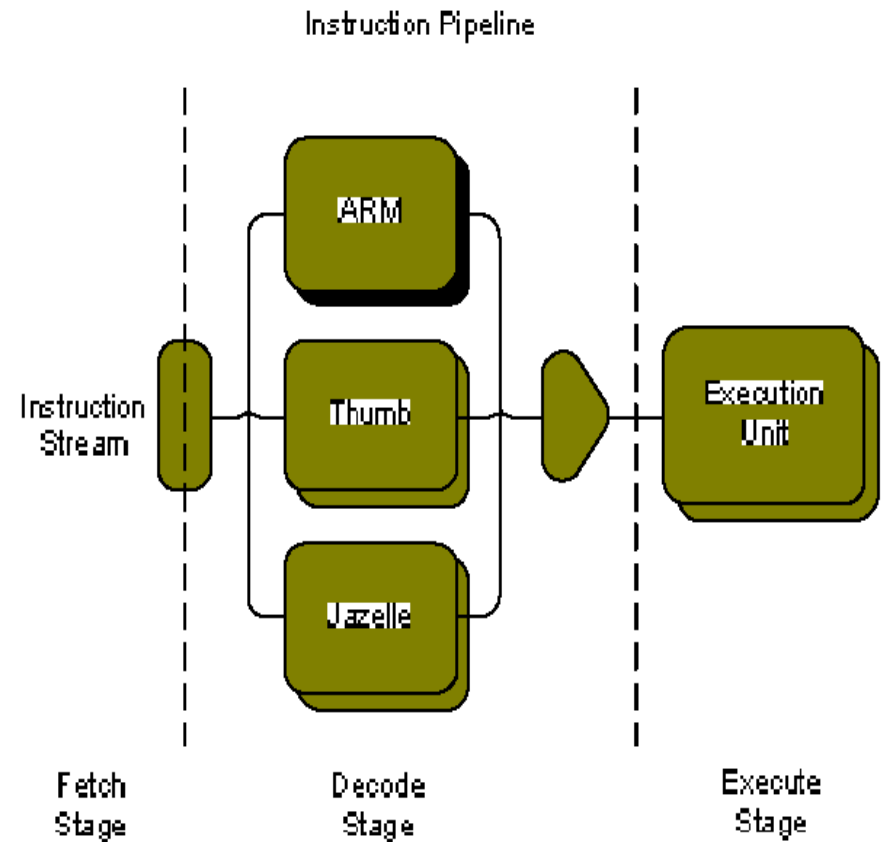
ARM & MIPS Similarities

- ARM: the most popular embedded core
- Similar basic set of instructions to MIPS

	ARM	MIPS
Date announced	1985	1985
Instruction size	32 bits	32 bits
Address space	32-bit flat	32-bit flat
Data alignment	Aligned	Aligned
Data addressing modes	9	3
Registers	15 × 32-bit	31 × 32-bit
Input/output	Memory mapped	Memory mapped

Instruction Set

- 32-bit ARM instruction
- 16-bit Thumb instruction
- 8-bit Java Instruction Set



Sources: white paper on [Jazelle Technology](#)

Main features of 32-bit Arm instruction

- Load-store architecture
- Fixed-length = 32 bits
- 3-address instruction format (2 sources, 1 result operands)
- DSP instruction
 - Support MAC & SIMD operations
- Conditional execution of ALL instructions

Registers

r0
r1
r2
r3
r4
r5
r6
r7

r8
r9
r10
r11
r12
r13 (SP)
r14 (LR)
r15 (PC)

Current program status register

31

0



Top four bits of CPSR

N : Negative

Z : Zero

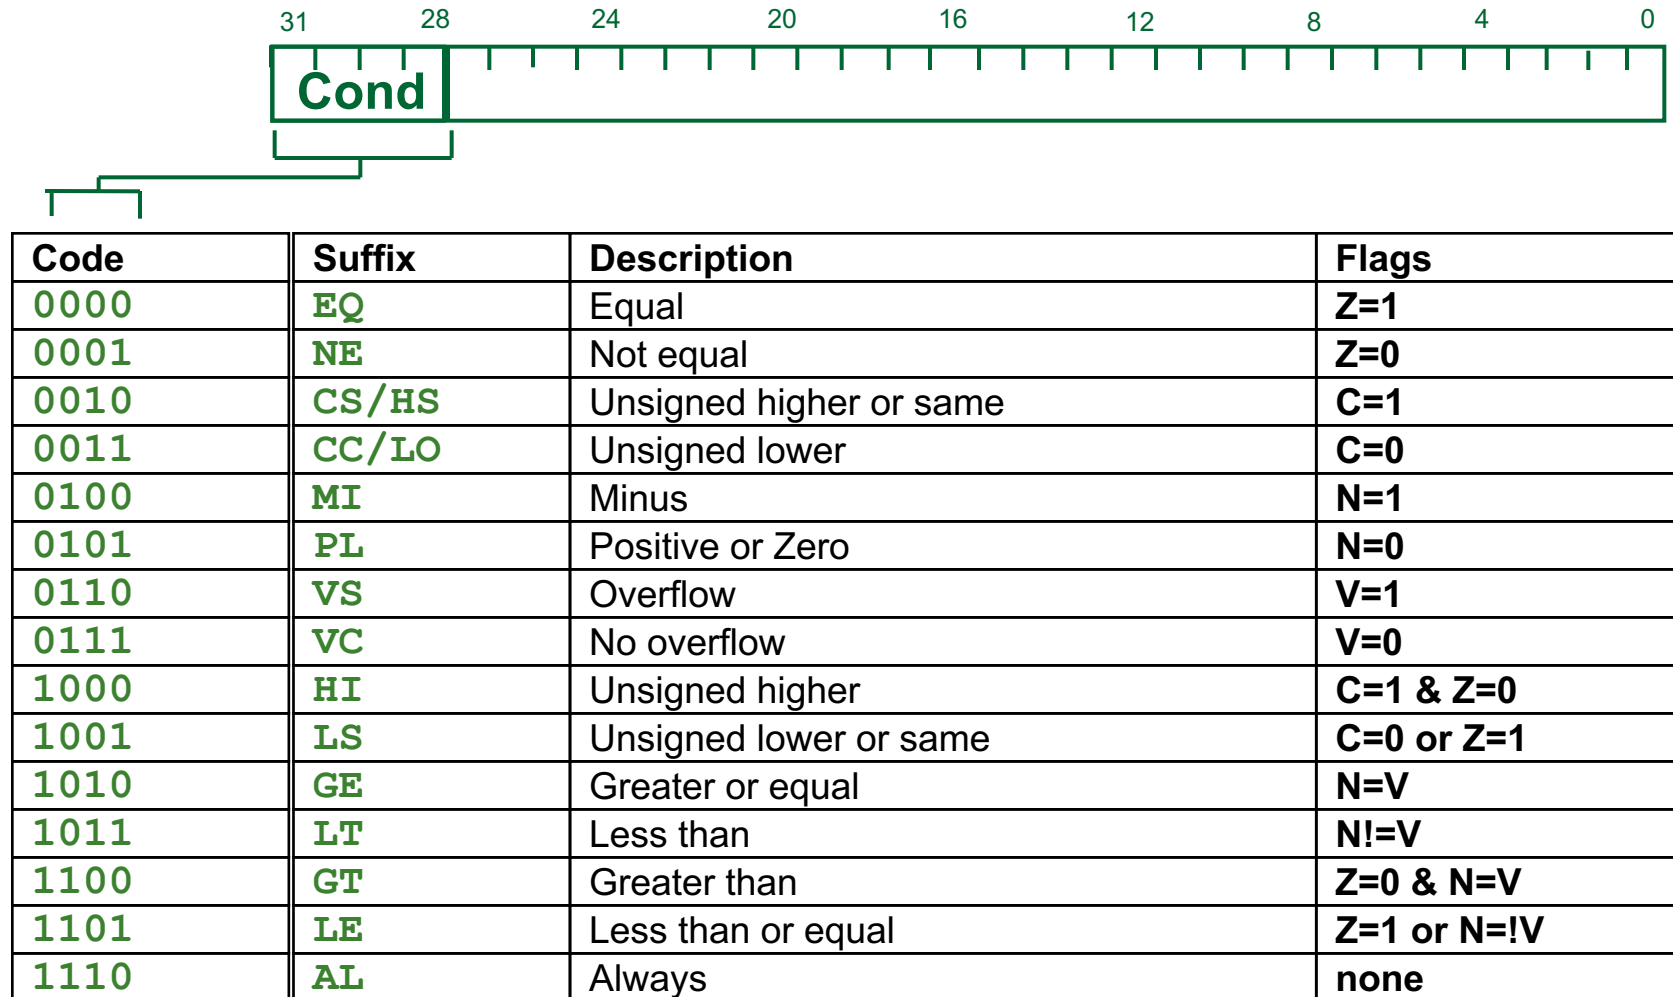
C : Carry

V : Overflow

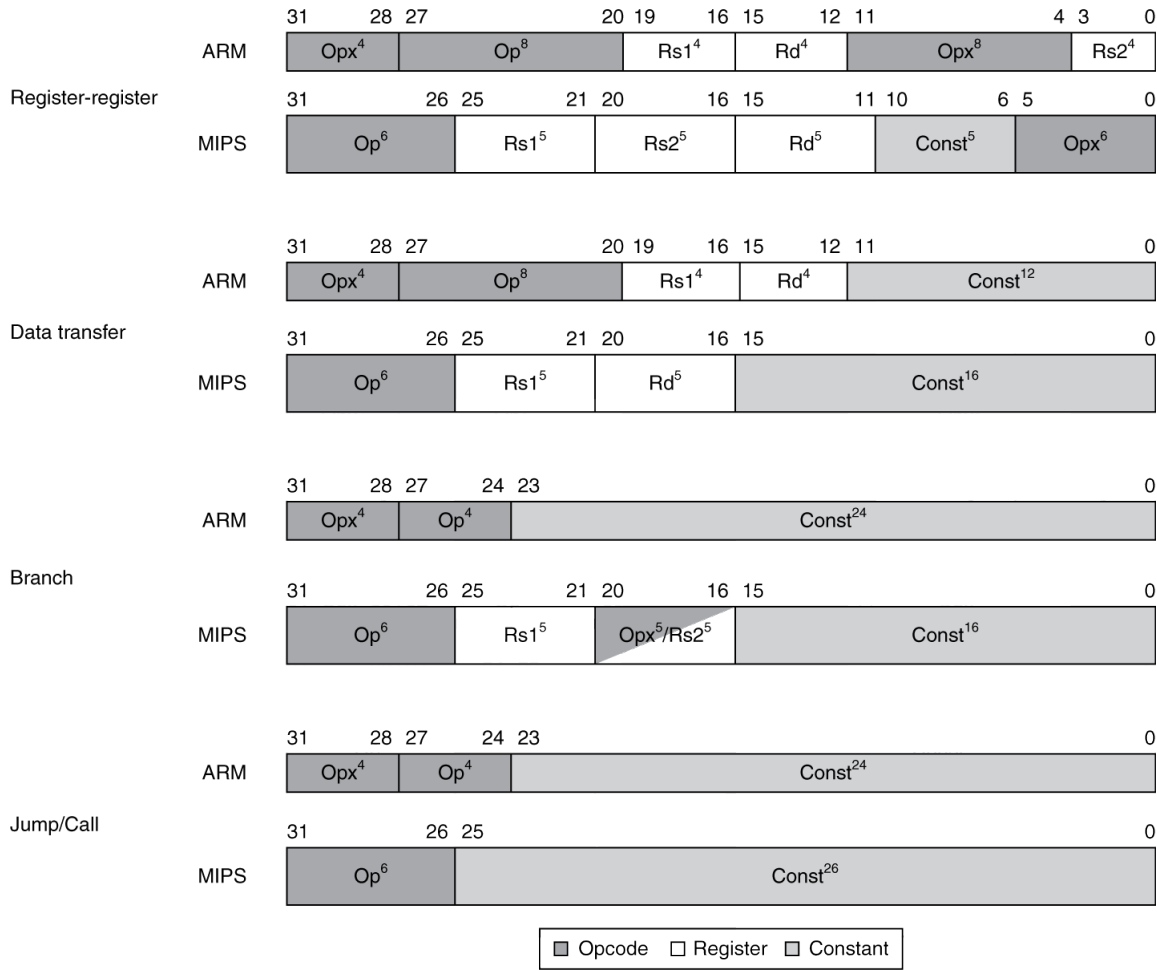
Conditional Execution

- All instructions contain a condition field which determines whether the CPU will execute them.
 - Allows very dense in-line code, without branches.
 - The time penalty of not executing several conditional instructions is frequently less than overhead of the branch or subroutine call that would otherwise be needed.
 - For example an add instruction takes the form:
 - `ADD r0,r1,r2` ; `r0 = r1 + r2`
 - To execute this only if the zero flag is set:
 - `ADDEQ r0,r1,r2` ; If zero flag set then...
; ... `r0 = r1 + r2`

The Condition Field



Instruction Encoding



The Intel x86 ISA

- Evolution with backward compatibility
 - ❑ 8080 (1974): 8-bit microprocessor
 - Accumulator, plus 3 index-register pairs
 - ❑ 8086 (1978): 16-bit extension to 8080
 - Complex instruction set (CISC)
 - ❑ 8087 (1980): floating-point coprocessor
 - Adds FP instructions and register stack
 - ❑ 80286 (1982): 24-bit addresses, MMU
 - Segmented memory mapping and protection
 - ❑ 80386 (1985): 32-bit extension (now IA-32)
 - Additional addressing modes and operations
 - Paged memory mapping as well as segments

The Intel x86 ISA

■ Further evolution...

- ❑ i486 (1989): pipelined, on-chip caches and FPU
 - Compatible competitors: AMD, Cyrix, ...
- ❑ Pentium (1993): superscalar, 64-bit datapath
 - Later versions added MMX (Multi-Media eXtension) instructions
 - The infamous FDIV bug
- ❑ Pentium Pro (1995), Pentium II (1997)
 - New microarchitecture (see Colwell, *The Pentium Chronicles*)
- ❑ Pentium III (1999)
 - Added SSE (Streaming SIMD Extensions) and associated registers (128 bits)
- ❑ Pentium 4 (2001)
 - New microarchitecture
 - Added SSE2 instructions

The Intel x86 ISA

■ And further...

- AMD64 (2003): extended architecture to 64 bits
- EM64T – Extended Memory 64 Technology (2004)
 - AMD64 adopted by Intel (with refinements)
 - Added SSE3 instructions
- Intel Core (2006)
 - Added SSE4 instructions, virtual machine support
- AMD64 (announced 2007): SSE5 instructions
 - Intel declined to follow, instead...
- Advanced Vector Extension (announced 2008)
 - Longer SSE registers (128bit->256 bits) , more instructions
- If Intel didn't extend with compatibility, its competitors would!
 - Technical elegance ≠ market success

IA-32 Registers

Name	31	0	Use
EAX			GPR 0
ECX			GPR 1
EDX			GPR 2
EBX			GPR 3
ESP			GPR 4
EBP			GPR 5
ESI			GPR 6
EDI			GPR 7
	CS		Code segment pointer
	SS		Stack segment pointer (top of stack)
	DS		Data segment pointer 0
	ES		Data segment pointer 1
	FS		Data segment pointer 2
	GS		Data segment pointer 3
EIP			Instruction pointer (PC)
EFLAGS			Condition codes

IA-32 has only 8 general purpose register vs. 32 in MIPS

IA-32 Typical Instructions

- Four major types of integer instructions:
 - ❑ Data movement including move, push, pop
 - ❑ Arithmetic and logical (destination register or memory)
 - ❑ Control flow (use of condition codes / flags)
 - ❑ String instructions, including string move and string compare

Instruction	Function
JE name	if equal(condition code) (EIP=name); EIP-128 ≤ name < EIP+128
JMP name	EIP=name
CALL name	SP=SP-4; M[SP]=EIP+5; EIP=name;
MOVW EBX,[EDI+45]	EBX=M[EDI+45]
PUSH ESI	SP=SP-4; M[SP]=ESI
POP EDI	EDI=M[SP]; SP=SP+4
ADD EAX,#6765	EAX= EAX+6765
TEST EDX,#42	Set condition code (flags) with EDX and 42
MOVSL	M[EDI]=M[ESI]; EDI=EDI+4; ESI=ESI+4

FIGURE 2.43 Some typical IA-32 instructions and their functions. A list of frequent operations appears in Figure 2.44. The CALL saves the EIP of the next instruction on the stack. (EIP is the Intel PC.)

1. IA-32: Two or three-operand operation vs. MIPS: three-operand operation
2. IA-32: Register-memory vs. MIPS: register-register

IA-32 Addressing Mode

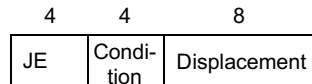
Mode	Description	Register restrictions	MIPS equivalent
Register Indirect	Address is in a register.	not ESP or EBP	lw \$s0,0(\$s1)
Based mode with 8- or 32-bit displacement	Address is contents of base register plus displacement.	not ESP or EBP	lw \$s0,100(\$s1) # ≤16-bit displacement
Base plus scaled Index	The address is $\text{Base} + (2^{\text{Scale}} \times \text{Index})$ where Scale has the value 0, 1, 2, or 3.	Base: any GPR Index: not ESP	mul \$t0,\$s2,4 add \$t0,\$t0,\$s1 lw \$s0,0(\$t0)
Base plus scaled Index with 8- or 32-bit displacement	The address is $\text{Base} + (2^{\text{Scale}} \times \text{Index}) + \text{displacement}$ where Scale has the value 0, 1, 2, or 3.	Base: any GPR Index: not ESP	mul \$t0,\$s2,4 add \$t0,\$t0,\$s1 lw \$s0,100(\$t0) # ≤16-bit displacement

Move EAX, [EBX+ EPI*4 + 100]

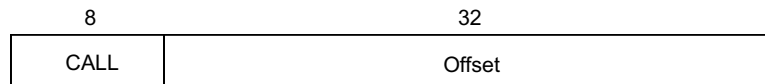
- IA-32 has more addressing modes than MIPS

IA-32 instruction Formats

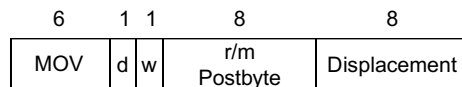
a. JE EIP + displacement



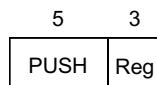
b. CALL



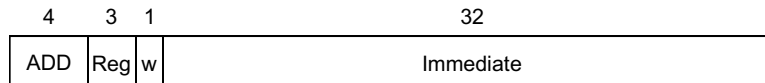
c. MOV EBX, [EDI + 45]



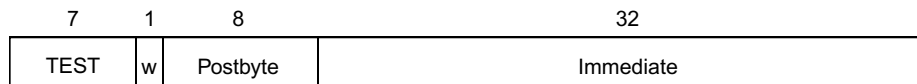
d. PUSH ESI



e. ADD EAX, #6765



f. TEST EDX, #42



IA-32 variable-length encoding vs. MIPS fixed-length encoding

Implementing IA-32

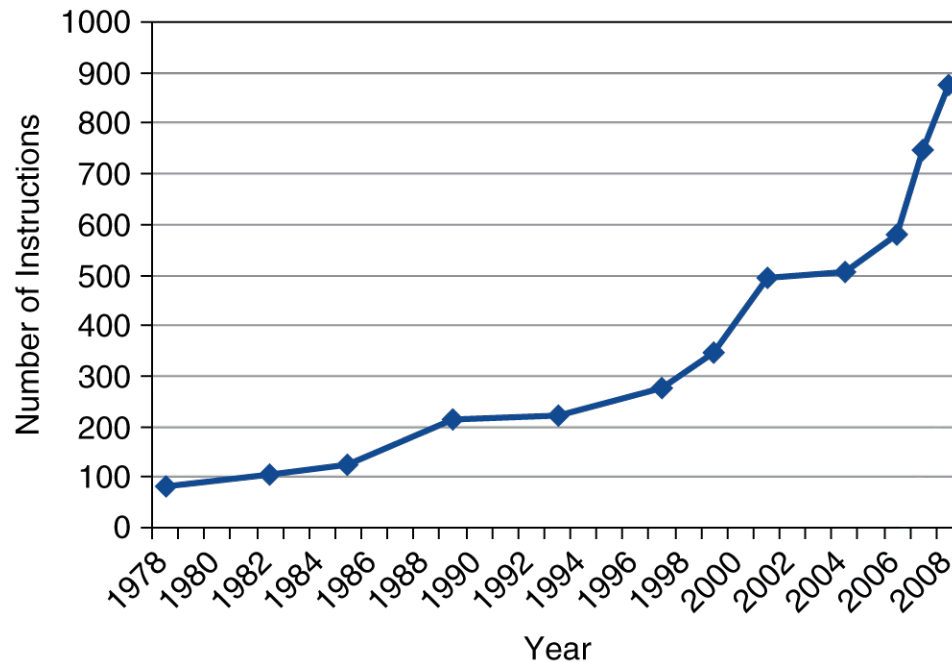
- Complex instruction set makes implementation difficult
 - Hardware translates instructions to simpler microoperations
 - Simple instructions: 1–1
 - Complex instructions: 1–many
 - Microengine similar to RISC
 - Market share makes this economically viable
- Comparable performance to RISC
 - Compilers avoid complex instructions

Fallacies

- Powerful instruction \Rightarrow higher performance
 - ❑ Fewer instructions required
 - ❑ But complex instructions are hard to implement
 - May slow down all instructions, including simple ones
 - ❑ Compilers are good at making fast code from simple instructions
- Use assembly code for high performance
 - ❑ But modern compilers are better at dealing with modern processors
 - ❑ More lines of code \Rightarrow more errors and less productivity

Fallacies

- Backward compatibility \Rightarrow instruction set doesn't change
 - But they do accrete more instructions



x86 instruction set

Pitfalls

- Sequential words are not at sequential addresses
 - Increment by 4, not by 1!
- Keeping a pointer to a local variable after procedure returns
 - Pointer becomes invalid when stack popped

Concluding Remarks

- Design principles
 1. Simplicity favors regularity
 2. Smaller is faster
 3. Make the common case fast
 4. Good design demands good compromises
- Layers of software/hardware
 - Compiler, assembler, hardware
- MIPS: typical of RISC ISAs
 - c.f. x86

Concluding Remarks

- Measure MIPS instruction executions in benchmark programs
 - Consider making the common case fast
 - Consider compromises

Instruction class	MIPS examples	SPEC2006 Int	SPEC2006 FP
Arithmetic	add, sub, addi	16%	48%
Data transfer	lw, sw, lb, lbu, lh, lhu, sb, lui	35%	36%
Logical	and, or, nor, andi, ori, sll, srl	12%	4%
Cond. Branch	beq, bne, slt, slti, sltiu	34%	8%
Jump	j, jr, jal	2%	0%