

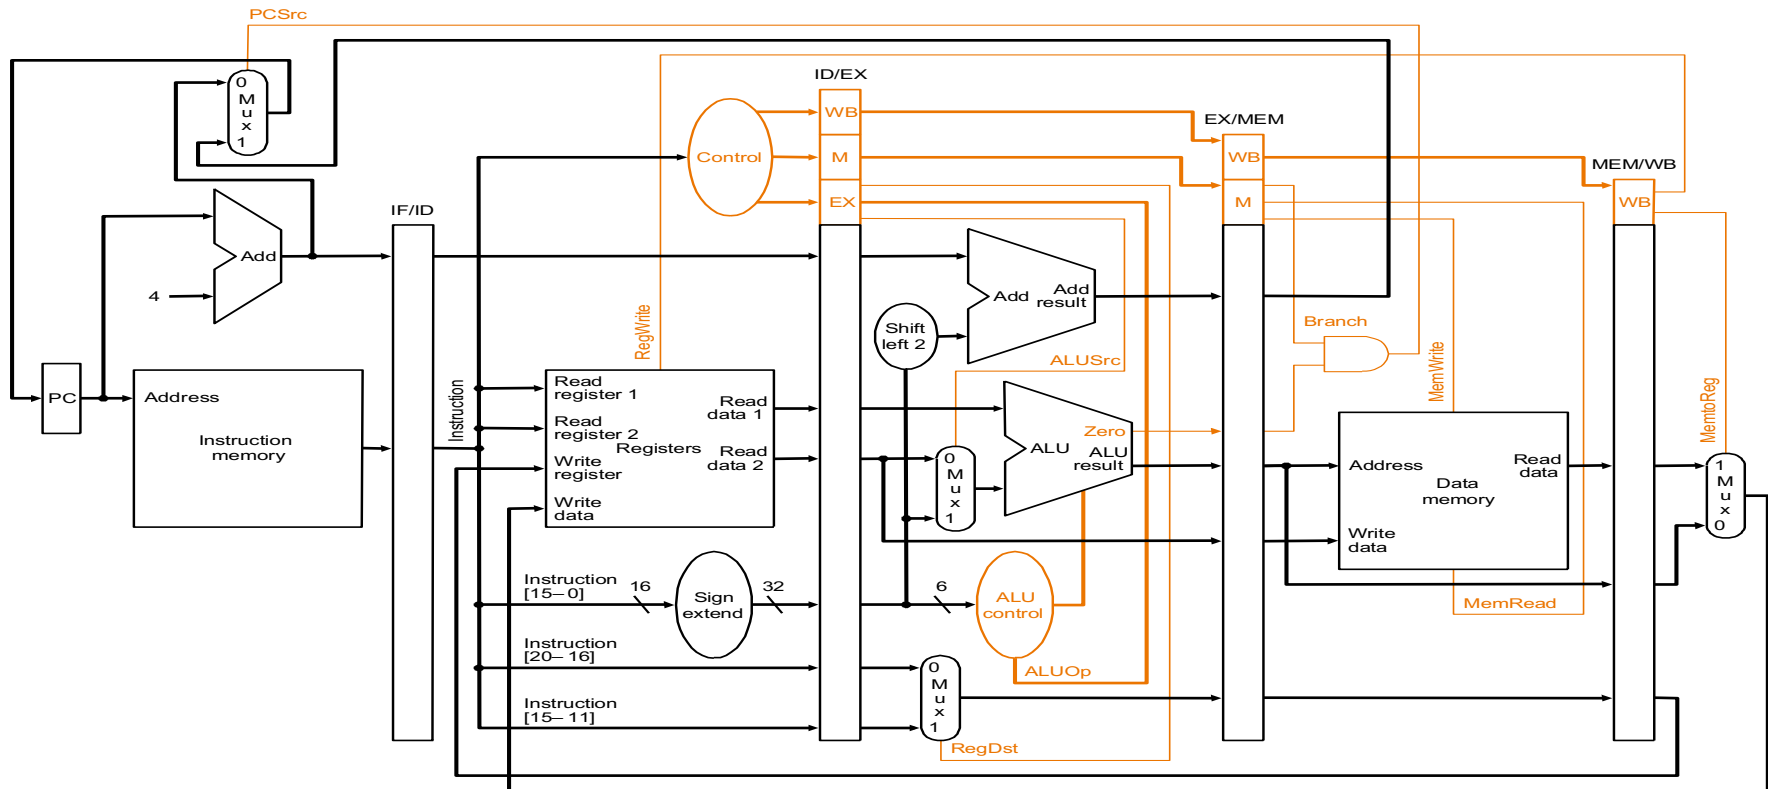
---

# Lecture 7: Pipelining (II)

---

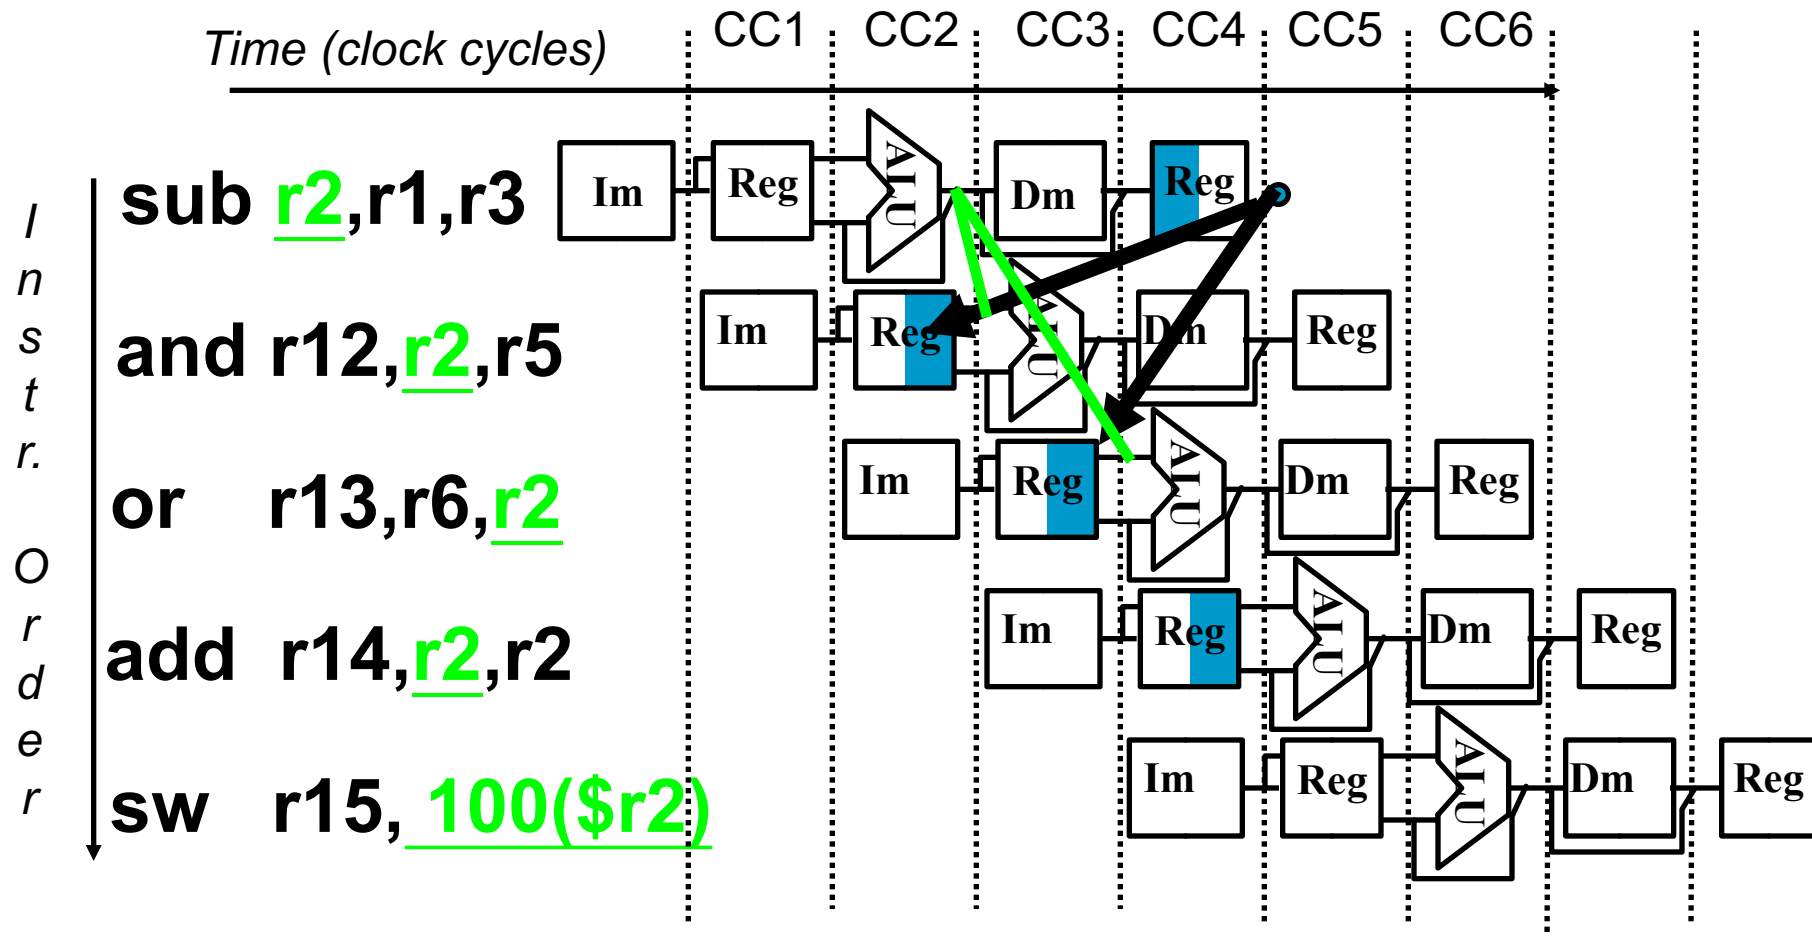
1. Data Hazards and Forwarding
2. Control Hazard Solution
3. How to Handle Exception

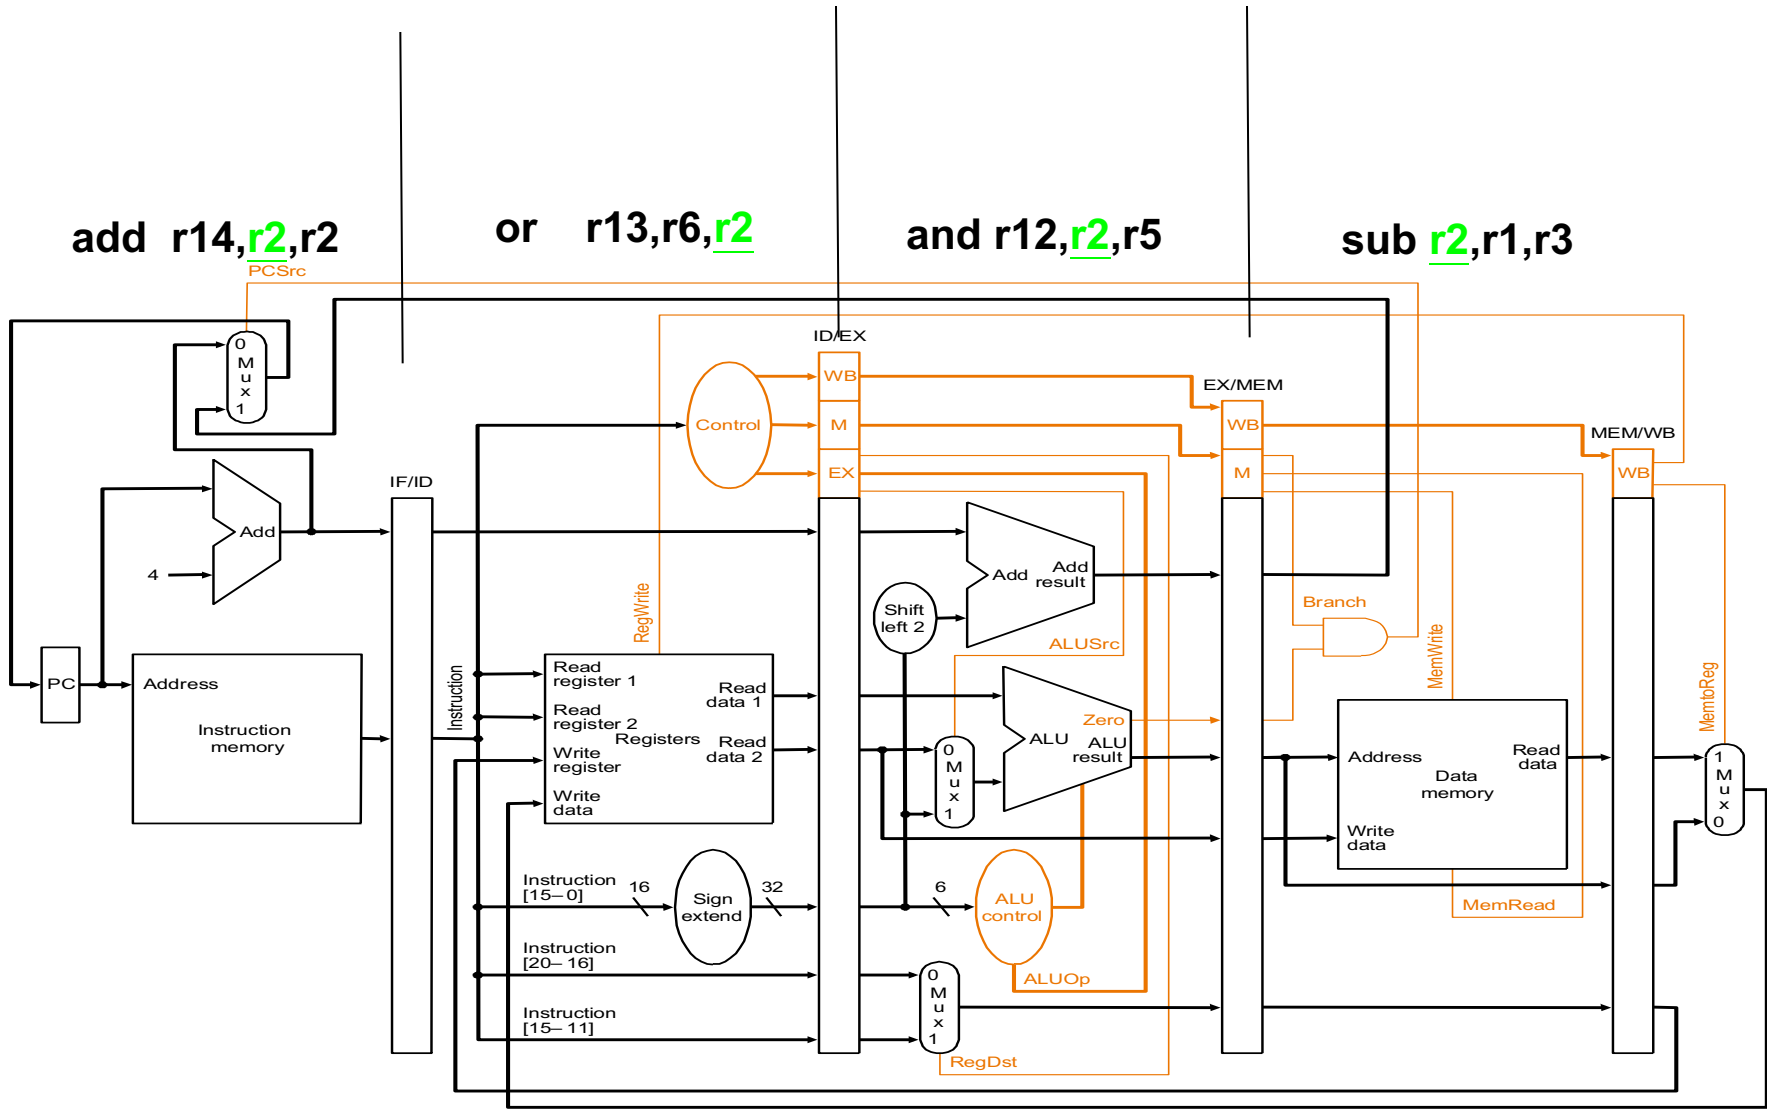
# Datapath with Control



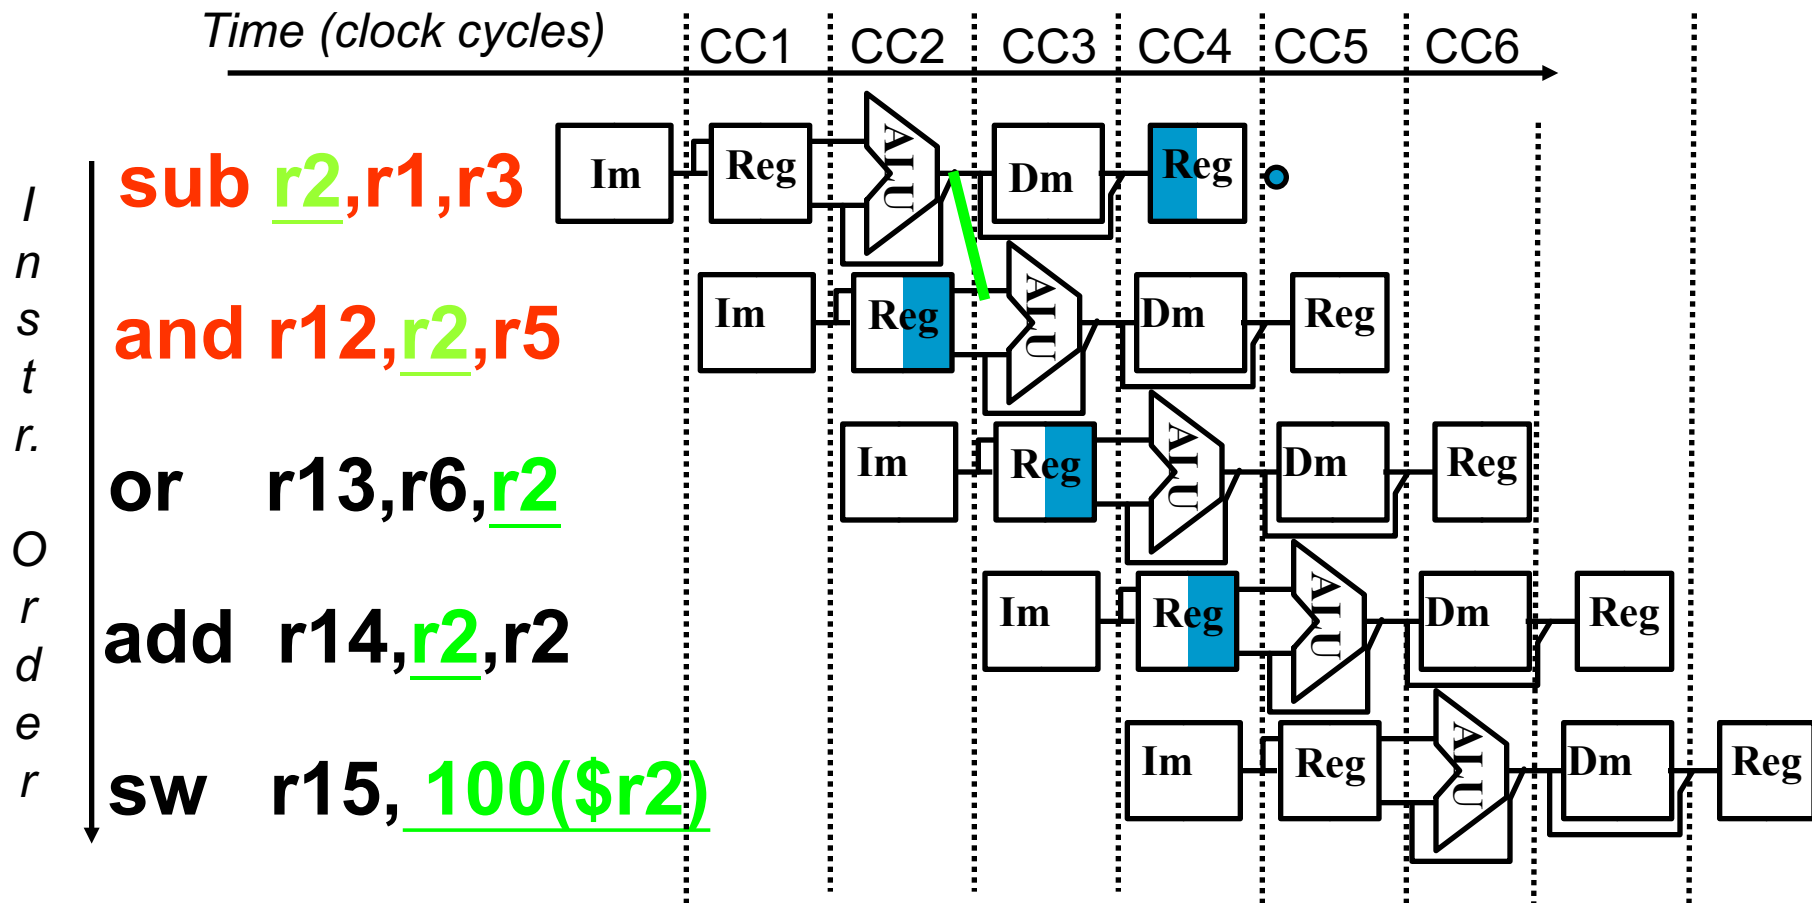
1. How to implement “data forwarding”?
2. How to detect load-use hazard? How to stall pipeline?
3. How to resolve branch in the decode stage?
4. How to flush pipeline?

# Data Dependence Detection & Forwarding





## How to detect dependency between (sub, and)?



1a: EX/MEM. RegisterRd = ID/EX.RegisterRs

1b: EX/MEM. RegisterRd = ID/EX.RegisterRt

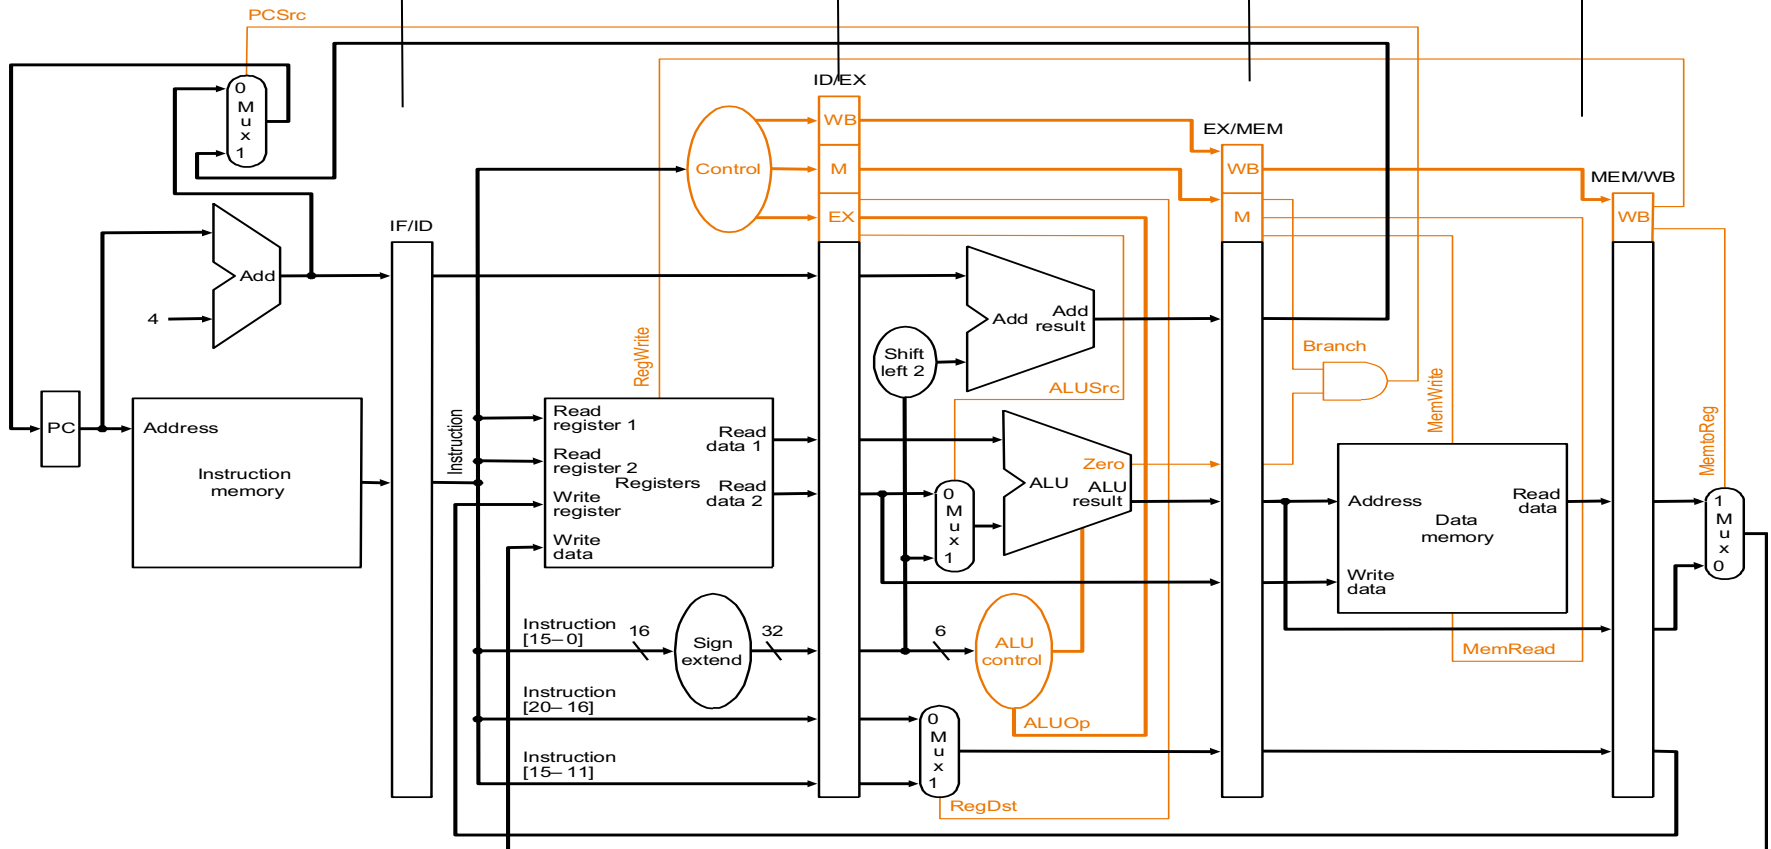
sw r15, 100(\$r2)

add r14, r2, r2

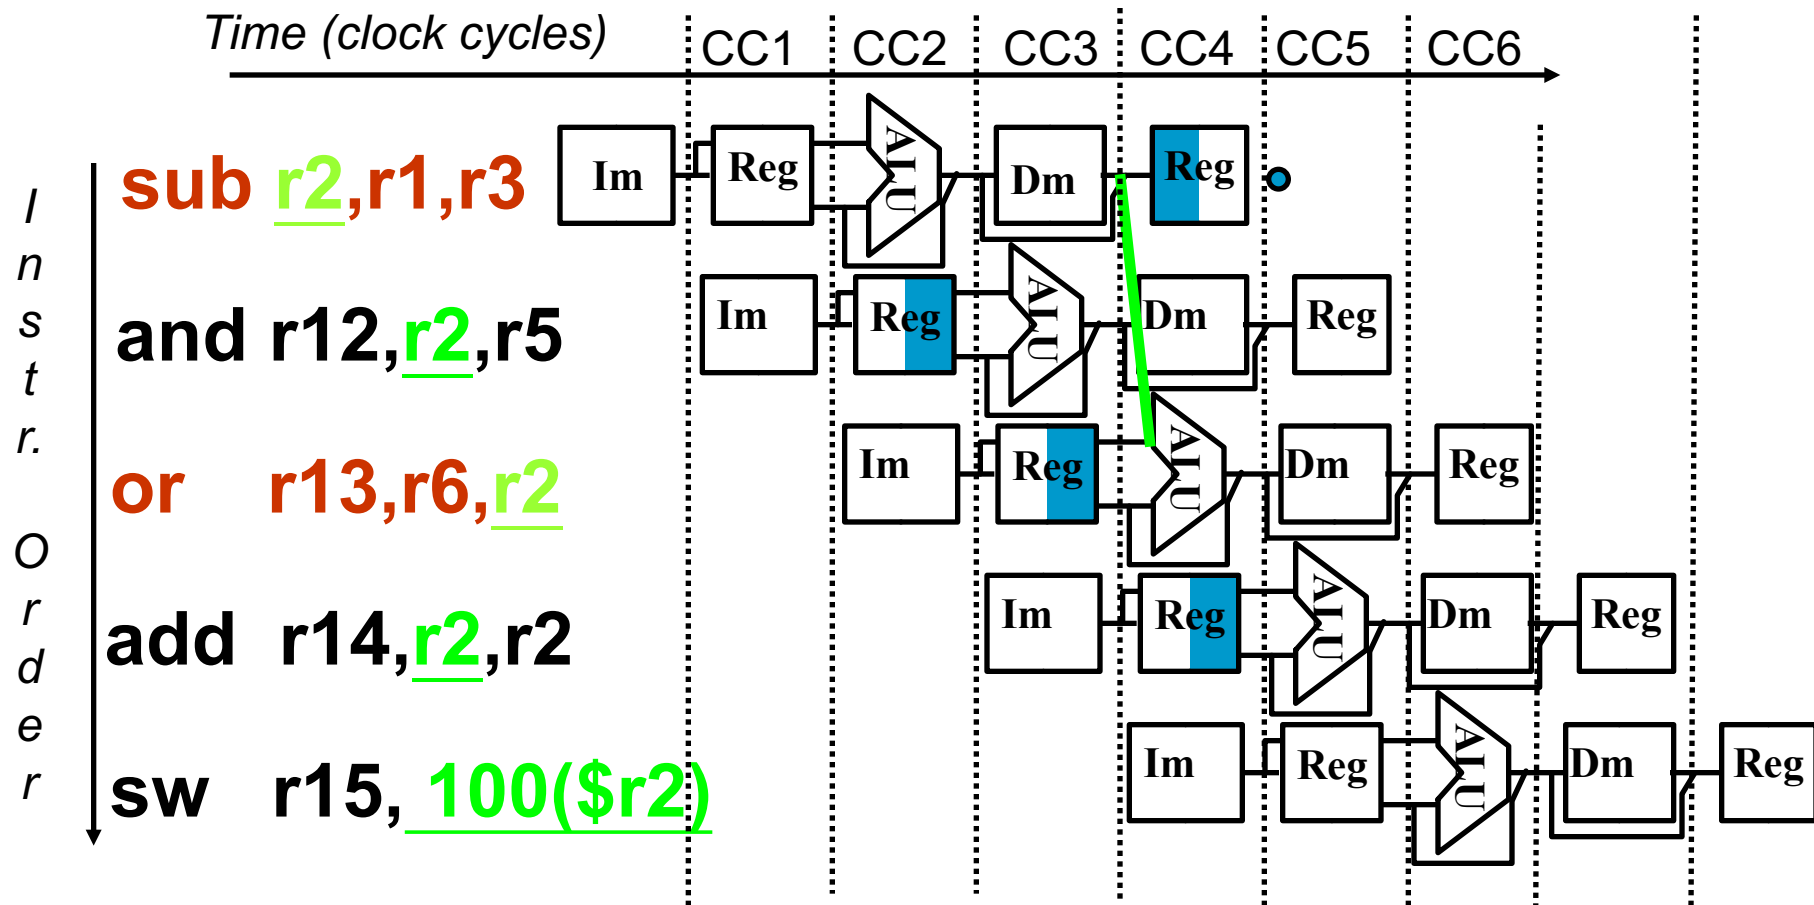
or r13, r6, r2

and r12, r2, r5

sub r2, r1, r3



# How to detect dependency between (sub, or)?



2a: MEM//WB.RegisterRd = ID/EX.RegisterRs (sub & or)

2b: MEM/WB.RegisterRd = ID/EX.RegisterRt

# Data Dependence Detection (cont.)

---

## ■ Hazard conditions:

- 1a: EX/MEM. RegisterRd = ID/EX.RegisterRs (sub & and) } EX hazard
- 1b: EX/MEM. RegisterRd = ID/EX.RegisterRt
- 2a: MEM/WB.RegisterRd = ID/EX.RegisterRs (sub & or) } MEM hazard
- 2b: MEM/WB.RegisterRd = ID/EX.RegisterRt
- RegWrite signal of WB Control field
  - EX/MEM.RegWrite, MEM/WB.RegWrite
- EX/MEM.RegisterRd  $\neq$  \$0
- MEM/WB.RegisterRd  $\neq$  \$0

How to forward data?

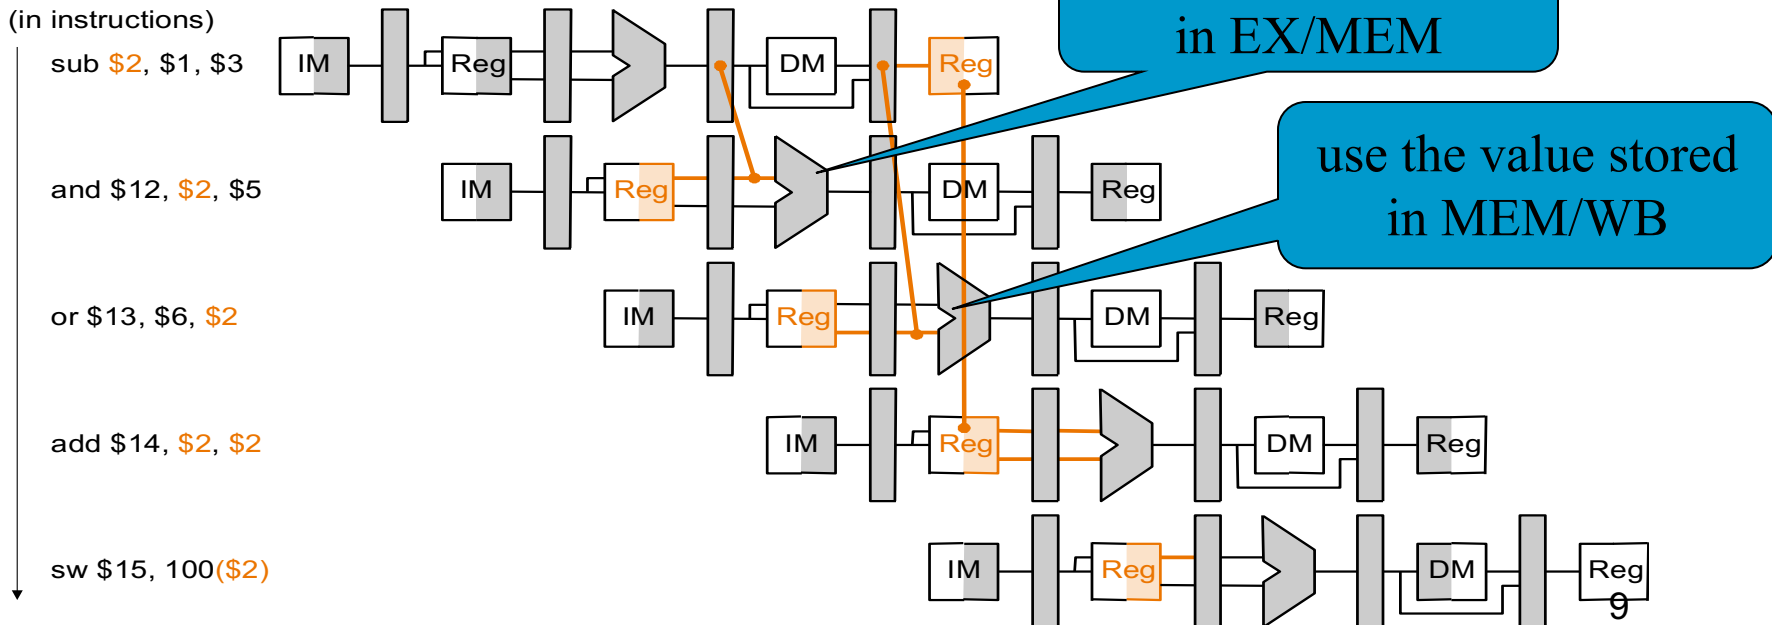


# Resolving Hazards by Forwarding

- Use the value in pipeline registers rather than waiting for the WB stage to write the register file.
  - EX/MEM.Aluout
  - MEM/WB.Aluout

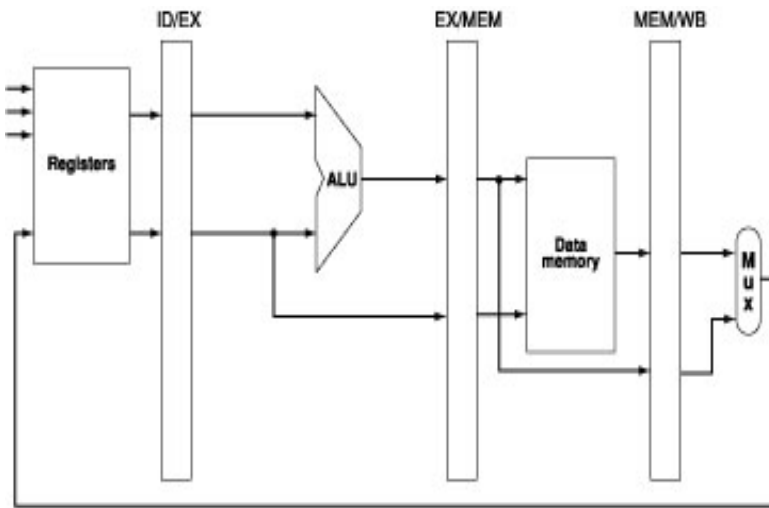
|                         | Time (in clock cycles) → |      |      |      |        |      |      |      |      |
|-------------------------|--------------------------|------|------|------|--------|------|------|------|------|
|                         | CC 1                     | CC 2 | CC 3 | CC 4 | CC 5   | CC 6 | CC 7 | CC 8 | CC 9 |
| Value of register \$2 : | 10                       | 10   | 10   | 10   | 10/-20 | -20  | -20  | -20  | -20  |
| Value of EX/MEM :       | X                        | X    | X    | -20  | X      | X    | X    | X    | X    |
| Value of MEM/WB :       | X                        | X    | X    | X    | -20    | X    | X    | X    | X    |

Program  
execution order  
(in instructions)

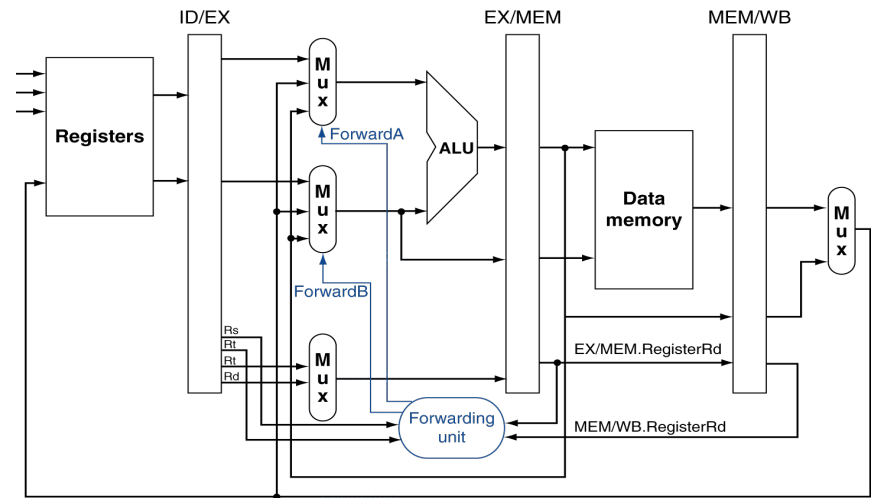


# Forwarding Logic

- Forwarding: input to ALU from any pipe reg.
  - Add multiplexors to ALU input
  - Forwarding Control will be in EX



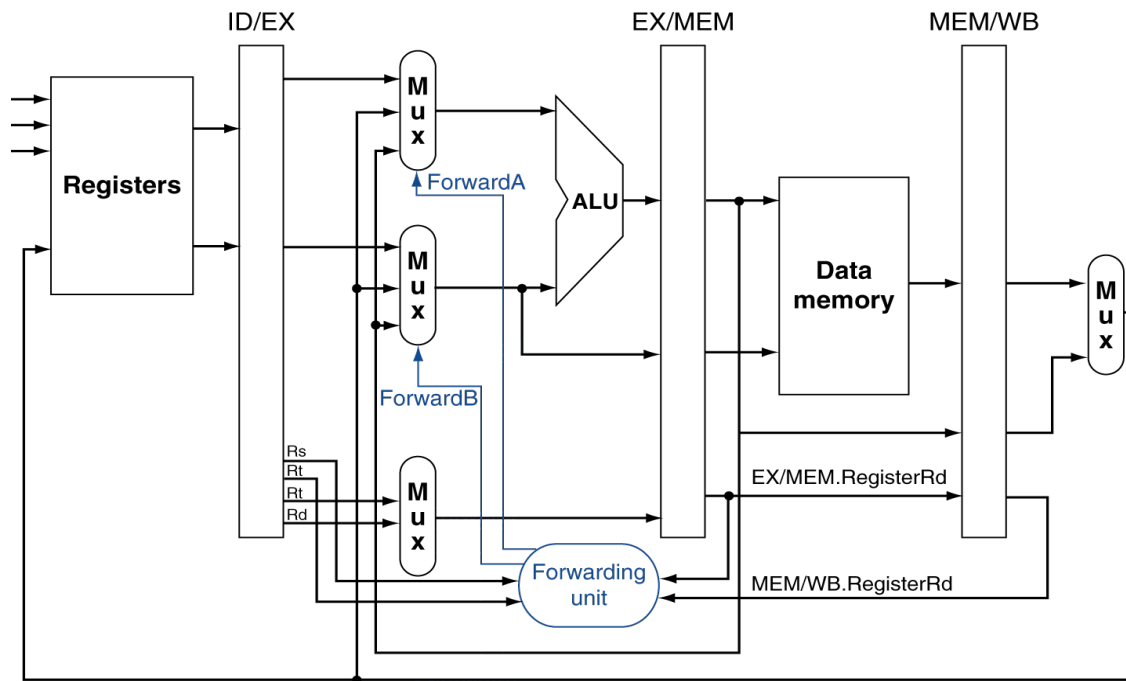
a. No forwarding



b. With forwarding

# Forwarding Control

| Mux control   | Source | Explanation  |
|---------------|--------|--|
| ForwardA = 00 | ID/EX  | The first ALU operand comes from the register file.                            |
| ForwardA = 10 | EX/MEM | The first ALU operand is forwarded from the prior ALU result.                  |
| ForwardA = 01 | MEM/WB | The first ALU operand is forwarded from data memory or an earlier ALU result.  |
| ForwardB = 00 | ID/EX  | The second ALU operand comes from the register file.                           |
| ForwardB = 10 | EX/MEM | The second ALU operand is forwarded from the prior ALU result.                 |
| ForwardB = 01 | MEM/WB | The second ALU operand is forwarded from data memory or an earlier ALU result. |



b. With forwarding

# Forwarding Control

| Mux control   | Source | Explanation  |
|---------------|--------|--|
| ForwardA = 00 | ID/EX  | The first ALU operand comes from the register file.                            |
| ForwardA = 10 | EX/MEM | The first ALU operand is forwarded from the prior ALU result.                  |
| ForwardA = 01 | MEM/WB | The first ALU operand is forwarded from data memory or an earlier ALU result.  |
| ForwardB = 00 | ID/EX  | The second ALU operand comes from the register file.                           |
| ForwardB = 10 | EX/MEM | The second ALU operand is forwarded from the prior ALU result.                 |
| ForwardB = 01 | MEM/WB | The second ALU operand is forwarded from data memory or an earlier ALU result. |

## 1. EX hazard

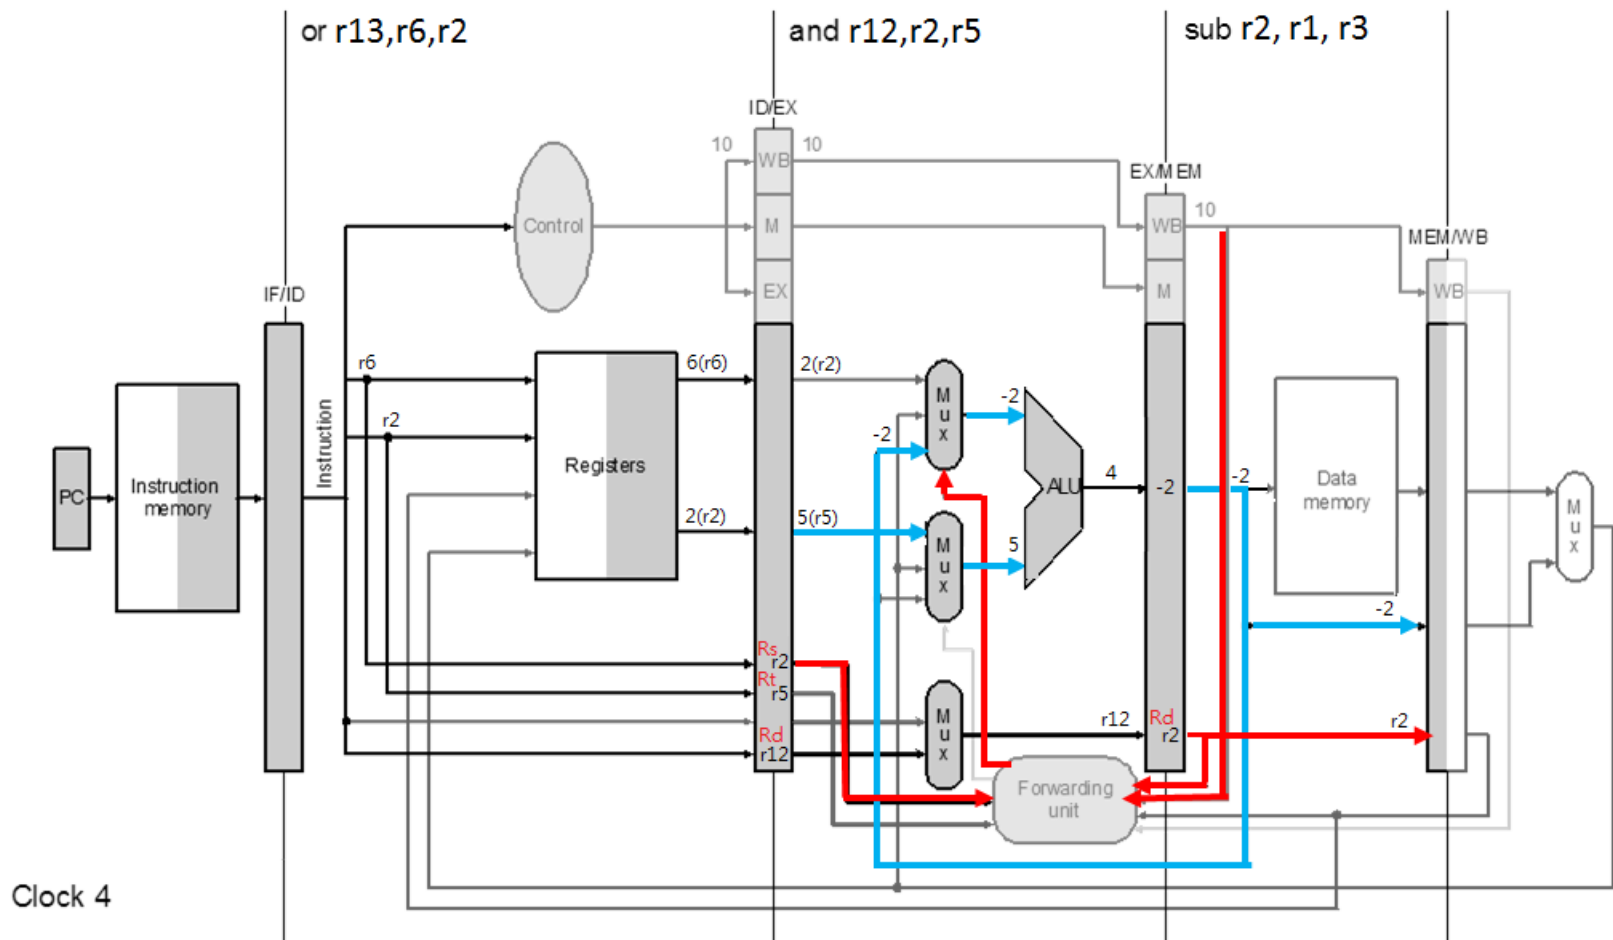
if (EX/MEM.RegWrite  
and (EX/MEM.RegisterRd  $\neq$  0)  
and (EX/MEM.RegisterRd=ID/EX.RegisterRs))  
ForwardA = 10

if (EX/MEM.RegWrite  
and (EX/MEM.RegisterRd  $\neq$  0)  
and (EX/MEM.RegisterRd=ID/Ex.RegisterRt))  
ForwardB = 10

## 2. MEM hazard

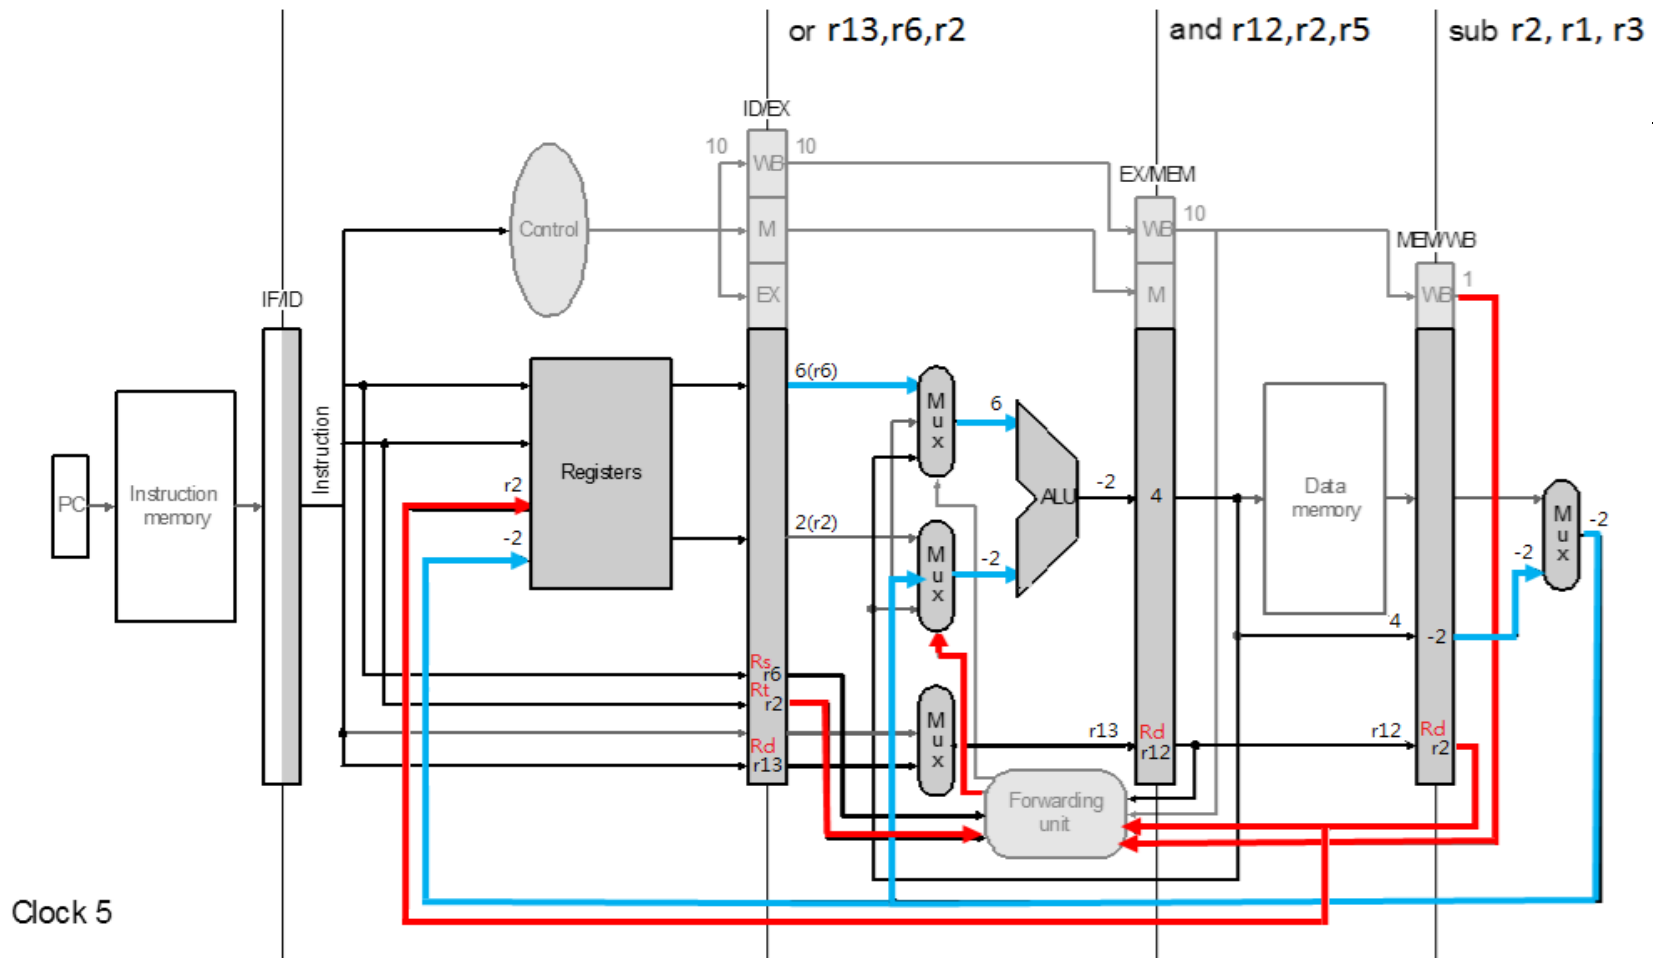
if (MEM/WB.RegWrite  
and (MEM/WB.RegisterRd  $\neq$  0)  
and (MEM/WB.RegisterRd=ID/Ex.RegisterRs))  
ForwardA = 01

if (MEM/WB.RegWrite  
and (MEM/WB.RegisterRd  $\neq$  0)  
and (MEM/WB.RegisterRd=ID/Ex.RegisterRt))  
ForwardB = 01



## 1. EX hazard

if (EX/MEM.RegWrite  
 and (EX/MEM.RegisterRd  $\neq$  0)  
 and (EX/MEM.RegisterRd=ID/EX.RegisterRs))  
 ForwardA = 10



## 2. MEM hazard

if (MEM/WB.RegWrite  
 and (MEM/WB.RegisterRd  $\neq$  0)  
 and (MEM/WB.RegisterRd=ID/Ex.RegisterRt))  
 ForwardB = 01

# Forwarding Control (cont.)

|       |                    |    |    |    |     |     |     |    |
|-------|--------------------|----|----|----|-----|-----|-----|----|
| inst1 | add \$1, \$1, \$2; | IF | ID | EX | MEM | WB  |     |    |
| inst2 | add \$1, \$1, \$3; |    | IF | ID | EX  | MEM | WB  |    |
| inst3 | add \$1, \$1, \$4; |    |    | IF | ID  | EX  | MEM | WB |

=> Which instruction should forward its results to instruction 3?

**MEM hazard condition becomes**

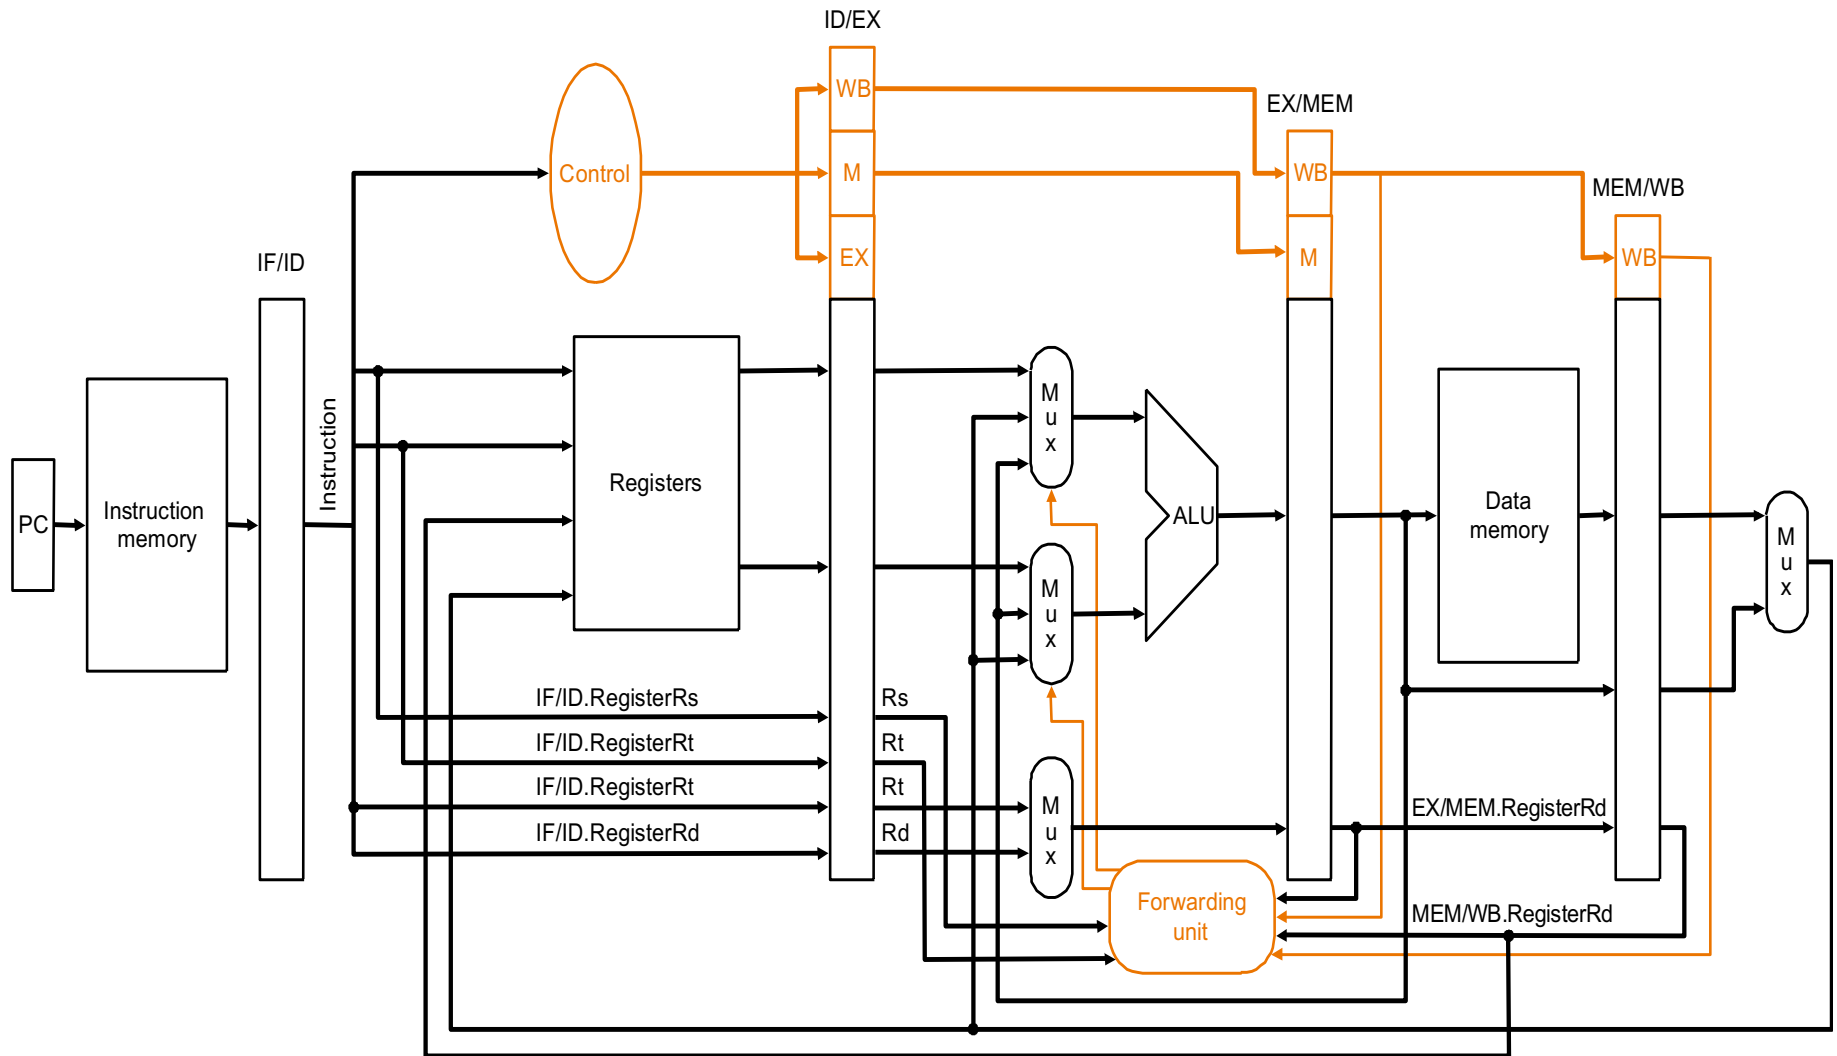
if (MEM/WB.RegWrite  
and (MEM/WB.RegisterRd  $\neq$  0)  
and (Ex/MEM.RegisterRd  $\neq$  ID/Ex.RegisterRs)  
and (MEM/WB.RegRd=ID/Ex.RegisterRs)) ForwardA = 01

if (MEM/WB.RegWrite  
and (MEM/WB.RegRd  $\neq$  0)  
and (Ex/MEM.RegisterRd  $\neq$  ID/Ex.RegisterRt)  
and (MEM/WB.RegRd=ID/Ex.RegisterRt)) ForwardB = 01

if EX/MEM  
Reg Rd  
= ID/EX  
Reg Rs(Rt)

會相衡

# Datapath with Forwarding





# Example

---

- Show how forwarding works with this instruction sequence (with dependencies highlighted):

sub ~~\$2~~, \$1, \$3  
and ~~\$4~~, ~~\$2~~, \$5  
or ~~\$4~~, ~~\$4~~, ~~\$2~~  
add \$9, ~~\$4~~, \$2

EX/MEM. Rd  
= IP/EX. Rs

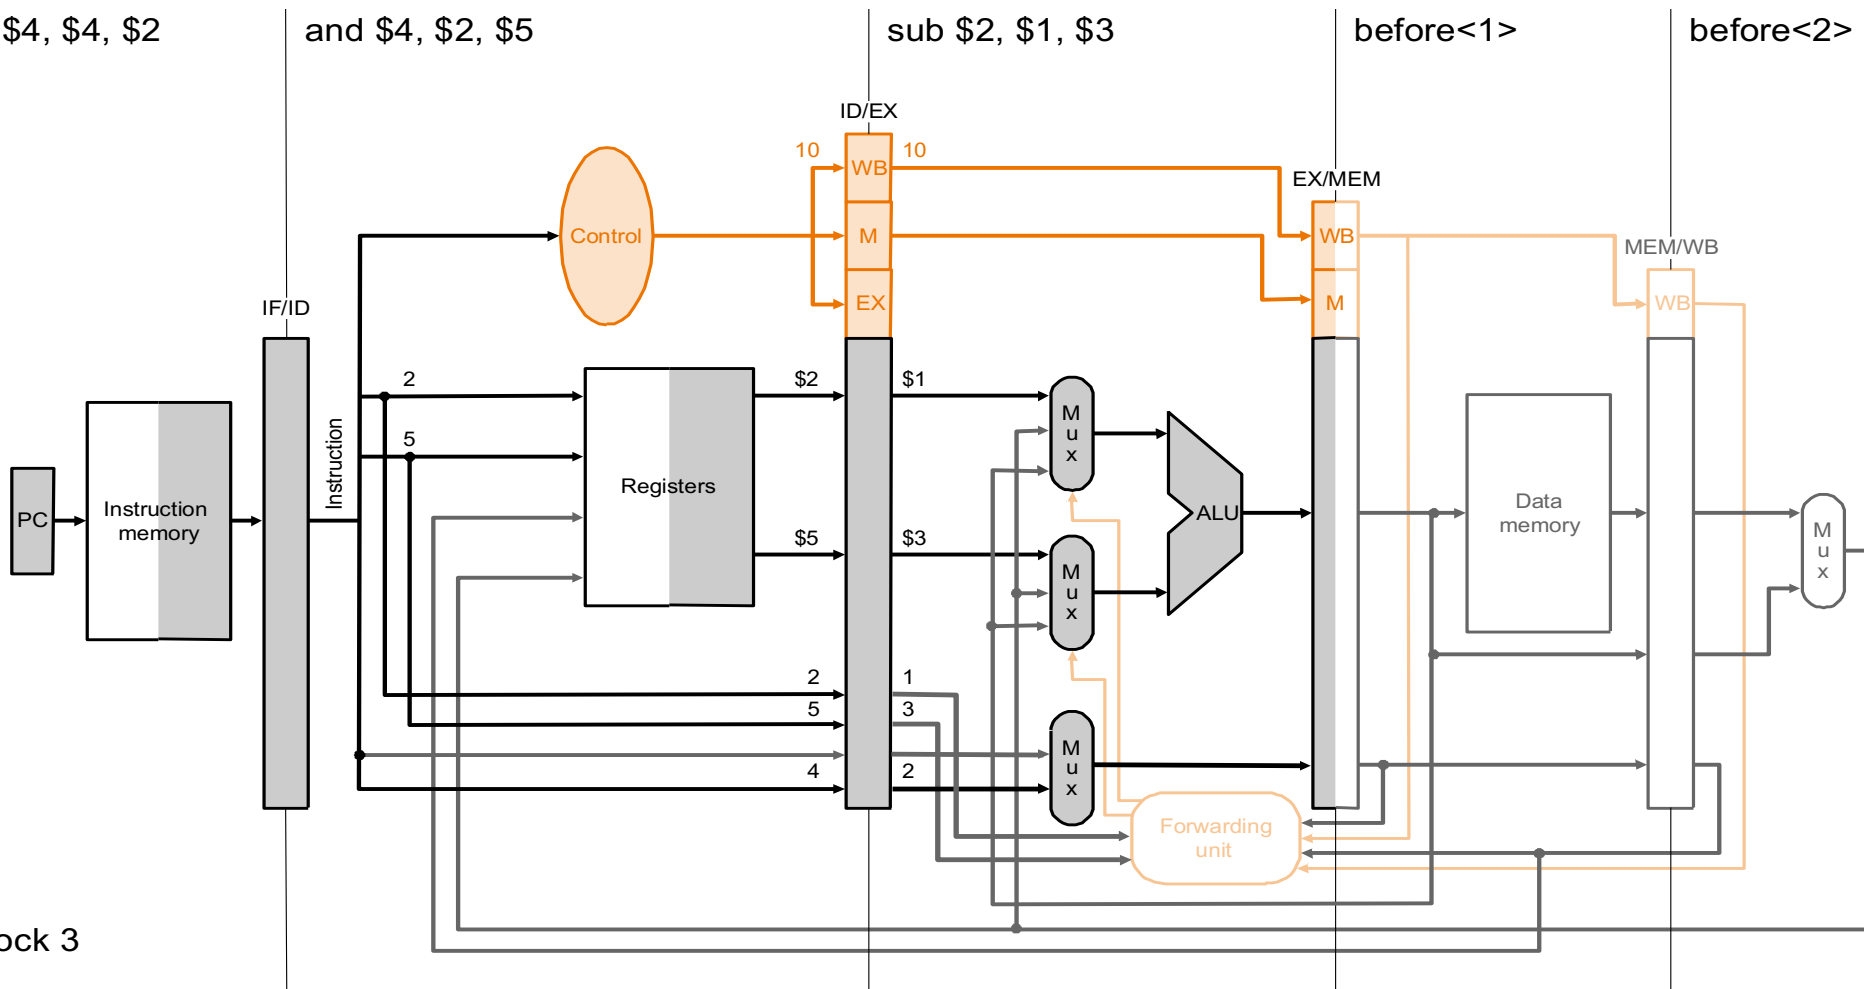
or \$4, \$4, \$2

and \$4, \$2, \$5

sub \$2, \$1, \$3

before<1>

before<2>



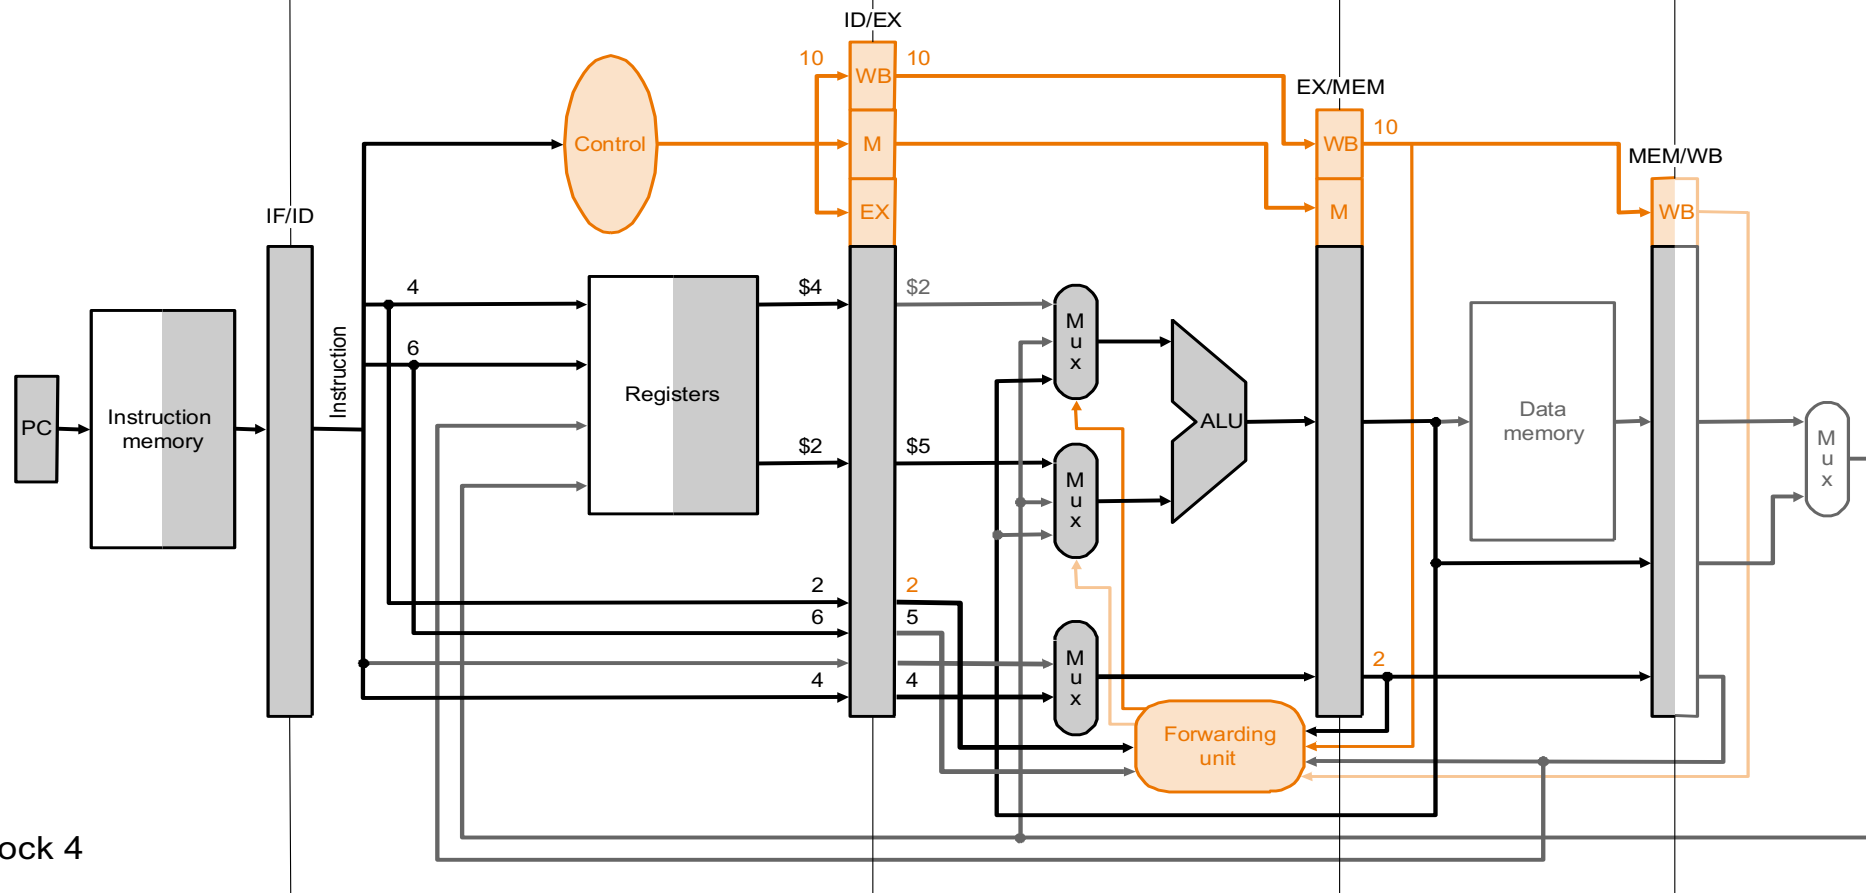
add \$9, \$4, \$2

or \$4, \$4, \$2

and \$4, \$2, \$5

sub \$2, ...

before<1>



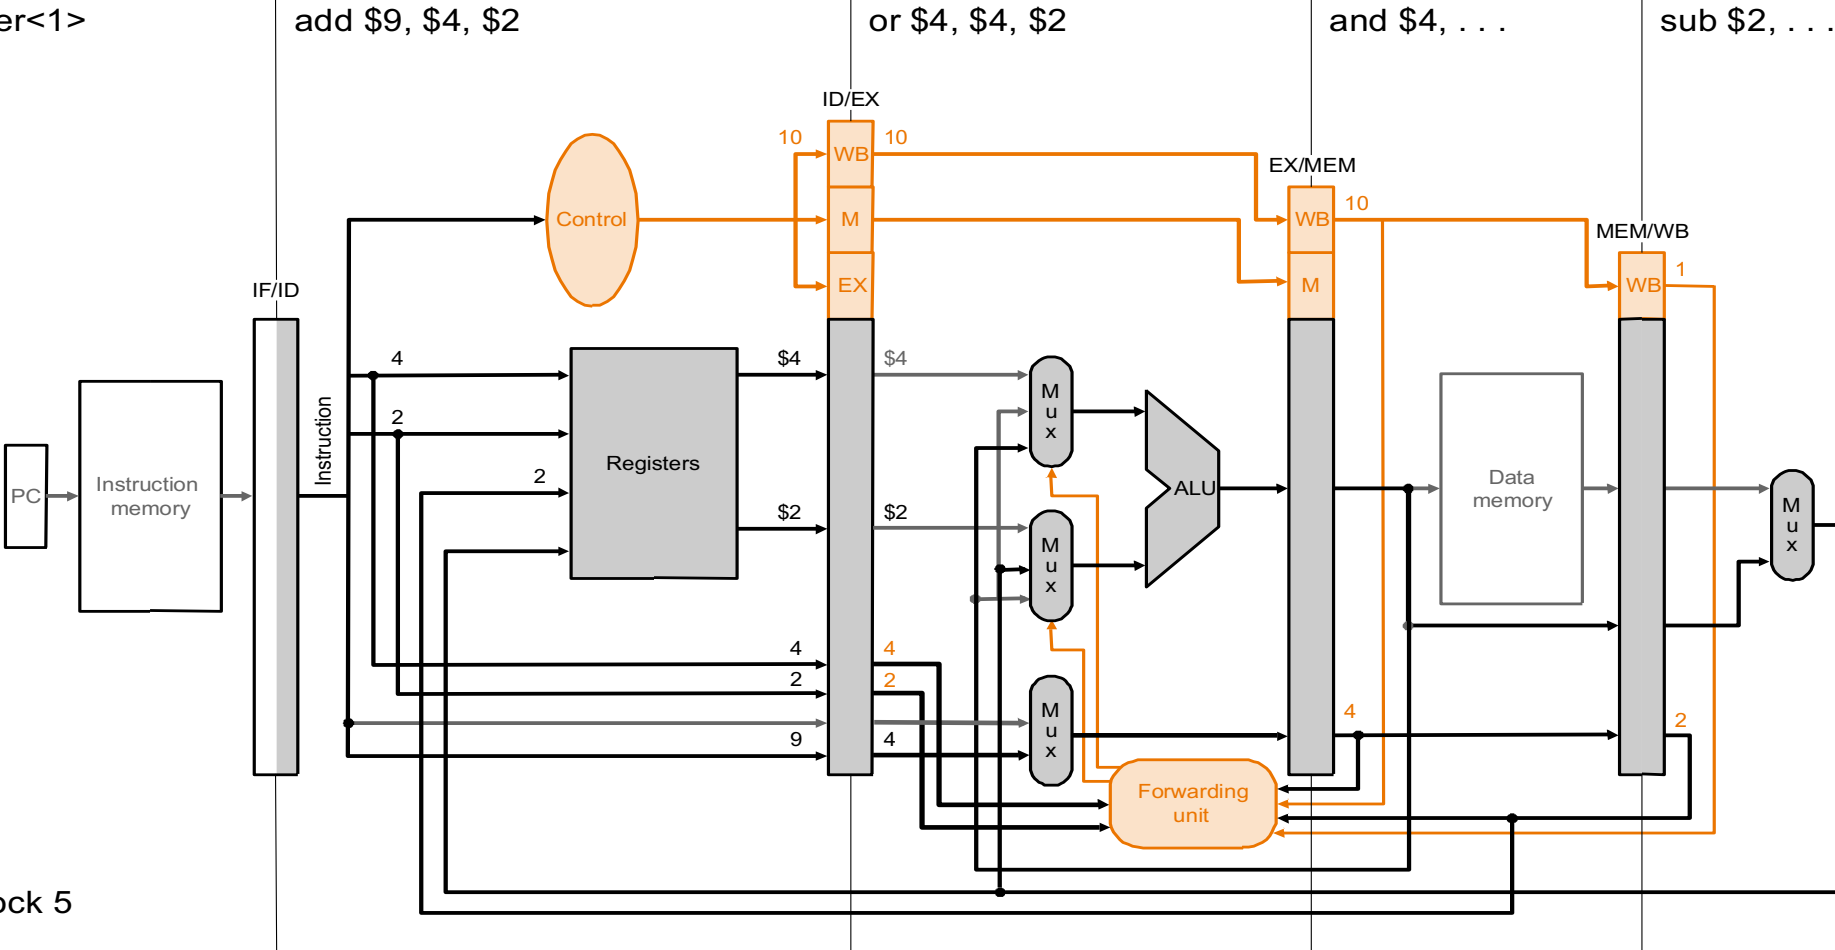
after<1>

add \$9, \$4, \$2

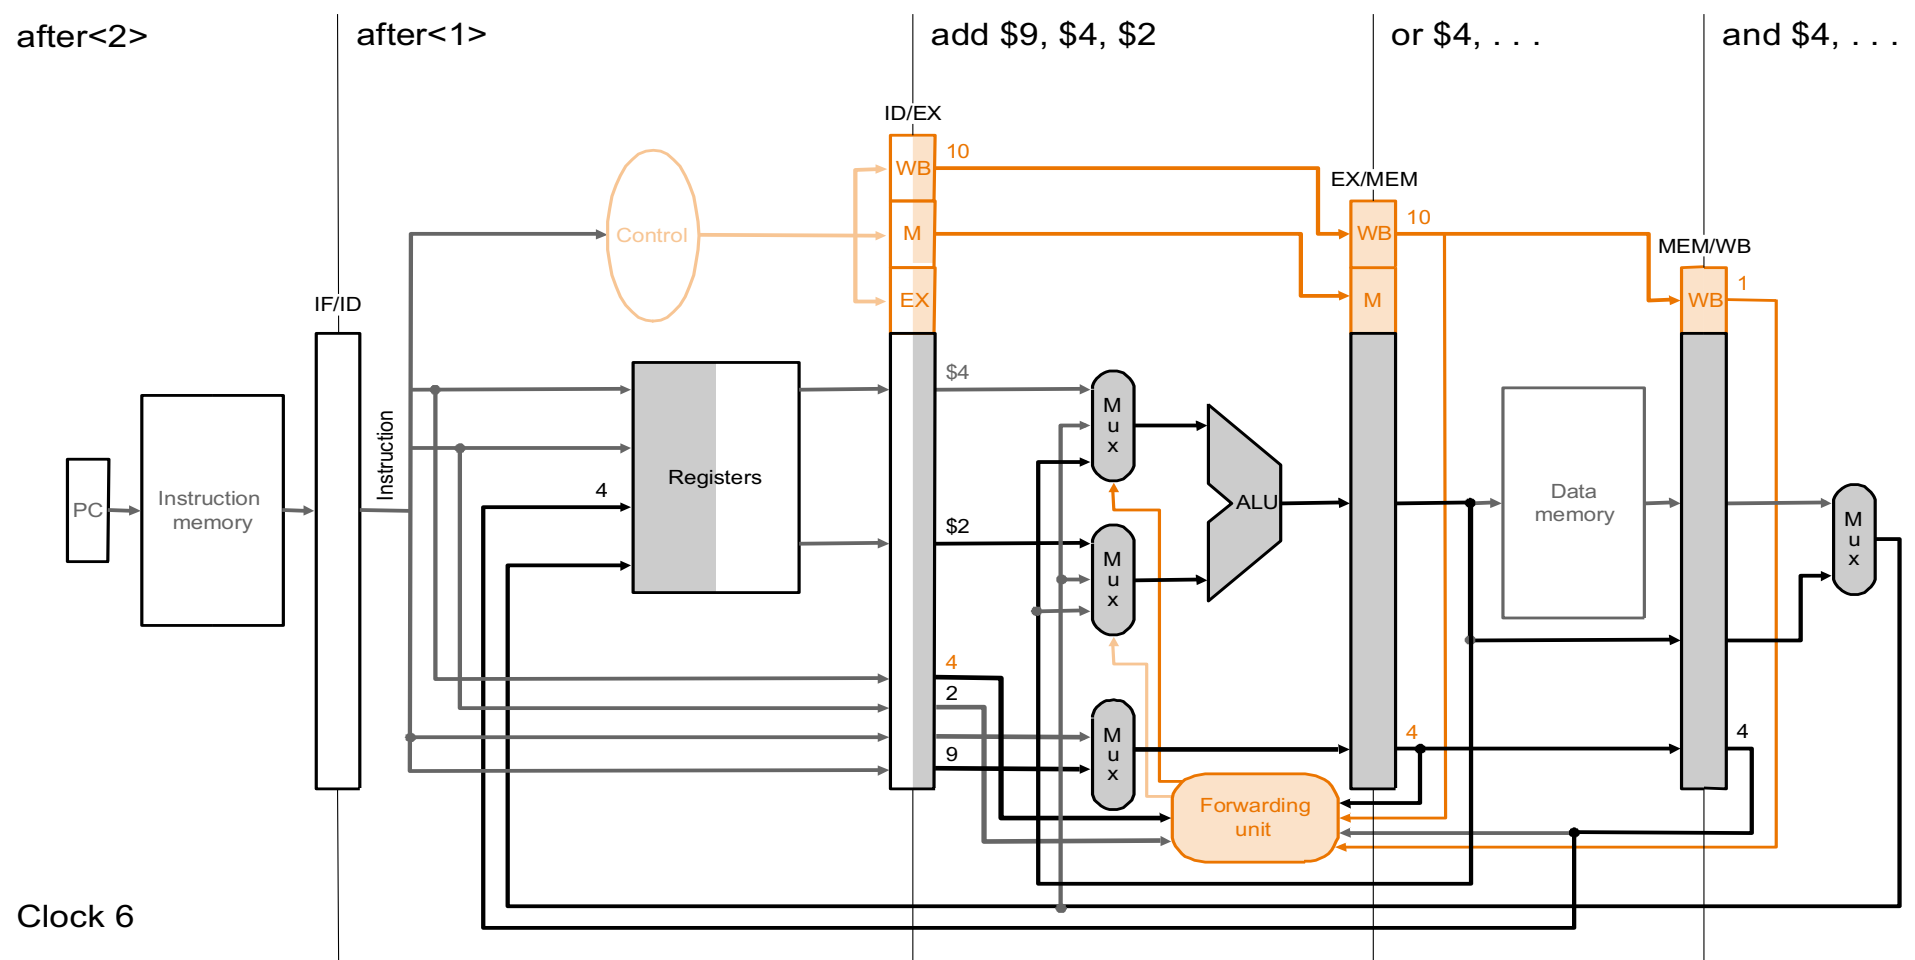
or \$4, \$4, \$2

and \$4, ...

sub \$2, ...



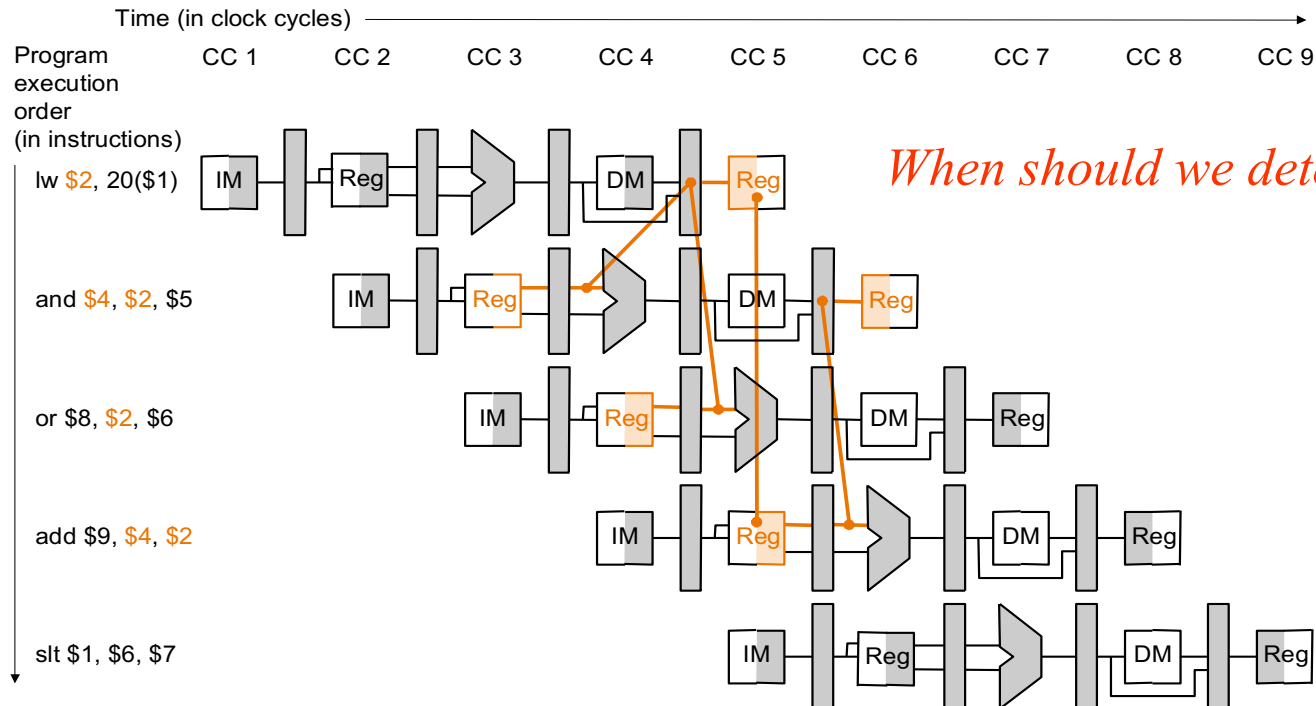
Clock 5



# Can't always forward

## Load word can still cause a hazard:

- an instruction tries to read a register following a load instruction that writes to the same register.

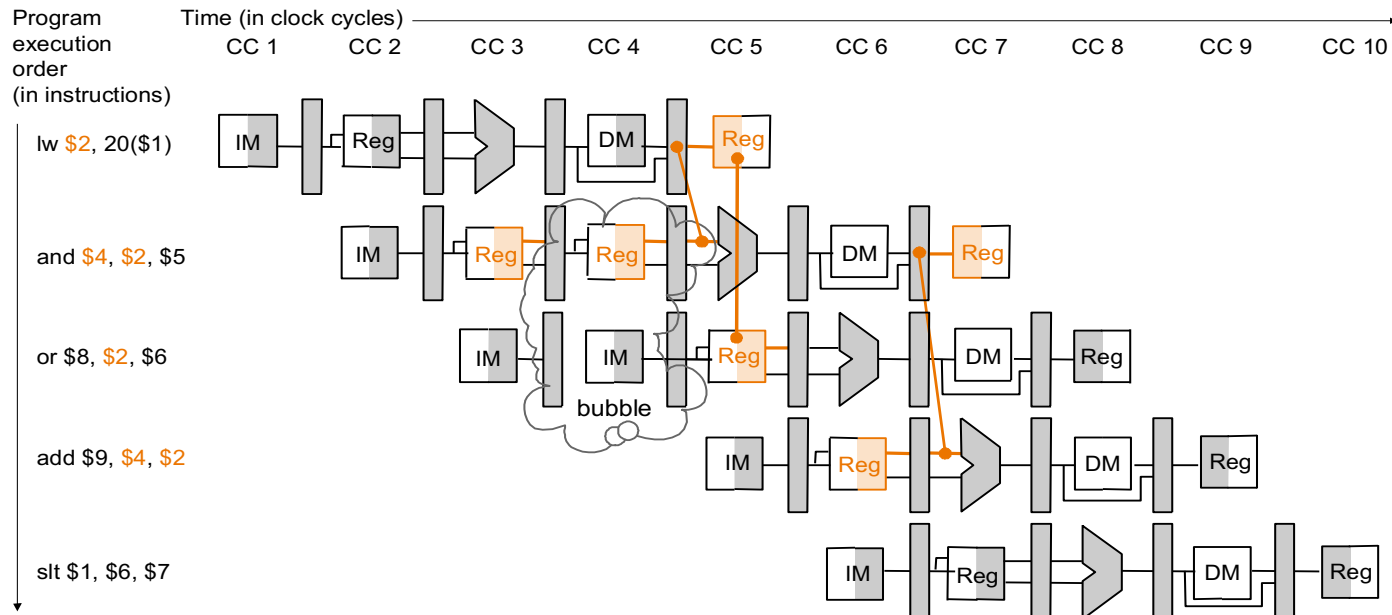


If (ID/EX.MemRead and  
((ID/EX.RegisterRt = IF/ID.RegisterRs) or  
(ID/EX.RegisterRt = IF/ID.RegisterRt)))

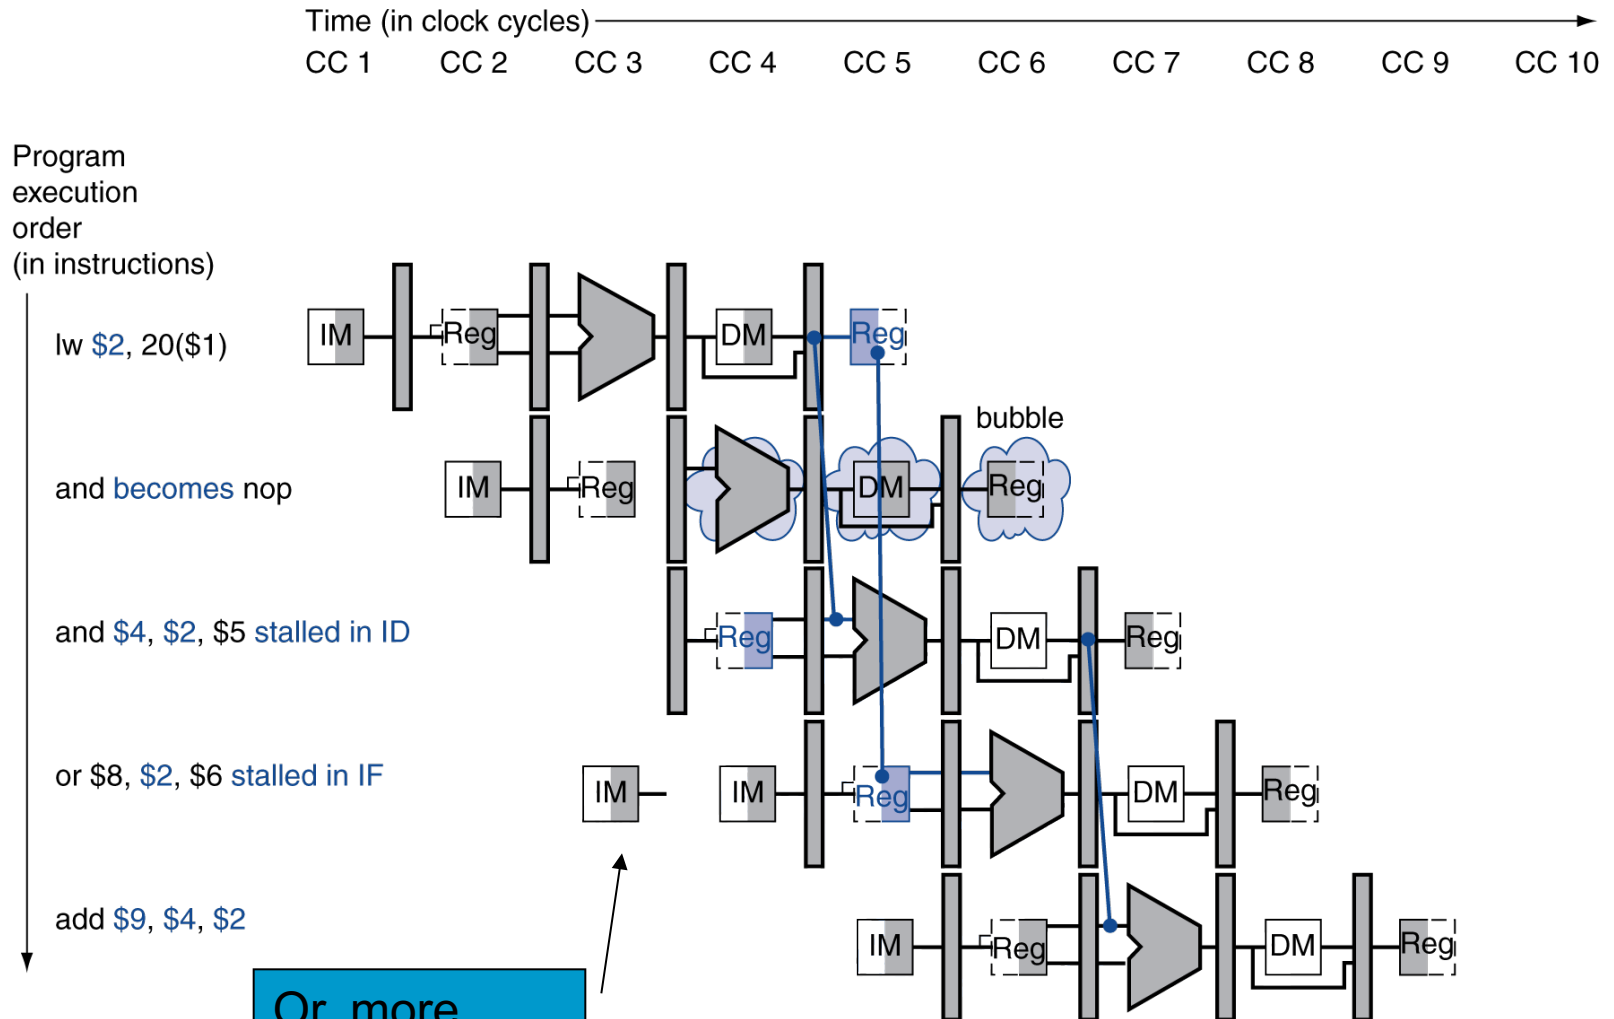
**stall the pipeline**

# Hazard Detection and Stall

- If (ID/EX.MemRead and  
((ID/EX.RegisterRt = IF/ID.RegisterRs) or  
(ID/EX.RegisterRt = IF/ID.RegisterRt)))  
**stall the pipeline**
- Stall the pipeline
  - Preventing instructions in the IF and ID stages from making progress
    - Preserve the PC and IF/ID pipeline registers
  - We need to do nothing in EX at CC4, MEM at CC5, WB at CC6
    - Deasserting all nine control signals in the EX, MEM and WB stage



# Stall/Bubble in the Pipeline

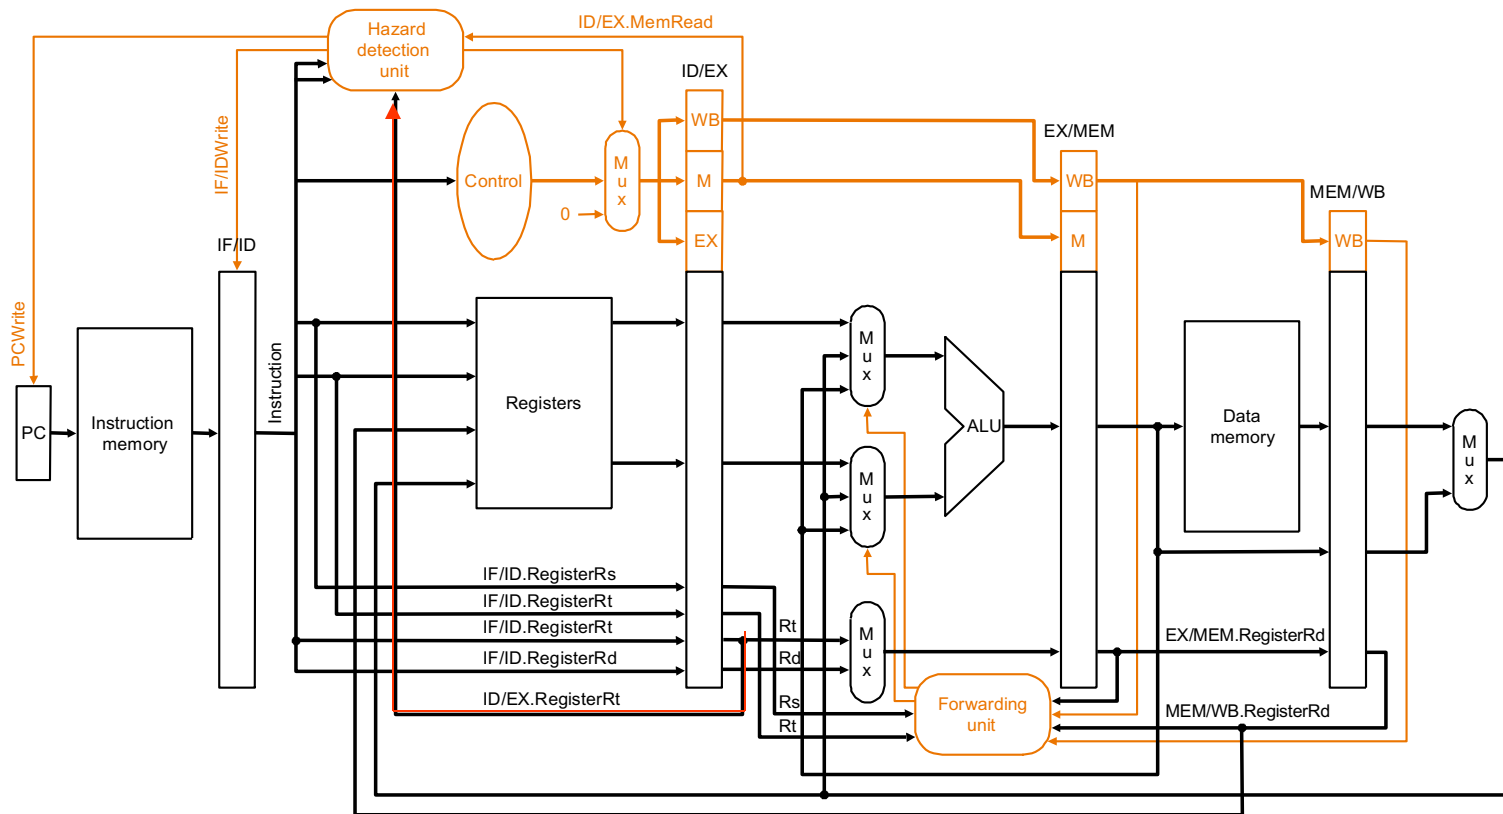


Or, more accurately...



# Hazard Detection Unit

- Stall by letting an instruction that won't write anything go forward



## Example

---

- Show how hazard detection unit works with this instruction sequence (with dependencies highlighted):

|     |     |            |   |                      |
|-----|-----|------------|---|----------------------|
| lw  | \$2 | , 20(\$1)  | } | load-use data hazard |
| and | \$4 | , \$2, \$5 |   |                      |
| or  | \$4 | , \$4, \$2 | } | Forwarding           |
| add | \$9 | , \$4, \$2 |   |                      |

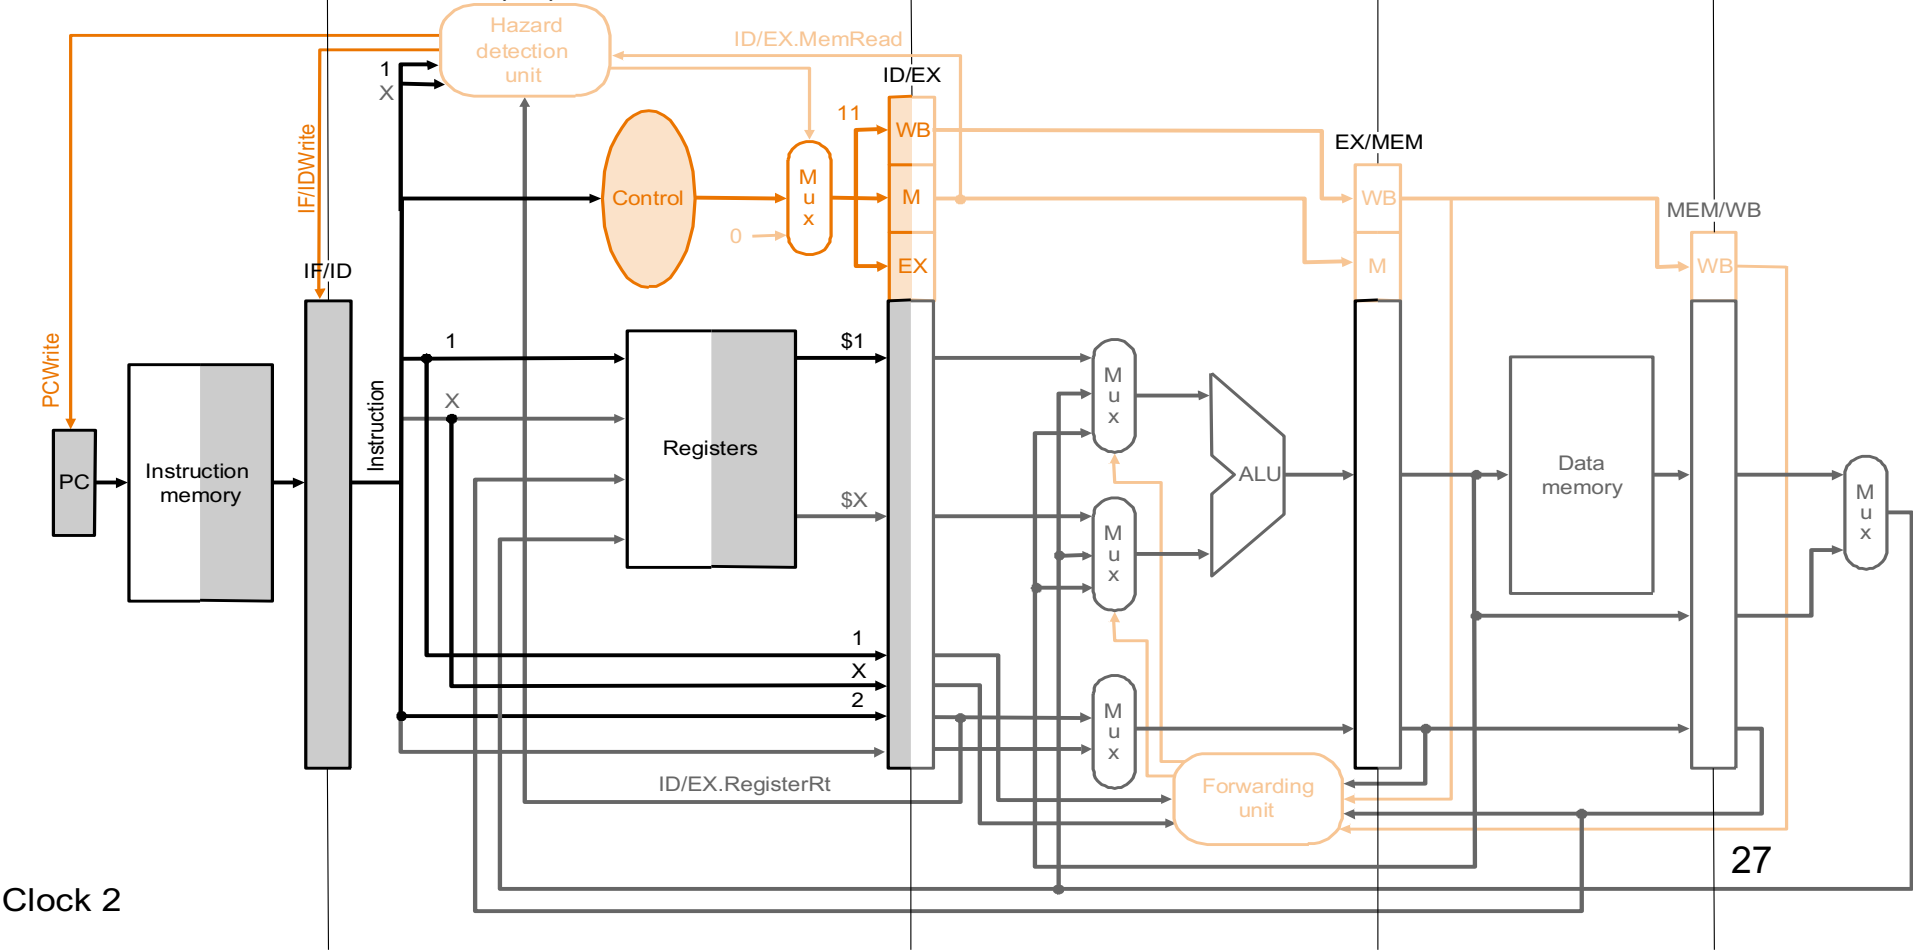
and \$4, \$2, \$5

lw \$2, 20(\$1)

before<1>

before<2>

before<3>



Clock 2

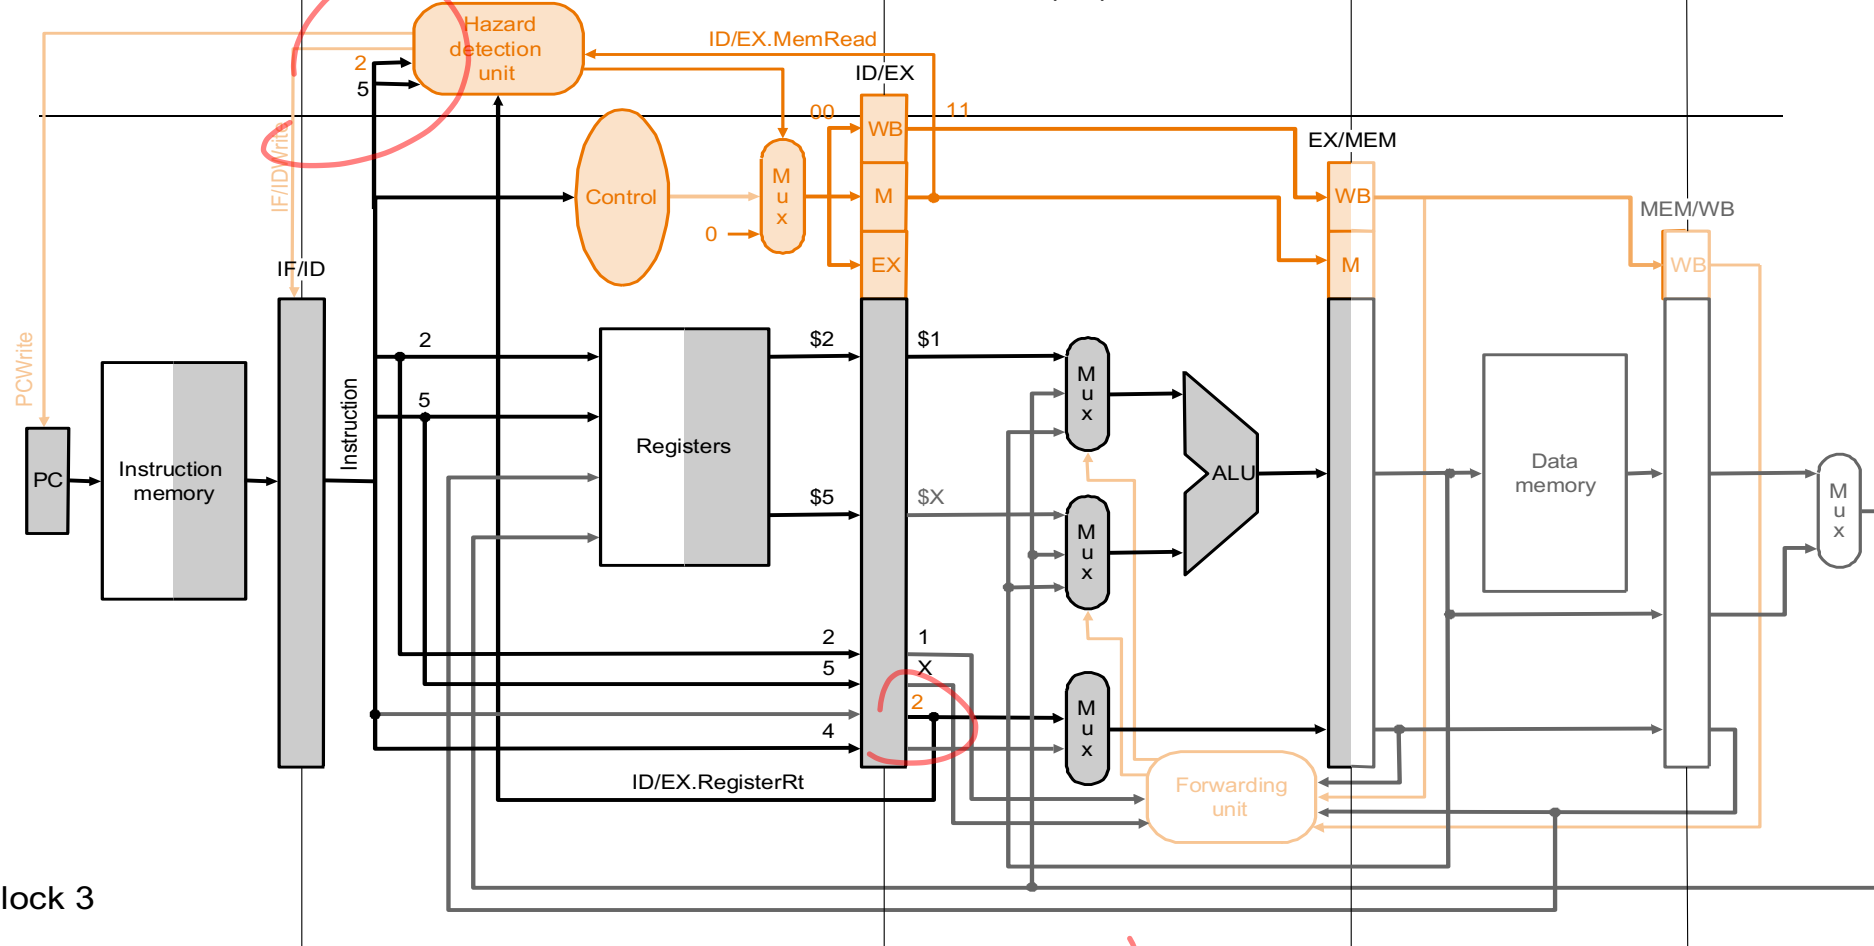
or \$4, \$4, \$2

and \$4, \$2, \$5

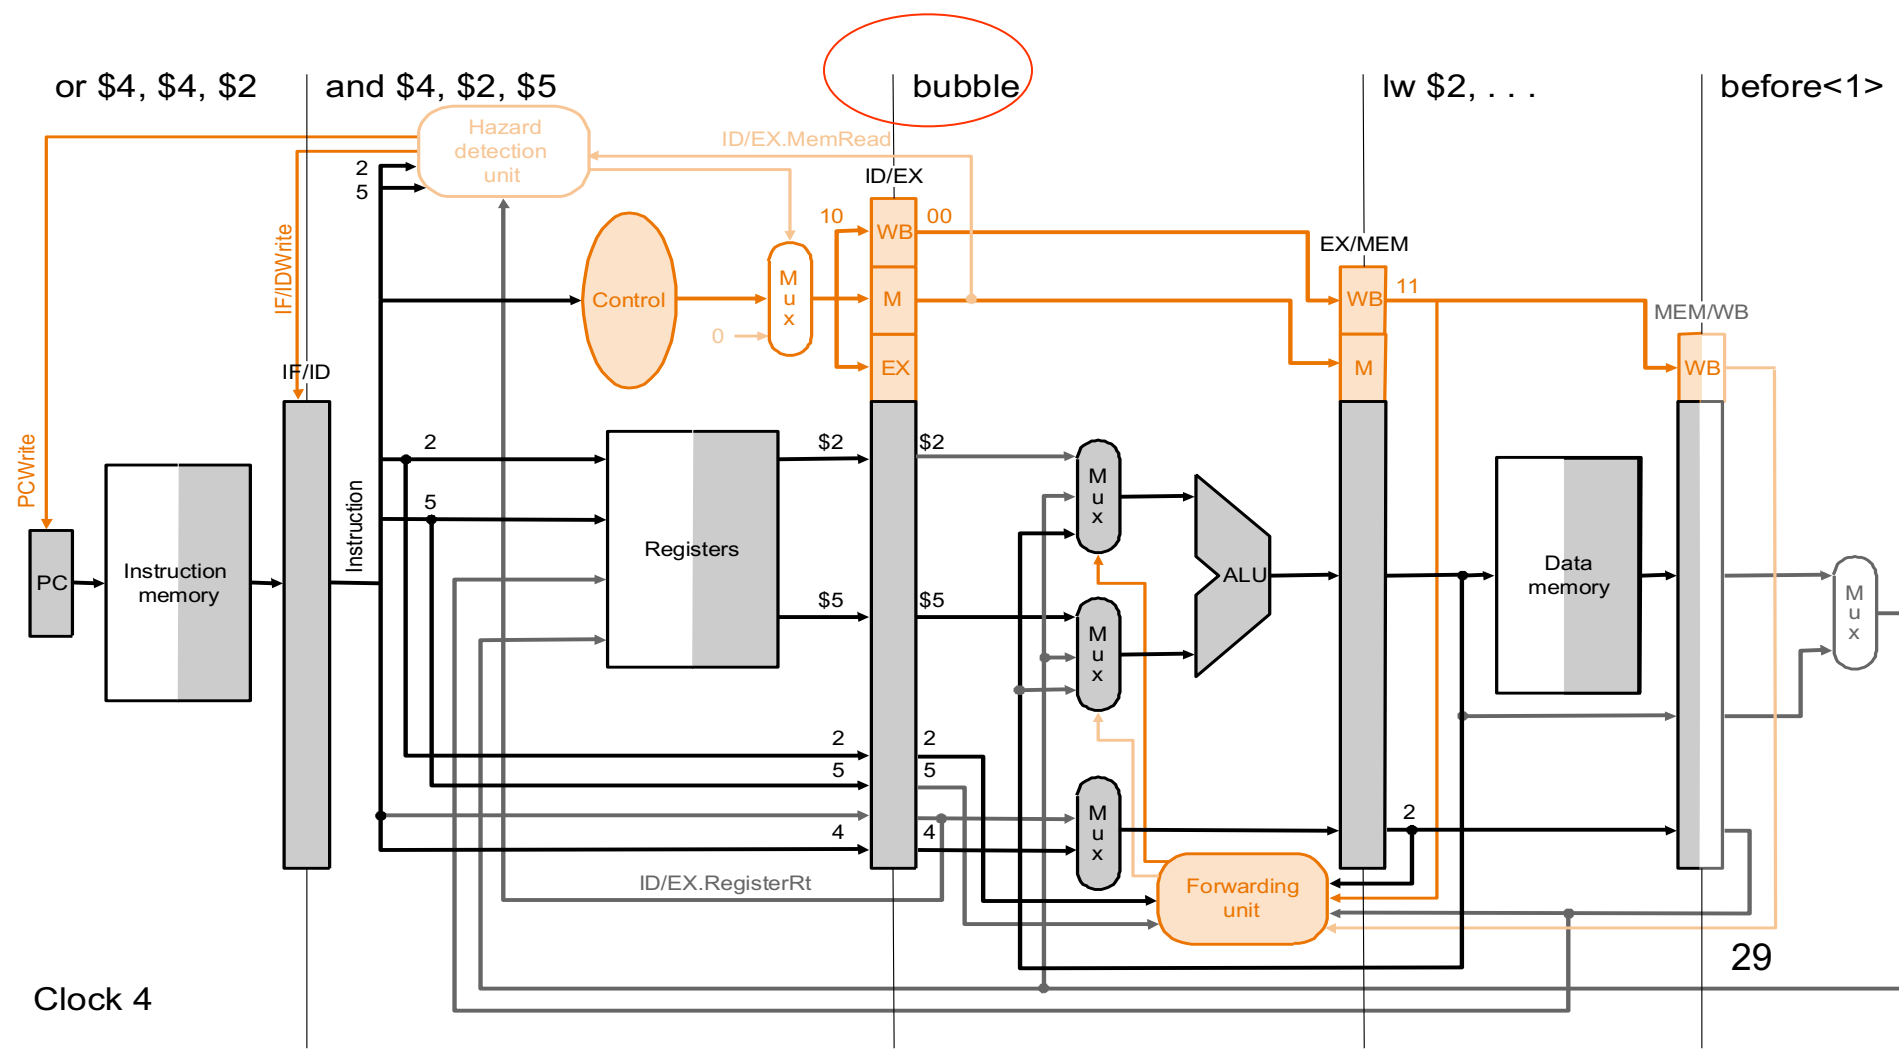
lw \$2, 20(\$1)

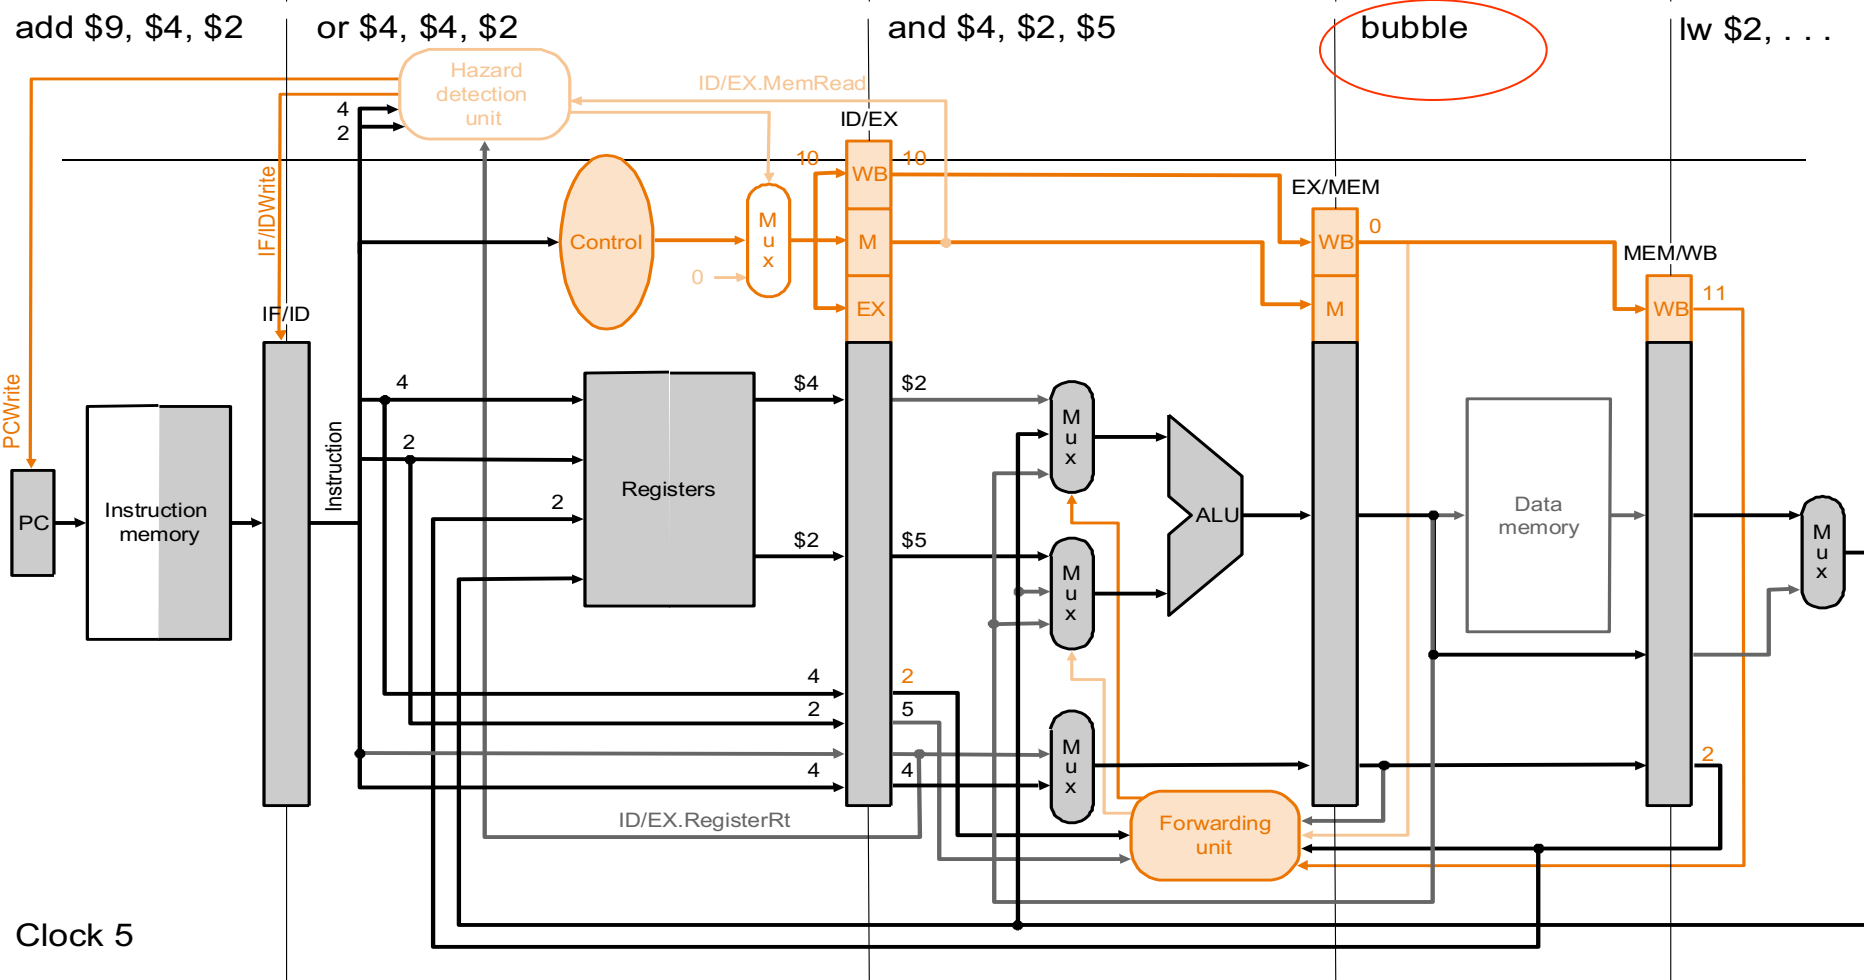
before<1>

before<2>



IF/ID.  $Rs(Rd) = 20, EX$   
 $\$2 \$5$   $Rt(\$2)$





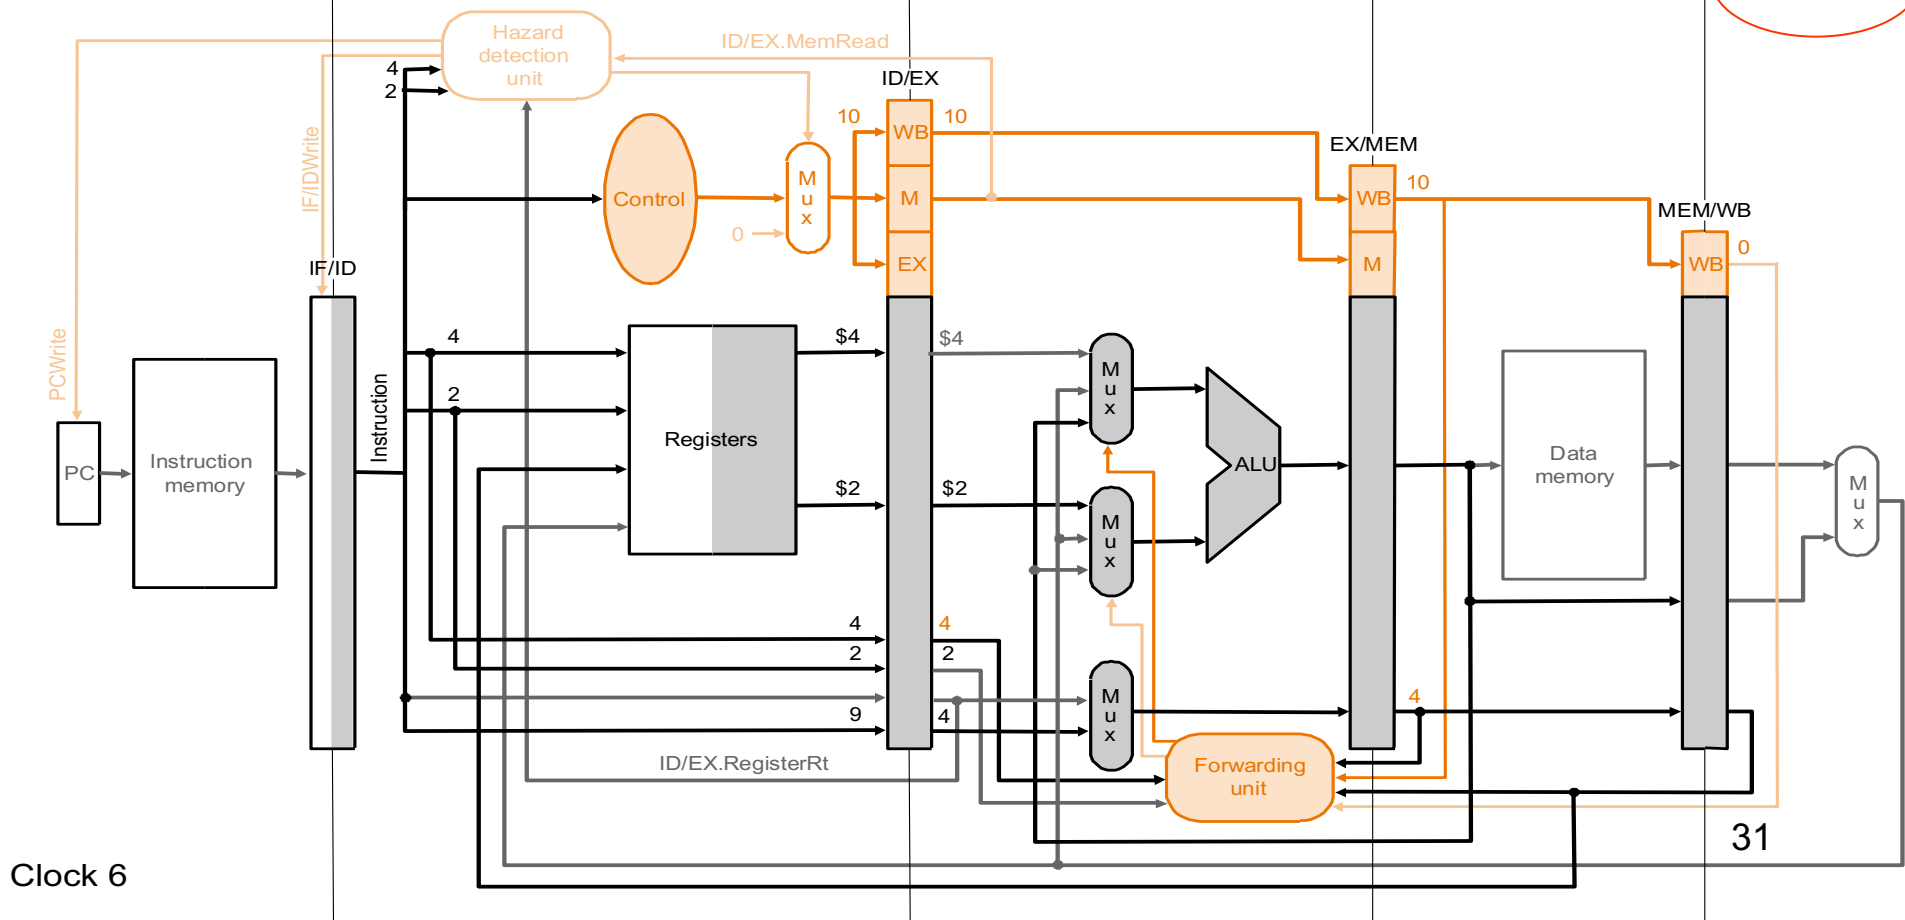
after<1>

add \$9, \$4, \$2

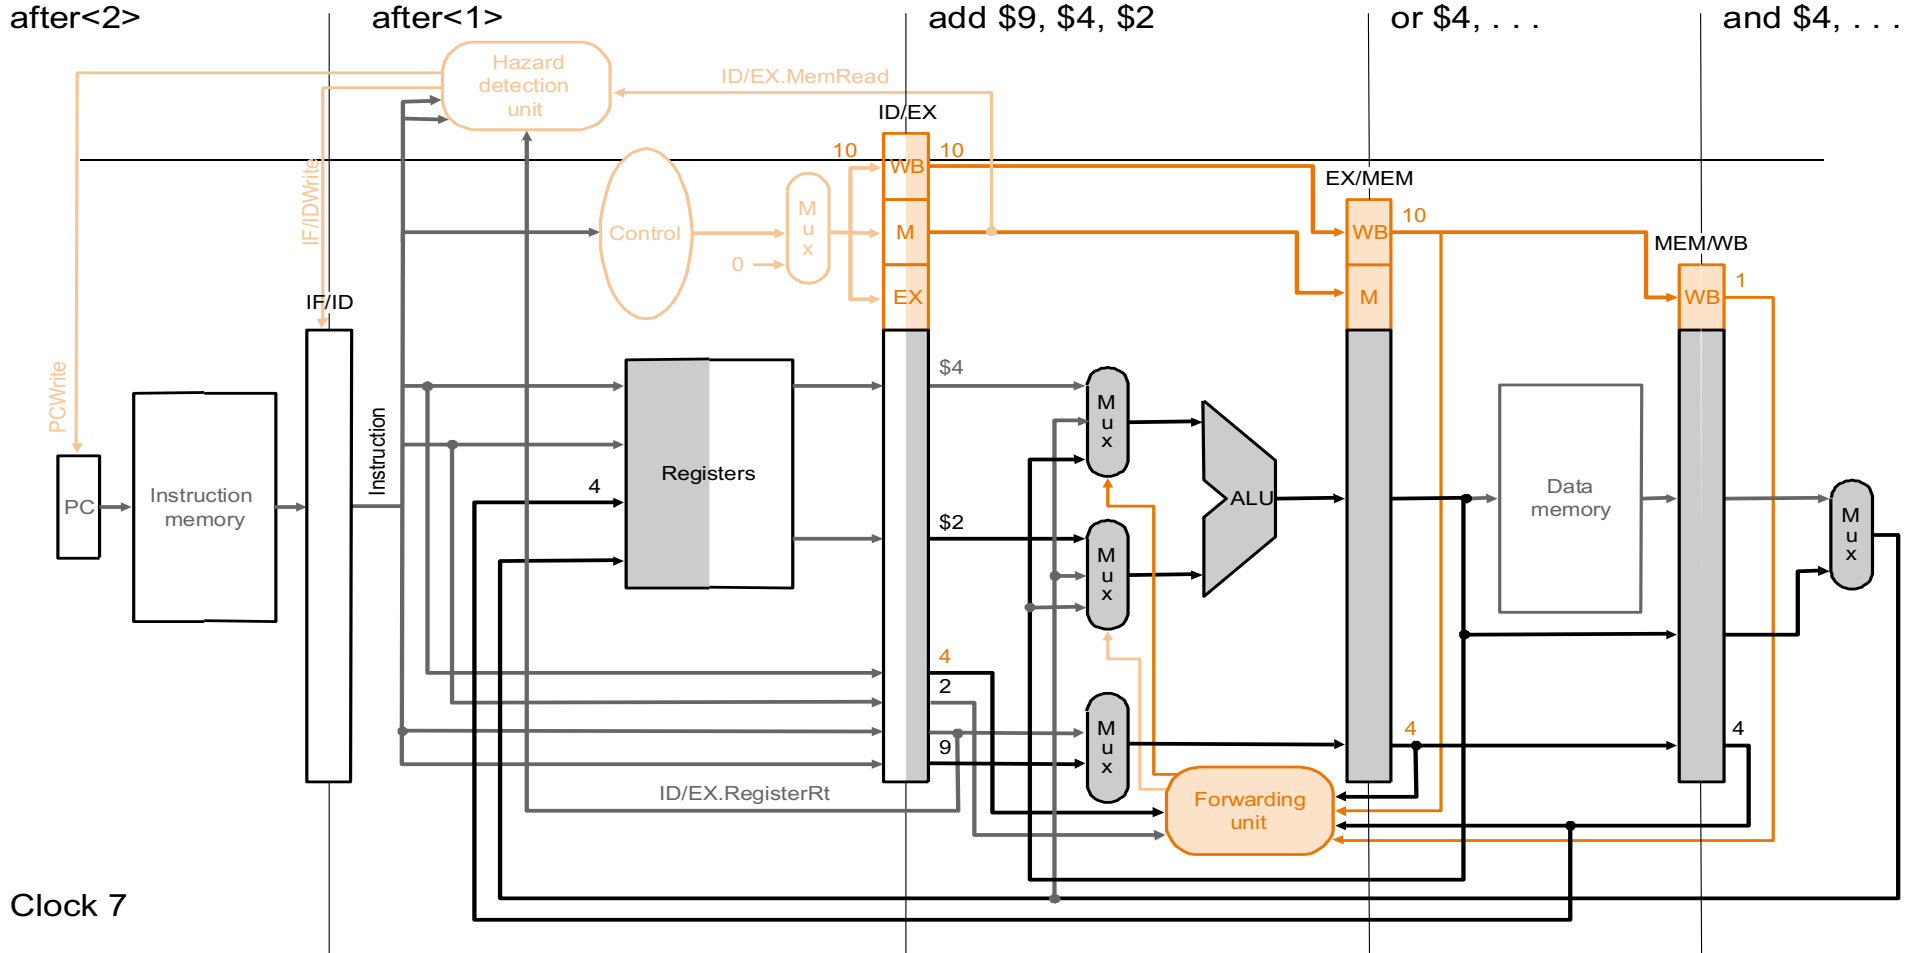
or \$4, \$4, \$2

and \$4, ...

bubble



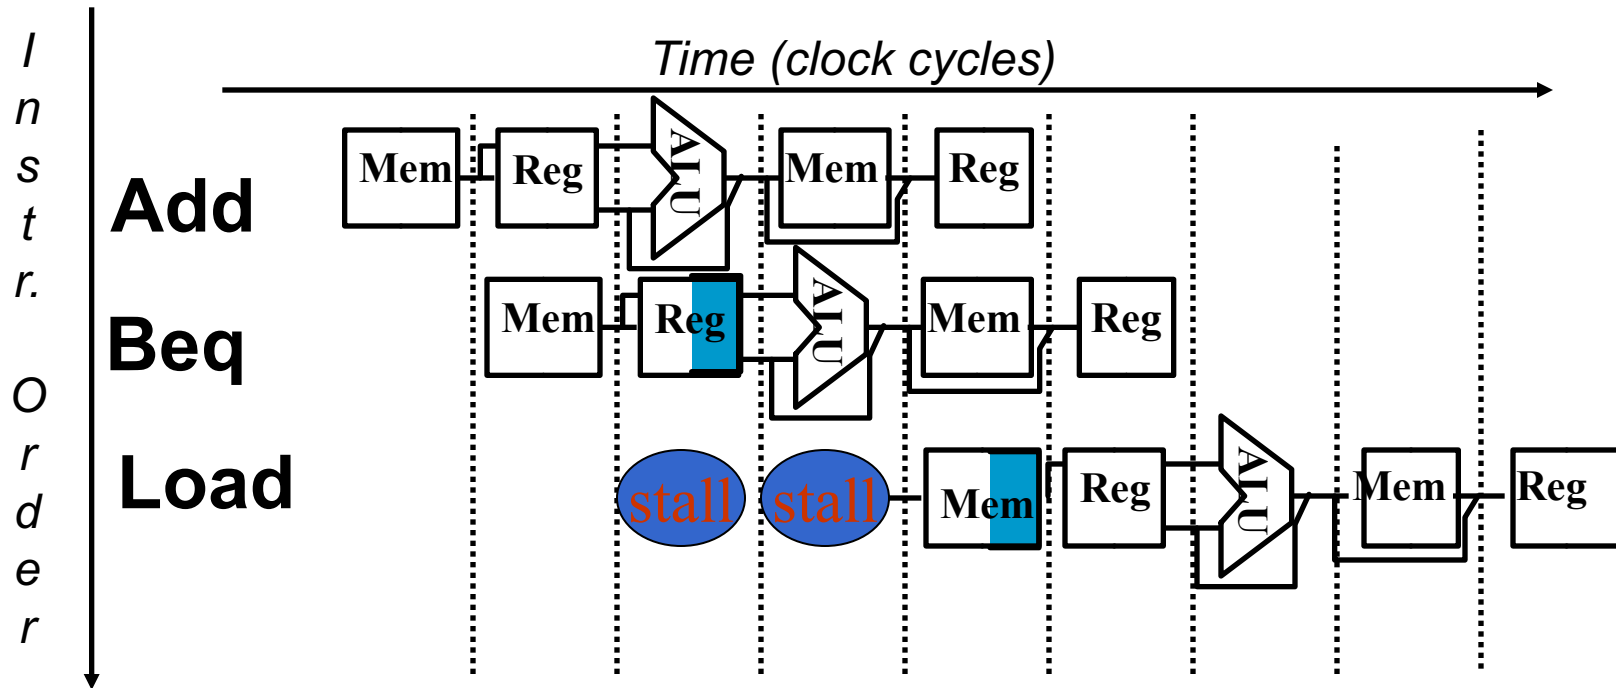
Clock 6





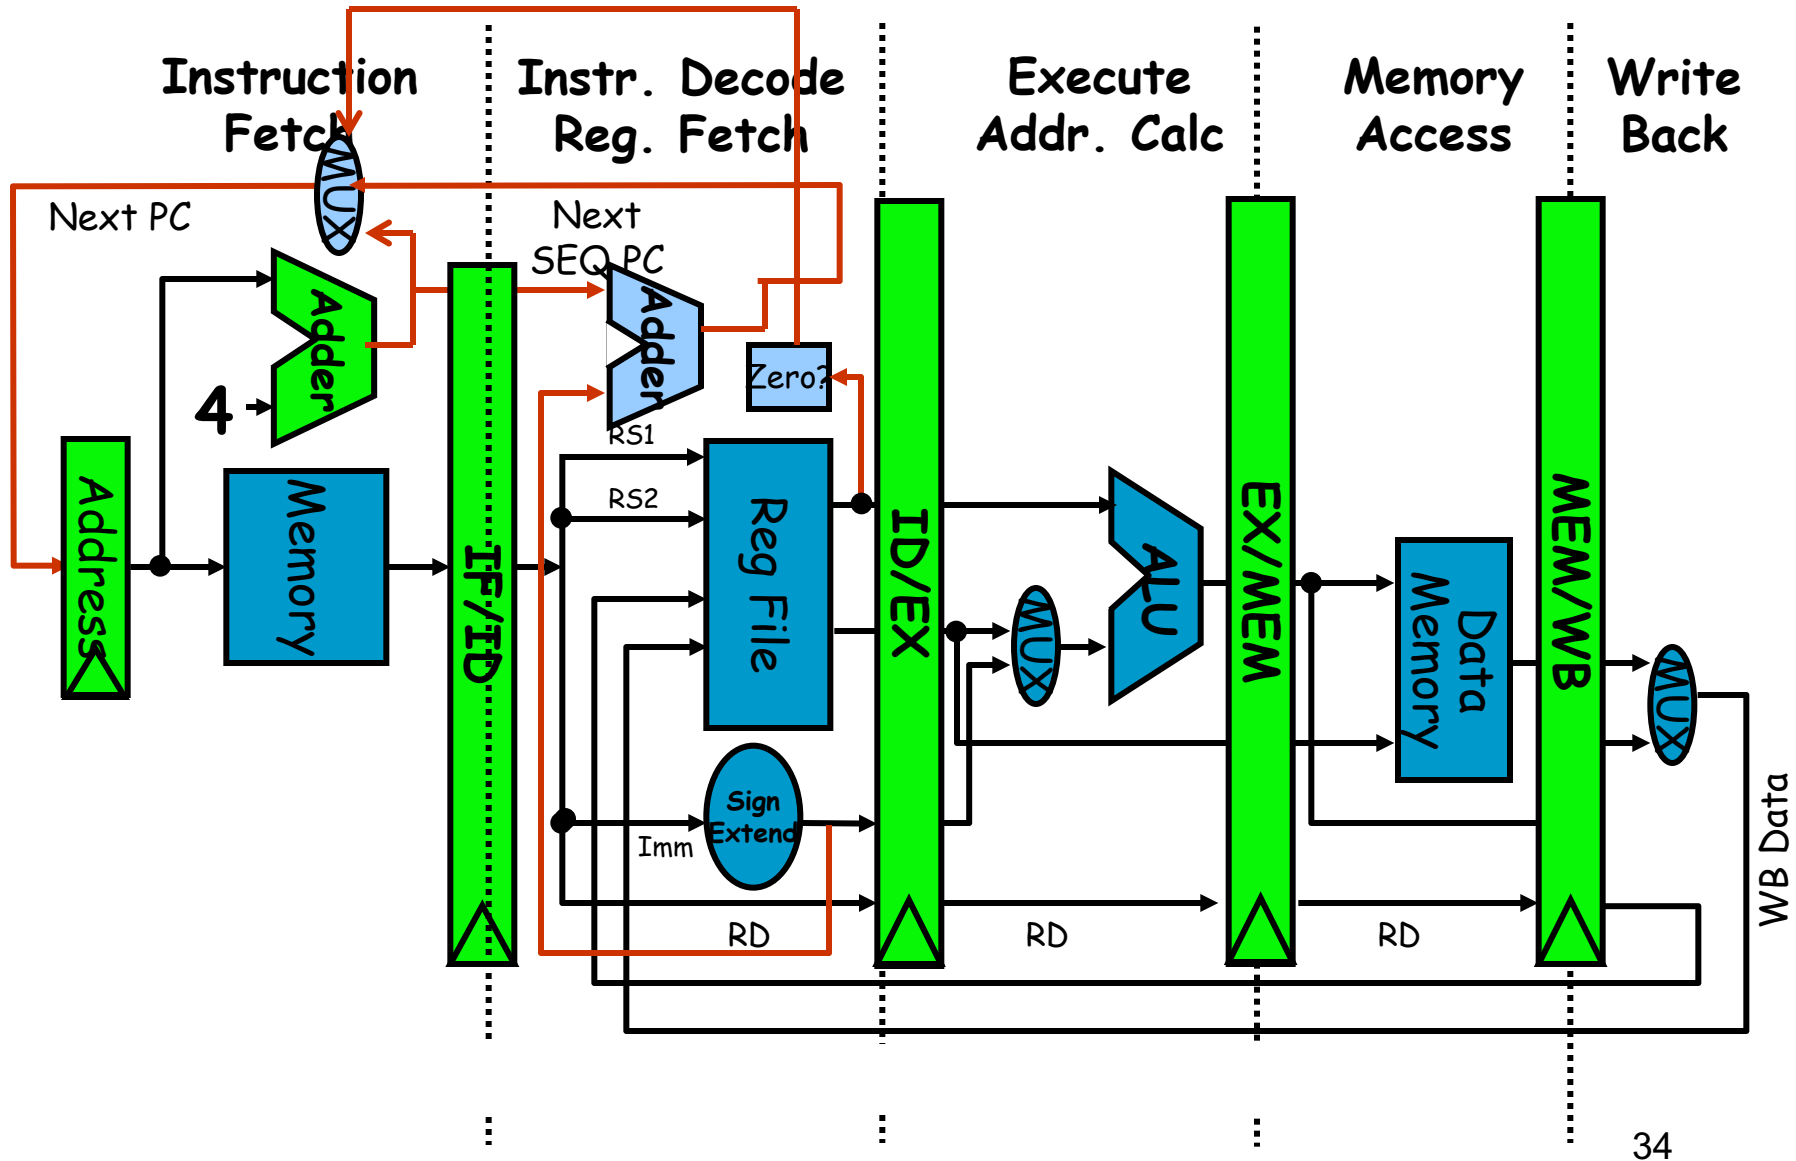
# Control Hazard Solutions

- Stall: wait until decision is clear



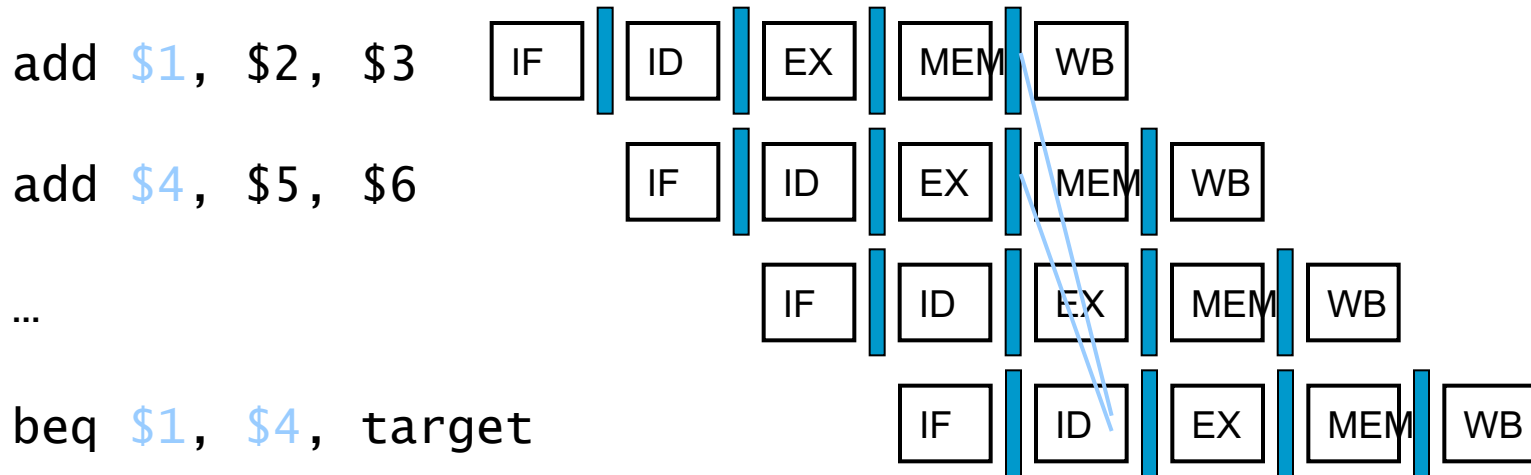
- Impact: 3 clock cycles per branch instruction  
=> slow

## Control Hazard Solution (1): Reducing the Delay of Branches (Example: BEQZ, BNEZ)



# Data Hazards for Branches

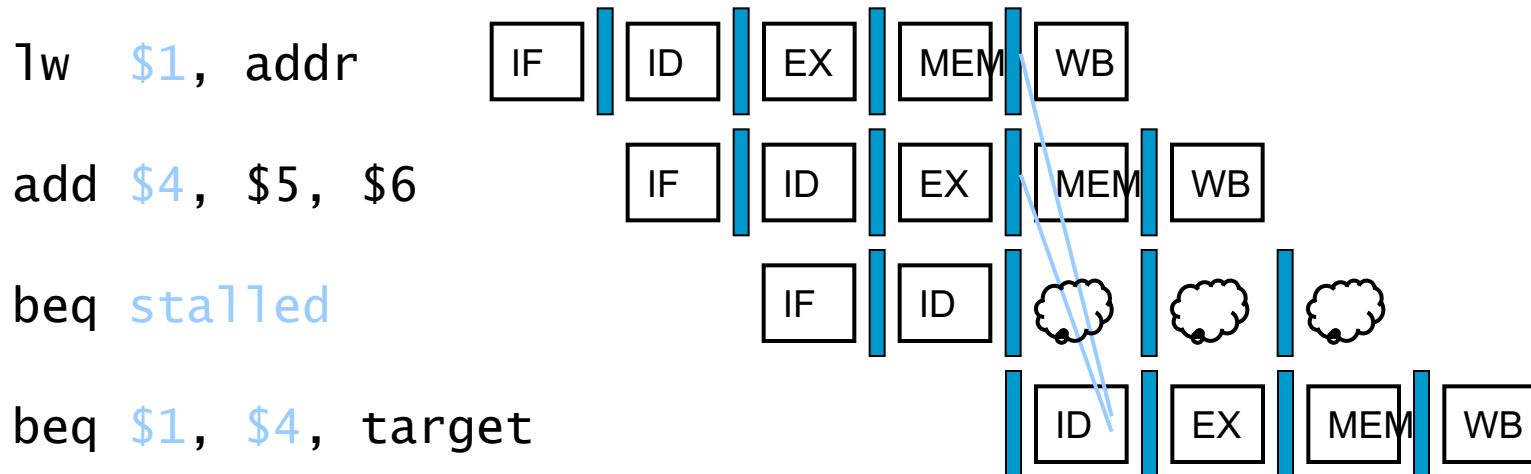
- If a comparison register is a destination of 2<sup>nd</sup> or 3<sup>rd</sup> preceding ALU instruction



- Can resolve using forwarding

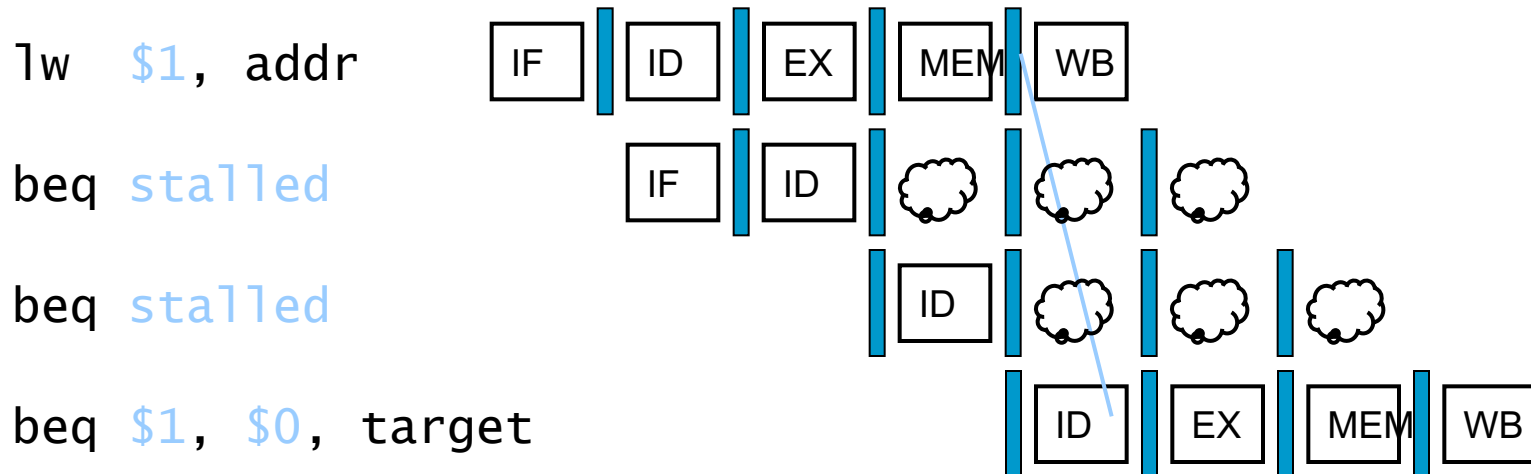
# Data Hazards for Branches

- If a comparison register is a destination of preceding ALU instruction or 2<sup>nd</sup> preceding load instruction
  - Need 1 stall cycle

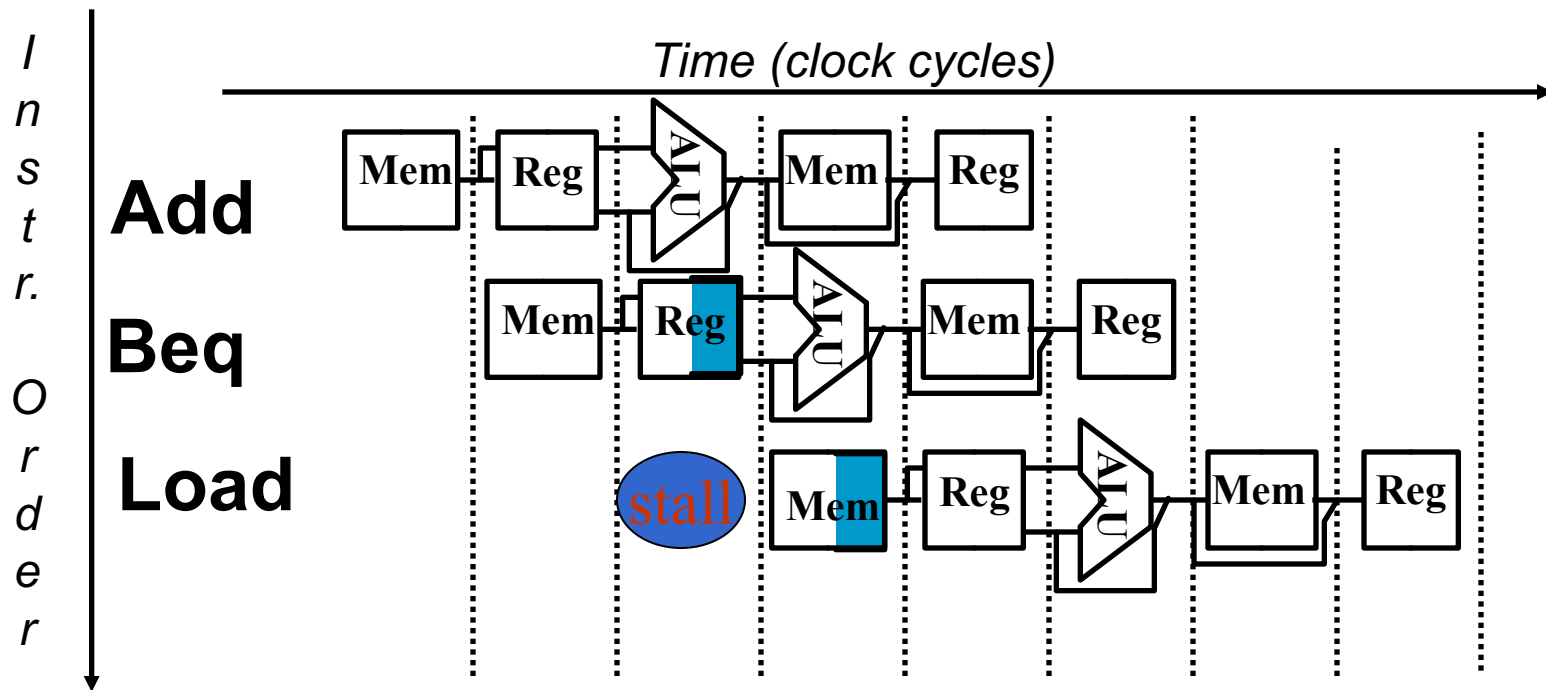


# Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
  - Need 2 stall cycles



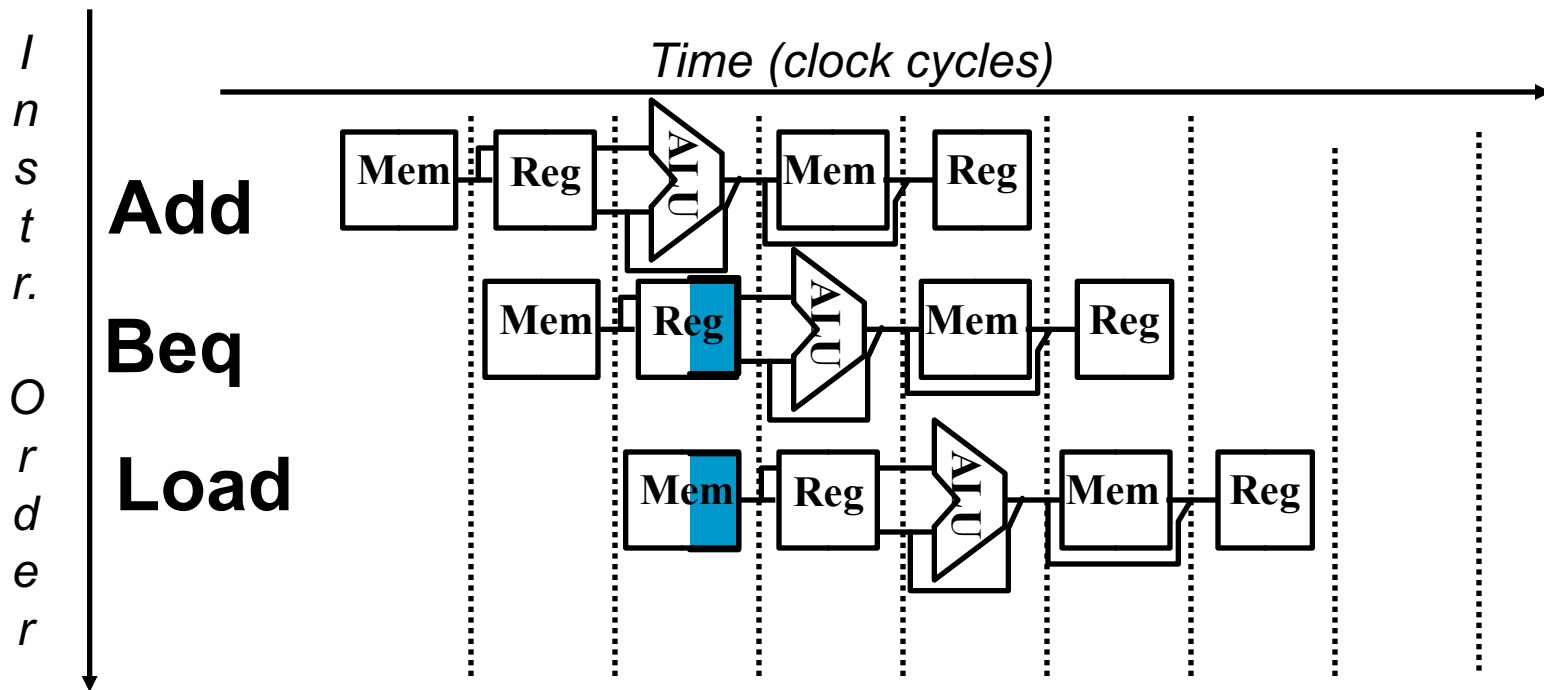
# Control Hazard Solutions (1)



- Impact: 2 clock cycles per branch instruction  
=> slow

## Control Hazard Solutions (2)

- Predict: guess one direction then back up if wrong
  - Predict not taken



- Impact: 1 clock cycles per branch instruction if right, 2 if wrong (right 50% of time)
- More dynamic scheme: history of 1 branch (90%)

## Predict branch not taken

---

| Branch Inst (i) | IF | ID | EX | MEM | WB  |     |     |     |
|-----------------|----|----|----|-----|-----|-----|-----|-----|
| Inst i+1        |    | IF | ID | EX  | MEM | WB  |     |     |
| Inst i+2        |    |    | IF | ID  | EX  | MEM | WB  |     |
| Inst i+3        |    |    |    | IF  | ID  | EX  | MEM | WB  |
| Inst i+4        |    |    |    |     | IF  | ID  | EX  | MEM |

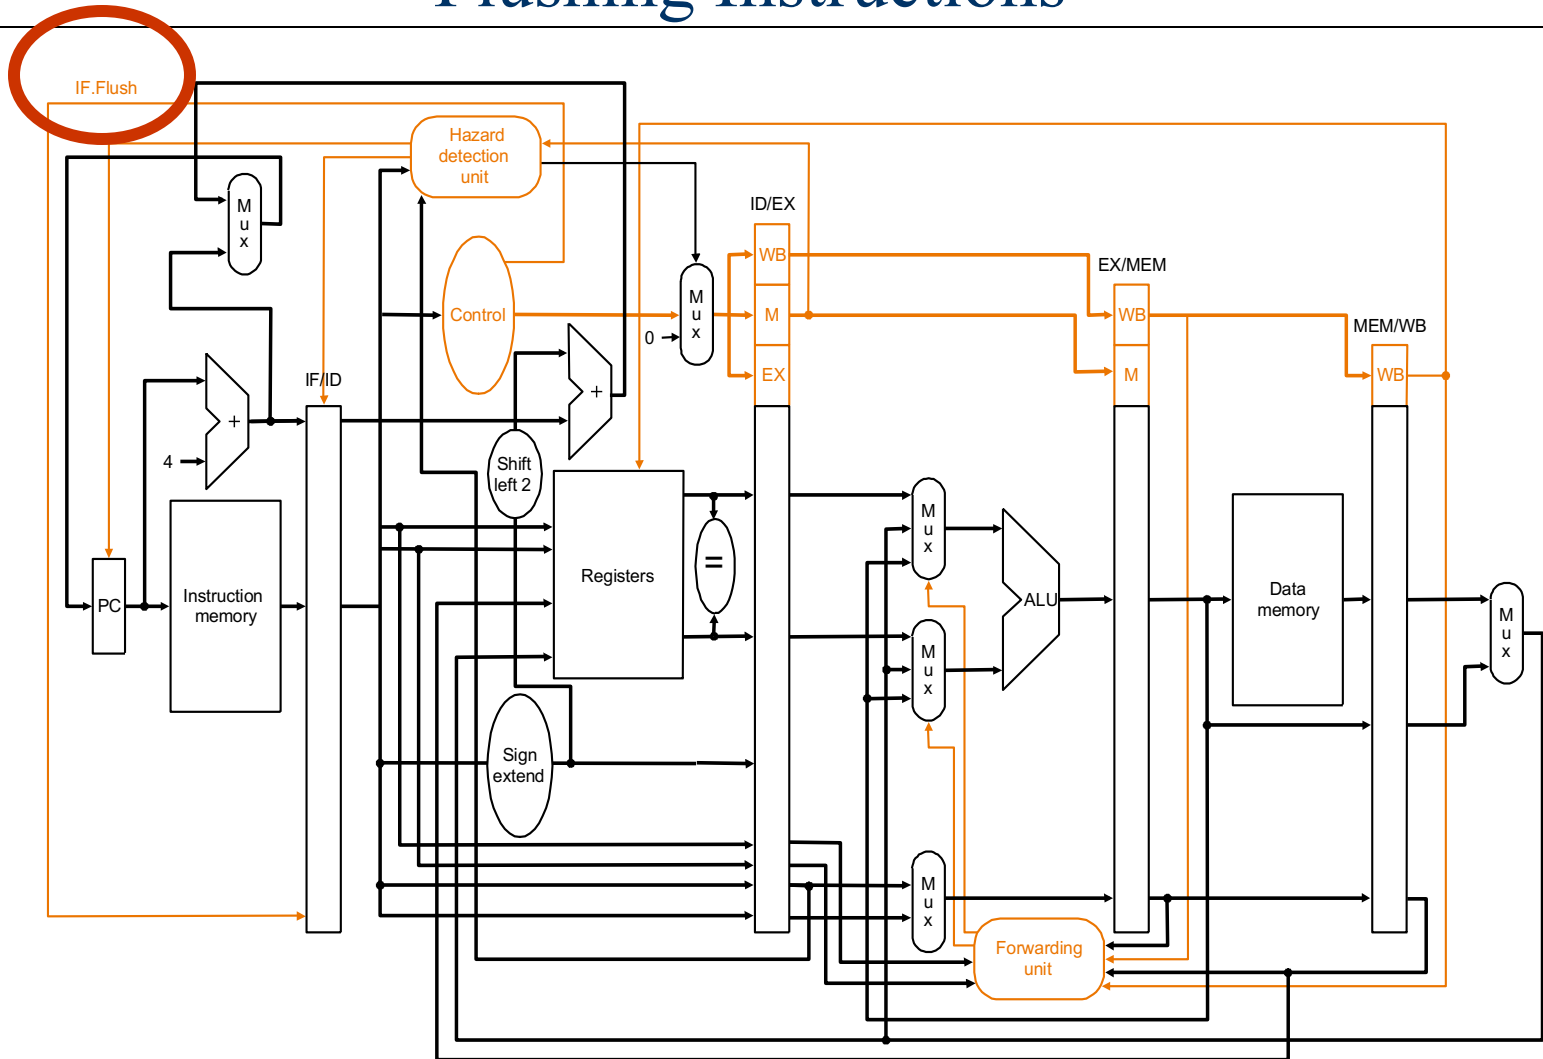
**Correct Prediction : Zero Cycle Branch Penalty!**

| Branch Inst (i) | IF | ID | EX  | MEM | WB  |     |    |  |
|-----------------|----|----|-----|-----|-----|-----|----|--|
| Inst i+1        |    | IF | nop | nop | nop | nop |    |  |
| Branch target   |    |    | IF  | ID  | EX  | MEM | WB |  |

**Incorrect Prediction - waste one cycle**  
**How to flush pipeline?**



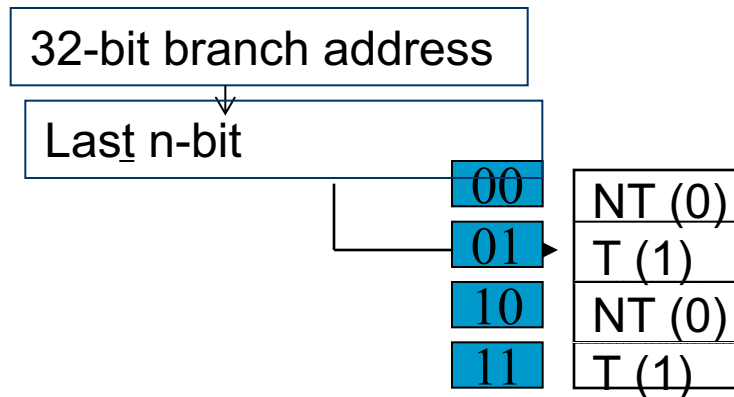
# Flushing Instructions



Zero the instruction field of the IF/ID pipeline register  
`sll $0, $0, 0`

# Dynamic Branch Prediction

- Branch History Table: Lower bits of PC address index table of 1-bit values
  - Says whether or not branch taken last time
  - No address check (saves HW, but may not be right branch)



branch history table =  $2^n$

Example: --00 T, --01 NT, --10 T, --11 T, --00 NT, --01 NT, --10 NT, --11 T,

# Dynamic Branch Prediction

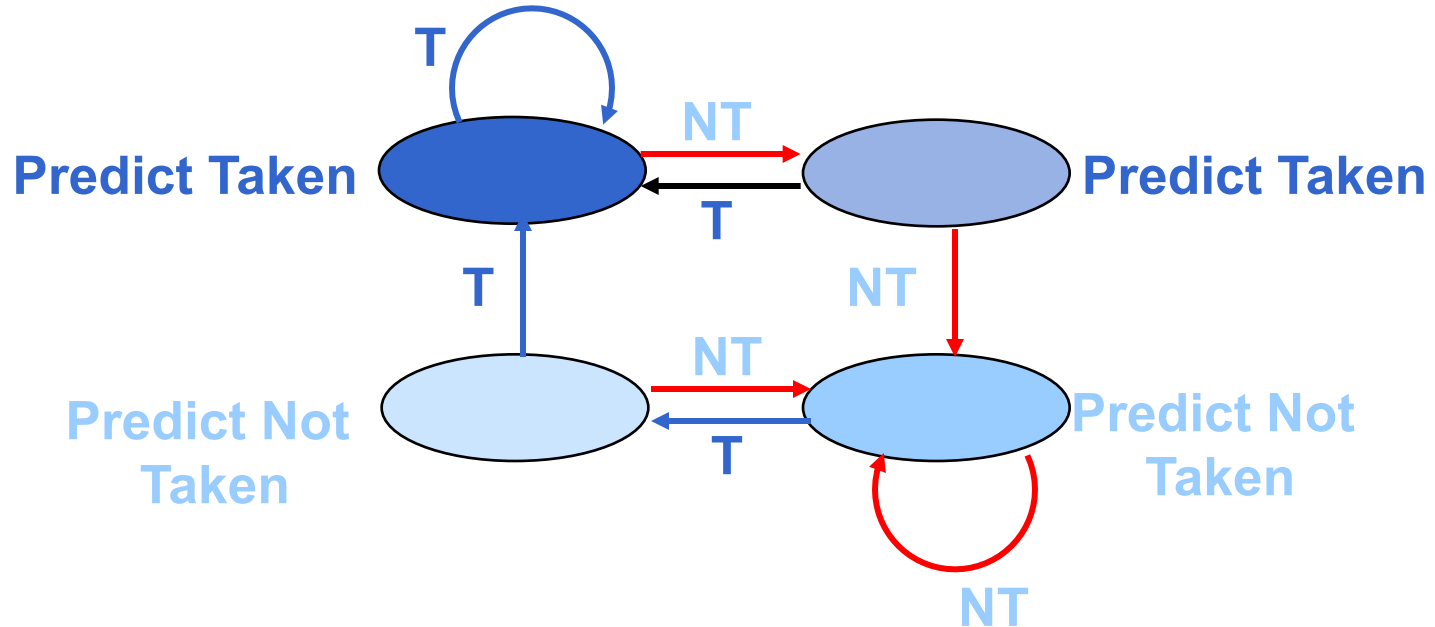
---

- Problem: in a loop, 1-bit BHT will cause 2 mispredictions (avg is 9 iterations before exit):
  - End of loop case, when it exits instead of looping as before
  - First time through loop on *next* time through code, when it predicts *exit* instead of looping
  - Only 80% accuracy even if loop 90% of the time

```
L1:  ...  
      ...  
      ...  
      SUBI    R1, #8  
      BNEZ    R1, L1
```

# Dynamic Branch Prediction

- Solution: 2-bit scheme which changes prediction only if get misprediction *twice*:



# Exercise

- Branch outcome of a single branch  
– T T T N N N T T T
- How many instances of this branch instruction are mis-predicted with a 1-bit predictor?
- How many instances of this branch instruction are mis-predicted with a 2-bit predictor?

The diagram illustrates the branch outcomes and predictions for a 1-bit and 2-bit predictor. The sequence of branch outcomes is T T T N N N T T T. For a 1-bit predictor, the predictions are X T T T N N N T T, with mis-predictions (X) occurring at the first, fourth, fifth, sixth, and seventh positions. For a 2-bit predictor, the predictions are X T T T N N N T T, with mis-predictions (X) occurring at the first, fourth, fifth, sixth, and seventh positions.

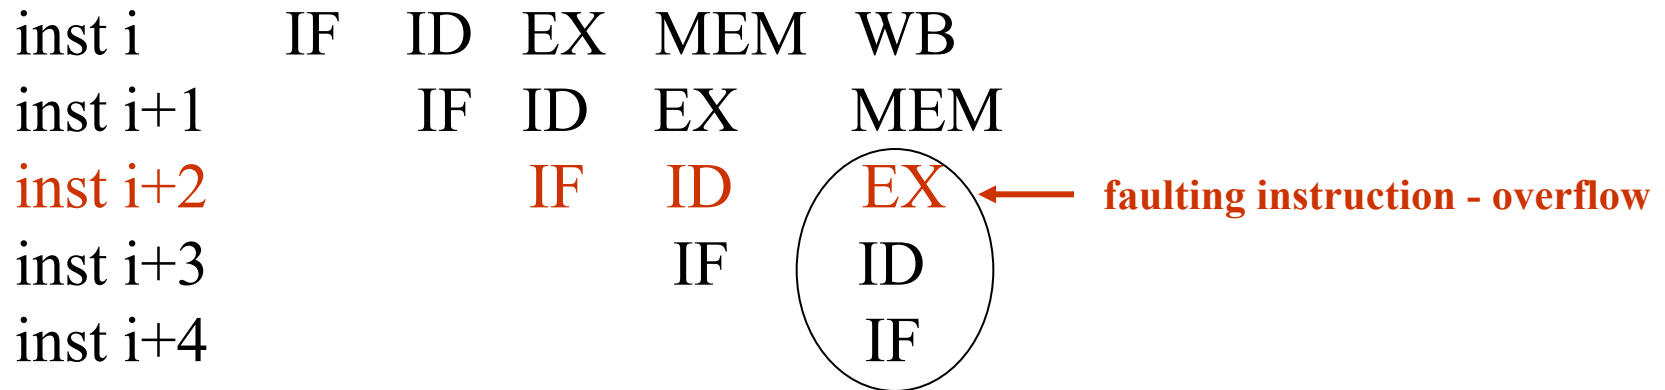
| Branch Outcome | 1-bit Predictor | 2-bit Predictor |
|----------------|-----------------|-----------------|
| T              | X               | X               |
| T              | T               | T               |
| T              | T               | T               |
| N              | T               | T               |
| N              | X               | X               |
| N              | X               | X               |
| T              | N               | N               |
| T              | N               | N               |
| T              | T               | T               |

# More on branch prediction

---

- Problems with predicted taken?
  - Need to calculate target address
  - Solution: branch target buffer
- Correlating predictor
  - A branch predictor that combines local behavior of a particular branch and global information about the behavior of some recent number of executed branches
- Tournament branch predictor
  - A branch predictor with multiple prediction for each branch and a selection mechanism that chooses which predictor to enable for a given branch

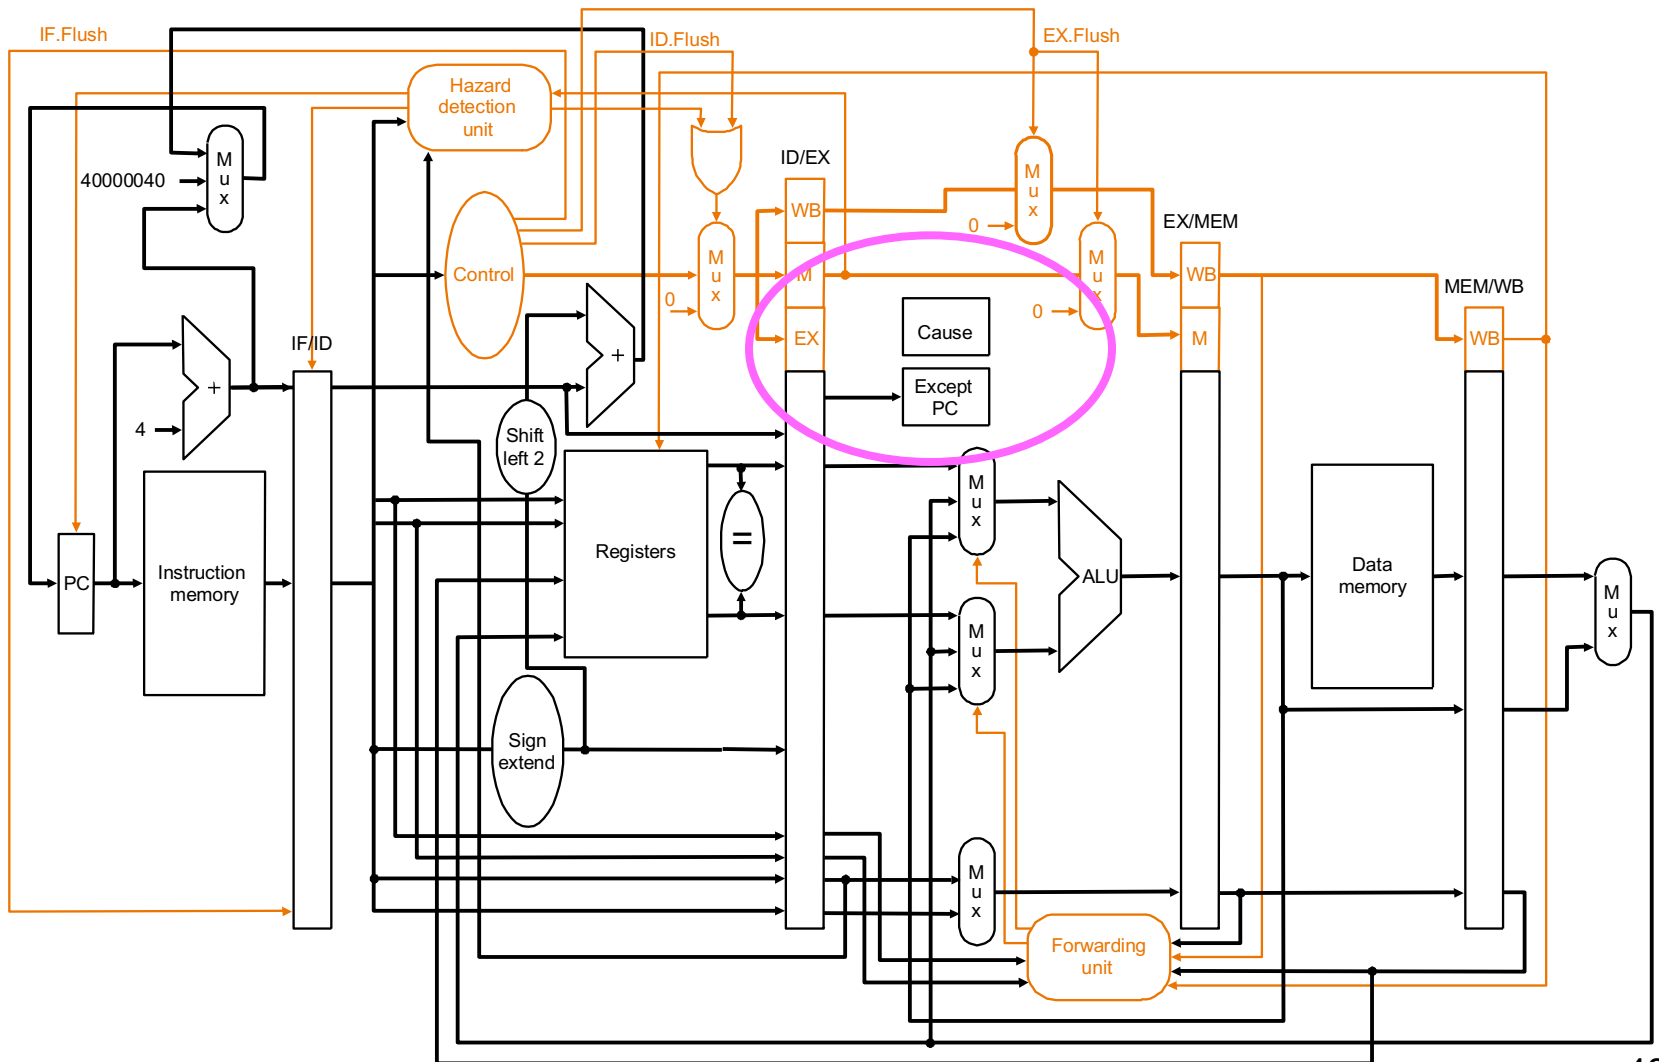
# Exceptions



## Steps to handle exceptions:

- Flush the instruction in the IF, ID and EX stages.
- Let all preceding instructions complete if they can
- EPC = address of (offending instruction) + 4
- Call the OS to handle the exception
  - PC = 0x40000040
- Return from the exception handler
  - PC = EPC -4

# Datapath with Controls to Handle Exceptions





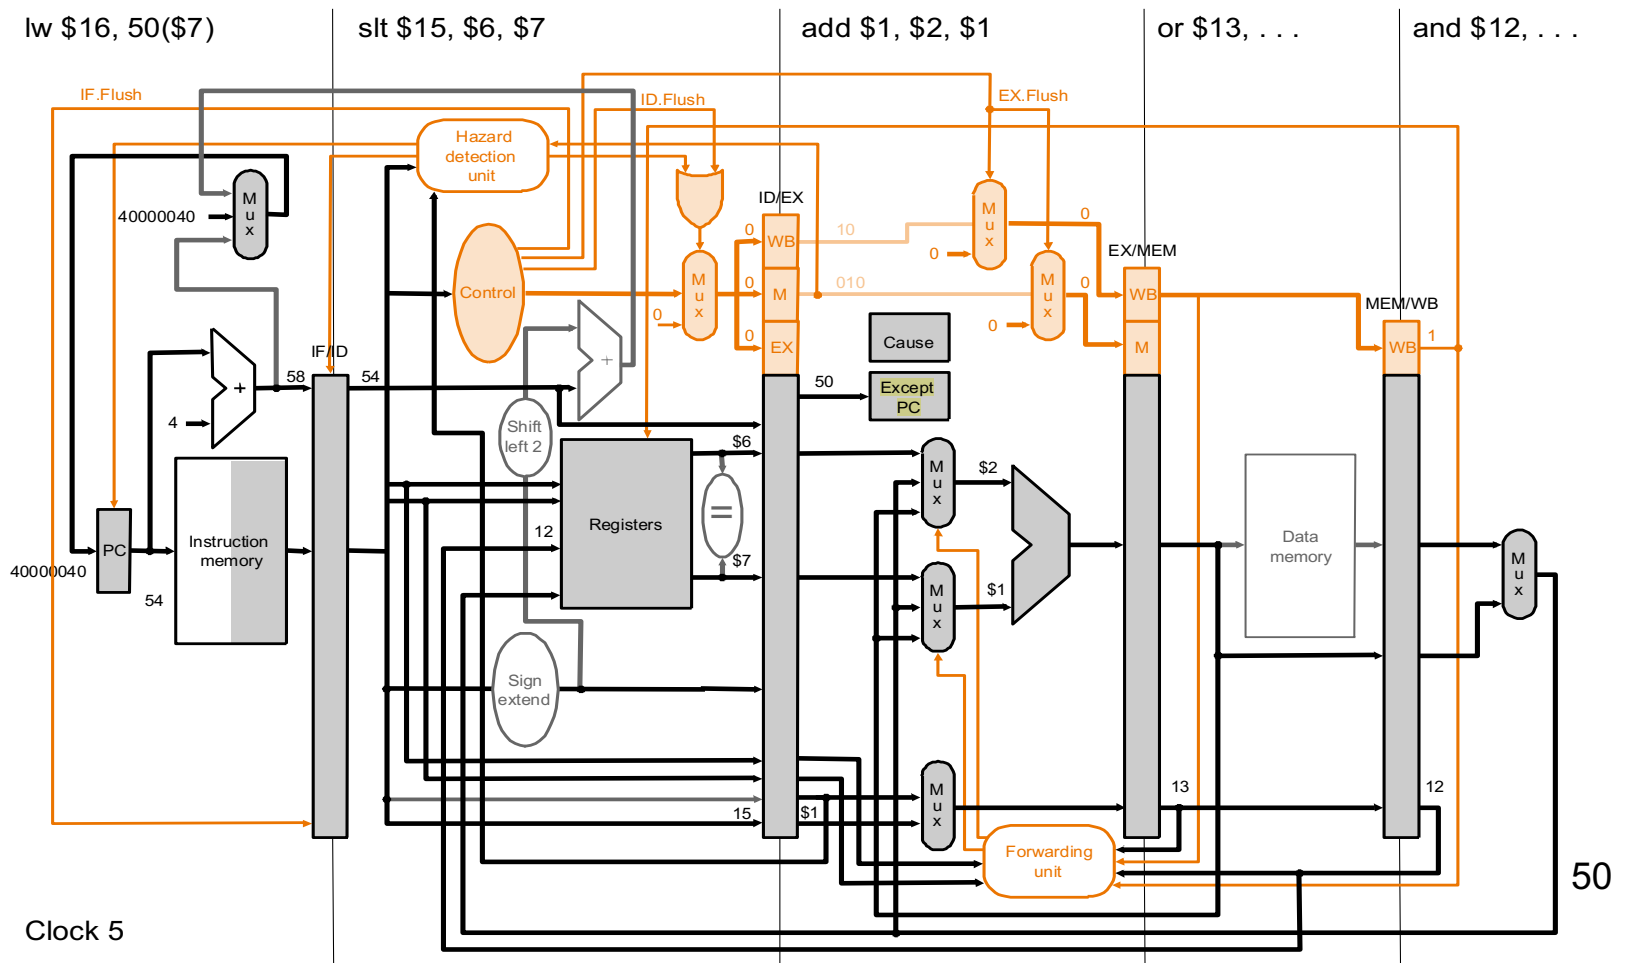
## Example: Handling Exception

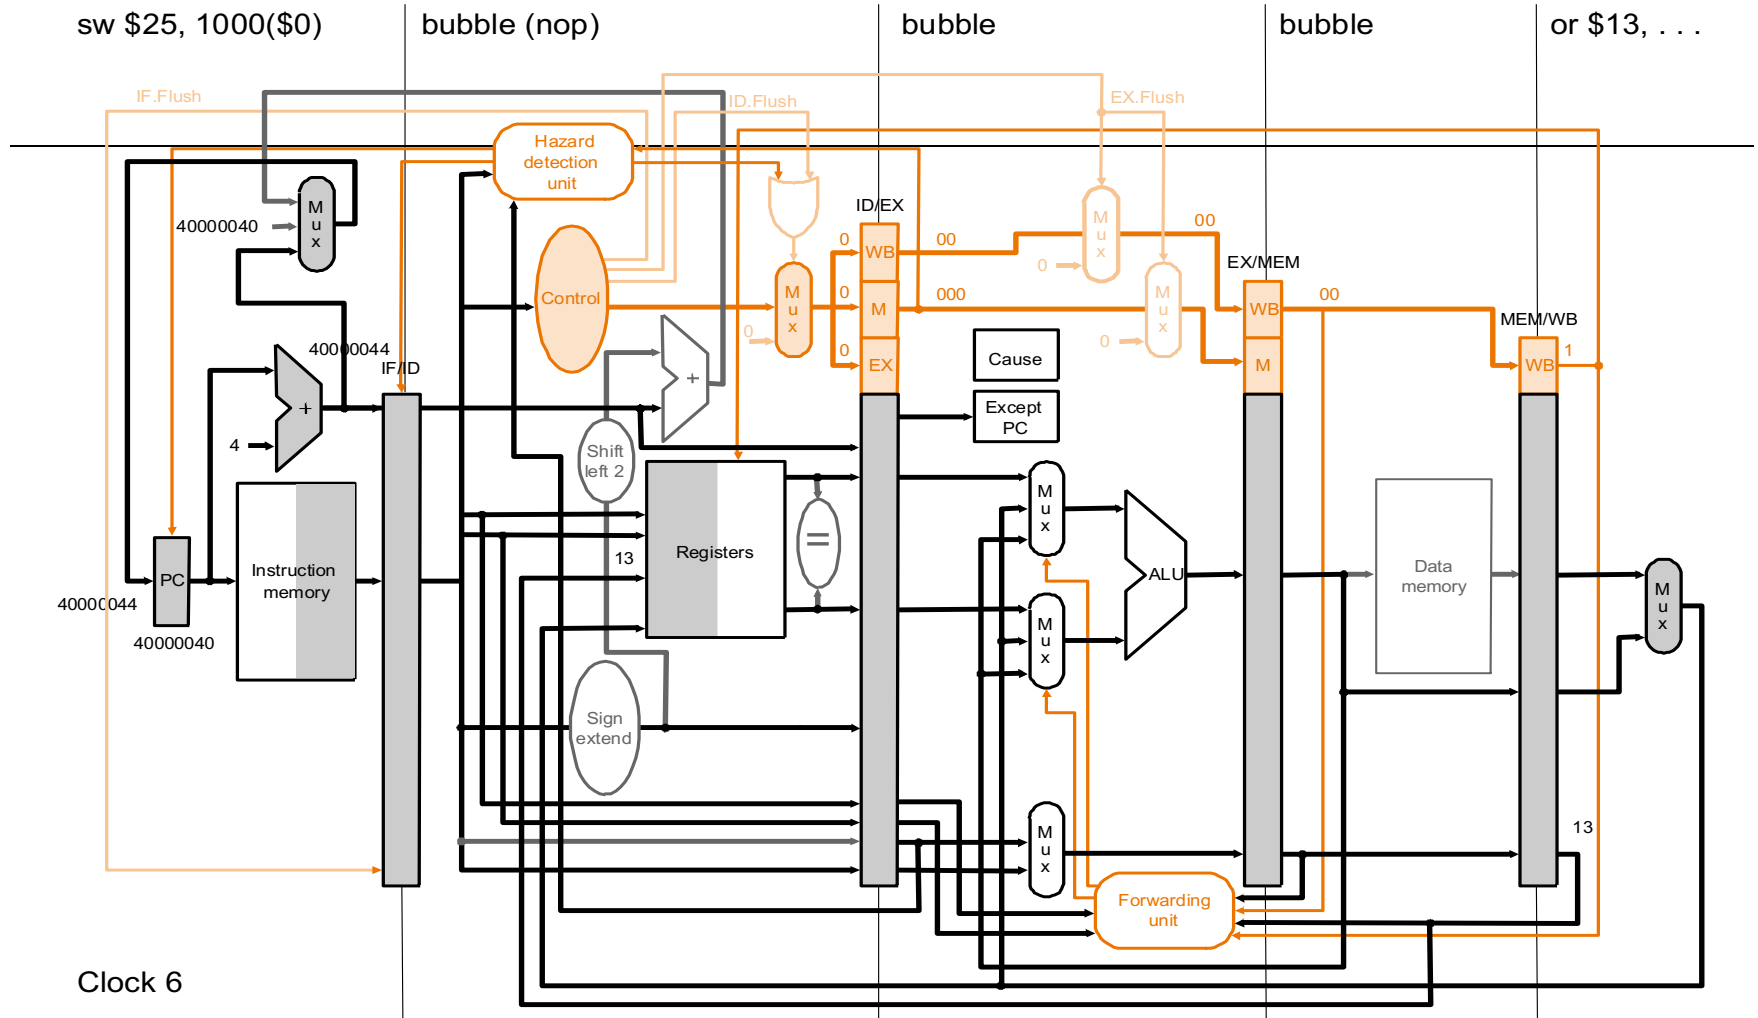
---

- Given the following instruction sequence:
  - 40hex      sub      \$11, \$2, \$4
  - 44hex      and      \$12, \$2, \$5
  - 48hex      or      \$13, \$2, \$6
  - 4Chex      add      \$1, \$2, \$1
  - 50hex      slt      \$15, \$6, \$7
  - 54hex      lw      \$16, 50(\$7)
- Assume the instruction to be invoked on an exception begin like this
  - 40000040hex      sw      \$25, 1000(\$0)
  - 40000044hex      sw      \$26, 1004(\$0)
- Show what happens in the pipeline if an overflow exception occurs in the **add** instruction.

# Example: Handling Exception

1. The overflow is detected when “add” is in the EXE stage.
2. Save PC+4 (50) in **EPC**
3. Assert IF.Flush, ID. Flush, and EX.Flush





1. Fetch stage: the first instruction of the exception routine
2. Instruction prior to the add instruction complete

# Multiple Issue

---

## ■ Static multiple issue

- Compiler groups instructions to be issued together
- Packages them into “issue slots”
- Compiler detects and avoids hazards

## ■ Dynamic multiple issue

- CPU examines instruction stream and chooses instructions to issue each cycle
- Compiler can help by reordering instructions
- CPU resolves hazards using advanced techniques at runtime

|    |    |    |     |     |     |    |
|----|----|----|-----|-----|-----|----|
| IF | ID | EX | MEM | WB  |     |    |
| IF | ID | EX | MEM | WB  |     |    |
|    | IF | ID | EX  | MEM | WB  |    |
|    | IF | ID | EX  | MEM | WB  |    |
|    |    | IF | ID  | EX  | MEM | WB |
|    |    | IF | ID  | EX  | MEM | WB |

# Speculation

---

- “Guess” what to do with an instruction
  - Start operation as soon as possible
  - Check whether guess was right
    - If so, complete the operation
    - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
  - Speculate on branch outcome
    - Roll back if path taken is different
  - Speculate on load
    - Roll back if location is updated

# Compiler/Hardware Speculation

---

- Compiler can reorder instructions
  - e.g., move load before branch
  - Can include “fix-up” instructions to recover from incorrect guess
- Hardware can look ahead for instructions to execute
  - Buffer results until it determines they are actually needed
  - Flush buffers on incorrect speculation

# Static Multiple Issue

---

- Compiler groups instructions into “issue packets”
  - Group of instructions that can be issued on a single cycle
  - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
  - Specifies multiple concurrent operations
  - $\Rightarrow$  Very Long Instruction Word (VLIW)
- Compiler must remove some/all hazards
  - Reorder instructions into issue packets
  - No dependencies with a packet
  - Possibly some dependencies between packets
    - Varies between ISAs; compiler must know!
  - Pad with nop if necessary

# MIPS with Static Dual Issue

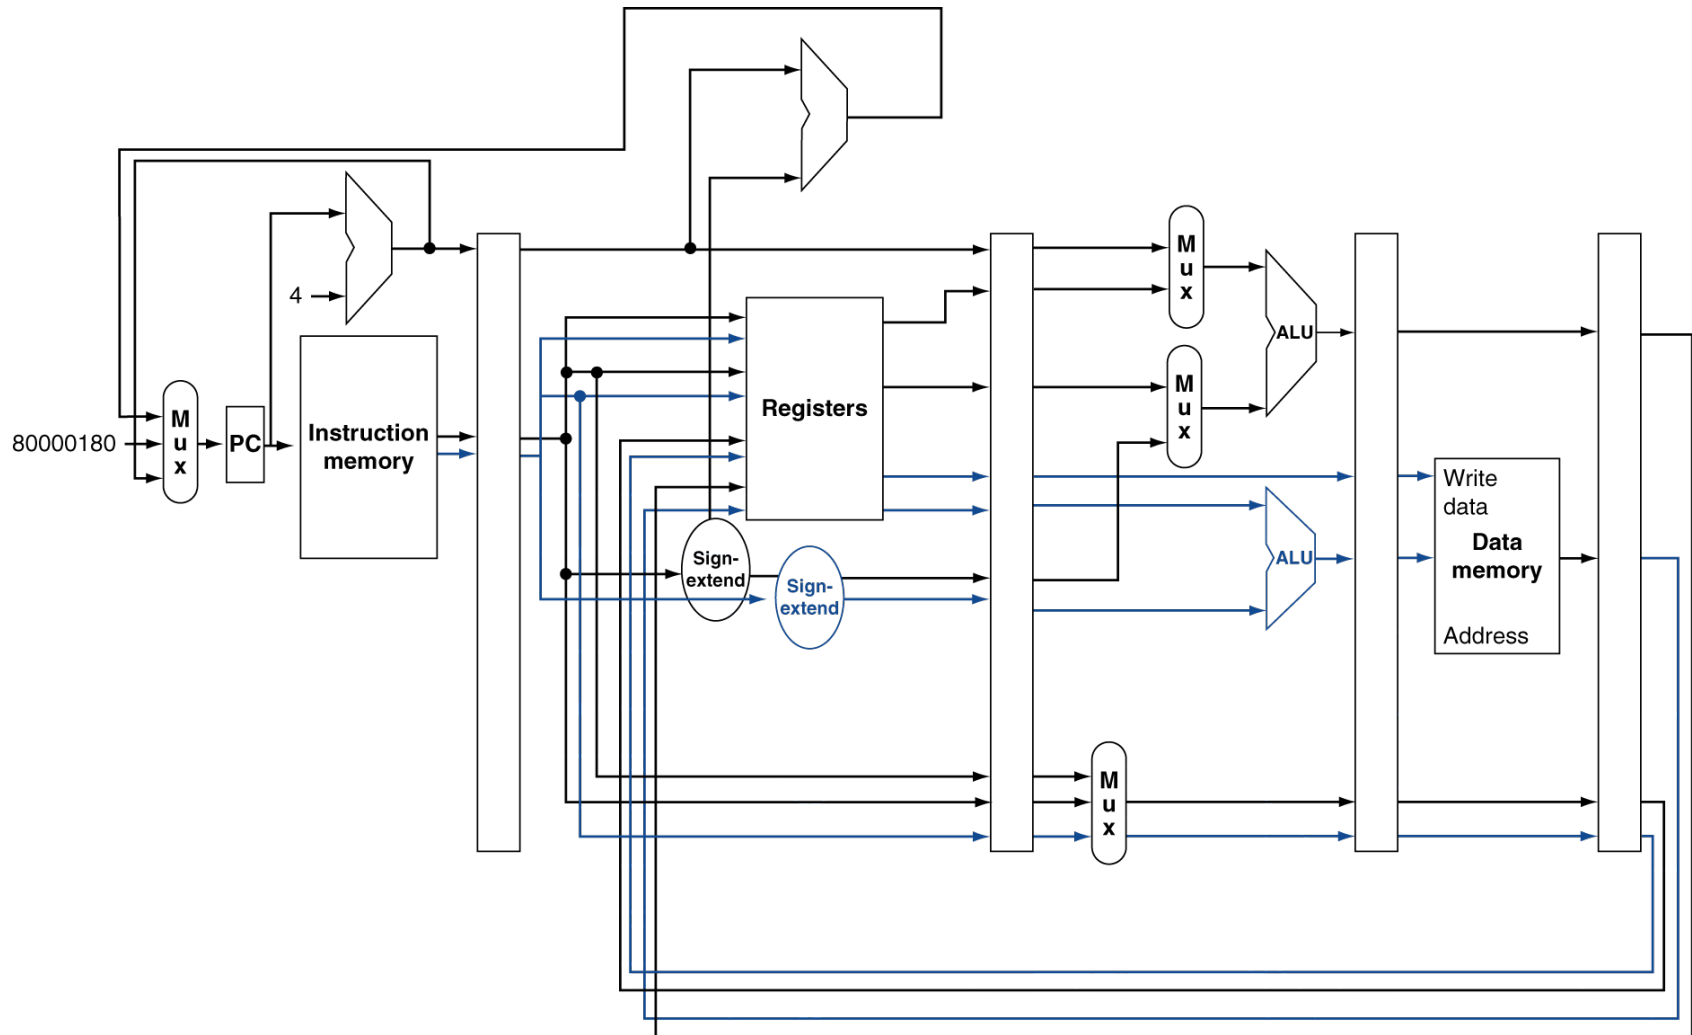
---

- Two-issue packets
  - One ALU/branch instruction
  - One load/store instruction
  - 64-bit aligned
    - ALU/branch, then load/store
    - Pad an unused instruction with nop

| Address | Instruction type | Pipeline Stages |    |    |     |     |     |    |
|---------|------------------|-----------------|----|----|-----|-----|-----|----|
| n       | ALU/branch       | IF              | ID | EX | MEM | WB  |     |    |
| n + 4   | Load/store       | IF              | ID | EX | MEM | WB  |     |    |
| n + 8   | ALU/branch       |                 | IF | ID | EX  | MEM | WB  |    |
| n + 12  | Load/store       |                 | IF | ID | EX  | MEM | WB  |    |
| n + 16  | ALU/branch       |                 |    | IF | ID  | EX  | MEM | WB |
| n + 20  | Load/store       |                 |    | IF | ID  | EX  | MEM | WB |



# MIPS with Static Dual Issue



# Hazards in the Dual-Issue MIPS

---

- More instructions executing in parallel
- EX data hazard
  - Forwarding avoided stalls with single-issue
  - Now can't use ALU result in load/store in same packet
    - add `$t0`, `$s0`, `$s1`  
load `$s2`, `0($t0)`
    - Split into two packets, effectively a stall
- Load-use hazard
  - Still one cycle use latency, but now two instructions
- More aggressive scheduling required

# Scheduling Example

## ■ Schedule this for dual-issue MIPS

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
```

|       | ALU/branch               | Load/store          | cycle |
|-------|--------------------------|---------------------|-------|
| Loop: | nop                      | lw    \$t0, 0(\$s1) | 1     |
|       | addi  \$s1, \$s1, -4     | nop                 | 2     |
|       | addu  \$t0, \$t0, \$s2   | nop                 | 3     |
|       | bne   \$s1, \$zero, Loop | sw    \$t0, 4(\$s1) | 4     |

■  $IPC = 5/4 = 1.25$  (c.f. peak  $IPC = 2$ )

# Loop Unrolling

---

- Replicate loop body to expose more parallelism
  - Reduces loop-control overhead
- Use different registers per replication
  - Called “register renaming”
  - Avoid loop-carried “anti-dependencies”
    - Store followed by a load of the same register
    - Aka “name dependence”
      - Reuse of a register name

---

■ Loop: lw \$t0, 0(\$s1)  
addu \$t0, \$t0, \$s2  
sw \$t0, 0(\$s1)  
addi \$s1, \$s1, -4  
bne \$s1, \$zero, Loop

# Unrolled Loop That Minimizes Stalls

|  | Renaming |  |    |
|--|----------|--|----|
| <b>Loop:</b>   | →        | <b>Loop:</b>   |    |
| lw <span style="color: blue;">\$t0</span> , 0(\$s1)  |          | lw <span style="color: blue;">\$t0</span> , 0(\$s1)  |    |
| addu <span style="color: green;">\$t0</span> , <span style="color: blue;">\$t0</span> , \$s2 |          | addu <span style="color: green;">\$t0</span> , <span style="color: blue;">\$t0</span> , \$s2 |    |
| sw <span style="color: green;">\$t0</span> , 0(\$s1)   |          | sw <span style="color: green;">\$t0</span> , 0(\$s1)   |    |
| <br>lw <span style="color: blue;">\$t0</span> , -4(\$s1)                                     |          | <br>lw <span style="color: blue;">\$t1</span> , -4(\$s1)                                     |    |
| addu <span style="color: green;">\$t0</span> , <span style="color: blue;">\$t0</span> , \$s2 |          | addu <span style="color: green;">\$t1</span> , <span style="color: blue;">\$t0</span> , \$s2 |    |
| sw <span style="color: green;">\$t0</span> , -4(\$s1)  |          | sw <span style="color: green;">\$t1</span> , -4(\$s1)  |    |
| <br>lw <span style="color: blue;">\$t0</span> , -8(\$s1)                                     |          | <br>lw <span style="color: blue;">\$t2</span> , -8(\$s1)                                     |    |
| addu <span style="color: green;">\$t0</span> , <span style="color: blue;">\$t0</span> , \$s2 |          | addu <span style="color: green;">\$t2</span> , <span style="color: blue;">\$t0</span> , \$s2 |    |
| sw <span style="color: green;">\$t0</span> , -8(\$s1)  |          | sw <span style="color: green;">\$t2</span> , -8(\$s1)  |    |
| <br>lw <span style="color: blue;">\$t0</span> , -12(\$s1)                                    |          | <br>lw <span style="color: blue;">\$t3</span> , -12(\$s1)                                    |    |
| addu <span style="color: green;">\$t0</span> , <span style="color: blue;">\$t0</span> , \$s2 |          | addu <span style="color: green;">\$t3</span> , <span style="color: blue;">\$t0</span> , \$s2 |    |
| sw <span style="color: green;">\$t0</span> , -12(\$s1)                                       |          | sw <span style="color: green;">\$t3</span> , -12(\$s1)                                       |    |
| <br>addi <span style="color: brown;">\$s1</span> , \$s1, -16                                 |          | <br>addi <span style="color: brown;">\$s1</span> , \$s1, -16                                 |    |
| bne <span style="color: brown;">\$s1</span> , \$zero, Loop                                   |          | bne <span style="color: brown;">\$s1</span> , \$zero, Loop                                   | 62 |

# Loop Unrolling Example

---

|       | ALU/branch                            | Load/store                | cycle |
|-------|---------------------------------------|---------------------------|-------|
| Loop: | addi <b>\$s1</b> , \$s1, -16          | lw <b>\$t0</b> , 0(\$s1)  | 1     |
|       | nop                                   | lw <b>\$t1</b> , 12(\$s1) | 2     |
|       | addu <b>\$t0</b> , <b>\$t0</b> , \$s2 | lw <b>\$t2</b> , 8(\$s1)  | 3     |
|       | addu <b>\$t1</b> , <b>\$t1</b> , \$s2 | lw <b>\$t3</b> , 4(\$s1)  | 4     |
|       | addu <b>\$t2</b> , <b>\$t2</b> , \$s2 | sw <b>\$t0</b> , 16(\$s1) | 5     |
|       | addu <b>\$t3</b> , <b>\$t4</b> , \$s2 | sw <b>\$t1</b> , 12(\$s1) | 6     |
|       | nop                                   | sw <b>\$t2</b> , 8(\$s1)  | 7     |
|       | bne <b>\$s1</b> , \$zero, Loop        | sw <b>\$t3</b> , 4(\$s1)  | 8     |

■  $IPC = 14/8 = 1.75$

– Closer to 2, but at cost of registers and code size

# Dynamic Multiple Issue

---

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, ... each cycle
  - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
  - Though it may still help
  - Code semantics ensured by the CPU



# Dynamic Pipeline Scheduling

---

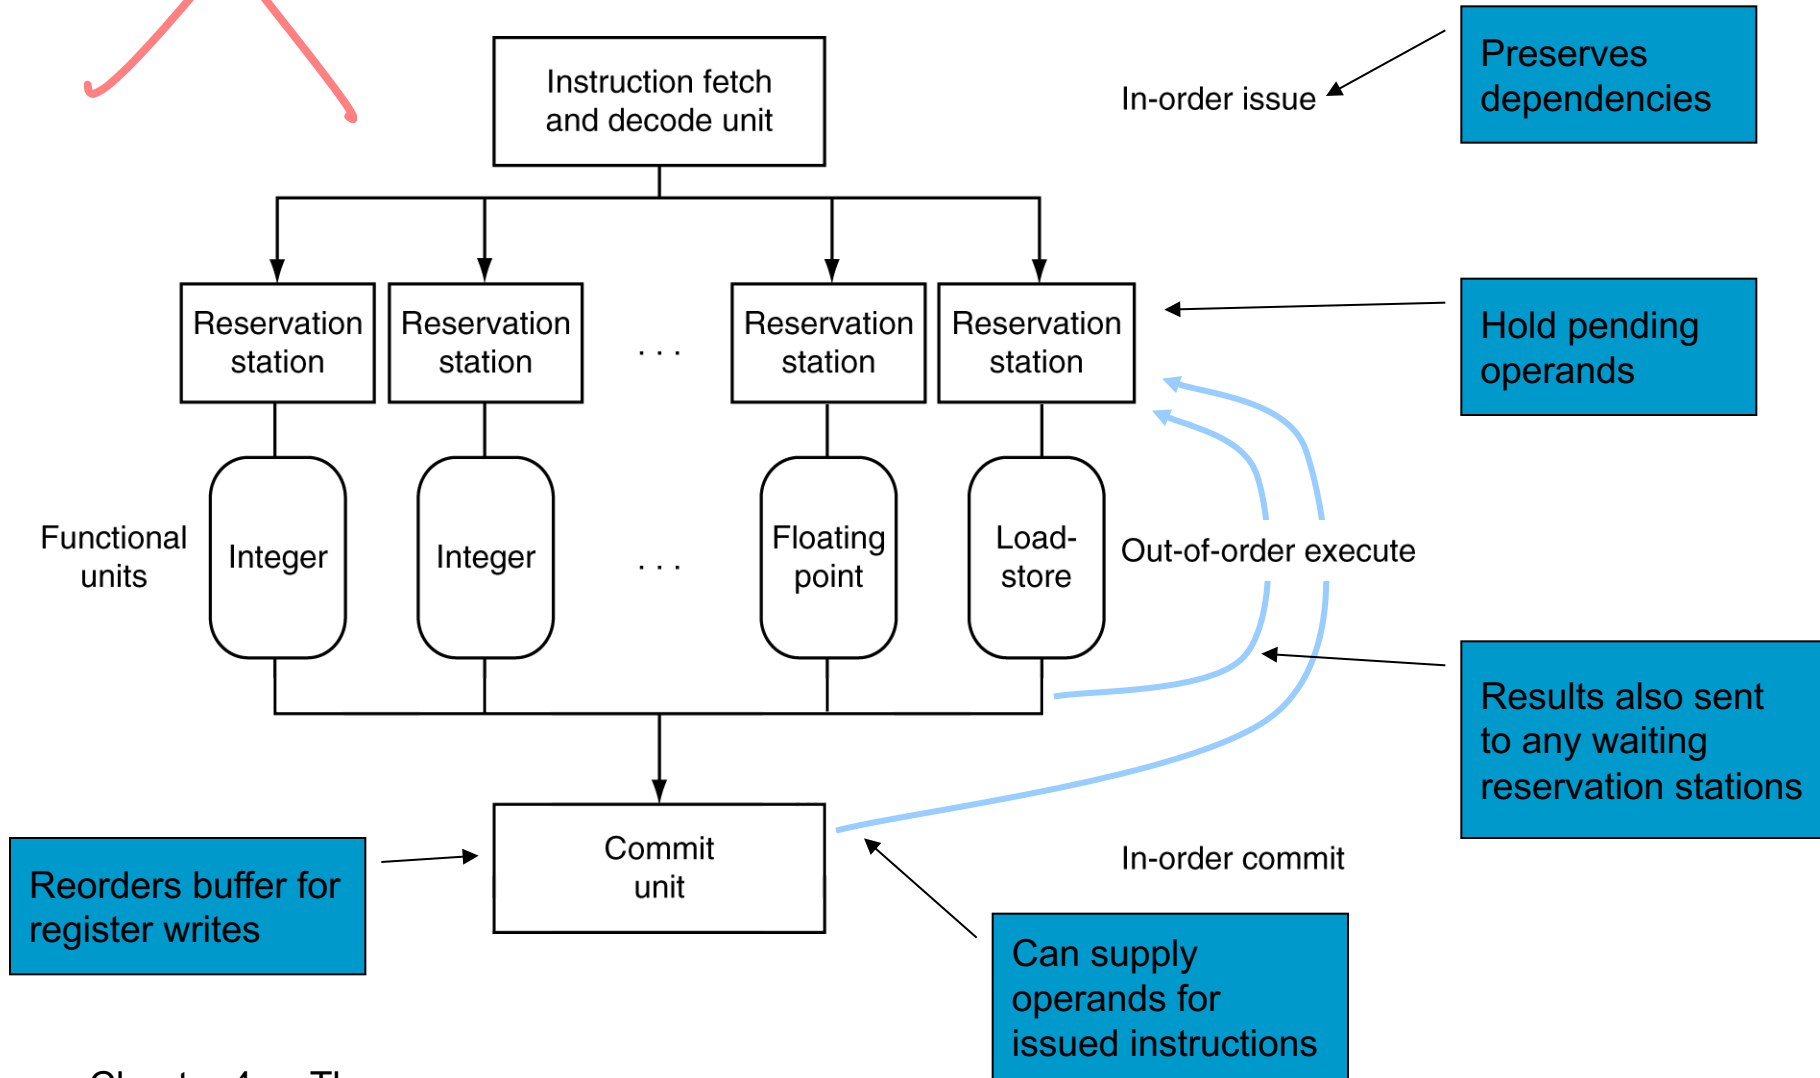
- Allow the CPU to execute instructions out of order to avoid stalls
  - But commit result to registers in order

- Example

```
lw      $t0, 20($s2)
addu    $t1, $t0, $t2
sub      $s4, $s4, $t3
slti    $t5, $s4, 20
```

- Can start sub while addu is waiting for lw

# Dynamically Scheduled CPU



# Why Do Dynamic Scheduling?

---

- Why not just let the compiler schedule code?
- Not all stalls are predictable
  - e.g., cache misses
- Can't always schedule around branches
  - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

# Does Multiple Issue Work?

---

## The BIG Picture

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Memory delays and limited bandwidth
  - Hard to keep pipelines full
- Speculation can help if done well

# Power Efficiency

---

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

| Microprocessor | Year | Clock Rate | Pipeline Stages | Issue width | Out-of-order/ Speculation | Cores | Power |
|----------------|------|------------|-----------------|-------------|---------------------------|-------|-------|
| i486           | 1989 | 25MHz      | 5               | 1           | No                        | 1     | 5W    |
| Pentium        | 1993 | 66MHz      | 5               | 2           | No                        | 1     | 10W   |
| Pentium Pro    | 1997 | 200MHz     | 10              | 3           | Yes                       | 1     | 29W   |
| P4 Willamette  | 2001 | 2000MHz    | 22              | 3           | Yes                       | 1     | 75W   |
| P4 Prescott    | 2004 | 3600MHz    | 31              | 3           | Yes                       | 1     | 103W  |
| Core           | 2006 | 2930MHz    | 14              | 4           | Yes                       | 2     | 75W   |
| UltraSparc III | 2003 | 1950MHz    | 14              | 4           | No                        | 1     | 90W   |
| UltraSparc T1  | 2005 | 1200MHz    | 6               | 1           | No                        | 8     | 70W   |

## Concluding Remarks

---

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
  - More instructions completed per second
  - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
  - Dependencies limit achievable parallelism
  - Complexity leads to the power wall