


# Multicores, Multiprocessor (Part I)

----- The Simple Components of a Computer -----



# Parallel Computers

- **Goal: connecting multiple computers to get higher performance**
  - Multiprocessors
  - Scalability, availability, power efficiency
- **Parallel Workloads**
  - Job-level (process-level) parallelism
    - » High throughput for independent jobs
  - Parallel processing program
    - » Single program run on multiple processors

# Challenge of Parallel Processing

- **Challenges**
  - Partitioning
  - Coordination
  - Communication overheads
- **Available Parallelism in applications**
  - Amdahl's Law (FracX: original % to be speed up)  
**Speedup =  $1 / [(FracX/SpeedupX + (1-FracX))]$**

What fraction of the original computation can be sequential to get 80X speedup from 100 processors?

Assume either 1 processor or 100 fully used

$$80 = 1 / [(FracX/100 + (1-FracX))]$$

$$0.8 * FracX + 80 * (1 - FracX) = 80 - 79.2 * FracX = 1$$

$$FracX = (80 - 1) / 79.2 = 0.9975$$

• Only 0.25% sequential!

# Scaling Example

- **Workload: sum of 10 scalars, and  $10 \times 10$  matrix sum**
  - Speed up from 10 to 100 processors
- **Single processor: Time =  $(10 + 100) \times t_{\text{add}}$**
- **10 processors**
  - Time =  $10 \times t_{\text{add}} + 100/10 \times t_{\text{add}} = 20 \times t_{\text{add}}$
  - Speedup =  $110/20 = 5.5$  (55% of potential)
- **100 processors**
  - Time =  $10 \times t_{\text{add}} + 100/100 \times t_{\text{add}} = 11 \times t_{\text{add}}$
  - Speedup =  $110/11 = 10$  (10% of potential)
- **Assumes load can be balanced across processors**

## Scaling Example (cont)


- What if matrix size is  $100 \times 100$ ?
- Single processor: Time =  $(10 + 10000) \times t_{\text{add}}$
- 10 processors
  - Time =  $10 \times t_{\text{add}} + 10000/10 \times t_{\text{add}} = 1010 \times t_{\text{add}}$
  - Speedup =  $10010/1010 = 9.9$  (99% of potential)
- 100 processors
  - Time =  $10 \times t_{\text{add}} + 10000/100 \times t_{\text{add}} = 110 \times t_{\text{add}}$
  - Speedup =  $10010/110 = 91$  (91% of potential)
- Assuming load balanced

# Strong vs Weak Scaling

- **Strong scaling: problem size fixed**
- **Weak scaling: problem size proportional to number of processors**

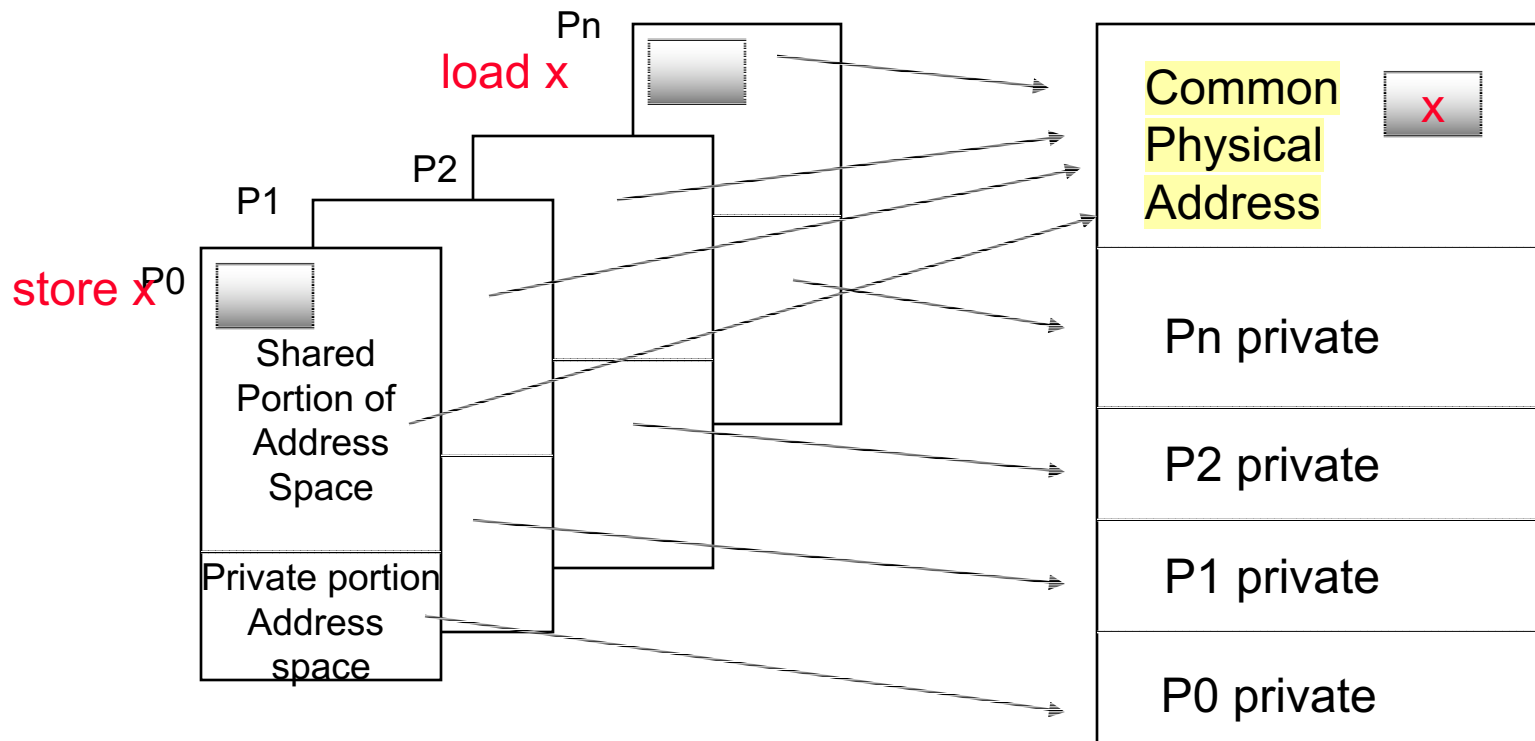
# Shared Memory

- **SMP: shared memory multiprocessor**
  - Hardware provides single physical address space for all processors
  - Synchronize shared variables using locks
  - Memory access time
    - » UMA (uniform) vs. NUMA (nonuniform)



# Communication Models

- Single Address Space: load/store






# Program Example – Single-Address Space

- sum 100,000 numbers & 100 processors (load & store)

**First Step: each processor ( $P_n$ ) sums his subset of numbers**

```
sum[Pn] = 0;      → Sum is shared variable  
for (i = 1000*Pn; i < 1000*(Pn+1); i = i + 1)  
    sum[Pn] = sum[Pn] + A[i]; /* sum the assigned areas*/
```




P0

# Program Example – Single-Address Space

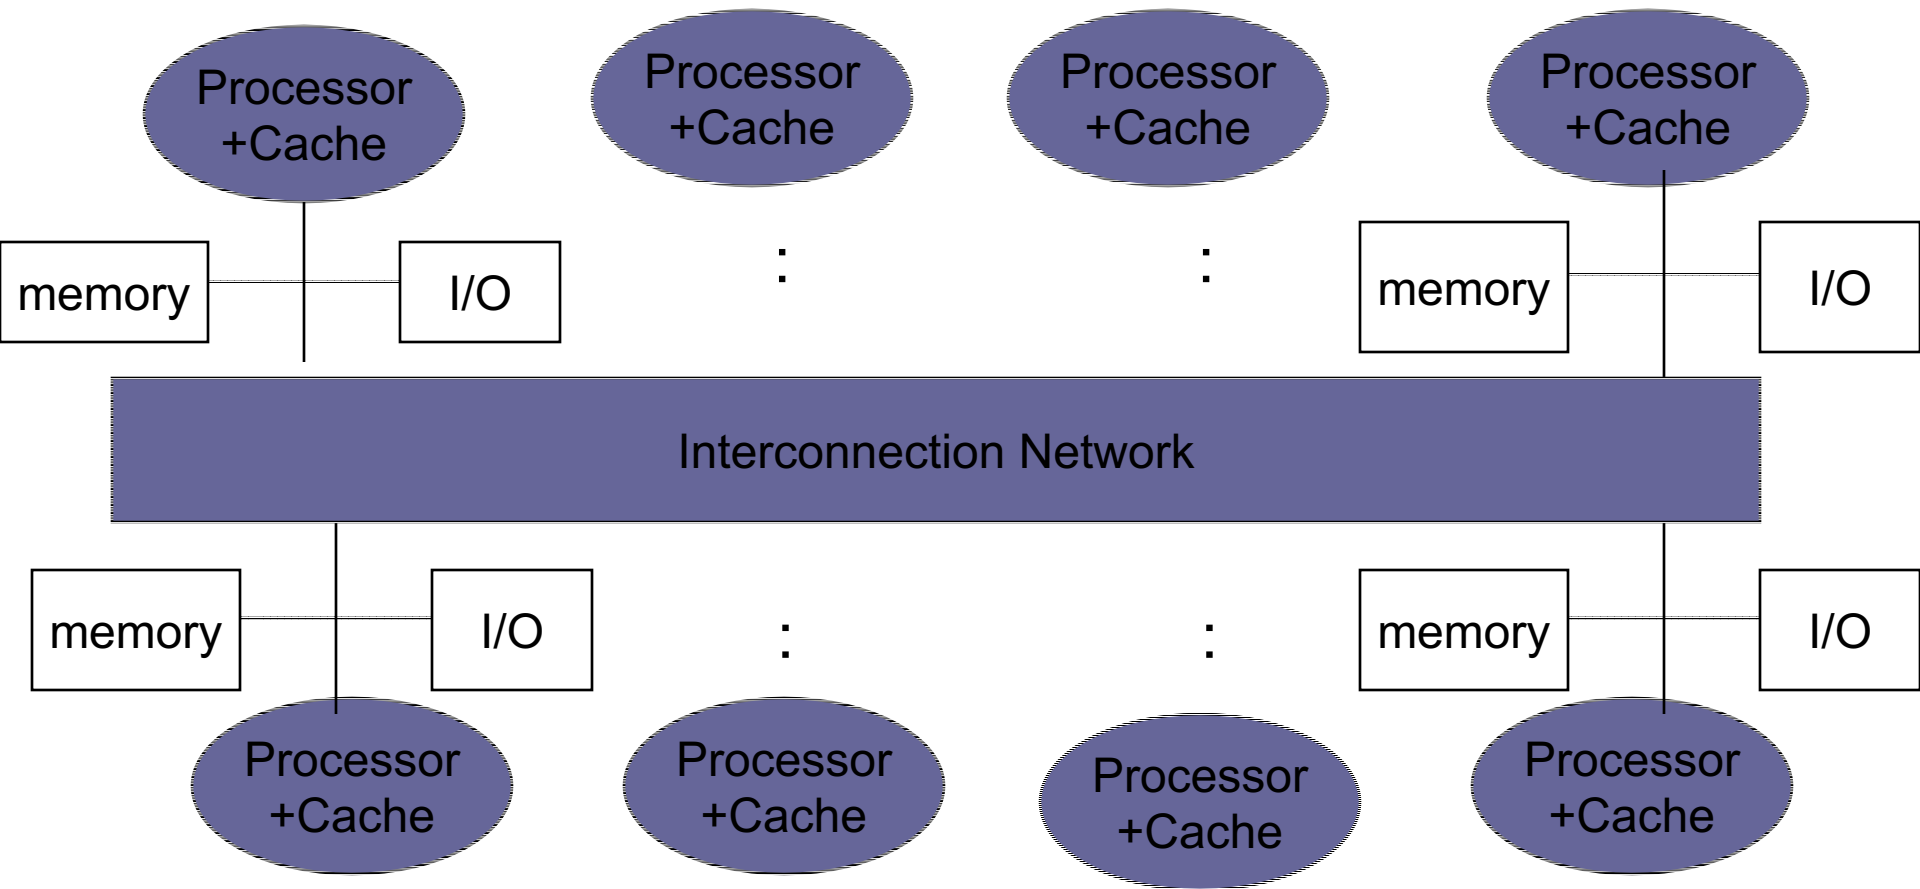
## Second Step: Add partial sums via divide-and-conquer

```
half = 100; /* 100 processors in multiprocessor*/  
repeat  
    synch(); /* wait for partial sum completion*/  
    if (half%2 != 0 && Pn == 0)  
        sum[0] = sum[0] + sum[half-1];  
        /* Conditional sum needed when half is  
        odd; Processor0 gets missing element */  
    half = half/2; /* dividing line on who sums */  
    if (Pn < half) sum[Pn] = sum[Pn] + sum[Pn+half];  
until (half == 1); /* exit with final sum in Sum[0] */
```



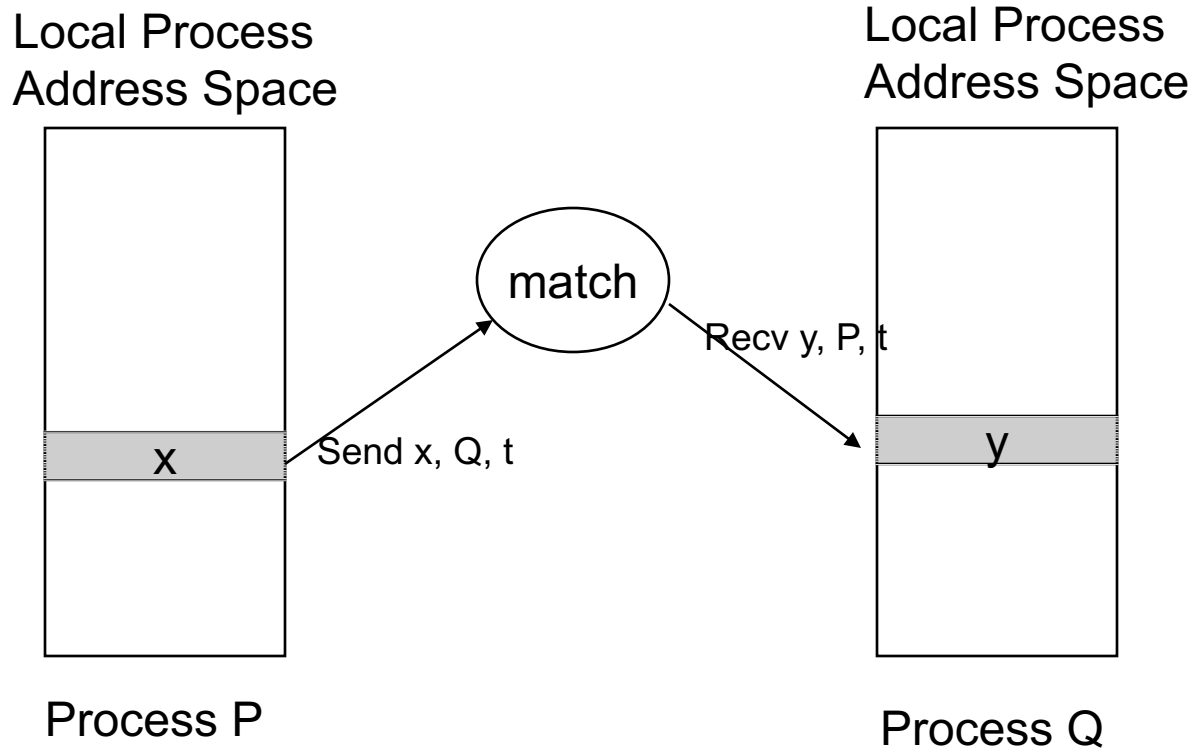
# Message Passing Multiprocessors

- **Clusters:** Collections of computers connected via I/O over standard network switches to form a message-passing multiprocessors



# Communication Models

- Multiple address spaces: **Message Passing**



# Parallel Program – Message Passing

- **sum 100,000 numbers & 100 processors (send & receive)**


**First Step: each processor (P<sub>n</sub>) sums his subset of numbers**

```
sum = 0;  
for (i = 0; i < 1000; i = i + 1) /* loop over each array */  
    sum = sum + A1[i]; /* sum the local arrays */
```

# Parallel Program – Message Passing


## Second Step: Add partial sums via divide-and-conquer

```
limit = 100; half = 100; /* 100 processors */
repeat
  half = (half+1)/2; /* send vs. receive dividing line*/
  if (Pn >= half && Pn < limit) send(Pn - half, sum);
  if (Pn < (limit/2)) sum = sum + receive();
  limit = half; /* upper limit of senders */
until (half == 1); /* exit with final sum */
```



# Bisection Bandwidth is Important

Bus Multicore




Total network bandwidth =

$$\text{bandwidth-per-link} \times \text{link\_no}$$


Bisection bandwidth =

the bandwidth between two equal parts of a multiprocessor


Ring Multicore



# Network Topology




Fully-connected



2D torus



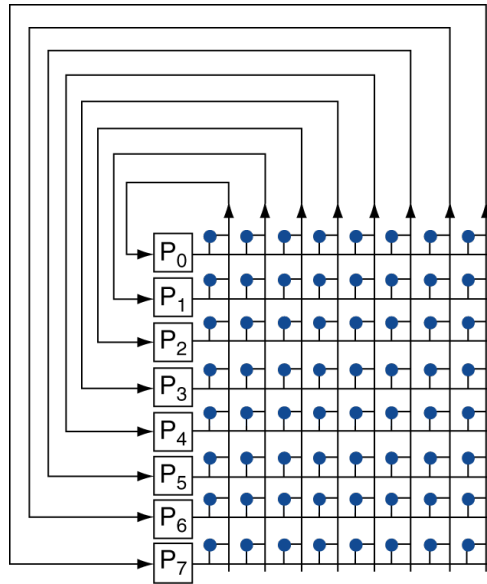
Ring




Cube




# Multistage Networks



a. Crossbar



b. Omega network




c. Omega network switch box

# Network Characteristics

- **Performance**
  - Latency per message (unloaded network)
  - Throughput
    - » Link bandwidth
    - » Total network bandwidth
    - » Bisection bandwidth
  - Congestion delays (depending on traffic)
- **Cost**
- **Power**
- **Routability in silicon**

# Cache Coherency Problem

- Traffic per processor and the bus bandwidth determine the # of processors
- Caches can lower bus traffic
  - Cache coherency problem



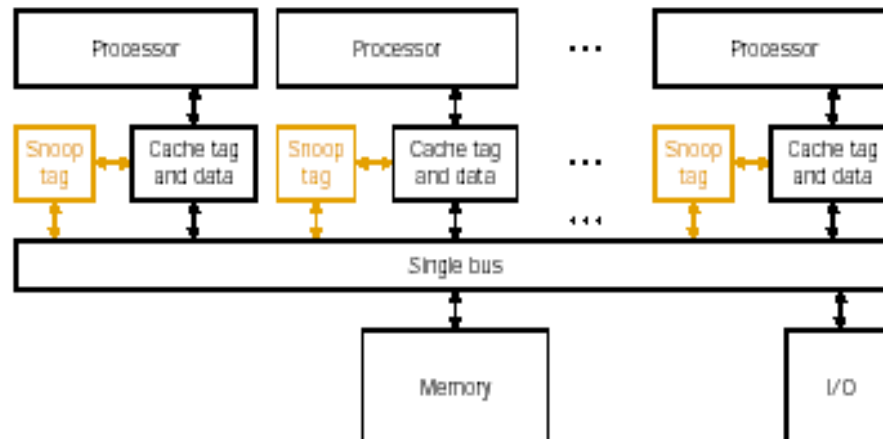
# Cache Coherency

Time	Event	\$ A	\$ B	X (memory)
0				1
1	CPU A reads X	1		1
2	CPU B reads X	1	1	1
3	CPU A stores 0 into X	0	1	0

# Cache Coherency Protocol

5.10

- **Snooping Solution (Snoopy Bus):**
  - Send all requests for data to all processors
  - Processors snoop to see if they have a copy and respond accordingly
  - Requires broadcast, since caching information is at processors
  - Works well with bus (natural broadcast medium)
  - Dominates for small scale machines (most of the market)



# Basic Snoopy Protocols

- **Write Invalidate Protocol:**
  - Multiple readers, single writer
  - Write to shared data: an invalidate is sent to all caches which snoop and *invalidate* any copy
  - Read Miss:
    - » Write-through: memory is always up-to-date
    - » Write-back: snoop in caches to find most recent copy
- **Write Update Protocol:**
  - Write to shared data: broadcast on bus, processors snoop, and *update* copies
  - Read miss: memory is always up-to-date
- What happens if two processors try to write to the same shared data word in the same clock cycle?
  - Write **serialization**: bus serializes requests

# Basic Snoopy Protocols

## • Invalidation write back

Processor activity	Bus activity	Contents of CPU A' cache	Contents of CPU B's cache	Contexts of memory location X
				0
CPU A reads X	Cache miss for X	0		
CPU B reads X	Cache miss for X	0	0	0
CPU A writes a 1 to X	<b>Invalidation for X</b>	1		0
CPU B reads X	Cache miss for X	1	1	<b>1</b>

## • Update

Processor activity	Bus activity	Contents of CPU A' cache	Contents of CPU B's cache	Contexts of memory location X
				0
CPU A reads X	Cache miss for X	0		
CPU B reads X	Cache miss for X	0	0	0
CPU A writes a 1 to X	<b>Write broadcast of X</b>	1	1	1
CPU B reads X		1	1	1

# Basic Snoopy Protocols

- **Write Invalidate versus Broadcast:**
  - Invalidate requires one transaction per write-run
  - Invalidate uses spatial locality: one transaction per block
  - Broadcast: BW (increased) vs. latency (decreased) tradeoff

**Invalidate protocol is more popular than update !**

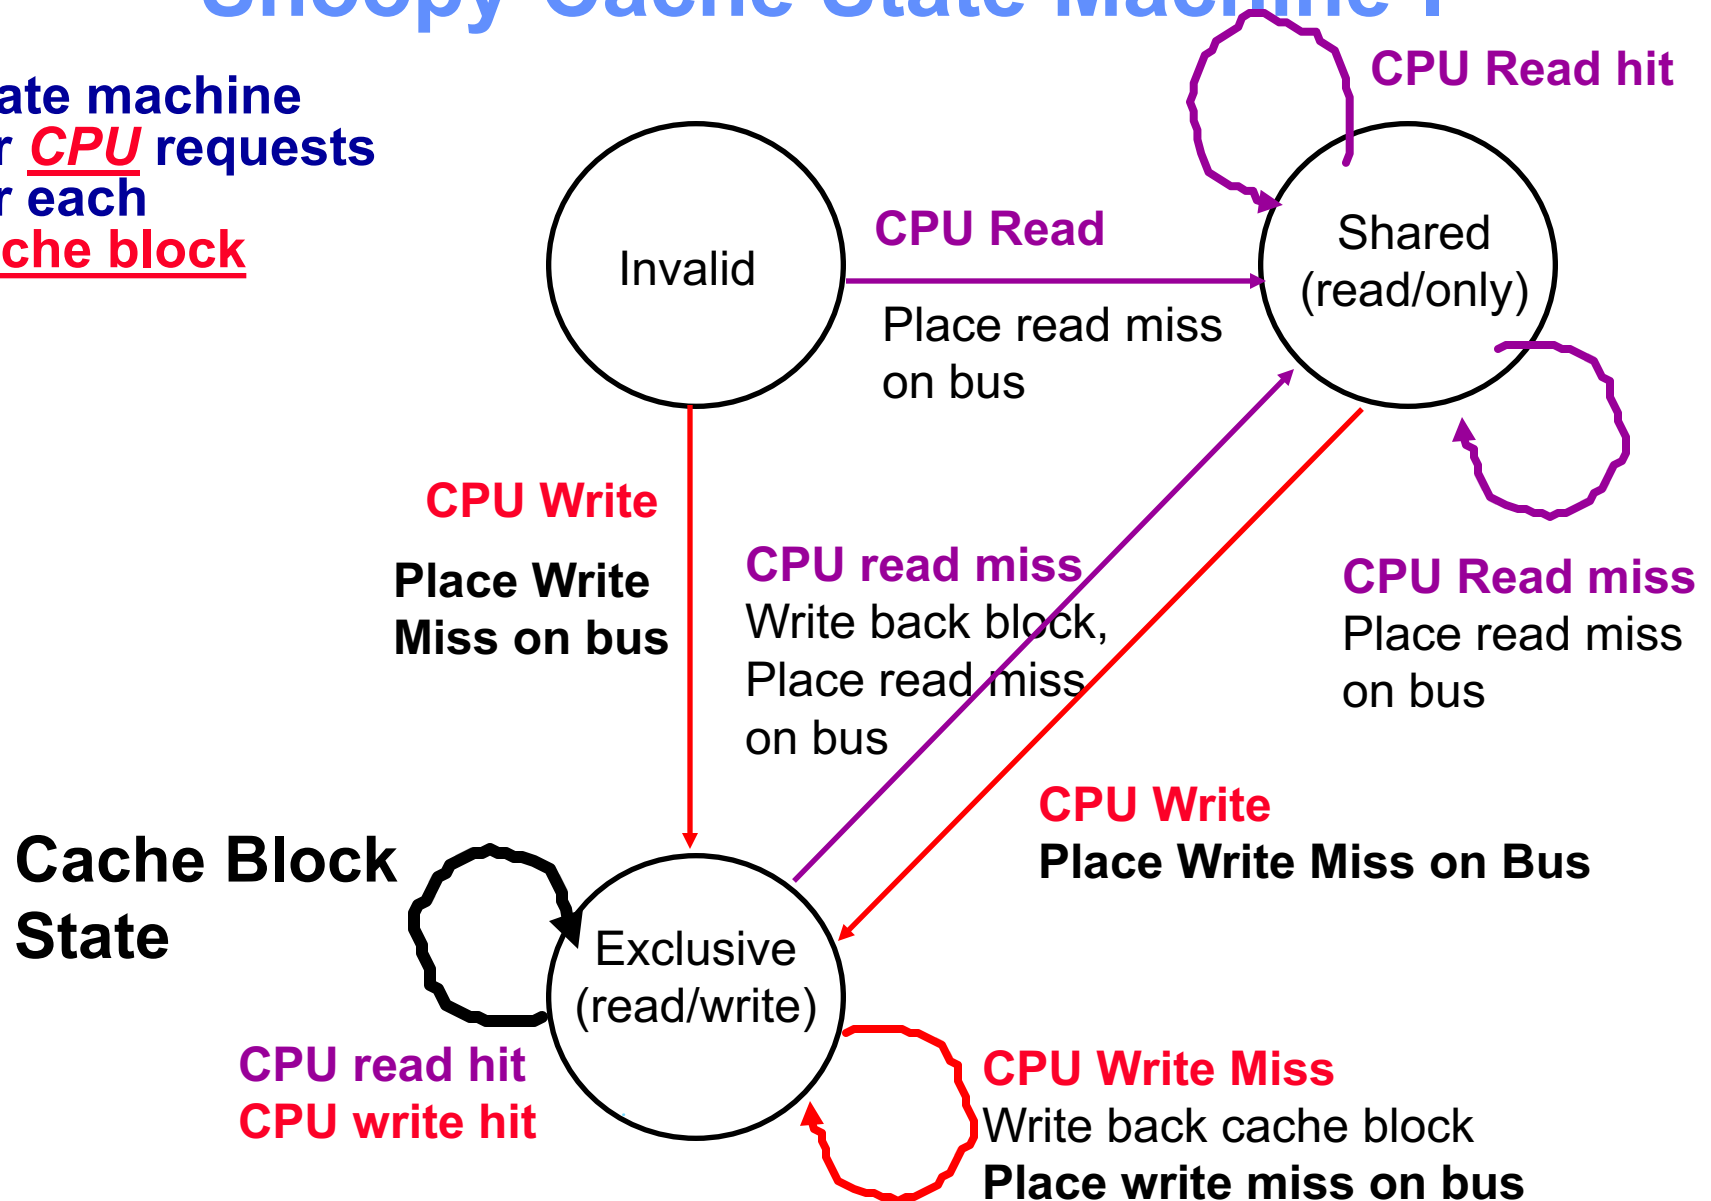


# An Example Snoopy Protocol

- Invalidation protocol, **write-back** cache
- Each block of memory is in one state:
  - **Clean** in all caches and up-to-date in memory
  - OR **Dirty** in exactly one cache
  - OR **Not** in any caches
- Each cache block is in one state:
  - **Shared**: block can be read
  - OR **Exclusive**: cache has only copy, its writeable, and dirty
  - OR Invalid: block contains no data
- **Read misses**: cause all caches to snoop
- Writes to clean line are treated as misses

# Snoopy-Cache State Machine-I

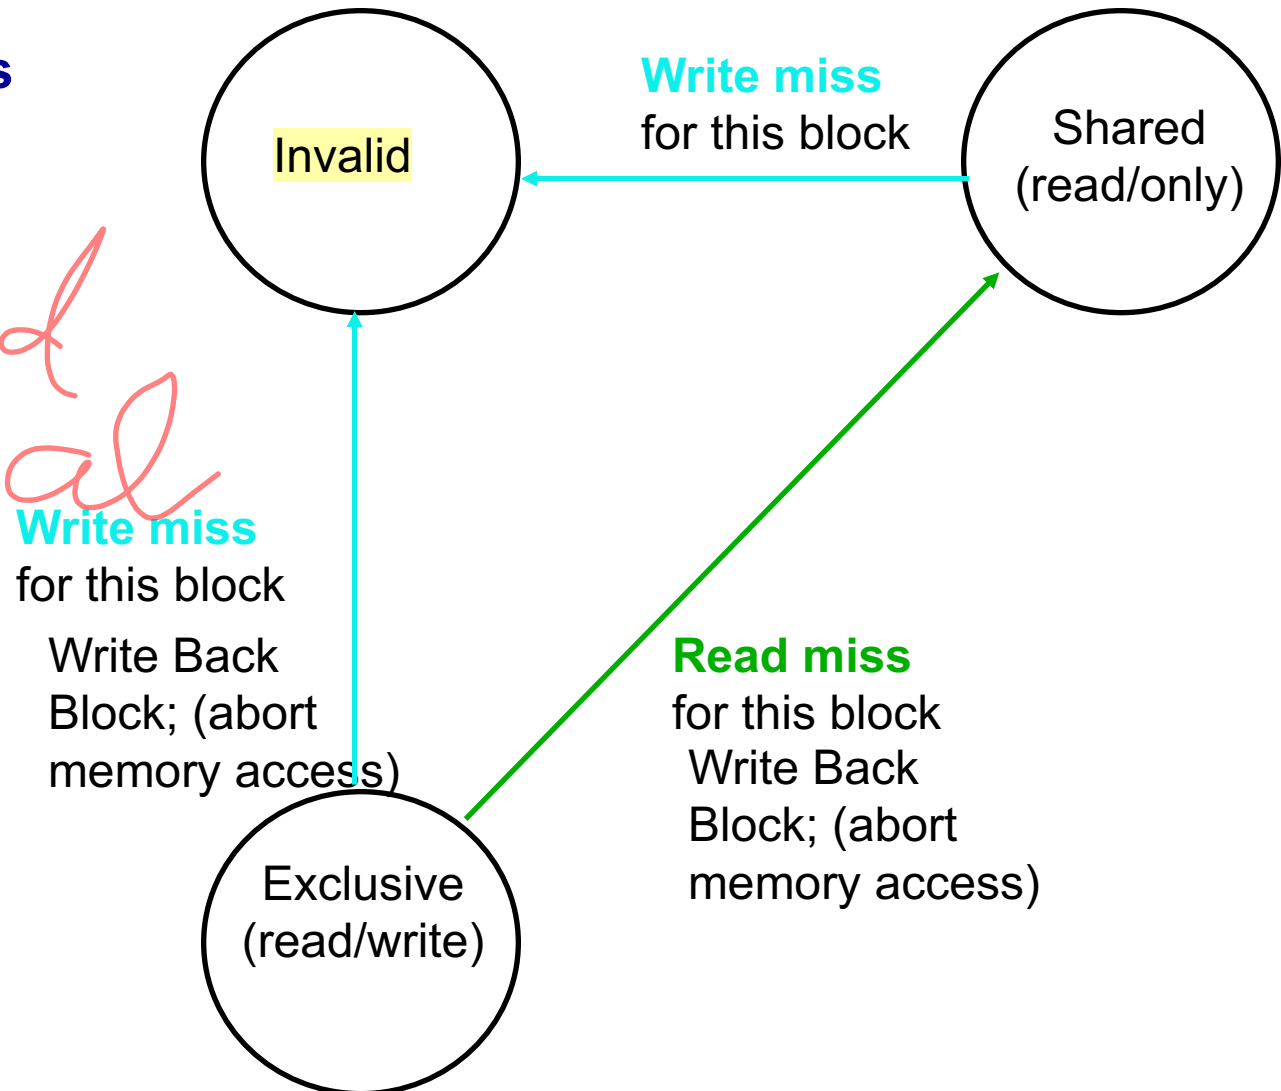
- State machine for **CPU** requests for each cache block



# Snoopy-Cache State Machine-II

- State machine for bus requests for each cache block

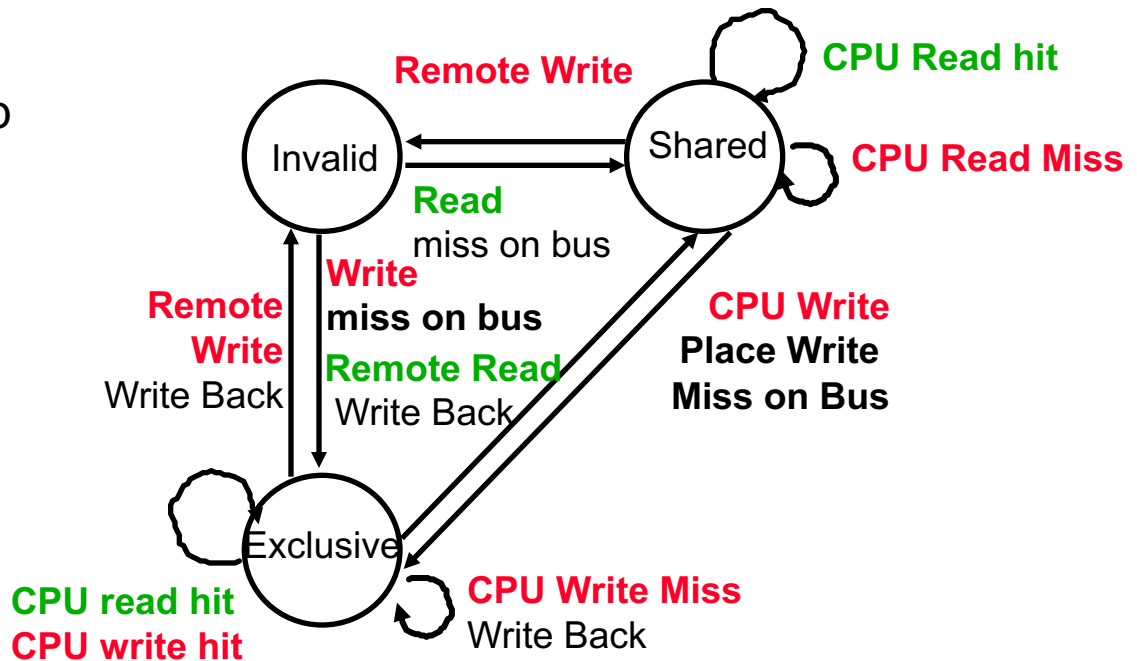
*invalid protocol*



# Example

	Processor 1			Processor 2			Bus			Memory		
	<i>P1</i>			<i>P2</i>			<i>Bus</i>				<i>Memory</i>	
<i>step</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>Action</i>	<i>Proc.</i>	<i>Addr</i>	<i>Value</i>	<i>Addr</i>	<i>Value</i>
P1: Write 10 to A1												
P1: Read A1												
P2: Read A1												
P2: Write 20 to A1												
P2: Write 40 to A2												

Assumes initial cache state is invalid and A1 and A2 map to same cache block, but A1 != A2

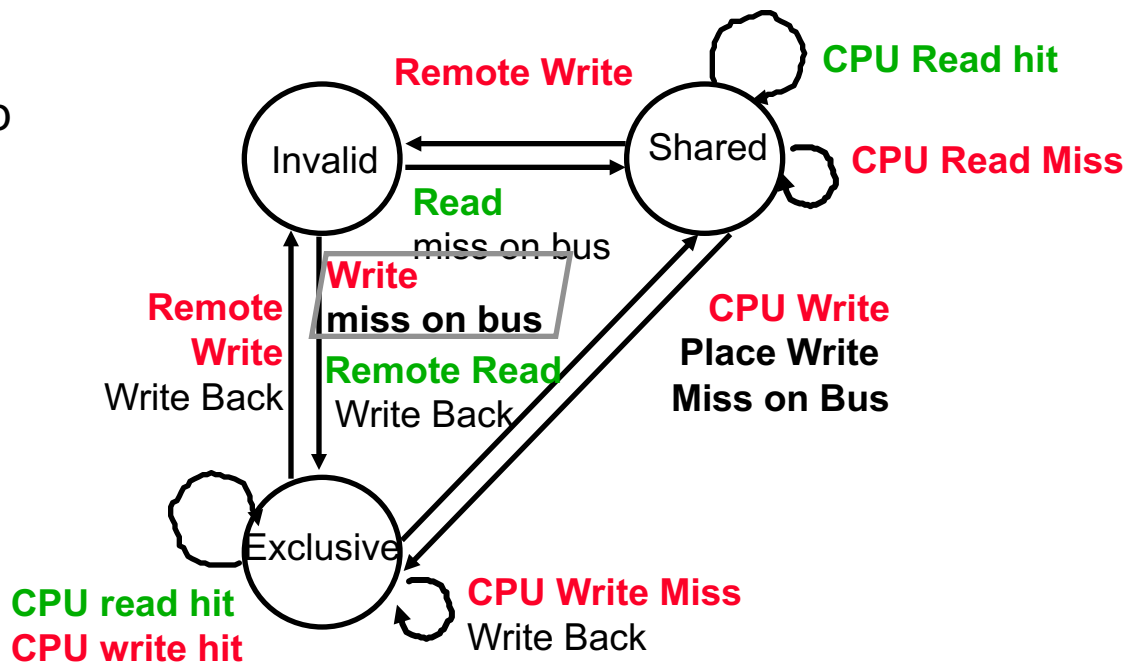


## Example: Step 1

[illegible]

Assumes initial cache state is invalid and A1 and A2 map to same cache block, but  $A1 \neq A2$ .

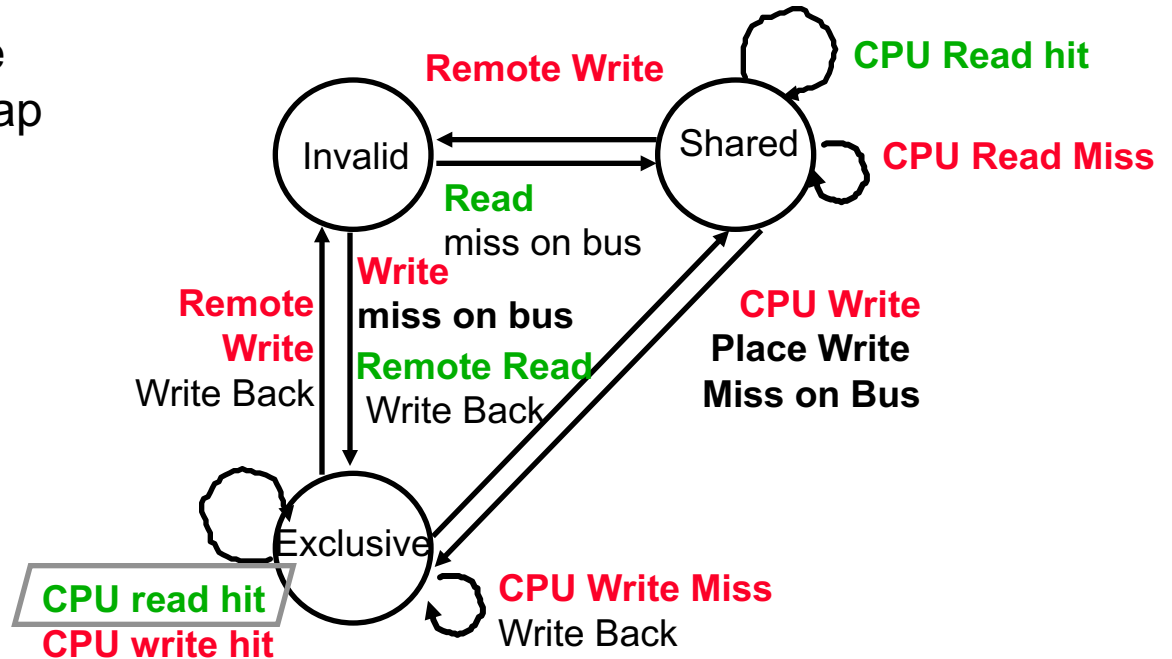
Active arrow =



## Example: Step 2

	<i>P1</i>			<i>P2</i>			<i>Bus</i>				<i>Memory</i>	
<i>step</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>Action</i>	<i>Proc.</i>	<i>Addr</i>	<i>Value</i>	<i>Addr</i>	<i>Value</i>
P1: Write 10 to A1	<u>Excl.</u>	<u>A1</u>	<u>10</u>				<u>WrMs</u>	P1	A1			
P1: Read A1	Excl.	A1	10									
P2: Read A1												
P2: Write 20 to A1												
P2: Write 40 to A2												

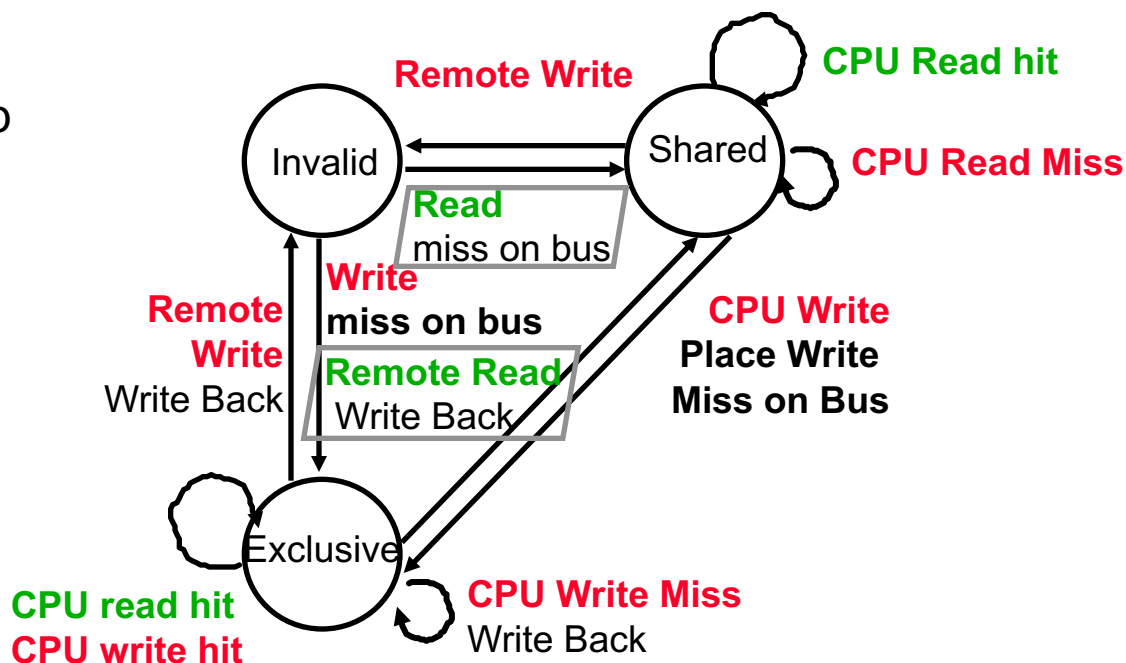
Assumes initial cache state is invalid and A1 and A2 map to same cache block, but A1 != A2



## Example: Step 3

	<i>P1</i>			<i>P2</i>			<i>Bus</i>			<i>Memory</i>		
<i>step</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>Action</i>	<i>Proc.</i>	<i>Addr</i>	<i>Value</i>	<i>Addr</i>	<i>Value</i>
P1: Write 10 to A1	<u>Excl.</u>	<u>A1</u>	<u>10</u>				<u>WrMs</u>	P1	A1			
P1: Read A1	Excl.	A1	10									
P2: Read A1				<u>Shar.</u>	<u>A1</u>		<u>RdMs</u>	P2	A1			
	<u>Shar.</u>	A1	10				<u>WrBk</u>	P1	A1	10	<u>A1</u>	<u>10</u>
				Shar.	A1	<u>10</u>	<u>RdDa</u>	P2	A1	10	A1	10
P2: Write 20 to A1												
P2: Write 40 to A2												

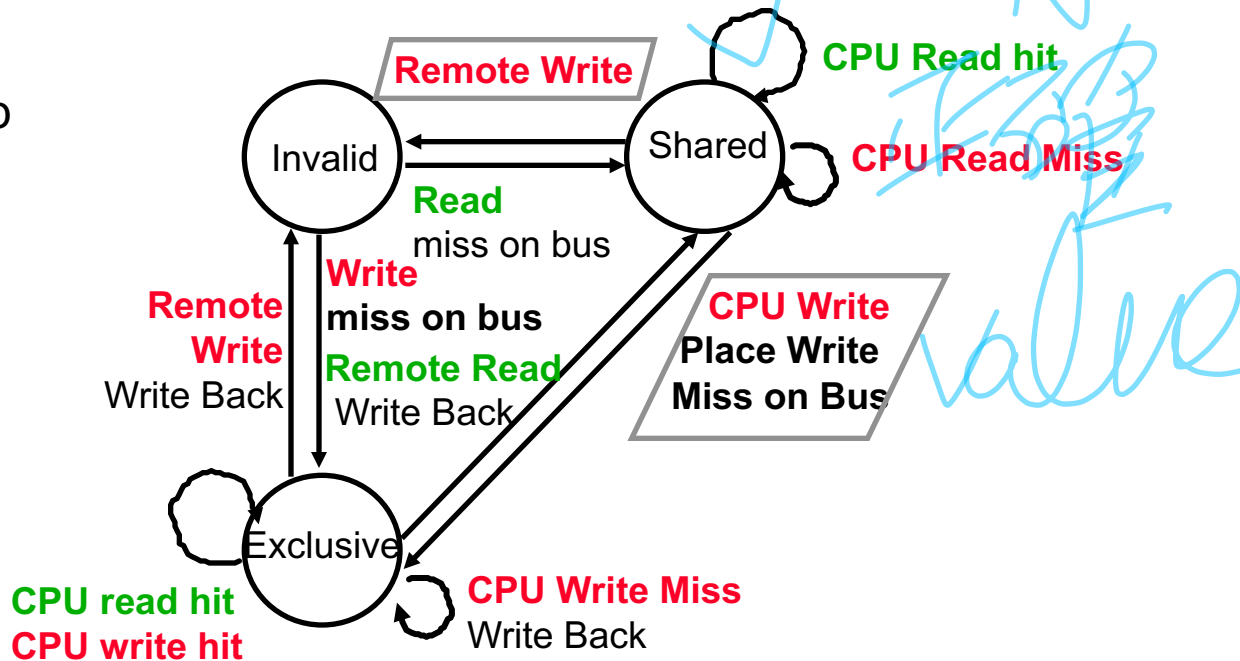
Assumes initial cache state is invalid and A1 and A2 map to same cache block, but A1 != A2.



# Example: Step 4

	<i>P1</i>			<i>P2</i>			<i>Bus</i>			<i>Memory</i>		
<i>step</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>State</i>	<i>Addr</i>	<i>Value</i>	<i>Action</i>	<i>Proc.</i>	<i>Addr</i>	<i>Value</i>	<i>Addr</i>	<i>Value</i>
P1: Write 10 to A1	<u>Excl.</u>	<u>A1</u>	<u>10</u>				<u>WrMs</u>	P1	A1			
P1: Read A1	Excl.	A1	10									
P2: Read A1				<u>Shar.</u>	<u>A1</u>		<u>RdMs</u>	P2	A1			
	<u>Shar.</u>	A1	10				<u>WrBk</u>	P1	A1	10	<u>A1</u>	<u>10</u>
				Shar.	A1	10	<u>RdDa</u>	P2	A1	10	A1	10
P2: Write 20 to A1	<u>Inv.</u>			<u>Excl.</u>	A1	<u>20</u>	<u>WrMs</u>	P2	A1		A1	10
P2: Write 40 to A2												..
												..

Assumes initial cache state is invalid and A1 and A2 map to same cache block, but A1 != A2

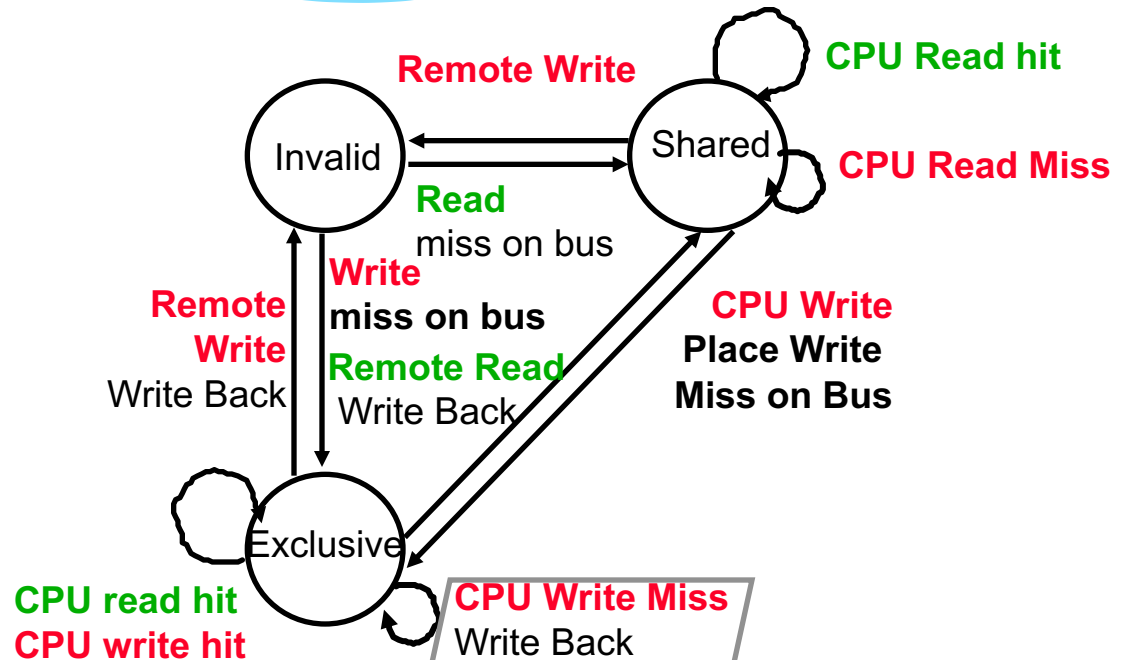





# Example: Step 5

	P1			P2			Bus					Memory	
step	State	Addr	Value	State	Addr	Value	Action	Proc.	Addr	Value		Addr	Value
P1: Write 10 to A1	<u>Excl.</u>	<u>A1</u>	<u>10</u>				<u>WrMs</u>	P1	<u>A1</u>				
P1: Read A1	Excl.	A1	10										
P2: Read A1				<u>Shar.</u>	<u>A1</u>		<u>RdMs</u>	P2	<u>A1</u>				
	<u>Shar.</u>	A1	10				<u>WrBk</u>	P1	<u>A1</u>	10		<u>A1</u>	<u>10</u>
				Shar.	A1	<u>10</u>	<u>RdDa</u>	P2	<u>A1</u>	10		A1	10
P2: Write 20 to A1	<u>Inv.</u>			<u>Excl.</u>	A1	<u>20</u>	<u>WrMs</u>	P2	<u>A1</u>			A1	10
P2: Write 40 to A2							<u>WrMs</u>	P2	<u>A2</u>			A1	10
				<u>Excl.</u>	<u>A2</u>	<u>40</u>	<u>WrBk</u>	P2	<u>A1</u>	20		<u>A1</u>	<u>20</u>


Assumes initial cache state is invalid and A1 and A2 map to same cache block, but A1 != A2



# What is multi-core?



Off-chip bus




On-chip bus


# Why Multi-core?

- Old Conventional Wisdom: Power is free, Transistors expensive
  - New Conventional Wisdom: “Power wall” Power expensive, Xtors free  
(Can put more on chip than can afford to turn on)
  - Old CW: Sufficiently increasing Instruction Level Parallelism via compilers, innovation (Out-of-order, speculation, VLIW, ...)
  - New CW: “ILP wall” law of diminishing returns on more HW for ILP
  - Old CW: Multiplies are slow, Memory access is fast
  - New CW: “Memory wall” Memory slow, multiplies fast  
(200 clock cycles to DRAM memory, 4 clocks for multiply)
  - Old CW: Uniprocessor performance 2X / 1.5 yrs
  - New CW: Power Wall + ILP Wall + Memory Wall = Brick Wall
    - Uniprocessor performance now 2X / 5(?) yrs
- ⇒ Sea change in chip design: multiple “cores”  
(2X processors per chip / ~ 2 years)
- » More simpler processors are more power efficient


# Parallelism for Energy Efficiency Present and Future



# From Multicore to Manycore

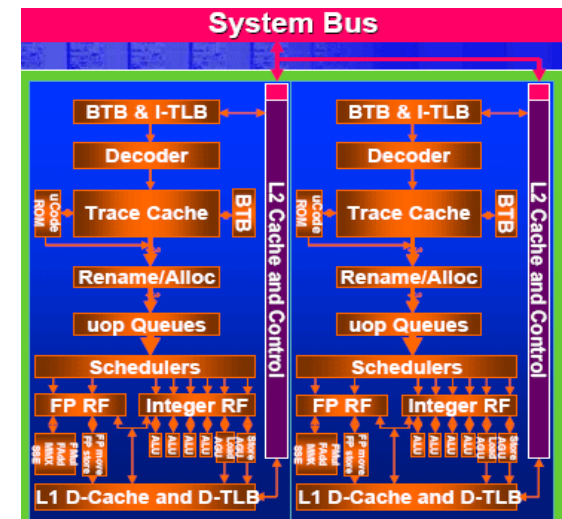



# Tera-Scale Computing



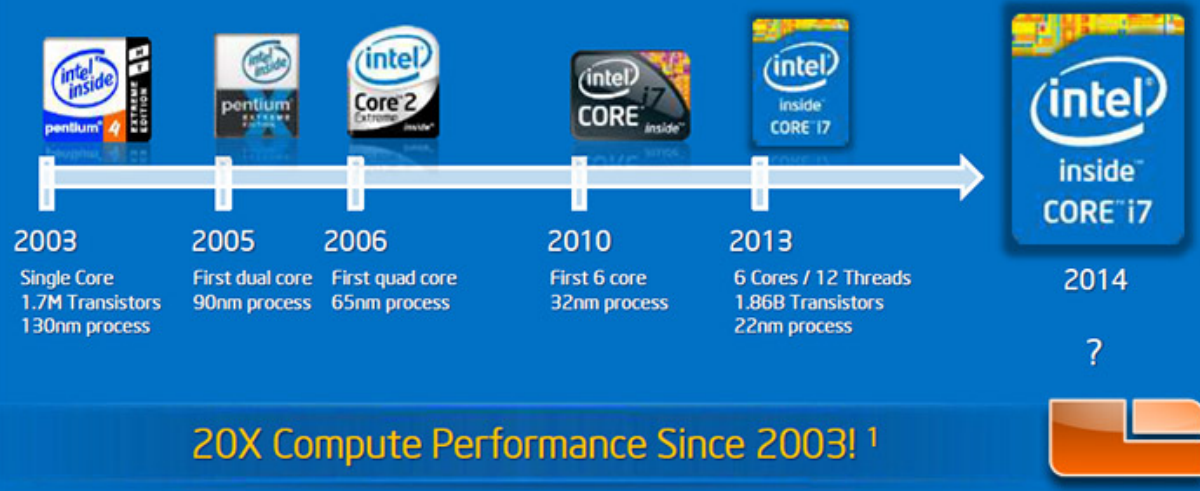
# Chip-Multiprocessor

- **Chip-Multiprocessor (2005 ~)**
  - A chip contains multiple cores
  - Example: Intel Xeon, AMD Opteron, Sony Playstation, Sun T1 Niagara




# Intel's Multi-Core Roadmap

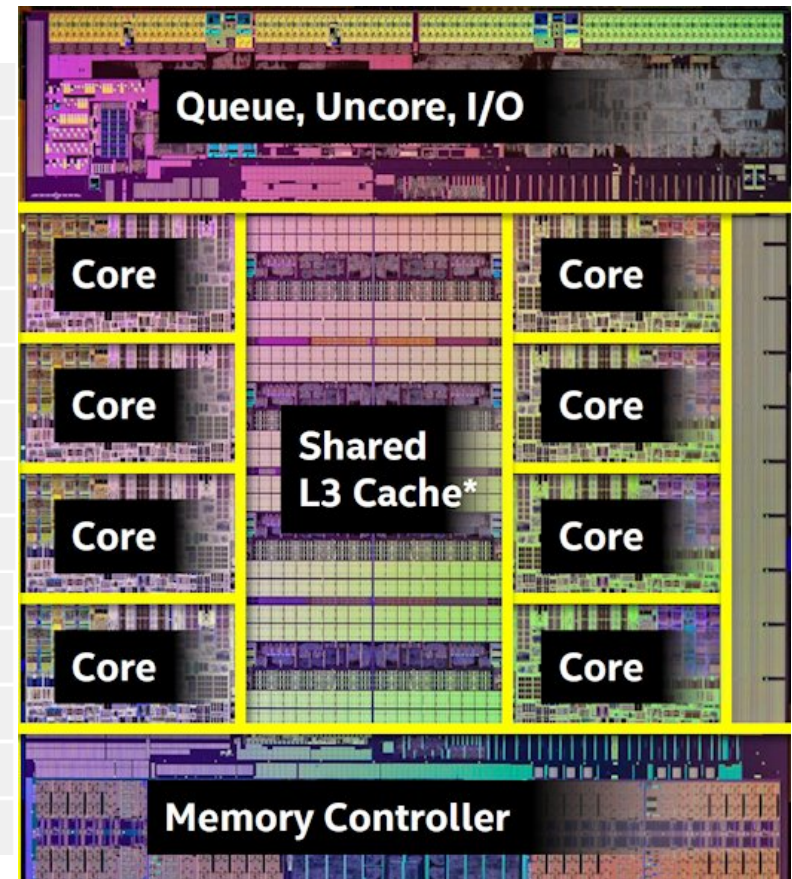
## Extreme Edition Innovation Continues





# Intel® Core™ i7-5960X Processor

# of Cores	8
# of Threads	16
Clock Speed	3 GHz
Max Turbo Frequency	3.5 GHz
Intel® Smart Cache	20 MB
Intel® QPI Speed	0 GT/s
# of QPI Links	0
Instruction Set	64-bit
Instruction Set Extensions	SSE4.2, AVX 2.0, AES
Max Memory Size (dependent on memory type)	64 GB
Memory Types	DDR4-1333/1600/2133
# of Memory Channels	4
Max Memory Bandwidth	68 GB/s
ECC Memory Supported ‡	 No



[http://ark.intel.com/products/82930/Intel-Core-i7-5960X-Processor-Extreme-Edition-20M-Cache-up-to-3\\_50-GHz#@specifications](http://ark.intel.com/products/82930/Intel-Core-i7-5960X-Processor-Extreme-Edition-20M-Cache-up-to-3_50-GHz#@specifications)

# Intel Core i9 ~ announced in 2017



## UNLOCKED INTEL® CORE™ X-SERIES PROCESSOR FAMILY

Processor number <sup>1</sup>	Base clock speed (GHz)	Intel® Turbo Boost Technology 2.0 frequency <sup>2</sup> (GHz)	Intel® Turbo Boost Max Technology 3.0 Frequency <sup>3</sup> (GHz)	Cores/threads	L3 cache	PCI express 3.0 lanes	Memory support	TDP	Socket (LGA)	RCP Pricing (1K USD)
i9-7980XE <b>NEW</b>	2.6	4.2	4.4	18/36	24.75 MB	44	Four channels DDR4-2666	165W	2066	\$1,999
i9-7960X <b>NEW</b>	2.8	4.2	4.4	16/32	22 MB	44	Four channels DDR4-2666	165W	2066	\$1,699
i9-7940X <b>NEW</b>	3.1	4.3	4.4	14/28	19.25 MB	44	Four channels DDR4-2666	165W	2066	\$1,399
i9-7920X <b>NEW</b>	2.9	4.3	4.4	12/24	16.5 MB	44	Four channels DDR4-2666	140W	2066	\$1,199
i9-7900X <b>NEW</b>	3.3	4.3	4.5	10/20	13.75 MB	44	Four channels DDR4-2666	140W	2066	\$999
i7-7820X <b>NEW</b>	3.6	4.3	4.5	8/16	11 MB	28	Four channels DDR4-2666	140W	2066	\$599
i7-7800X <b>NEW</b>	3.5	4.0	NA	6/12	8.25 MB	28	Four channels DDR4-2400	140W	2066	\$389
i7-7740X <b>NEW</b>	4.3	4.5	NA	4/8	8 MB	16	Two channels DDR4-2666	112W	2066	\$339
i5-7640X <b>NEW</b>	4.0	4.2	NA	4/4	6 MB	16	Two channels DDR4-2666	112W	2066	\$242

1. Intel processor numbers are not a measure of performance. Processor numbers differentiate features within each processor family, not across different processor families.

See [intel.com/products/processor\\_number](https://www.intel.com/products/processor_number) for details.

2. Refers to the maximum dual-core frequency that can be achieved with Intel® Turbo Boost Technology 2.0.

3. Refers to the maximum dual-core frequency that can be achieved with Intel® Turbo Boost Max Technology 3.0

# IBM: Power Processor

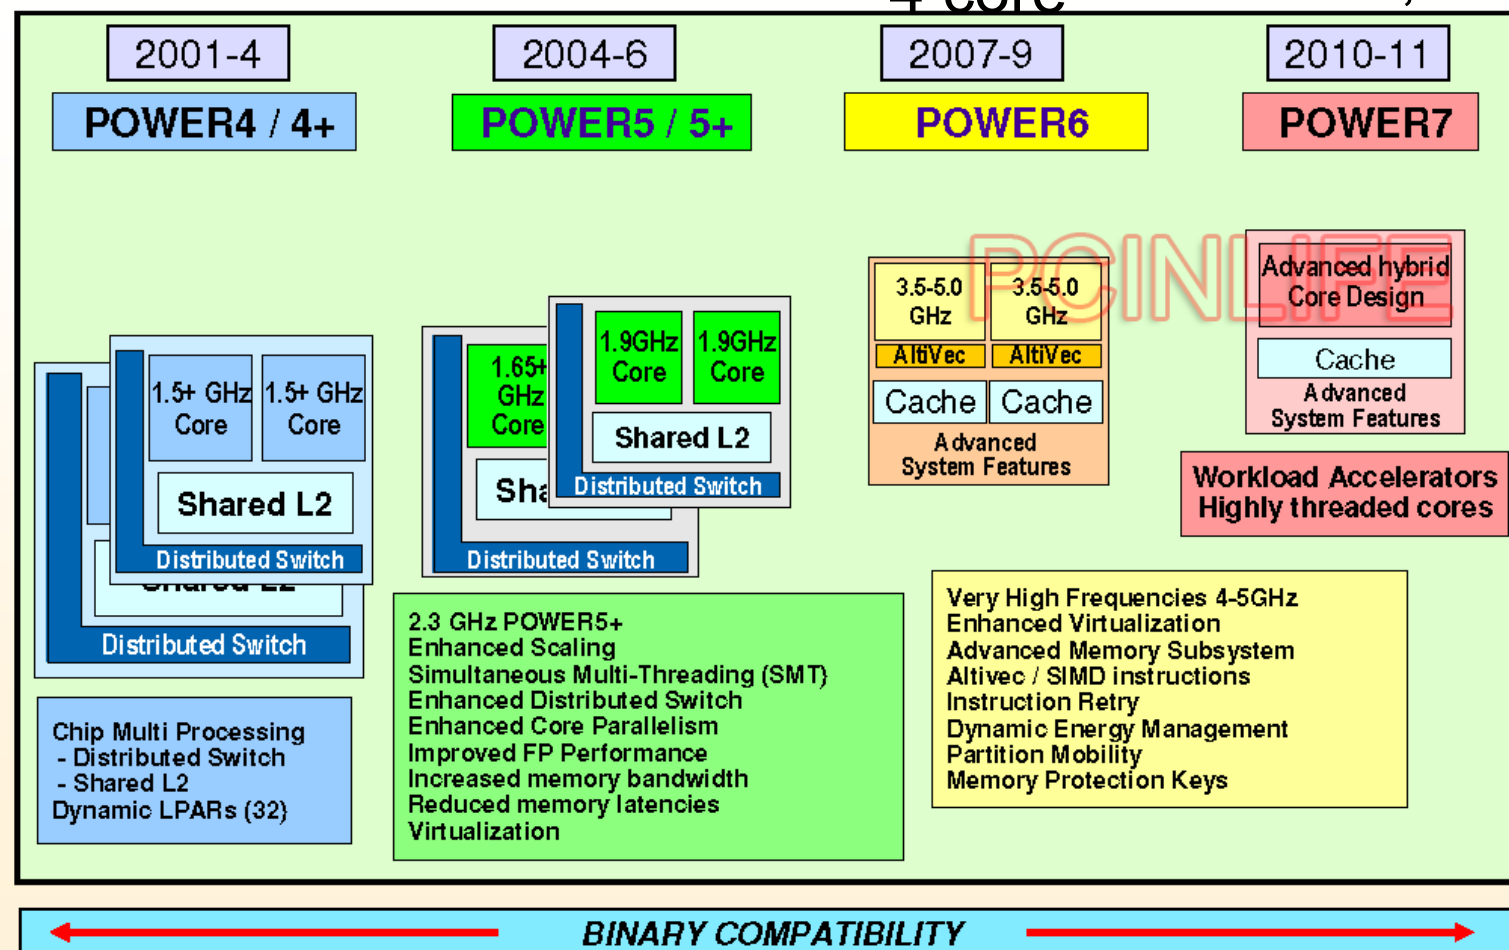
IBM System p



## POWER Processor Roadmap

4 core

8 core, 32 threads





# POWER8 Processor

## Technology

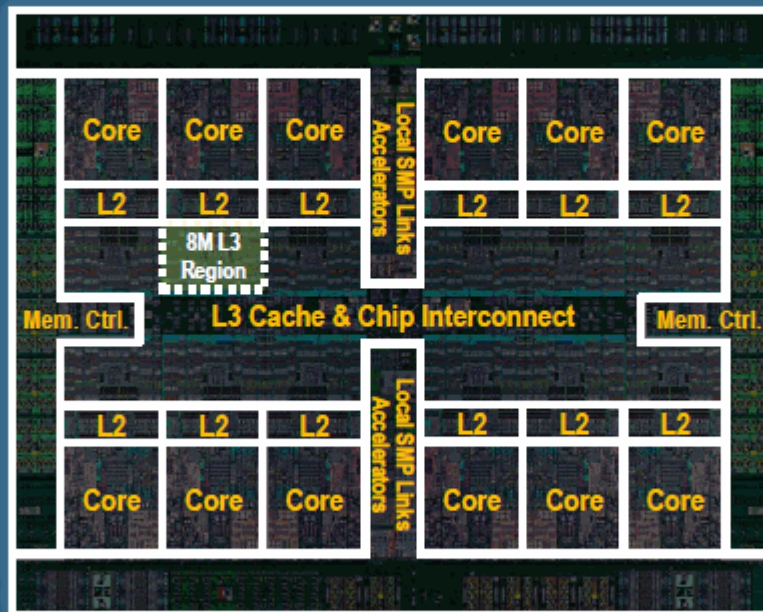
- 22nm SOI, eDRAM, 15 ML 650mm<sup>2</sup>

## Cores

- 12 cores (SMT8)
- 8 dispatch, 10 issue, 16 exec pipe
- 2X internal data flows/queues
- Enhanced prefetching
- 64K data cache, 32K instruction cache

## Accelerators

- Crypto & memory expansion
- Transactional Memory
- VMM assist
- Data Move / VM Mobility



## Energy Management

- On-chip Power Management Micro-controller
- Integrated Per-core VRM
- Critical Path Monitors

## Caches

- 512 KB SRAM L2 / core
- 96 MB eDRAM shared L3
- Up to 128 MB eDRAM L4 (off-chip)


## Memory

- Up to 230 GB/s sustained bandwidth

## Bus Interfaces

- Durable open memory attach interface
- Integrated PCIe Gen3
- SMP Interconnect
- CAPI (Coherent Accelerator Processor Interface)

# IBM Power 9 (announced in Dec. 2017)



## POWER9 Processor – Common Features



### New Core Microarchitecture

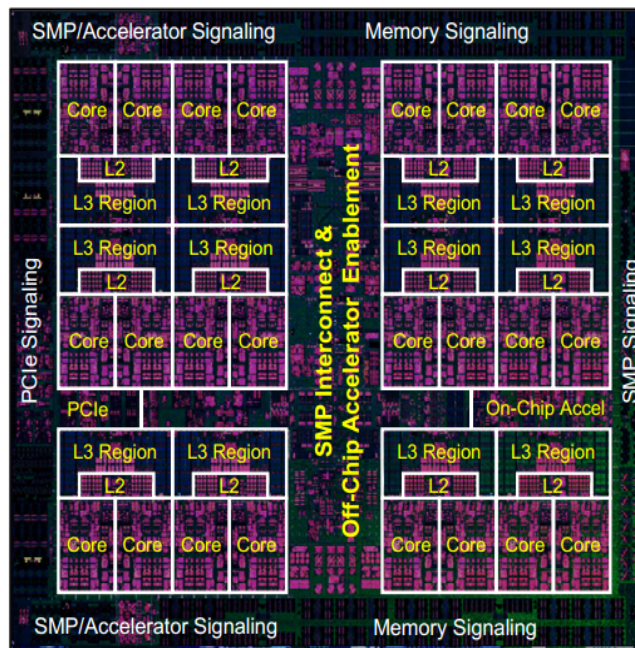
- Stronger thread performance
- Efficient agile pipeline
- POWER ISA v3.0

### Enhanced Cache Hierarchy

- 120MB NUCA L3 architecture
- 12 x 20-way associative regions
- Advanced replacement policies
- Fed by 7 TB/s on-chip bandwidth

### Cloud + Virtualization Innovation

- Quality of service assists
- New interrupt architecture
- Workload optimized frequency
- Hardware enforced trusted execution



### 14nm finFET Semiconductor Process

- Improved device performance and reduced energy
- 17 layer metal stack and eDRAM

### Leadership

#### Hardware Acceleration Platform

- Enhanced on-chip acceleration
- Nvidia NVLink 2.0: High bandwidth and advanced new features (25G)
- CAPI 2.0: Coherent accelerator and storage attach (PCIe G4)
- New CAPI: Improved latency and bandwidth, open interface (25G)

#### State of the Art I/O Subsystem



- PCIe Gen4 – 48 lanes

#### High Bandwidth Signaling Technology

- 16 Gb/s interface
  - Local SMP
- 25 Gb/s Common Link interface
- Accelerator remote SMP

# ARM Big.Little Technology

- ARM big.LITTLE processing is designed to deliver the vision of the right processor for the right job. In current big.LITTLE system implementations a **'big' ARM Cortex™-A15 processor** is paired with a **'LITTLE' Cortex™-A7 processor** to create a system that can accomplish both high intensity and low intensity tasks in the most energy efficient manner. For example, the performance capabilities of the Cortex-A15 processor can be utilized for heavy workloads, while the Cortex-A7 can take over to process most efficiently majority of smartphone workloads. These include operating system activities, user interface and other always on, always connected tasks.




# Eight-core in mobile device



## Samsung's Exynos 5 Octa will become a "true" eight-core chip in Q4

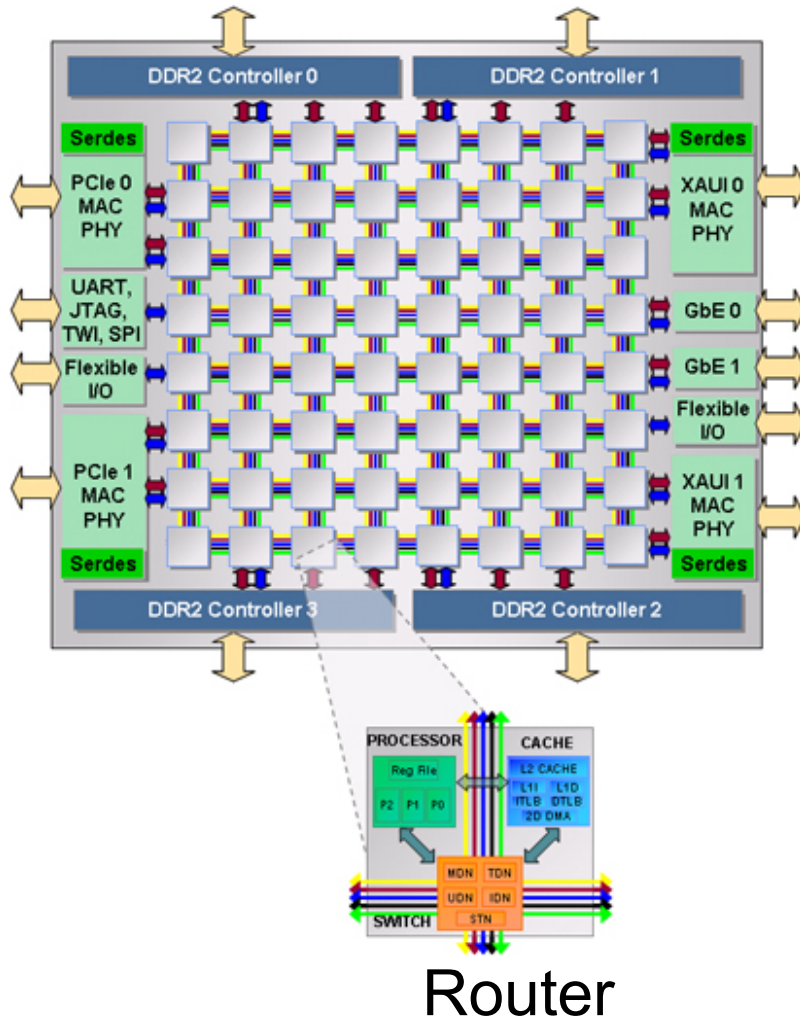
Currently, only four of its eight CPU cores can be used at one time.

by Andrew Cunningham - Sept 10 2013, 2:20pm +0800





# Tilera 64-core processor



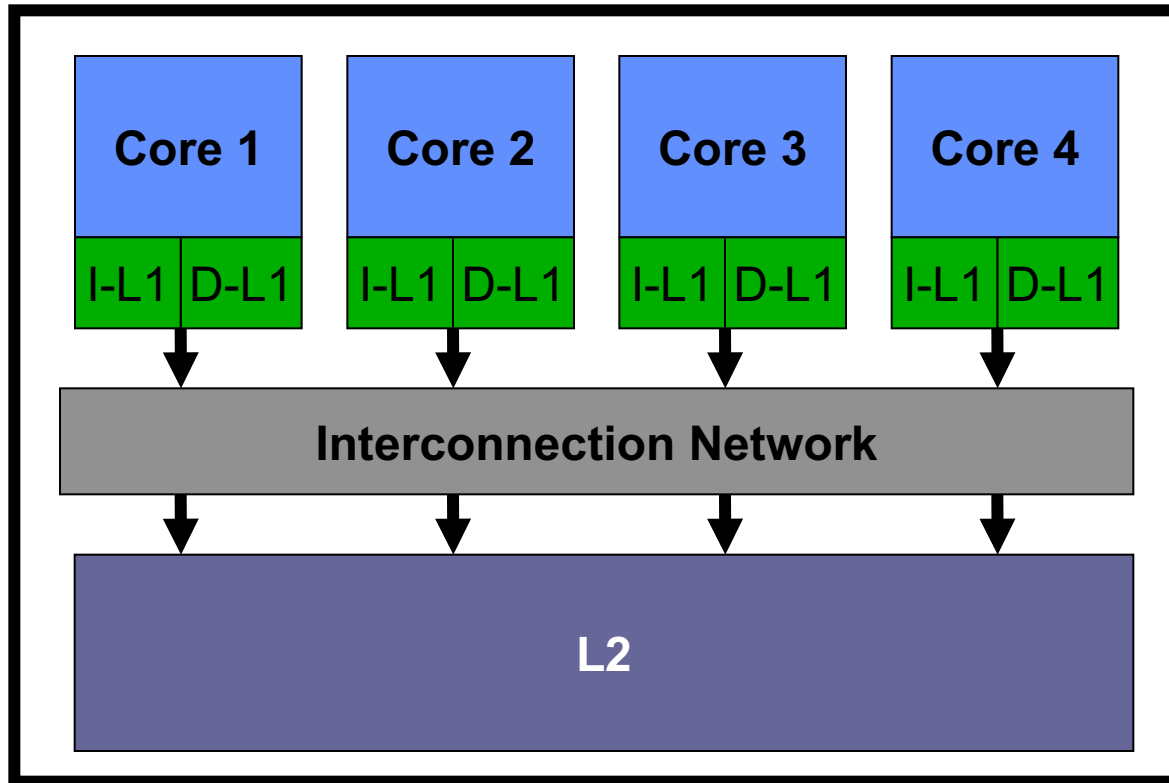
"Tilera® Corporation provides multicore processors that deliver the world's highest performance computing for a broad range of applications such as networking, wireless infrastructure, digital multimedia, and cloud computing. Tilera's processors are based on an intelligent mesh (iMesh™) architecture and provide far greater scalability than any other processor on the market."

***"The Processor is  
the new Transistor"  
[Rowen]***




# Basic CMP Architecture

- Shared last level cache



# Scalable CMP Architecture

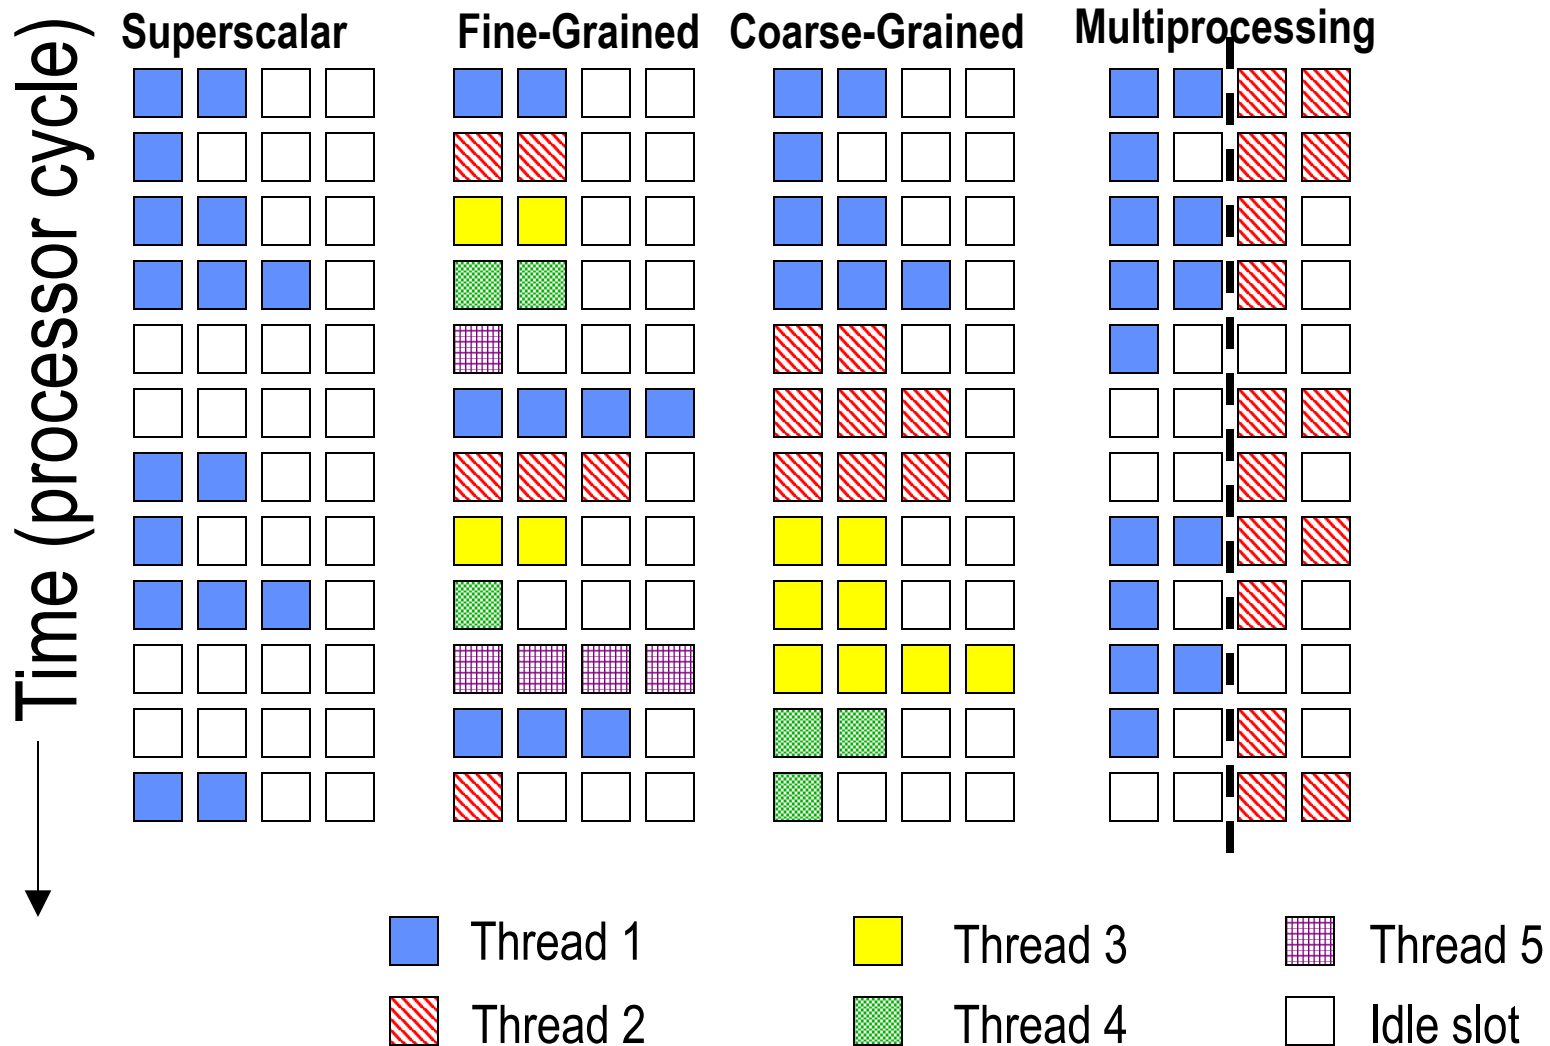
- **Tiled CMP**
  - Each tile includes processor, L1, L2, and router
  - Physically distributed last level cache



# Multithreading

- **Performing multiple threads of execution in parallel**
  - Replicate registers, PC, etc.
  - Fast switching between threads
- **Fine-grain multithreading**
  - Switch threads after each cycle
  - Interleave instruction execution
  - If one thread stalls, others are executed
- **Coarse-grain multithreading**
  - Only switch on long stall (e.g., L2-cache miss)
  - Simplifies hardware, but doesn't hide short stalls (eg, data hazards)

# Multithreaded Categories



# Simultaneous Multithreading

- **In multiple-issue dynamically scheduled processor**
  - Schedule instructions from multiple threads
  - Instructions from independent threads execute when function units are available
  - Within threads, dependencies handled by scheduling and register renaming
- **Example: Intel Pentium-4 HT**
  - Two threads: duplicated registers, shared function units and caches

# Multithreaded Categories

