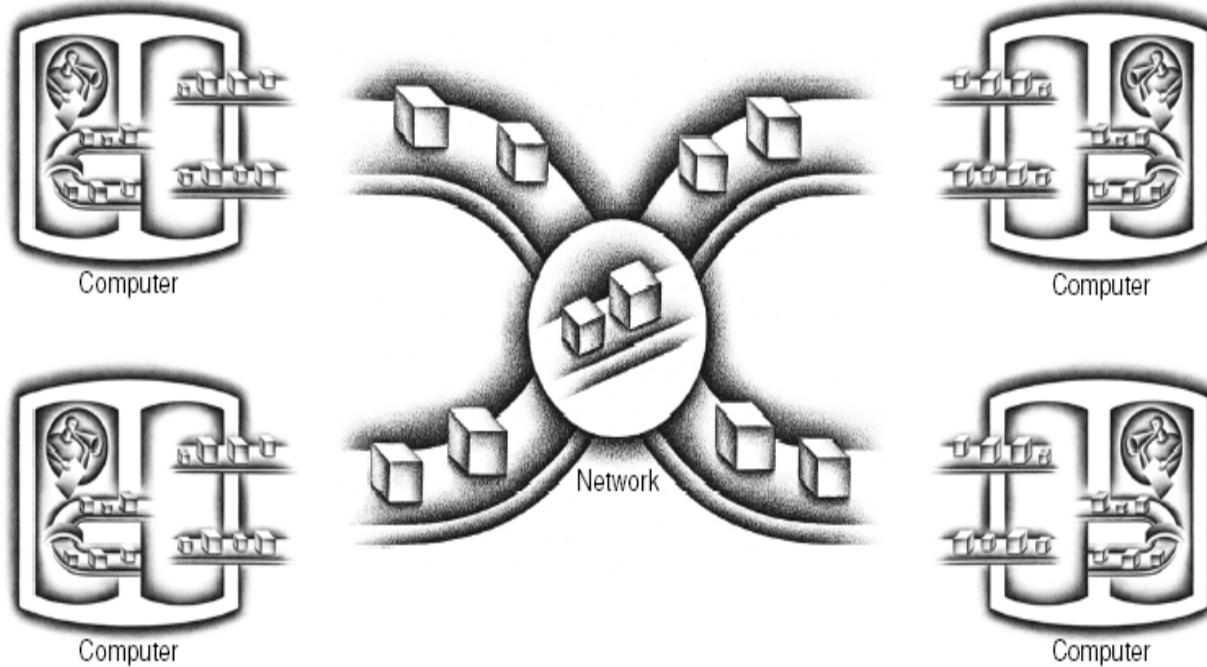


Lecture 14 Multicores, Multiprocessor and Cluster (Part II)

..... the three components of a computer



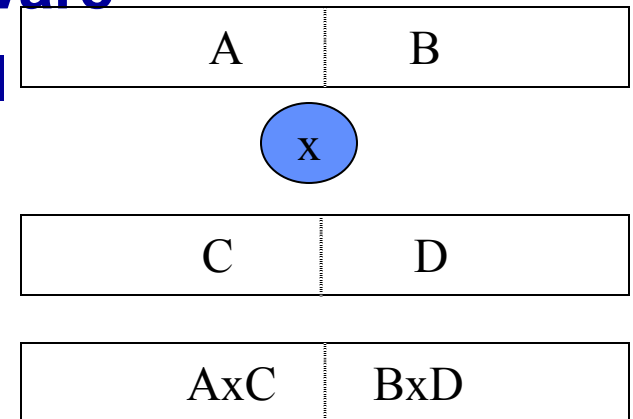
Computing Device Classification: Instruction and Data Streams

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345

- **SPMD: Single Program Multiple Data**
 - A parallel program on a MIMD computer
 - Conditional code for different processors

SIMD

- **Operate elementwise on vectors of data**
 - E.g., MMX and SSE instructions in x86
 - » Multiple data elements in 128-bit wide registers
- **All processors execute the same instruction at the same time**
 - Each with different data address, etc.
- **Simplifies synchronization**
- **Reduced instruction control hardware**
- **Works best for highly data-parallel applications**



Vector Processors

- Highly pipelined function units
- **Stream data from/to vector registers to units**
 - Data collected from memory into registers
 - Results stored from registers to memory
- **Example: Vector extension to MIPS**
 - 32×64 -element registers (64-bit elements)
 - **Vector instructions**
 - » `lv, sv`: load/store vector
 - » `addv.d`: add vectors of double
 - » `addvs.d`: add scalar to each element of vector of double
- **Significantly reduces instruction-fetch bandwidth**

Example: DAXPY ($Y = a \times X + Y$)

double

- Conventional MIPS code

```
load
loop: l.d    $f0,a($sp)      ;load scalar a
      addiu  r4,$s0,#512     ;upper bound of what to
      l.d    $f2,0($s0)     ;load x(i)
      mul.d  $f2,$f2,$f0     ;a x x(i)
      l.d    $f4,0($s1)     ;load y(i)
      add.d  $f4,$f4,$f2     ;a x x(i) + y(i)
      s.d    $f4,0($s1)     ;store into y(i)
      addiu  $s0,$s0,#8      ;increment index to x
      addiu  $s1,$s1,#8      ;increment index to y
      subu   $t0,r4,$s0      ;compute bound
      bne    $t0,$zero,loop ;check if done
```

- Vector MIPS code

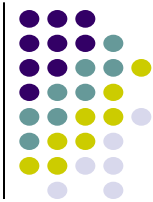
```
l.d    $f0,a($sp) ;load scalar a
lv      $v1,0($s0) ;load vector x
mulvs.d $v2,$v1,$f0 ;vector-scalar multiply
lv      $v3,0($s1) ;load vector y
addv.d  $v4,$v2,$v3 ;add y to product
sv      $v4,0($s1) ;store the result
```

aX

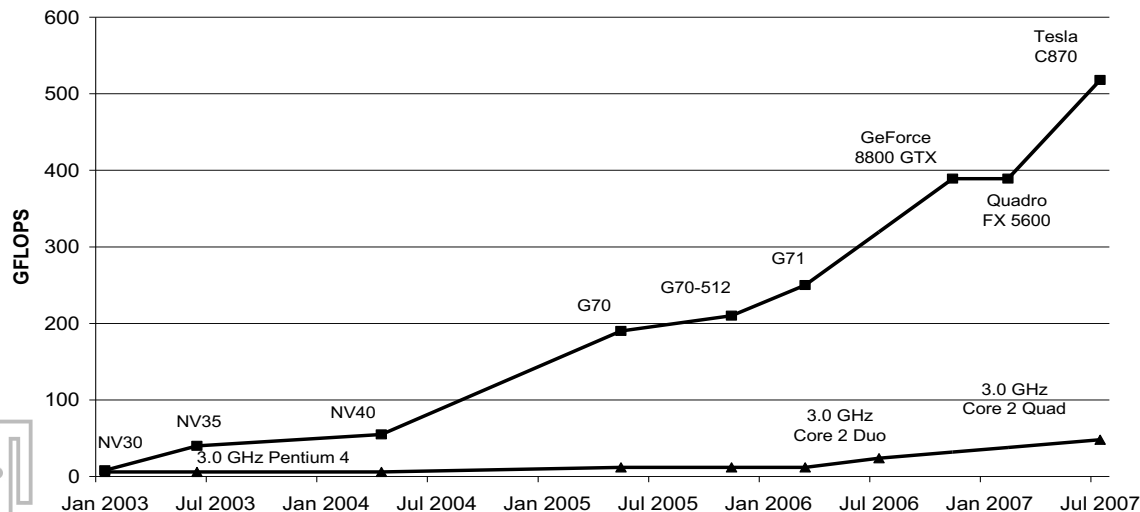
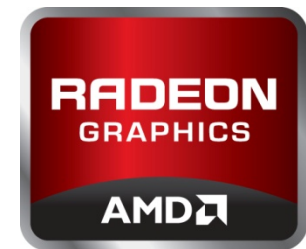
Vector vs. Scalar

- **Vector architectures and compilers**
 - Simplify data-parallel programming
 - Explicit statement of absence of loop-carried dependences
 - » Reduced checking in hardware
 - Regular access patterns benefit from interleaved and burst memory
 - Avoid control hazards by avoiding loops
- **More general than ad-hoc media extensions (such as MMX, SSE)**
 - Better match with compiler technology

What is GPU?



- Graphics Processing Units
- GPU is a device to compute massive vertices, pixels, and general purpose data
- Feature
 - High availability
 - High computing performance
 - Low price of computing capability



GPU's History and Evolution



- **Early History**

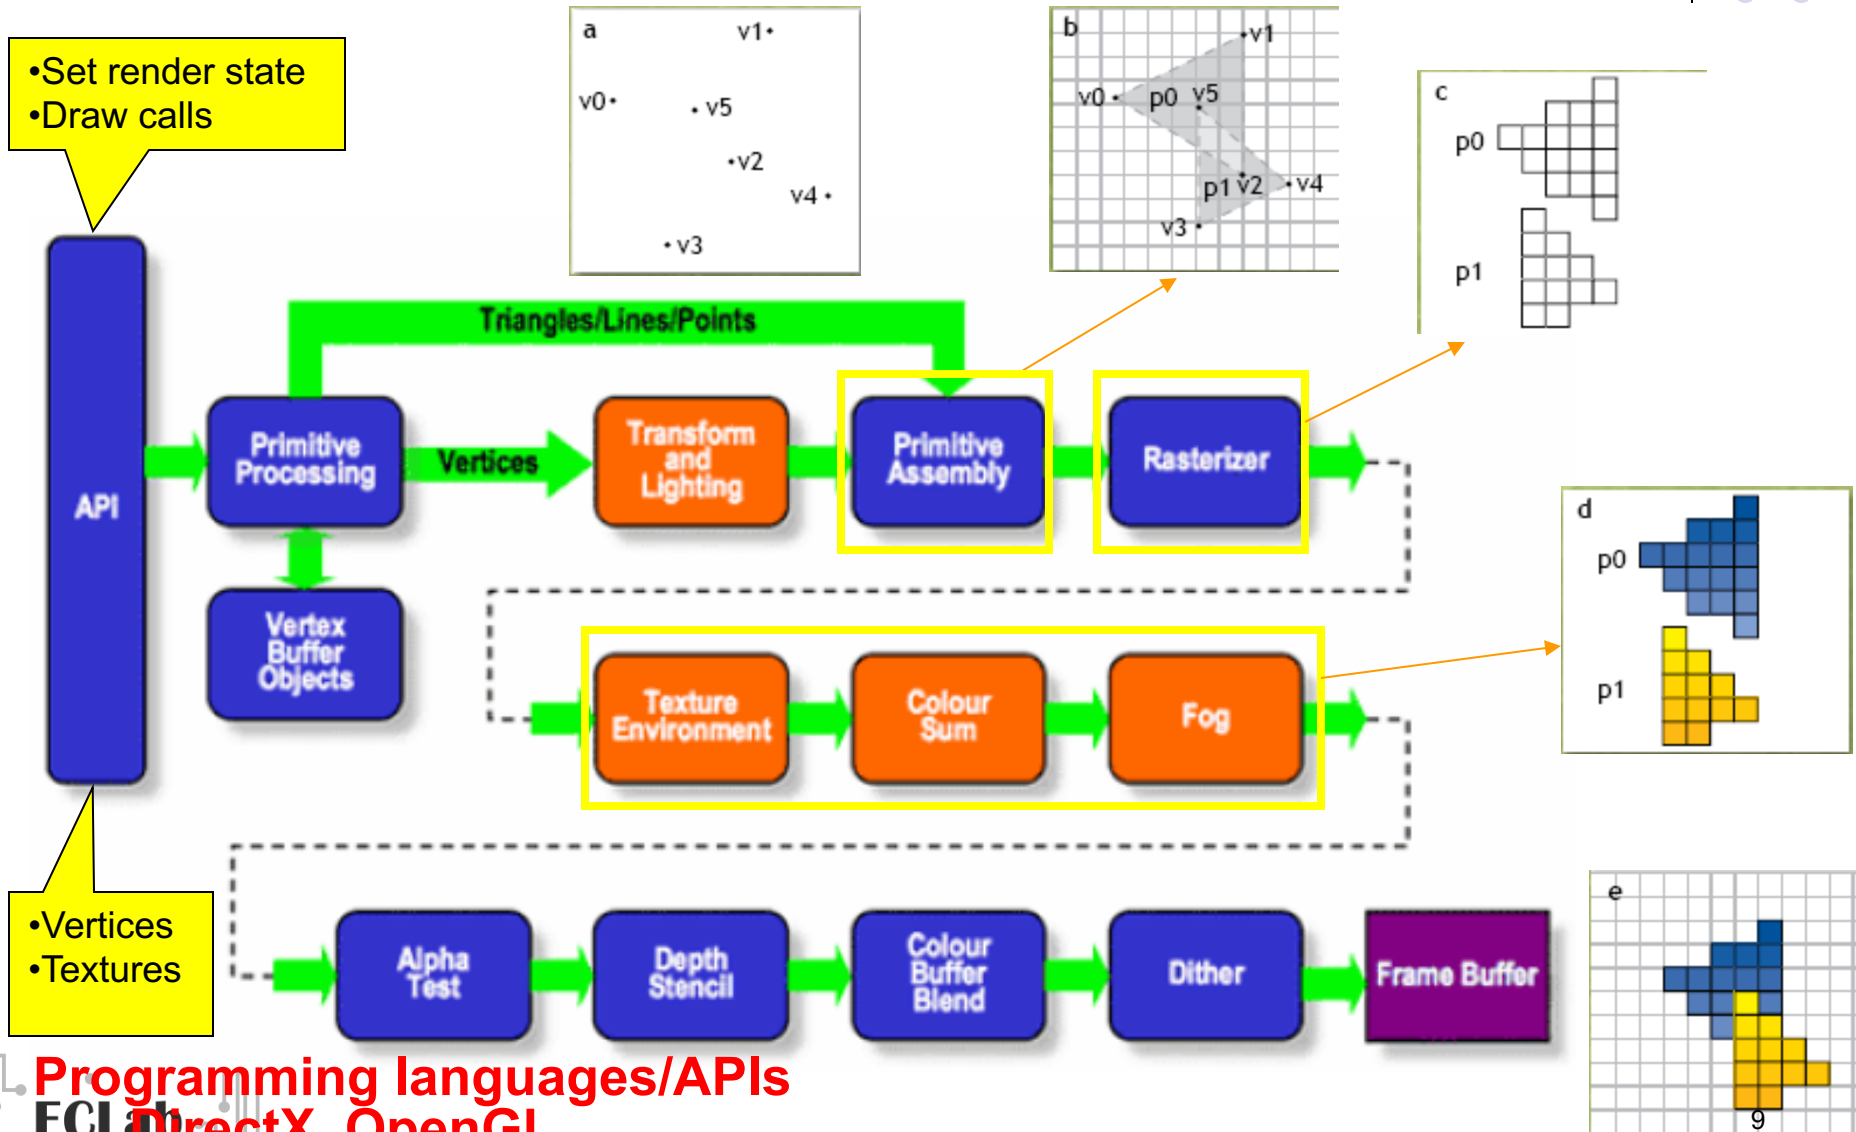
- In early 90's, graphics are only performed by a video graphics array (VGA) controller
- In 1997, VGA controllers start to incorporate 3D acceleration functions
- In 2000, the term GPU is coined to denote that the graphics devices had become a processor

- **GPU Evolution**

- **Fixed-Function → Programmable**
 - 1999, NVIDIA *Geforce 256*, Fixed-function vertex transform and pixel pipeline
 - 2001, NVIDIA *Geforce 3*, 1st programmable vertex processor
 - 2002, ATI *Radeon 9700*, 1st programmable pixel(fragment) processor
- **Non-unified Processor → Unified Processor**
 - 2005, Microsoft *XBOX 360*, 1st unified shader architecture
- **Tesla GPU series released in 2007**
- **Fermi Architecture released in 2009**
- **Kepler Architecture released in 2012**
- **Maxwell Architecture released in 2014**
- **Pascal Architecture released in 2016 (16 nm FinFET process)**
- **Volta Architecture released in 2017 – Tensor Core for AI (12 nm process)**

energy
2

Fixed-function 3D Graphics Pipeline

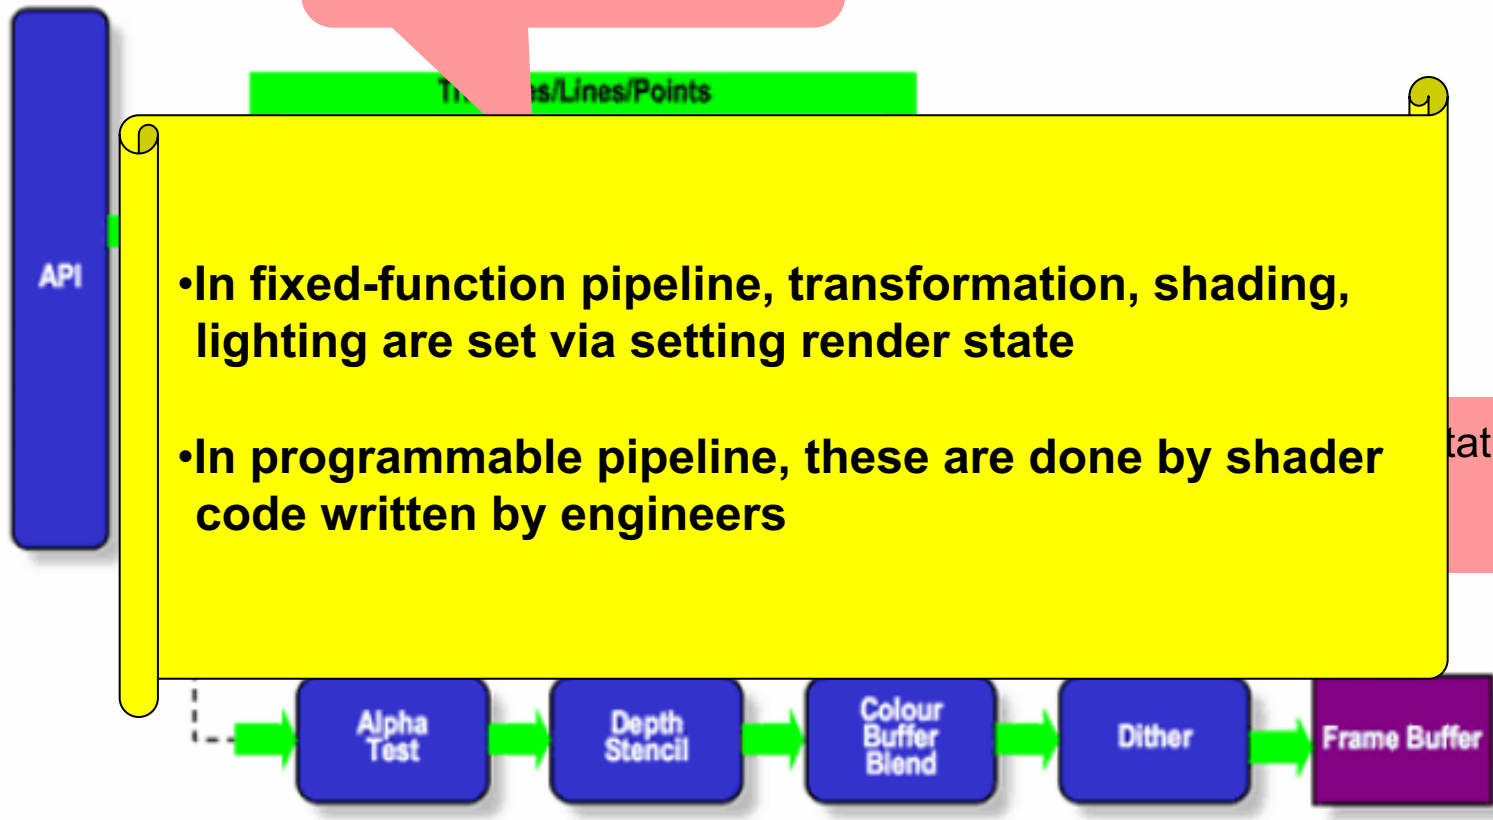


Programmable 3D Graphics Pipeline



High Level Shader Language
(HLSL)

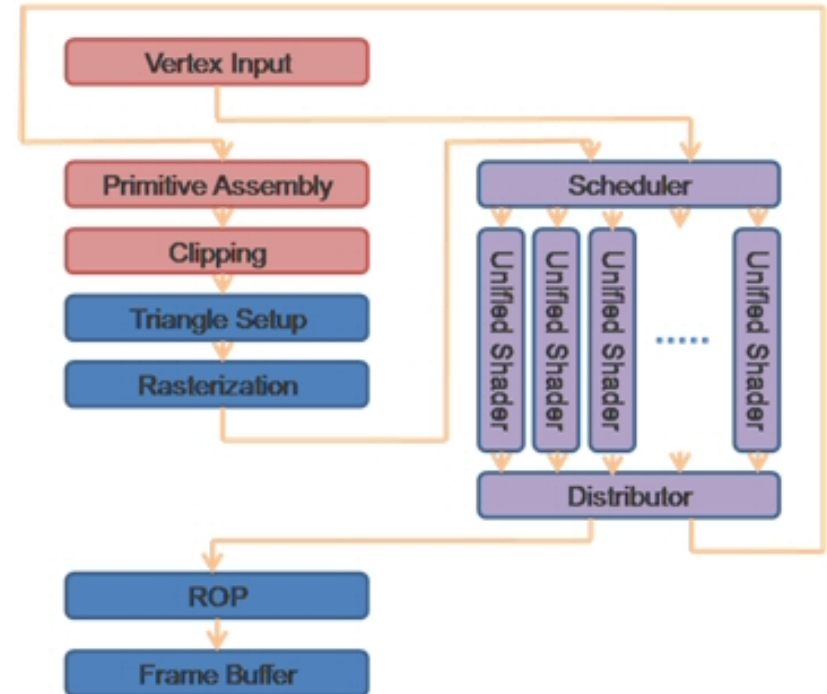
- Vertex computation



Unified Shader Architecture



- Use the same shader processors for all types of computation
 - Vertex threads
 - Pixel threads
 - Computation threads
- Advantage
 - Better resource utilization
 - Lower hardware complexity



Modern GPUs: A Computing Device



- **GPUs** have orders of magnitudes more computing power than CPU
- General-purpose tasks with **high-degree** of data level parallelism running on GPU outperform those on CPU
=> **General-Purpose** computing on GPU (**GPGPU**)
- GPGPU programming models
 - NVIDIA's CUDA
 - AMD's StreamSDK
 - OpenCL

GPGPU Performance	
Medical Imaging	300X
Molecular Dynamics	150X
SPICE	130X
Fourier Transform	130X
Fluid Dynamics	100X

Fundamental Architectural Differences between CPU & GPU

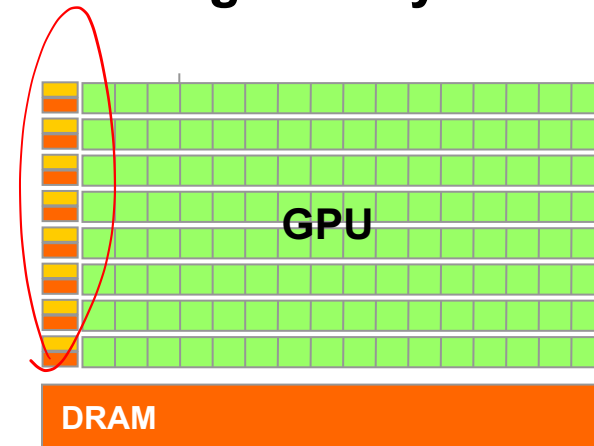
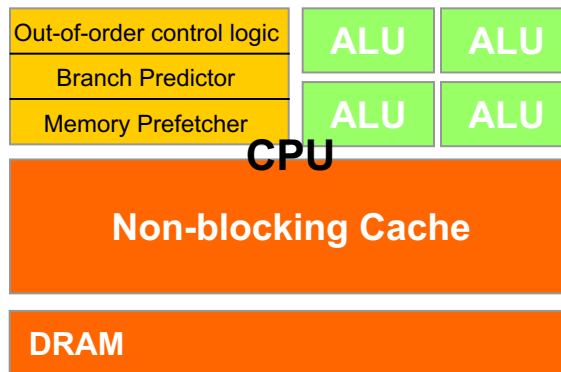


- **Multi-core CPU**

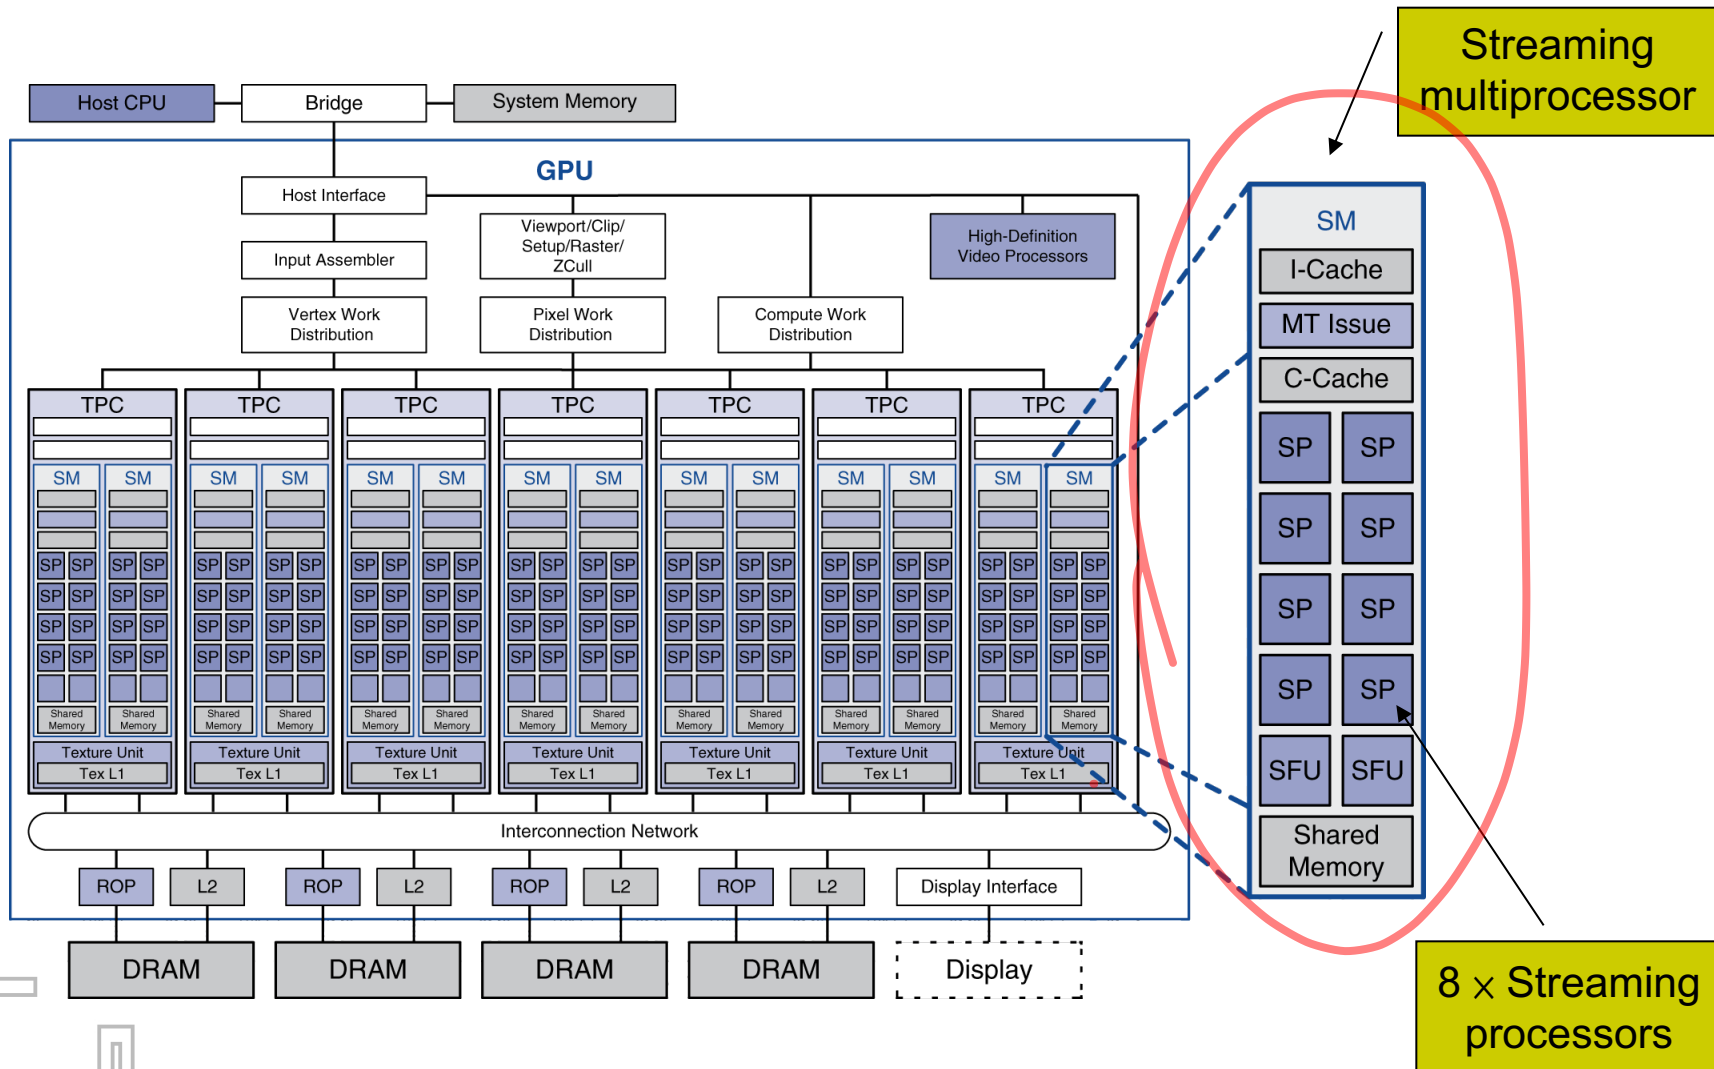
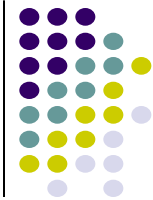
- Coarse-grain, heavyweight threads
- Memory latency is resolved through large on-chip caches & out-of-order execution

- **Modern GPU**

- Fine-grain, lightweight threads
- Exploit thread-level parallelism for hiding latency



Example: NVIDIA Tesla



SIMT Execution Model of GPUs

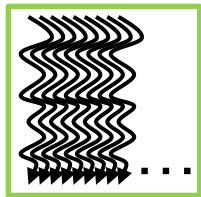


- **SIMT (Single Instruction Multiple Threads)**

- **Warp**

- A group of threads (pixel, vertex, compute...)
- Basic scheduling/execution unit
- Common PC value

Thread block



1 ~ x thread ID

1 ~ 32 thread ID



Warp 1

33 ~ 64 thread ID



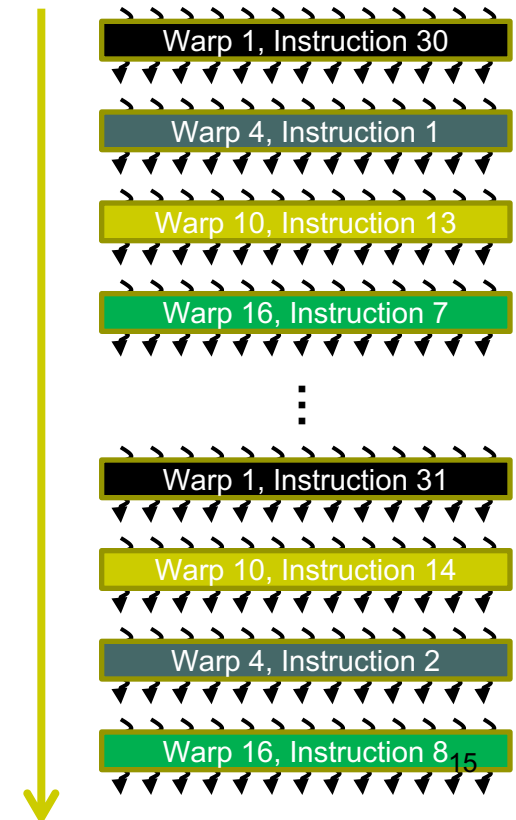
Warp 2

...



Warp n

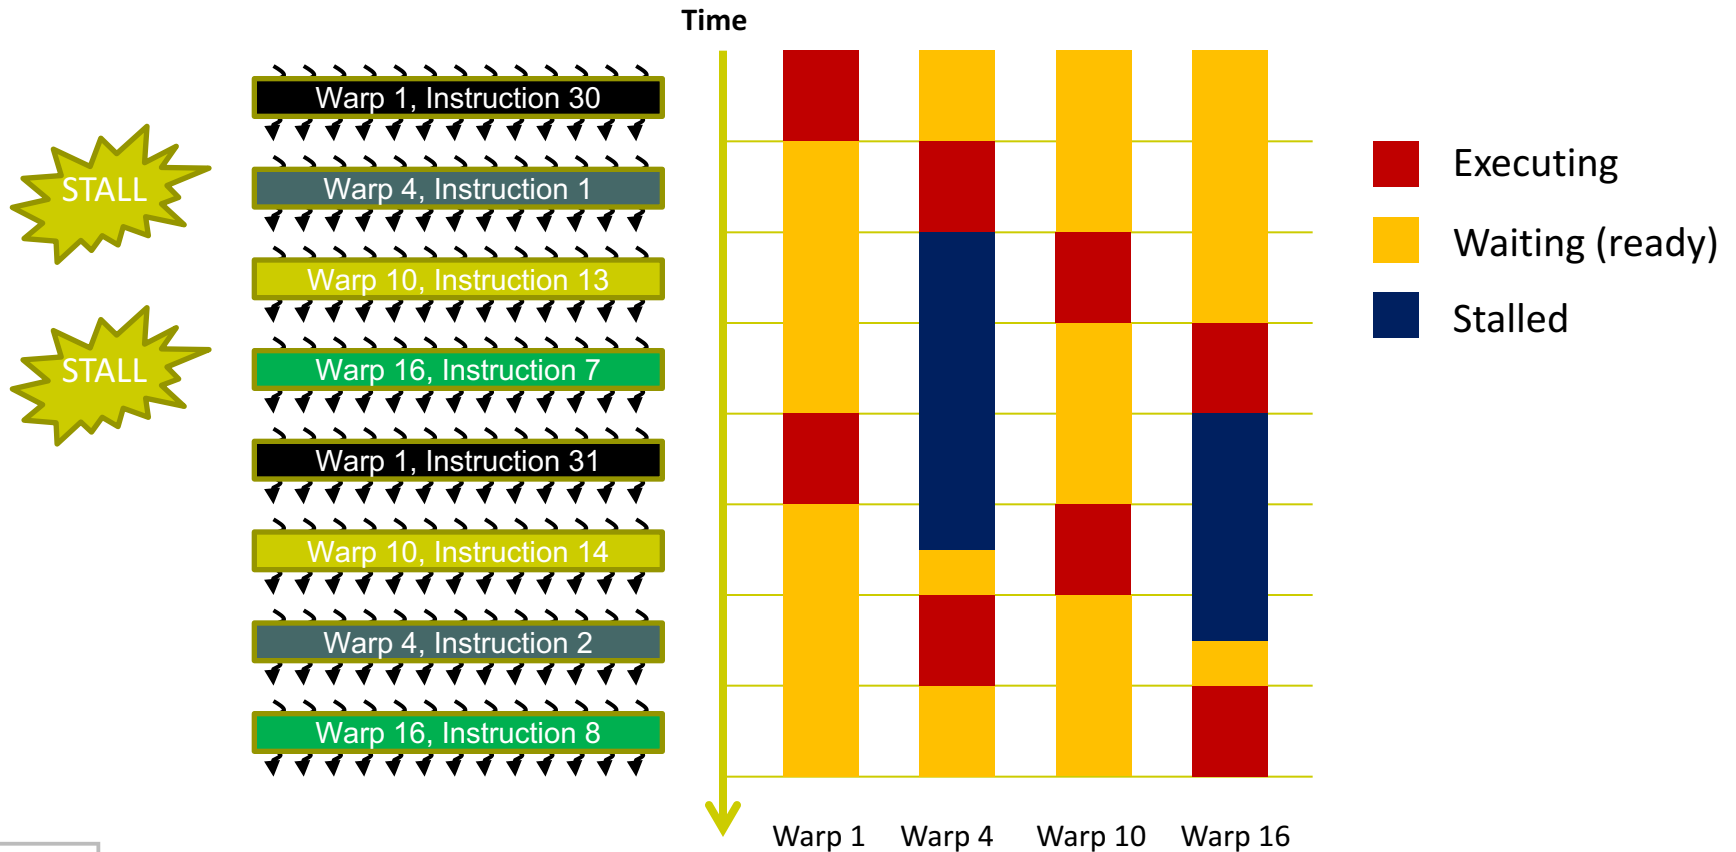
Time



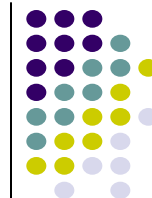
Latency Hiding



- Interleaved warp execution



Volta



$D =$

$$\begin{pmatrix} A_{0,0} & A_{0,1} & A_{0,2} & A_{0,3} \\ A_{1,0} & A_{1,1} & A_{1,2} & A_{1,3} \\ A_{2,0} & A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,0} & A_{3,1} & A_{3,2} & A_{3,3} \end{pmatrix}$$

FP16 or FP32

FP16

$$\begin{pmatrix} B_{0,0} & B_{0,1} & B_{0,2} & B_{0,3} \\ B_{1,0} & B_{1,1} & B_{1,2} & B_{1,3} \\ B_{2,0} & B_{2,1} & B_{2,2} & B_{2,3} \\ B_{3,0} & B_{3,1} & B_{3,2} & B_{3,3} \end{pmatrix}$$

FP16

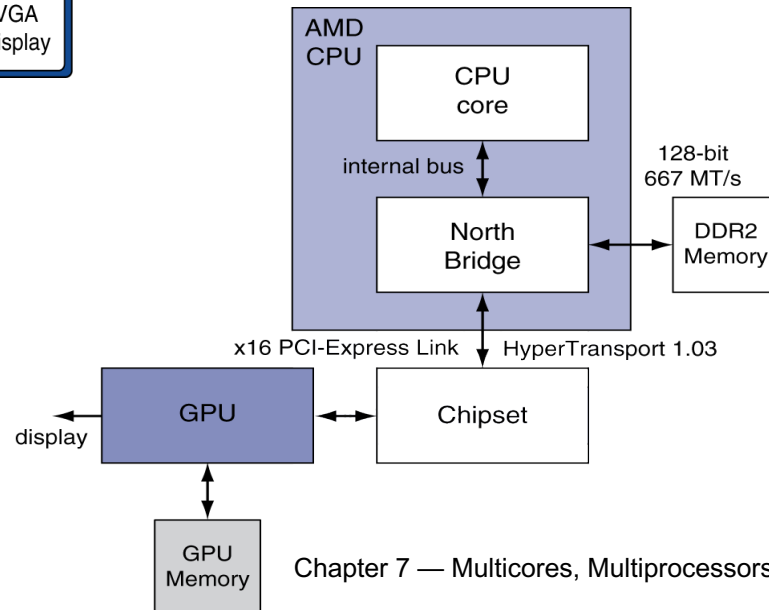
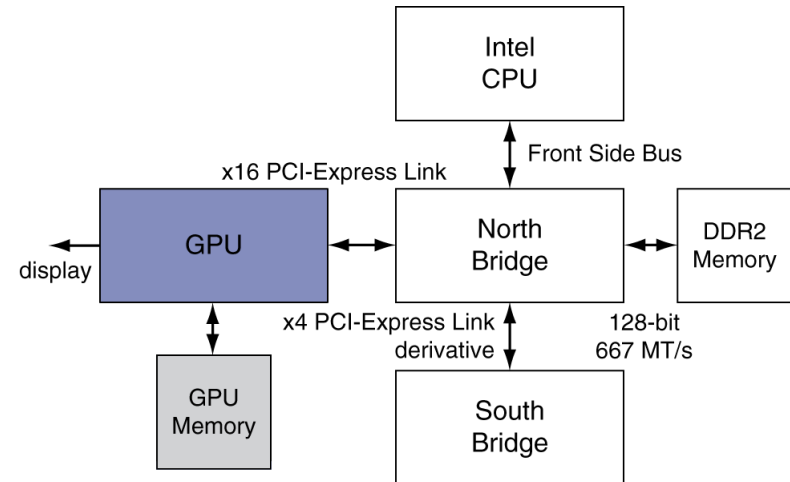
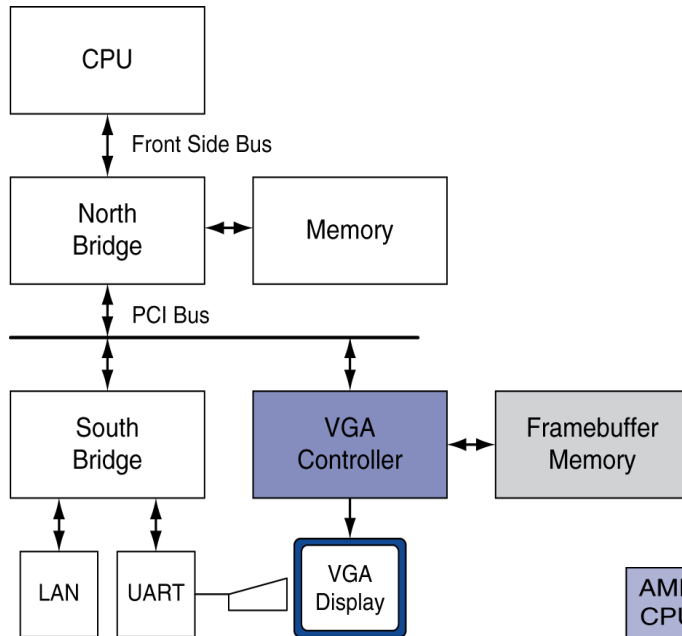
+

$$\begin{pmatrix} C_{0,0} & C_{0,1} & C_{0,2} & C_{0,3} \\ C_{1,0} & C_{1,1} & C_{1,2} & C_{1,3} \\ C_{2,0} & C_{2,1} & C_{2,2} & C_{2,3} \\ C_{3,0} & C_{3,1} & C_{3,2} & C_{3,3} \end{pmatrix}$$

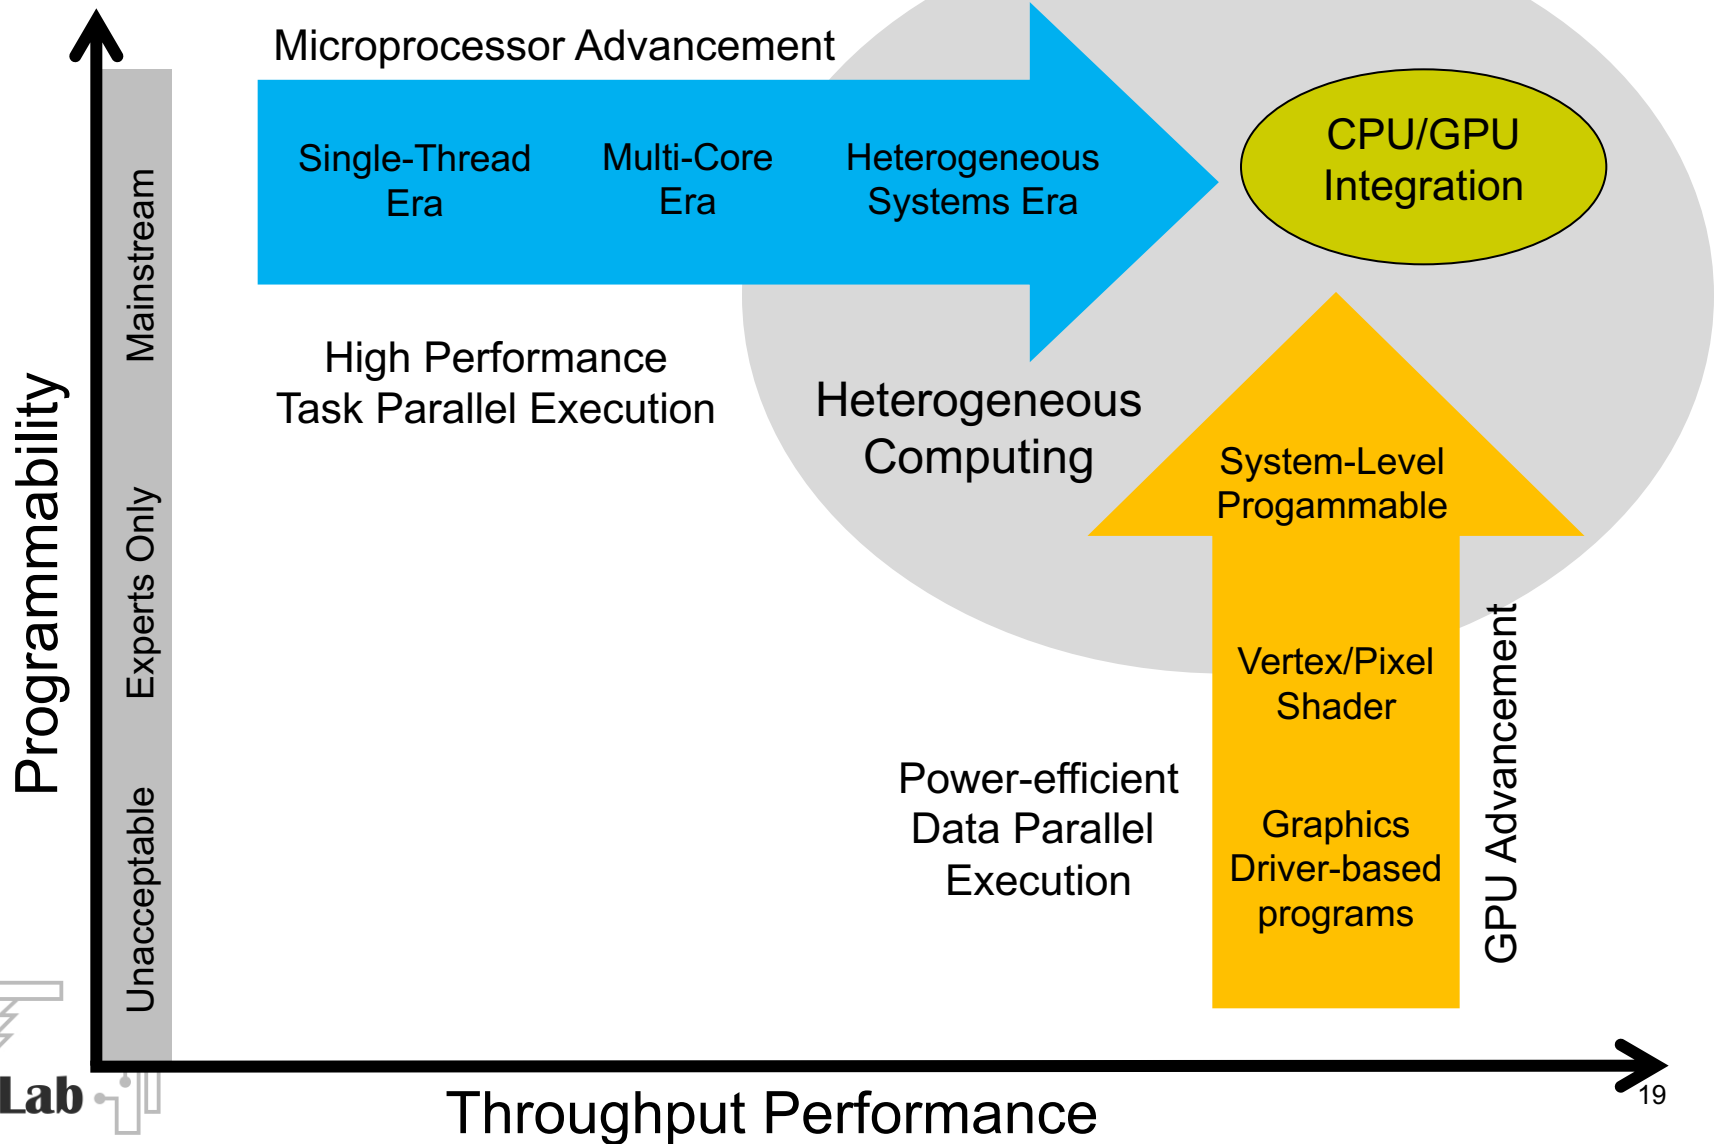
FP16 or FP32



Graphics in the System



CPU/GPU Integration: CPU's Advancement Meets GPU's



Heterogeneous Computing ~ 2011

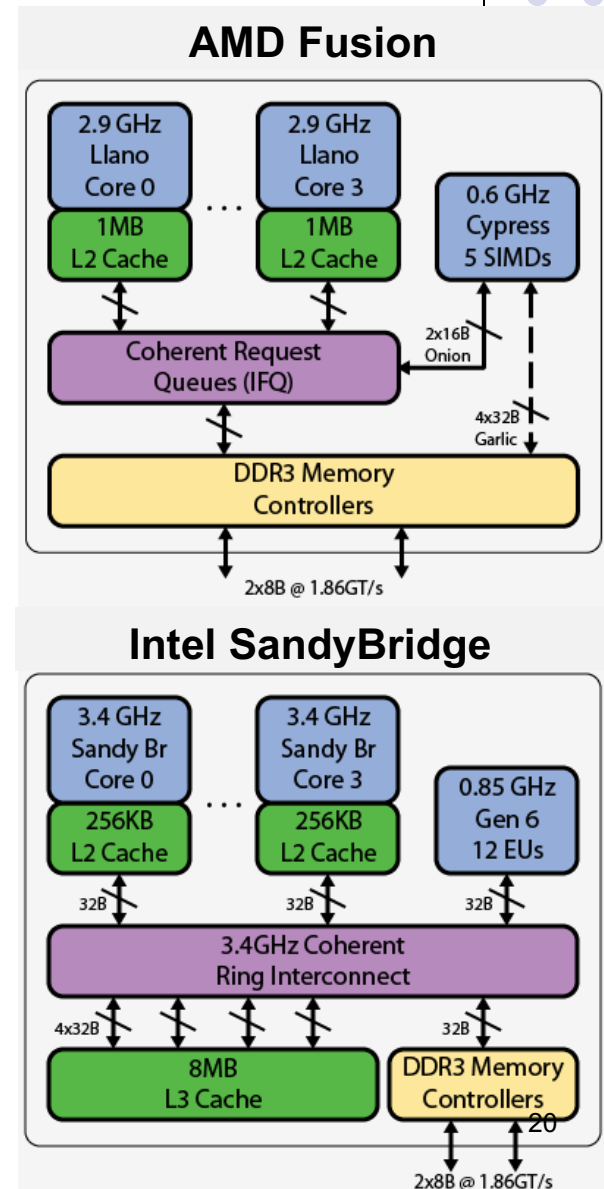


- **Intel SandyBridge**

- Shared last-level cache (LLC) and main memory

- **AMD Fusion APU (Accelerated Processing Unit)**

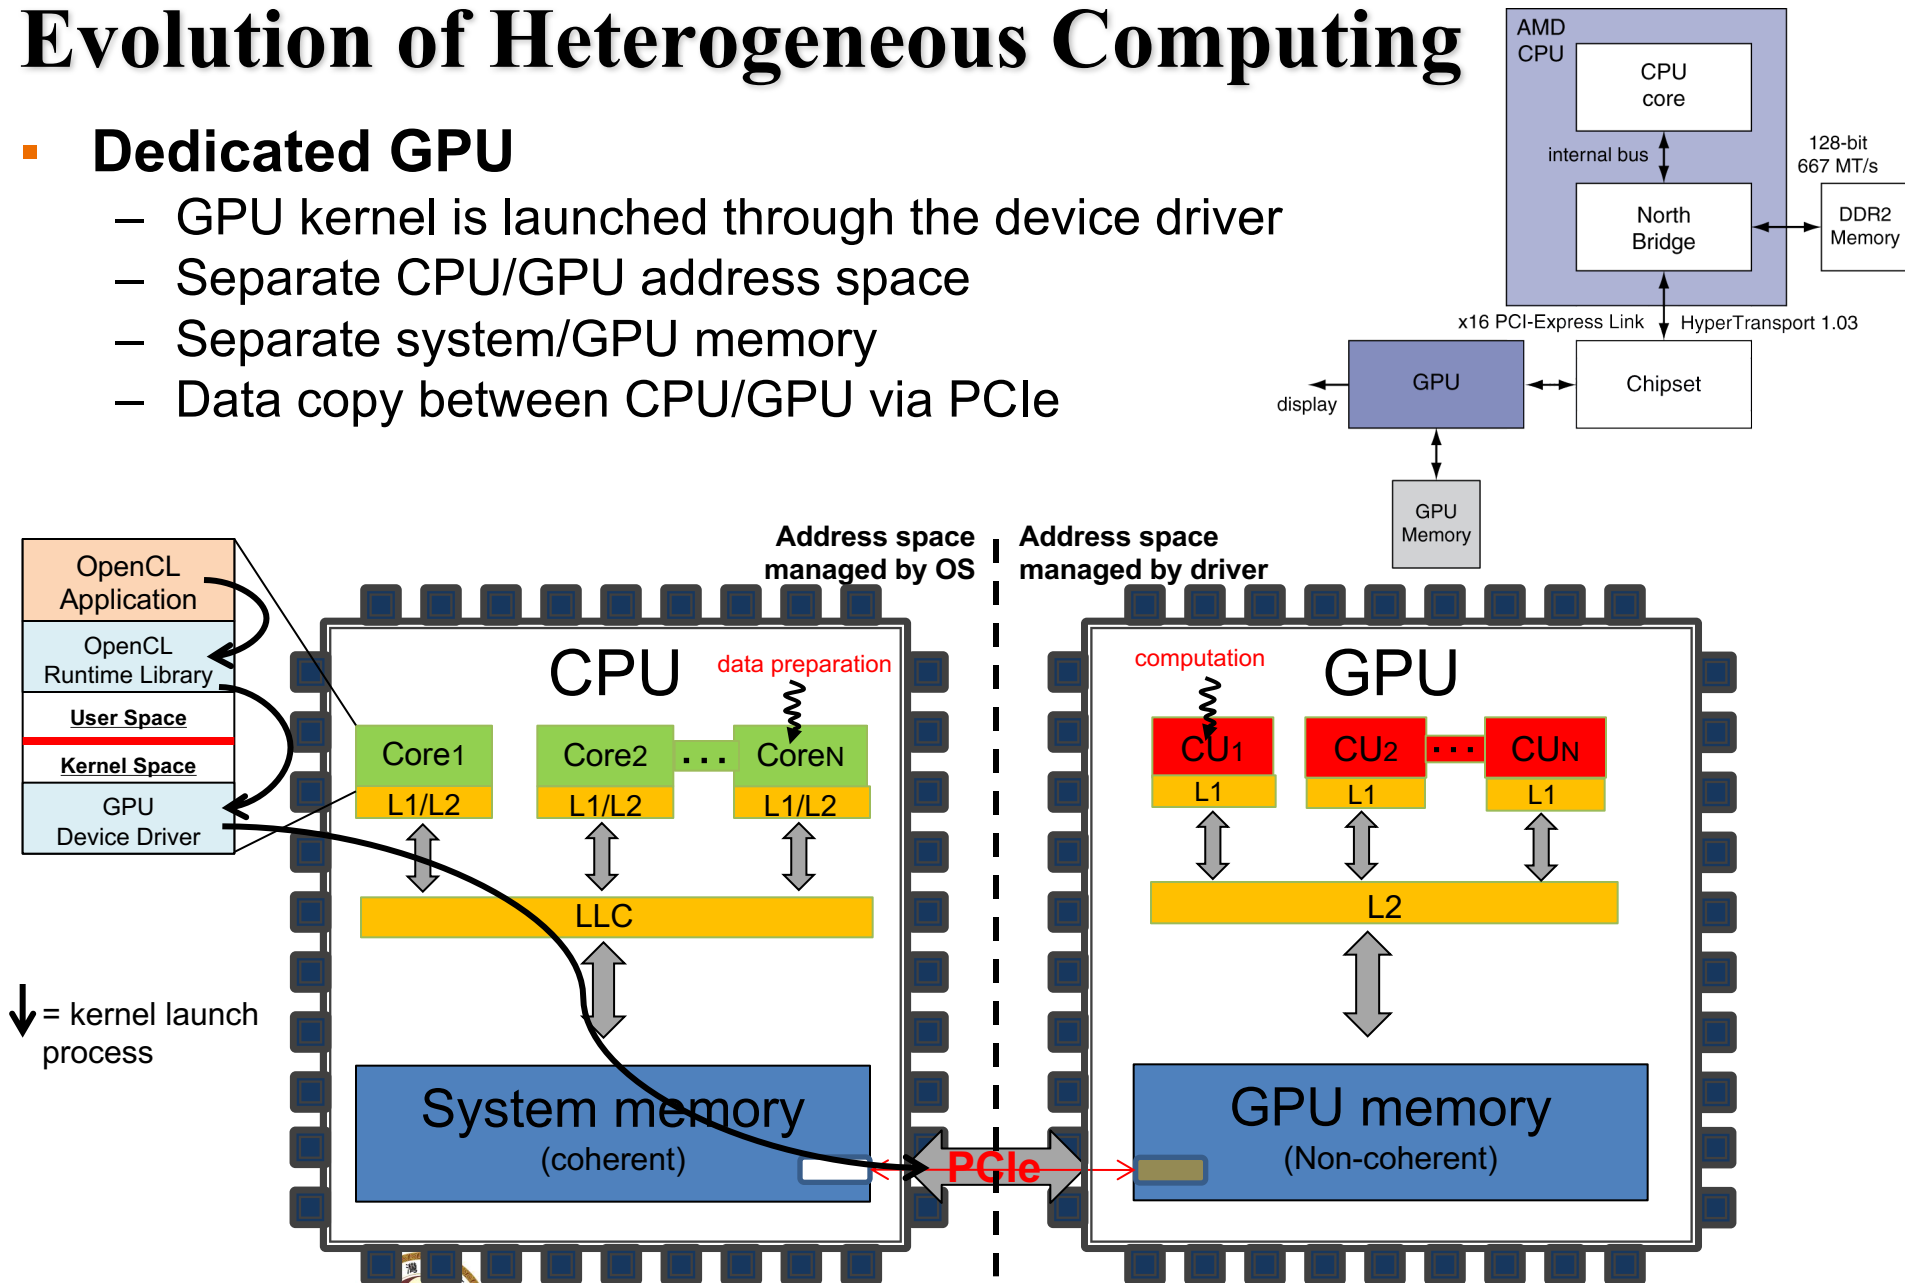
- Shared main memory



Evolution of Heterogeneous Computing

■ Dedicated GPU

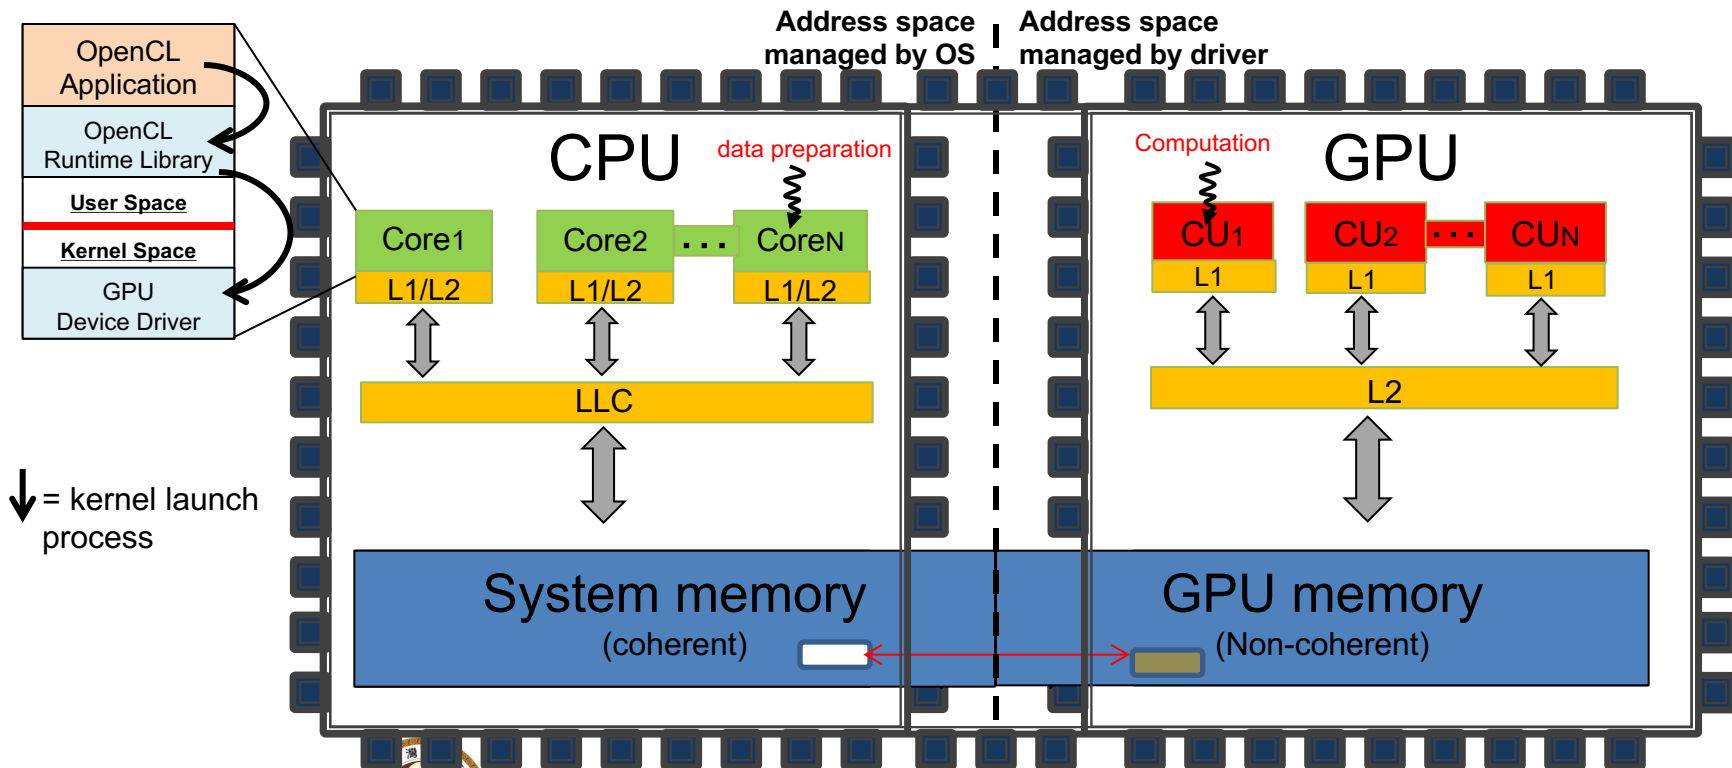
- GPU kernel is launched through the device driver
- Separate CPU/GPU address space
- Separate system/GPU memory
- Data copy between CPU/GPU via PCIe



Evolution of Heterogeneous Computing

■ Integrated GPU architecture

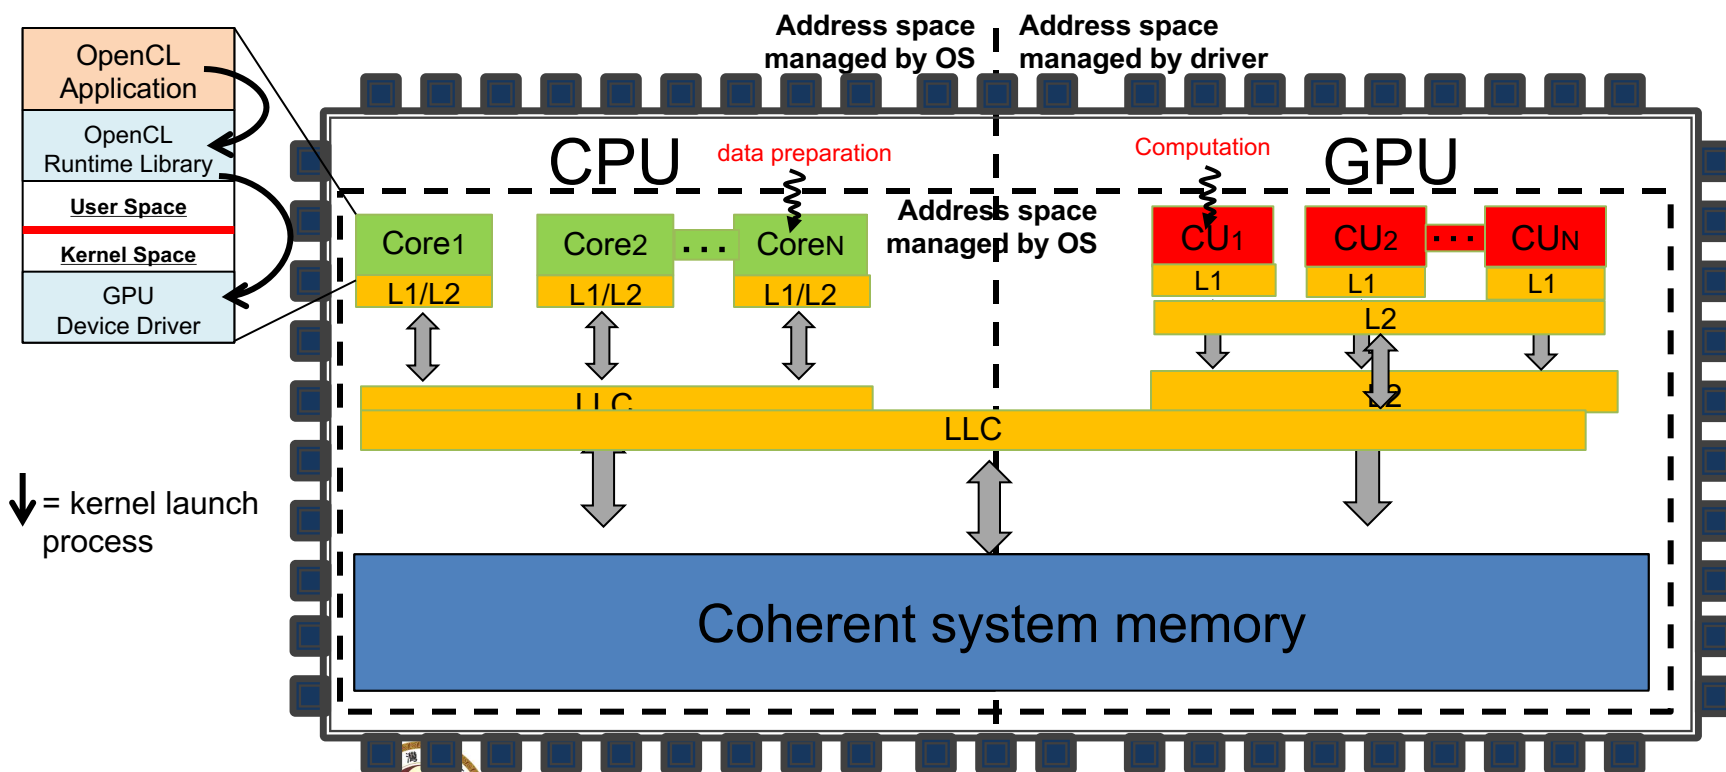
- GPU kernel is launched through the device driver
- Separate CPU/GPU address space
- Separate system/GPU memory
- Data copy between CPU/GPU *via memory bus*



Evolution of Heterogeneous Computing

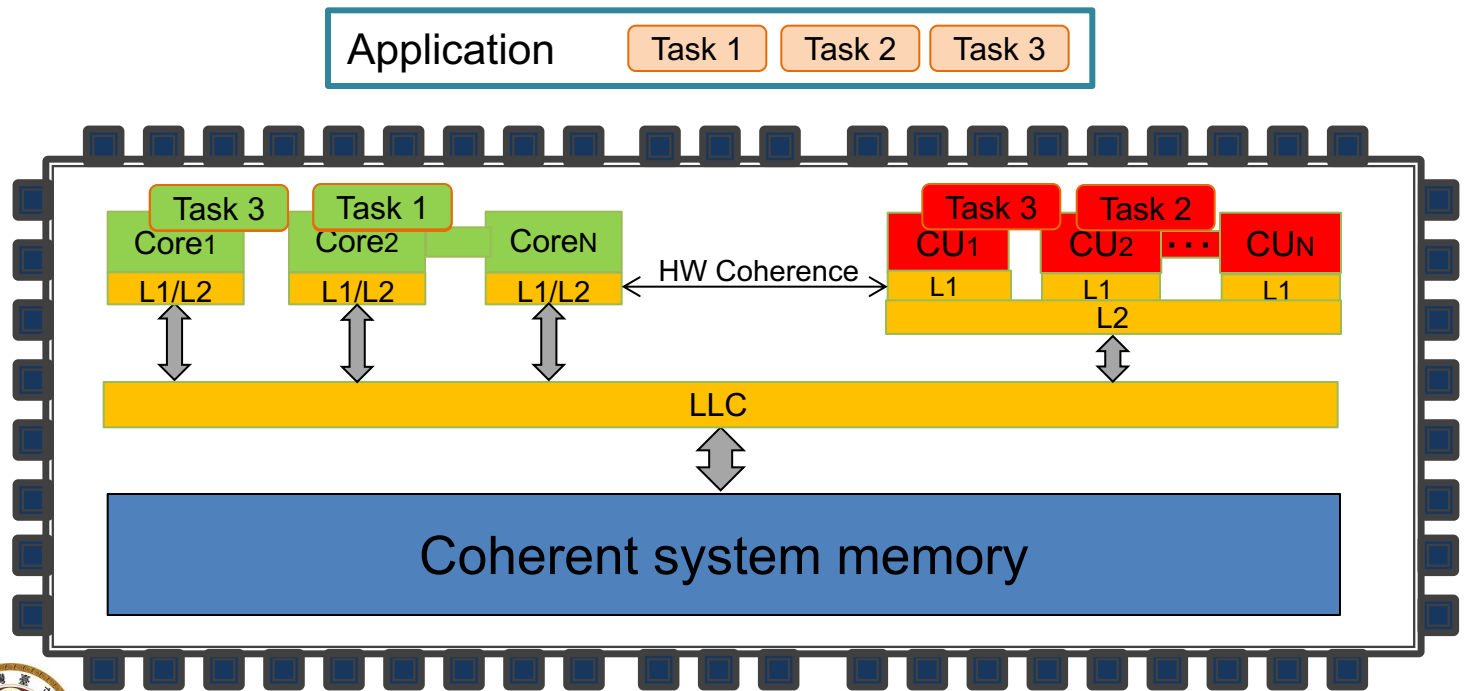
■ Integrated GPU architecture

- GPU kernel is launched through the device driver
- Unified CPU/GPU address space (managed by OS)
- Unified system/GPU memory
- No data copy - data can be retrieved by pointer passing



Utopia World of Heterogeneous Computing

- Processors are architected to operate cooperatively
 - Tasks in an application are executed on different types of core
 - Unified coherent memory enables data sharing across all processors
- Designed to enable the applications to run on different processors at different time
 - Capability to translate from high-level language to target binary at run-time
 - User-level task dispatch
 - Decision making module



Intel joins forces with AMD to battle Nvidia

By Paul Lilly November 06, 2017

A Core processor with custom Radeon graphics could be a game changer.

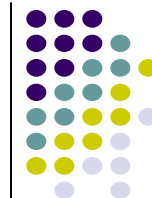


COMMENTS



https://www.youtube.com/watch?time_continue=84&v=gaHs_guCp2o

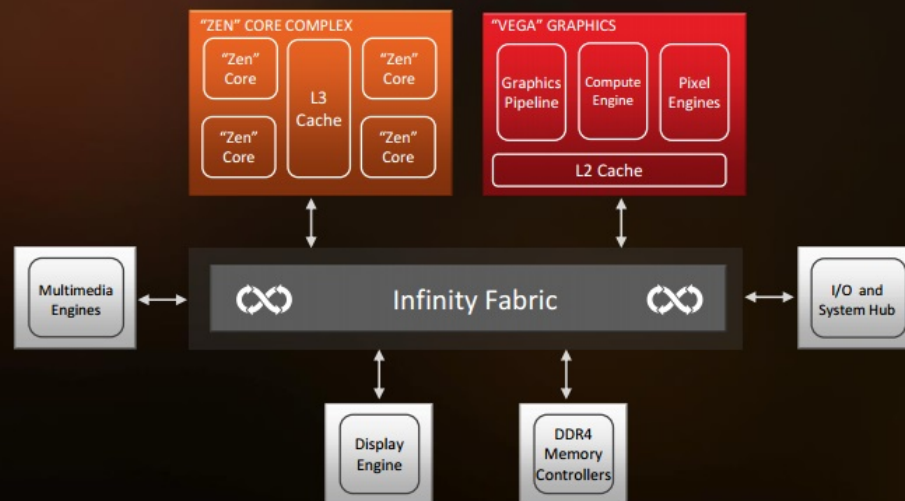
AMD Ryzen



Uniting “Zen” and “Vega”

With Infinity Fabric

- Our Fabric is the backbone of the AMD Ryzen™ APU
- One coherent control and data interface to integrate and manage the full SoC
- The Fabric services six clients in the SoC
- Telemetry monitoring in Infinity Fabric also governs Precision Boost 2 and mXFR*
- A project 4 years in the making



38 AMD Ryzen™ Processors with Vega Graphics | Confidential – Under Embargo until 10/26/2017 @ 9 AM EASTERN.

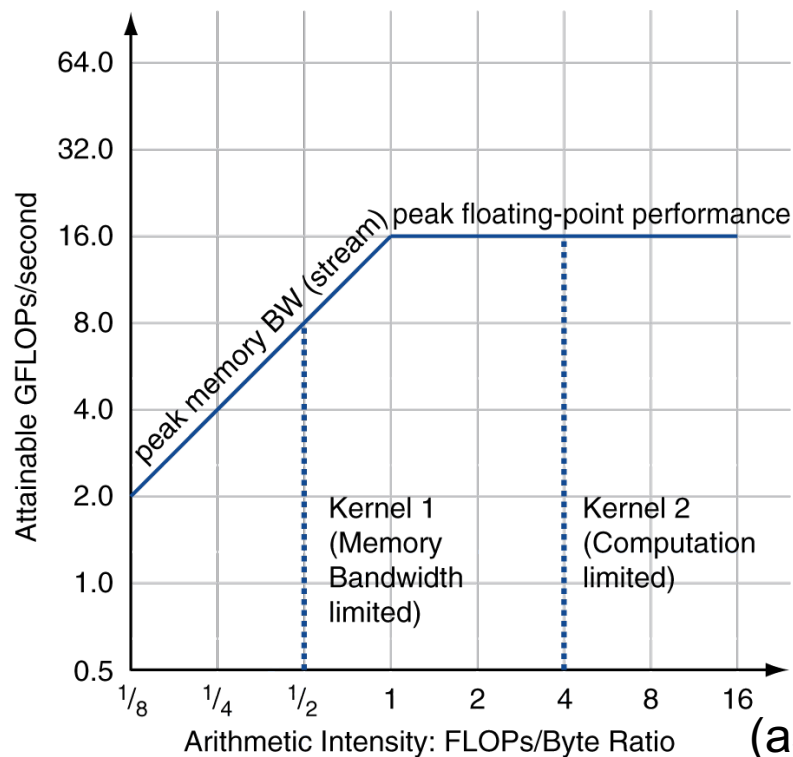
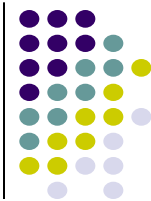
* mXFR enablement must meet AMD requirements. Not enabled on all notebook designs. Check with manufacturer to confirm “amplified mXFR performance” support.



Parallel Benchmarks

- **Linpack:** matrix linear algebra
- **SPECrate:** parallel run of SPEC CPU programs
 - Job-level parallelism
- **SPLASH: Stanford Parallel Applications for Shared Memory**
 - Mix of kernels and applications, strong scaling
- **NAS (NASA Advanced Supercomputing) suite**
 - computational fluid dynamics kernels
- **PARSEC (Princeton Application Repository for Shared Memory Computers) suite**
 - Multithreaded applications using Pthreads and OpenMP

Roofline : A Simple Performance Model



$$\frac{\text{Floating-Point Ops/sec}}{\text{Floating-Point Ops/ byte}} = \text{Bytes/sec}$$

Arithmetic Intensity

$$\text{Attainable GFLOPs/sec} = \text{Min} (\text{Peak Memory BW} \times \text{Arithmetic Intensity}, \text{Peak FP Performance})$$