

Operating System, Spring 2018

Project 3

DUE DATE: June 20, 2018

Team 16: B03505053 曾彦青 B03902129 陳鵬宇


Code reading

- How readahead is called when page faults occur?
 - mmap()
 - filemap_fault()

The structure **task_struct** (defined in "linux/sched.h") contains a substructure **mm_struct** (defined in "linux/mm_types.h") which is the memory descriptor, storing some useful information about the usage of memory.

In **mm_struct**, the first member is **vm_area_struct**, the *memory region* (defined in "linux/mm_types.h") which is a linked list of virtual memory area.

Thus we can get the following diagram:



In "linux/filemap.c", we can find following operations structure:

```
const struct vm_operations_struct generic_file_vm_ops = {  
    .fault = filemap_fault,  
};
```

Therefore, when a page fault occurs, it'll consequently invoke **filemap_fault()**, which will then check whether the required page is in the page cache by function **find_get_page()** first.

There are two conditions:

- (a) The page is in the page cache.
- (b) The page is not in the page cache.

For (a), we'll execute **async_readahead** to read pages.

For (b), we'll execute **sync_readahead** to read the required page and *readahead* other pages to cache.

Both **do_async_mmap_readahead()** and **do_sync_mmap_readahead()** will check whether VMA is randomly reading by **VM_RandomReadHint()**.

If **VM_RandomReadHint()** returns true, it's no need to do *readahead*, therefore both functions will return; otherwise, both functions will keep executing.

Finally, we'll find the page by **async_readahead** and **sync_readahead** if **MADV_RANDOM** has no effect (VMA isn't reading randomly). If we found the page (checking by **find_get_page()**), we'll lock the page and check whether it is truncated and up-to-date. After checking its size under page lock, we return the required page.

If **MADV_RANDOM** has no effect, we goto **no_cached_page**, and it'll execute **page_cache_read()** to read the required page and go back to **find_get_page()**.

Revise the readahead algorithm for smaller response time

We modify the original readahead function in "mm/readahead.c" and find that **struct file_ra_state** defined in "include/linux/fs.h" controls the readahead state of the file.

We also change **get_next_ra_size()** to decide the size of readahead.

```
static unsigned long get_next_ra_size(struct file_ra_state *ra, unsigned long max) {
    unsigned long cur = ra->size;
    unsigned long newsize;

    int ORIGINAL = false;

    if (ORIGINAL) {
        // original algorithm
        if (cur < max / 16)
            newsize = 4 * cur;
        else
            newsize = 2 * cur;
    } else {
        // revised algorithm
        if (cur < max / 32)
            newsize = 16 * cur;
        else if (cur < max / 16)
            newsize = 8 * cur;
        else
            newsize = 4 * cur;
    }

    return min(newsize, max);
}
```

There are two cases, and we test each case for 5 times:

1. Original readahead.c

	1	2	3	4	5	average
# of major pagefault	4158	4158	4158	4158	4158	4158
# of minor pagefault	2639	2641	2640	2640	2641	2640.2
# of resident set size	26620	26636	26624	26632	26632	26628.8
real(sec)	1.845	1.818	1.833	1.904	1.817	1.8434
user(sec)	0	0.036	0	0.016	0	0.0104
sys(sec)	0.018	0.16	0.196	0.0192	0.176	0.11384

2. Revised readahead.c

	1	2	3	4	5	average
# of major pagefault	178	178	178	178	178	178
# of minor pagefault	6618	6618	6619	6620	6620	6619
# of resident set size	26476	26476	26476	26476	26476	26476
real(sec)	0.35	0.371	0.378	0.359	0.354	0.3624
user(sec)	0	0	0.012	0.012	0.004	0.0056
sys(sec)	0.064	0.064	0.06	0.06	0.068	0.0632

Screenshots

```

jay@jay-VirtualBox: ~/hw3
831873783
1954997726
-802210018
-571331550
687073953
-1702804749
-2118621831
871876224
-186862404
-1924523629
1808037273
1187597919
361373333
102010677
671903510
1608468492
# of major pagefault: 4158
# of minor pagefault: 2641
# of resident set size: 26632 KB

real    0m1.817s
user    0m0.000s
sys     0m0.176s
jay@jay-VirtualBox:~/hw3$

```

(a) Original algorithm

```

jay@jay-VirtualBox: ~/hw3
831873783
1954997726
-802210018
-571331550
687073953
-1702804749
-2118621831
871876224
-186862404
-1924523629
1808037273
1187597919
361373333
102010677
671903510
1608468492
# of major pagefault: 178
# of minor pagefault: 6618
# of resident set size: 26476 KB

real    0m0.350s
user    0m0.000s
sys     0m0.064s
jay@jay-VirtualBox:~/hw3$

```

(b) Revised algorithm

```

[ 551.924525] page fault test program starts !
[ 553.643371] page fault test program ends !
[ 556.018349] page fault test program starts !
[ 556.090100] page fault test program ends !
[ 618.422999] page fault test program starts !
[ 620.233822] page fault test program ends !
[ 769.383010] page fault test program starts !
[ 771.152309] page fault test program ends !
[ 799.227301] page fault test program starts !
[ 801.010537] page fault test program ends !
[ 909.680981] page fault test program starts !
[ 911.536768] page fault test program ends !
[ 933.321050] page fault test program starts !
[ 934.995666] page fault test program ends !
jay@jay-VirtualBox:~/hw3$

```

Bonus

We use the following command to change the size of **VM_MAX_READHEAD** and test the revised algorithm, and we don't revise any code.

```
$ sudo /sbin/blockdev --setra 512 /dev/sda
$ sudo /sbin/blockdev --setra 2048 /dev/sda
$ sudo /sbin/blockdev --setra 8192 /dev/sda
```

Different size of **VM_MAX_READHEAD**:

	128(original)	512	2048	8192
average time(sec)	1.8434	1.593	0.3624	0.281

Screenshots

```
jay@jay-VirtualBox: ~/hw3
831873783
1954997726
-802210018
-571331550
687073953
-1702804749
-2118621831
871876224
-186862404
-1924523629
1808037273
1187597919
361373333
102010677
671903510
1608468492
# of major pagefault: 4158
# of minor pagefault: 2641
# of resident set size: 26632 KB

real    0m1.817s
user    0m0.000s
sys     0m0.176s
jay@jay-VirtualBox:~/hw3$
```

(a) **VM_MAX_READHEAD** = 128

```
jay@jay-VirtualBox: ~/hw3
831873783
1954997726
-802210018
-571331550
687073953
-1702804749
-2118621831
871876224
-186862404
-1924523629
1808037273
1187597919
361373333
102010677
671903510
1608468492
# of major pagefault: 3444
# of minor pagefault: 3353
# of resident set size: 26632 KB

real    0m1.593s
user    0m0.020s
sys     0m0.156s
jay@jay-VirtualBox:~/hw3$
```

(b) **VM_MAX_READHEAD** = 512

```
jay@jay-VirtualBox: ~/hw3
831873783
1954997726
-802210018
-571331550
687073953
-1702804749
-2118621831
871876224
-186862404
-1924523629
1808037273
1187597919
361373333
102010677
671903510
1608468492
# of major pagefault: 178
# of minor pagefault: 6618
# of resident set size: 26476 KB

real    0m0.350s
user    0m0.000s
sys     0m0.064s
jay@jay-VirtualBox:~/hw3$
```

(c) **VM_MAX_READHEAD** = 2048

```
jay@jay-VirtualBox: ~/hw3
831873783
1954997726
-802210018
-571331550
687073953
-1702804749
-2118621831
871876224
-186862404
-1924523629
1808037273
1187597919
361373333
102010677
671903510
1608468492
# of major pagefault: 47
# of minor pagefault: 6751
# of resident set size: 26472 KB

real    0m0.281s
user    0m0.012s
sys     0m0.064s
jay@jay-VirtualBox:~/hw3$
```

(d) **VM_MAX_READHEAD** = 8192

It's obvious that after enlarging **VM_MAX_READHEAD**, the real time is reduced. The reason is that when the parameter is enlarged, it'll make the number of pages we can pre-read each time more, thus letting the times of major page fault lesser, and improving the efficiency.