

# Contents

1. Introduction
2. System Structures
3. Process Concept
4. Multithreaded Programming
5. Process Scheduling
6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Mass-Storage Structures
12. I/O Systems
13. Protection, Security, Distributed Systems



# Chapter 6 Synchronization

# Process Synchronization

- Why Synchronization?
  - To ensure data consistency for concurrent access to shared data!
- Contents:
  - Various mechanisms to ensure the orderly execution of cooperating processes

# Process Synchronization

## ■ A Consumer-Producer Example

### ■ Producer

```
while (1) {  
    while (counter == BUFFER_SIZE)  
        ;  
    produce an item in nextp;  
    ....  
    buffer[in] = nextp;  
    in = (in+1) % BUFFER_SIZE;  
    counter++;  
}
```

### ■ Consumer:

```
while (1) {  
    while (counter == 0)  
        ;  
    nextc = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    consume an item in nextc;  
}
```

# Process Synchronization

- counter++ vs counter—  
r1 = counter      r2 = counter  
r1 = r1 + 1      r2 = r2 - 1  
counter = r1      counter = r2
- Initially, let counter = 5.
  1. P: r1 = counter
  2. P: r1 = r1 + 1
  3. C: r2 = counter
  4. C: r2 = r2 - 1
  5. P: counter = r1
  6. C: counter = r2


⇒ A Race Condition!

# Process Synchronization

- A Race Condition:
  - A situation where the outcome of the execution depends on the particular order of process scheduling.
- The Critical-Section Problem:
  - Design a protocol that processes can use to cooperate.
    - Each process has a segment of code, called a critical section, whose execution must be mutually exclusive.

# Process Synchronization

- A General Structure for the Critical-Section Problem



# The Critical-Section Problem

- Three Requirements

- 1. Mutual Exclusion

- a. Only one process can be in its critical section.

- 2. Progress

- a. Only processes not in their remainder section can decide which will enter its critical section.
- b. The selection cannot be postponed indefinitely.

- 3. Bounded Waiting

- a. A waiting process only waits for a bounded number of processes to enter their critical sections.



# The Critical-Section Problem – Peterson's Solution

## ■ Notation

- Processes  $P_i$  and  $P_j$ , where  $j=1-i$ ;

## ■ Assumption

- Every basic machine-language instruction is atomic.

## ■ Algorithm 1

- Idea: Remember which process is allowed to enter its critical section, That is, process  $i$  can enter its critical section if  $turn = i$ .

do {

while ( $turn \neq i$ ) ;

critical section


turn=j;

remainder section

} while (1);

# The Critical-Section Problem – Peterson's Solution

- Algorithm 1 fails the progress requirement:



# The Critical-Section Problem – Peterson's Solution

## ■ Algorithm 2

- Idea: Remember the state of each process.
- $\text{flag}[i] == \text{true} \rightarrow P_i$  is ready to enter its critical section.
- Algorithm 2 fails the progress requirement when  $\text{flag}[0] == \text{flag}[1] == \text{true}$ ;
  - the exact timing of the two processes?

Initially,  $\text{flag}[0] = \text{flag}[1] = \text{false}$

do {

$\text{flag}[i] = \text{true};$

while ( $\text{flag}[j]$ ) ;

critical section

$\text{flag}[i] = \text{false};$

remainder section

} while (1);

# The Critical-Section Problem – Peterson's Solution

- Algorithm 3
  - Idea: Combine the ideas of Algorithms 1 and 2
  - When  $(\text{flag}[i] \ \&\& \ \text{turn}=i)$ ,  $P_j$  must wait.
  - Initially,  $\text{flag}[0]=\text{flag}[1]=\text{false}$ , and  $\text{turn} = 0$  or  $1$

do {

```
flag[i]=true;
```

```
turn=j;
```

```
while (flag[j] && turn==j) ;
```

critical section

```
flag[i]=false;
```

remainder section

```
} while (1);
```

# The Critical-Section Problem – Peterson's Solution

- Properties of Algorithm 3
  - Mutual Exclusion
    - The eventual value of *turn* determines which process enters the critical section.
  - Progress
    - A process can only be stuck in the while loop, and the process which can keep it waiting must be in its critical sections.
  - Bounded Waiting
    - Each process wait at most one entry by the other process.

# The Critical-Section Problem – A Multiple-Process Solution

- Bakery Algorithm
  - Originally designed for distributed systems
  - Processes which are ready to enter their critical section must take a number and wait till the number becomes the lowest.
  - **int number[i]:  $P_i$ 's number if it is nonzero.**
  - **boolean choosing[i]:  $P_i$  is taking a number.**

# The Critical-Section Problem – A Multiple-Process Solution

do {

```
    choosing[i]=true;  
    number[i]=max(number[0], ...number[n-1])+1;  
    choosing[i]=false;  
    for (j=0; j < n; j++)  
        while choosing[j] ;  
        while (number[j] != 0 && (number[j],j)<(number[i],i)) ;
```

critical section

```
    number[i]=0;
```

remainder section

} while (1);

- An observation: If  $P_i$  is in its critical section, and  $P_k$  ( $k \neq i$ ) has already chosen its number  $k$ , then  $(\text{number}[i], i) < (\text{number}[k], k)$ .

# Synchronization Hardware

- Motivation:
  - Hardware features make programming easier and improve system efficiency.
- Approach:
  - Disable Interrupt → No Preemption
    - Infeasible in multiprocessor environment where message passing is used.
    - Potential impacts on interrupt-driven system clocks.
  - Atomic Hardware Instructions
    - Test-and-set, Swap, etc.



# Synchronization Hardware

```
boolean TestAndSet(boolean *target) {  
    boolean rv = *target;  
    *target=true;  
    return rv;  
}
```

---

```
do {  
    while (TestAndSet(&lock)) ;  
    critical section  
    lock=false;  
    remainder section  
} while (1);
```

# Synchronization Hardware

```
void Swap(boolean *a, boolean *b) {  
    boolean temp = *a;  
    *a=*b;  
    *b=temp;  
}
```

---

```
do {  
    key=true;  
    while (key == true)  
        Swap(&lock, &key);  
    critical section  
    lock=false;  
    remainder section  
} while (1);
```

# Synchronization Hardware

do {

```
waiting[i]=true;
key=true;
while (waiting[i] && key)
    key=TestAndSet(&lock);
waiting[i]=false;
```

critical section;

```
j= (i+1) % n;
while(j != i) && (not waiting[j])
    j= (j+1) % n;
If (j=i) lock=false;
else waiting[j]=false;
```

remainder section

} while (1);

- Mutual Exclusion
  - Pass if key == F or waiting[i] == F
- Progress
  - Exit process sends a process in.
- Bounded Waiting
  - Wait at most n-1 times
- Atomic TestAndSet is hard to implement in a multiprocessor environment.

# Mutex Locks

- Motivation:
  - A high-level software solution to provide protect critical sections with mutual exclusion.
- Implementation:
  - Atomic execution of acquire() and release().
  - Spinlock.
    - Disadvantage: Busy waiting.
    - Advantage: No context switch for multiprocessor systems.

```
acquire() {  
    while (!available) ;  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

# Semaphores

- Motivation:
  - A high-level solution for more complex problems.
- Semaphore
  - A variable **S** only accessible by two atomic operations:

```
wait(S) {      /* P */           signal(S) {    /* V */
    while (S <= 0) ;              S++;
    S—;                           }
}
```

•Indivisibility for “(S<=0)”, “S—”, and “S++”

# Semaphores – Usages

## ■ Critical Sections

```
do {  
    wait(mutex);  
    critical section  
    signal(mutex);  
    remainder section  
} while (1);
```

## ■ Precedence Enforcement

```
P1:  
    S1;  
    signal(synch);
```

```
P2:  
    wait(synch);  
    S2;
```

# Semaphores

- Implementation
  - Spinlock – A Busy-Waiting Semaphore
    - “while ( $S \leq 0$ )” causes the wasting of CPU cycles!
    - Advantage:
      - When locks are held for a short time, spinlocks are useful since no context switching is involved.
  - Semaphores with Block-Waiting
    - No busy waiting from the entry to the critical section!

# Semaphores

- Semaphores with Block Waiting

```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore ;
```

```
void wait(semaphore S) {  
    S.value--;  
    if (S.value < 0) {  
        add this process to S.L;  
        block();  
    }  
}
```

```
void signal(semaphore S);  
    S.value++;  
    if (S.value <= 0) {  
        remove a process P form S.L;  
        wakeup(P);  
    }  
}
```



# Semaphores

- The queueing strategy can be arbitrary, but there is a restriction for the bounded-waiting requirement.
- Mutual exclusion in wait() & signal()
  - Uniprocessor Environments
    - Interrupt Disabling
    - TestAndSet, Swap
    - Software Methods, e.g., the Bakery Algorithm
  - Multiprocessor Environments
- Remarks: Busy-waiting is limited to only the critical sections of the wait() & signal()!

# Deadlocks and Starvation

- Deadlock

- A set of processes is in a deadlock state when every process in the set is waiting for an event that can be caused only by another process in the set.

P0: wait(S);

wait(Q);

...

signal(S);

signal(Q);

P1: wait(Q);

wait(S);

...

signal(Q);


signal(S);

- Starvation (or Indefinite Blocking)

- E.g., a LIFO queue

# Priority Inversion

- Definition: A higher-priority task is blocked by a lower-priority task due to some resource access conflict.
  - Examples such as system calls



# Binary Semaphore

- Binary Semaphores versus Counting Semaphores
  - The value ranges from 0 to 1 → easy implementation!

## WAIT(S)

```
wait(S3);  
wait(S1); /* protect C */  
C--;  
if (C < 0) {  
    signal(S1);  
    wait(S2);  
} else signal(S1);  
signal(S3);
```

## SIGNAL(S)

```
wait(S1);  
C++;  
if (C <= 0)  
    signal (S2); /* wakeup */  
signal (S1);
```

\* S1 & S2: binary semaphores

# Classical Synchronization Problems – The Bounded Buffer

Producer:

do {

    produce an item in nextp;

    .....

Initialized to  $n$   $\Rightarrow$  wait(empty); /\* control buffer availability \*/

Initialized to 1  $\Rightarrow$  wait(mutex); /\* mutual exclusion \*/

    .....

    add nextp to buffer;

    signal(mutex);

Initialized to 0  $\Rightarrow$  signal(full); /\* increase item counts \*/

} while (1);

# Classical Synchronization Problems – The Bounded Buffer

Consumer:

do {

Initialized to 0  $\Rightarrow$  wait(full); /\* control buffer availability \*/

Initialized to 1  $\Rightarrow$  wait(mutex); /\* mutual exclusion \*/

.....

remove an item from buffer to nextp;

.....

signal(mutex);

Initialized to  $n$   $\Rightarrow$  signal(empty); /\* increase item counts \*/  
consume nextp;

} while (1);

# Classical Synchronization Problems – Readers and Writers

- The Basic Assumption:
  - Readers: shared locks
  - Writers: exclusive locks
- The first reader-writers problem
  - No readers will be kept waiting unless a writer has already obtained permission to use the shared object → potential hazard to writers!
- The second reader-writers problem:
  - Once a writer is ready, it performs its write asap! → potential hazard to readers!

# Classical Synchronization Problems – Readers and Writers

First R/W  
Solution



```
semaphore wrt, mutex;  
    (initialized to 1);  
int readcount=0;
```

Writer:

Queueing  
mechanism



```
wait(wrt);  
.....  
writing is performed  
.....  
signal(wrt)
```

Reader:

```
wait(mutex);  
readcount++;  
if (readcount == 1)  
    wait(wrt);  
signal(mutex);  
..... reading .....  
wait(mutex);  
readcount--;  
if (readcount== 0)  
    signal(wrt);  
signal(mutex);
```





Which is awoken?





# Classical Synchronization Problems – Dining-Philosophers


- Each philosopher must pick up one chopstick beside him/her at a time
- When two chopsticks are picked up, the philosopher can eat.



# Classical Synchronization Problems – Dining-Philosophers

```
semaphore chopstick[5];  
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i + 1) % 5]);  
    ... eat ...  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    ...think ...  
} while (1);
```

# Classical Synchronization Problems – Dining-Philosophers



- Deadlock or Starvation?!
- Solutions to Deadlocks:
  - At most four philosophers appear.
  - Pick up two chopsticks “simultaneously”.
  - Order their behaviors, e.g., odds pick up their right one first, and evens pick up their left one first.
- Solutions to Starvation:
  - No philosopher will starve to death.
    - A deadlock could happen??

# Critical Regions/Monitor

- Motivation:
  - Various programming errors in using low-level constructs, e.g., semaphores
    - Interchange the order of wait and signal operations
    - Miss some waits or signals
    - Replace waits with signals
    - etc
- The needs of high-level language constructs to reduce the possibility of errors!

# Critical Regions

- Region  $v$  when  $B$  do  $S$ ;
  - Variable  $v$  – shared among processes and only accessible in the region

```
struct buffer {  
    item pool[n];  
    int count, in, out;  
};
```
  - $B$  – condition
    - $\text{count} < 0$
  - $S$  – statements

**Example: Mutual Exclusion**  
region  $v$  when (true)  $S1$ ;  
region  $v$  when (true)  $S2$ ;

# Critical Regions – Consumer-Producer

```
struct buffer {  
    item pool[n];  
    int count, in, out;  
};
```

## Producer:

```
region buffer when  
(count < n) {  
    pool[in] = nextp;  
    in = (in + 1) % n;  
    count++;  
}
```

## Consumer:

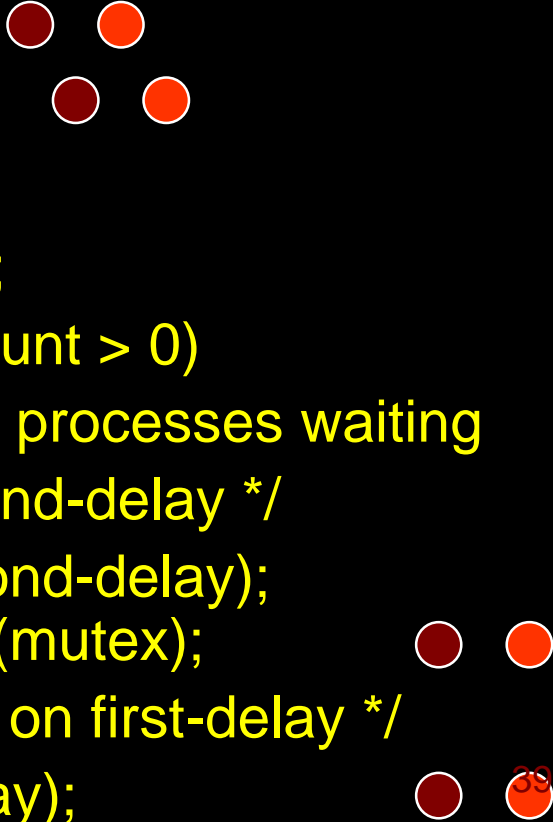
```
region buffer when  
(count > 0) {  
    nextc = pool[out];  
    out = (out + 1) % n;  
    count--;  
}
```

# Critical Regions – Implementation by Semaphores

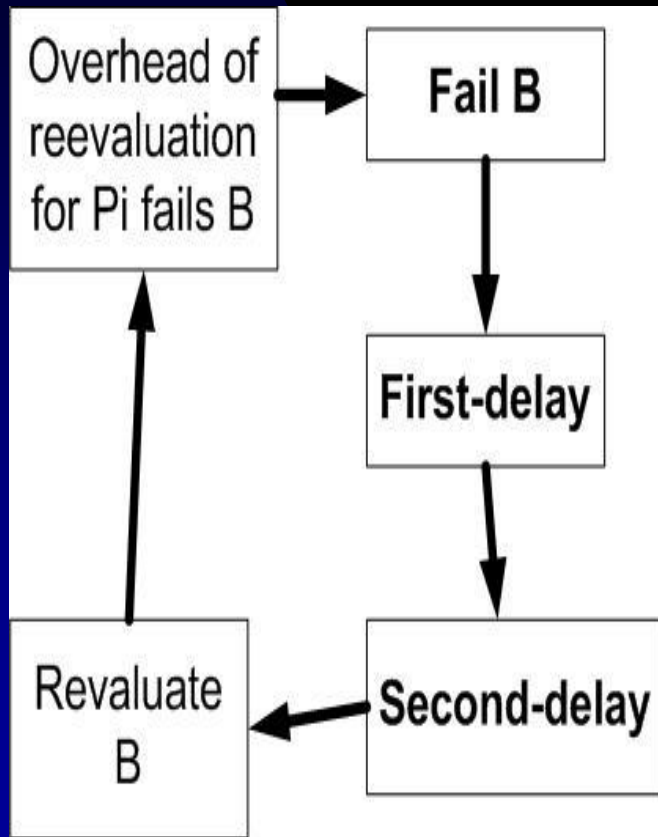
Region x when B do S;

```
/* to protect the region */  
semaphore mutex;  
/* to (re-)test B */  
semaphore first-delay;  
int first-count=0;  
/* to retest B */  
semaphore second-delay;  
int second-count=0;
```

```
wait(mutex);  
while (!B) {  
    /* fail B */  
    first-count++;  
    if (second-count > 0)  
        /* try other processes waiting  
         on second-delay */  
        signal(second-delay);  
    else signal(mutex);  
    /* block itself on first-delay */  
    wait(first-delay);
```



# Critical Regions – Implementation by Semaphores



```
first-count--;  
second-count++;  
if (first-count > 0)  
    signal(first-delay);  
    else signal(second-delay);  
/* block itself on first-delay */  
wait(second-delay);  
second-count--;
```


```
}  
S;
```

```
if (first-count > 0)  
    signal(first-delay);  
else if (second-count > 0)  
    signal(second-delay);  
else signal(mutex);
```



# Monitor

- Components
  - Variables – monitor state
  - Procedures
    - Only access local variables or formal parameters
  - Condition variables
    - Tailor-made sync
    - `x.wait()` or `x.signal`



```
monitor name {  
    variable declaration  
    void proc1(...) {  
    }  
    ...  
    void procn(...) {  
    }  
}
```

# Monitor

- Semantics of signal & wait
  - $x.\text{signal}()$  resumes one suspended process. If there is none, no effect is imposed.
  - $P\ x.\text{signal}()$  a suspended process  $Q$ 
    - $P$  either waits until  $Q$  leaves the monitor or waits for another condition
    - $Q$  either waits until  $P$  leaves the monitor, or waits for another condition.

# Monitor – Dining-Philosophers

```
monitor dp {  
    enum {thinking, hungry, eating} state[5];  
    condition self[5];  
    void pickup(int i) {  
        stat[i]=hungry;  
        test(i);  
        if (stat[i] != eating)  
            self[i].wait;  
    }  
    void putdown(int i) {  
        stat[i] = thinking;  
        test((i+4) % 5);  
        test((i + 1) % 5);  
    }  
}
```

Pi:

```
dp.pickup(i);  
... eat ...  
dp.putdown(i);
```

# Monitor – Dining-Philosophers

No deadlock!  
But starvation could occur!

```
void test(int i) {  
    if (stat[(i+4) % 5] != eating &&  
        stat[i] == hungry &&  
        state[(i+1) % 5] != eating) {  
        stat[i] = eating;  
        self[i].signal();  
    }  
}  
  
void init() {  
    for (int i=0; i < 5; i++)  
        state[i] = thinking;  
}
```

# Monitor – Implementation by Semaphores

- Semaphores
  - *mutex* – to protect the monitor
  - *next* – being initialized to zero, on which processes may suspend themselves
    - *nextcount*
- For each external function  $F$ 
  - wait(mutex);*
  - ...*
  - body of F;*
  - ...*
  - if (next-count > 0)*
    - signal(next);*
  - else signal(mutex);*

# Monitor – Implementation by Semaphores

- For every condition  $x$ 
  - A semaphore  $x\text{-sem}$
  - An integer variable  $x\text{-count}$
  - Implementation of  $x.\text{wait}()$  and  $x.\text{signal}$  :
    - $x.\text{wait}()$ 

```
x-count++;  
if (next-count > 0)  
    signal(next);  
else signal(mutex);  
wait(x-sem);  
x-count--;
```
    - $x.\text{signal}$ 

```
if (x-count > 0) {  
    next-count++;  
    signal(x-sem);  
    wait(next);  
    next-count--;  
}
```

# Monitor

```
monitor ResAllc {  
  boolean busy;  
  condition x;  
  void acquire(int time) {  
    if (busy)  
      x.wait(time);  
    busy=true;  
  }  
  ...  
}
```

- Process-Resumption Order
  - Queuing mechanisms for a monitor and its condition variables.
  - A solution:  
    x.wait(c);
    - where the expression c is evaluated to determine its process's resumption order.  
        R.acquire(t);  
        ...  
        access the resource;  
        R.release;

# Monitor

- Concerns:
  - Processes may access resources without consulting the monitor.
  - Processes may never release resources.
  - Processes may release resources which they never requested.
  - Process may even request resources twice.



# Monitor

- Remark: Whether the monitor is correctly used?
  - => Requirements for correct computations
    - Processes always make their calls on the monitor in correct order.
    - No uncooperative process can access resource directly without using the access protocols.
- Note: Scheduling behavior should consult the built-in monitor scheduling algorithm if resource access RPC are built inside the monitor.

# Synchronization – Windows

- General Mechanism
  - Spin-locking for short code segments in a multiprocessor platform.
  - Interrupt disabling when the kernel accesses global variables in a uniprocessor platform.
- Dispatcher Object
  - State: signaled or non-signaled
  - *Mutex* – select one process from its waiting queue to the ready queue.
    - Critical-section object – user-mode mutex
  - *Events* – select all processes waiting for the event.


# Synchronization – Linux

- Preemptive Kernel After Version 2.6
  - Atomic integer

```
atomic_t counter;  
...  
atomic_set(&counter, 5); atomic_add(10, &counter);
```
  - Semaphores for long code segments. Mutex locks for the kernel code.
  - Spin-locking for short code segments in a multiprocessor platform.
  - Preemption disabling and enabling in a uniprocessor platform.
    - `preempt_disable()` and `preempt_enable()`
    - `Preempt_count` for each task in the system.

# Synchronization – Solaris

- Semaphores and Condition Variables
- Adaptive Mutex
  - Spin-locking if the lock-holding thread is running; otherwise, blocking is used.
- Readers-Writers Locks
  - Expensive in implementations.
- Turnstile
  - A queue structure containing threads blocked on a lock.
  - Priority inversion → priority inheritance protocol for kernel threads



# Synchronization – Pthreads

- General Mechanism
  - Mutex locks – mutual exclusion
  - Condition variables – Monitor
  - Read-write locks
- Extensions
  - POSIX SEM extension: semaphores
    - Named and unnamed semaphores

```
sem_t sem;  
...  
sem_init(&sem, 0 ,1); /* sharing mode and initial value */  
sem_wait(&sem); .... sem_post(&sem);
```
  - Spinlocks – portability?

# Alternative Approaches

- Motivation:
  - As the number of cores increases, it becomes hard to avoid race conditions and deadlocks.
- Transactional Memory
  - Memory Transaction: A sequence of memory read-write operations that are atomic.
  - Committed or being rolled back.

```
void update() {  
    atomic {  
        /* modify shared data */  
    }  
}
```

# Alternative Approaches

- Advantages:
  - No deadlock
  - Identification of potentially concurrently executing statements in atomic blocks.
- Implementations (with features added to program language):
  - Software Transactional Memory
    - Code is inserted by the compiler.
  - Hardware Transactional Memory
    - Hardware cache hierarchies and cache coherency protocols are used

# Alternative Approaches

- OpenMP

- A set of compiler directives and an API
  - The critical-section compiler directive behaves like a binary semaphore or mutex.

```
void update() {  
    #pragma omp critical {  
        counter += value;  
    }  
}
```

- Advantage: Easy to use
- Disadvantage: Identification of protected code and potentials of deadlocks.

- Functional Programming Language