

Contents

1. Introduction
2. System Structures
3. Process Concept
4. Multithreaded Programming
5. Process Scheduling
6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Mass-Storage Structures
12. I/O Systems
13. Distributed Systems




Distributed Systems

- Definition: A collection of loosely coupled processors interconnected by a communication network.
- Design Issues:
 - A Transparent Interface to Access of Local/Remote Resources
 - Fault Tolerance – Communication, Machines, Services, etc.
 - Scalability

Distributed Systems

- Distributed Synchronization
 - Challenges: No Common Clocks and Memory
 - A Happened-Before Relation – A Partial Order of Events



Distributed Systems

- Why ordering matters?
 - Correctness in concurrent process executions!
- How to achieve the goal?
 - Locking – Two-Phase Locking
 - Timestamp – Synchronization of Logical Counters
 - E.g., $A \rightarrow B \Rightarrow LC(B) = LC(A) + X$
- Issues:
 - Atomicity – Two-Phase Commit Protocol
 - Deadlocks

Q & A

Atomic Transactions

- Why Atomic Transactions?
 - Critical sections ensure mutual exclusion in data sharing, but the relationship between critical sections and the logical unit of work might also be important!
→ Atomic Transactions
- Operating systems can be viewed as manipulators of data!

Atomic Transactions – System Model

- Transaction – a logical unit of computation
 - A sequence of read and write operations followed by a commit or an abort.
- Beyond “critical sections”
 1. Atomicity: All or Nothing
 - An aborted transaction must be *rolled back*.
 - The effect of a committed transaction must persist and be imposed as a logical unit of operations.

Atomic Transactions – System Model

2. Serializability:

- The order of transaction executions must be equivalent to a serial schedule.

T0	T1
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

Two operations O_i & O_j conflict if

1. Access the same object
2. One of them is write

Atomic Transactions – System Model


- Conflict Serializable:
 - S is conflict serializable if S can be transformed into a serial schedule by swapping nonconflicting operations.

T0	T1		T0	T1
R(A)			R(A)	
W(A)			W(A)	
	R(A)		R(B)	
	W(A)	⇒	W(B)	
R(B)				R(A)
W(B)				W(A)
	R(B)			R(B)
	W(B)			W(B)

Atomic Transactions – Concurrency Control

- Locking Protocols
 - Lock modes (A general approach!)
 - 1. Shared-Mode: “Reads”.
 - 2. Exclusive-Mode: “Reads” & “Writes”
 - General Rule
 - A transaction must receive a lock of an appropriate mode of an object before it accesses the object. The lock may not be released until the last access of the object is done.

Atomic Transactions – Concurrency Control




Atomic Transactions – Concurrency Control

- When to release locks w/o violating serializability

$R0(A)$ $W0(A)$ $R1(A)$ $R1(B)$ $R0(B)$ $W0(B)$

- Two-Phase Locking Protocol (2PL) –
Not Deadlock-Free




- How to improve 2PL?
 - Semantics, Order of Data, Access Pattern, etc.

Atomic Transactions – Concurrency Control



- Timestamp-Based Protocols
 - A time stamp for each transaction $TS(T_i)$
 - Determine transactions' order in a schedule in advance!
 - A General Approach:
 - $TS(T_i)$ – System Clock or Logical Counter
 - Unique?
 - Scheduling Scheme – deadlock-free & serializable
 -
 -

Atomic Transactions – Concurrency Control

- $R(Q)$ requested by $T_i \rightarrow$ check $TS(T_i)$!



- $W(Q)$ requested by $T_i \rightarrow$ check $TS(T_i)$!




- Rejected transactions are rolled back and restated with a new time stamp.

Failure Recovery – A Way to Achieve Atomicity

- Failures of Volatile and Nonvolatile Storages!
 - Volatile Storage: Memory and Cache
 - Nonvolatile Storage: Disks, Magnetic Tape, etc.
 - Stable Storage: Storage which never fail.
- Log-Based Recovery
 - Write-Ahead Logging
 - Log Records
 - < Ti starts >
 - < Ti commits >
 - < Ti aborts >
 - < Ti, Data-Item-Name, Old-Value, New-Value>

Failure Recovery

- Two Basic Recovery Procedures:



- $\text{undo}(T_i)$: restore data updated by T_i
- $\text{redo}(T_i)$: reset data updated by T_i
- Operations must be idempotent!
- Recover the system when a failure occurs:
 - “Redo” committed transactions, and “undo” aborted transactions.

Failure Recovery

- Why Checkpointing?
 - The needs to scan and rerun all log entries to redo committed transactions.
- CheckPoint
 - Output all log records, Output DB, and Write <check point> to stable storage!
 - Commit: A Force Write Procedure

