

Deep Learning for Computer Vision

Spring 2019

<http://vllab.ee.ntu.edu.tw/dlcv.html> (primary)

<https://ceiba.ntu.edu.tw/1072CommE5052> (grade, etc.)

FB: [DLCV Spring 2019](#)

Yu-Chiang Frank Wang 王鈺強, Associate Professor
Dept. Electrical Engineering, National Taiwan University

What's to Be Covered Today...


- Intro to Neural Networks & CNN
 - Linear Classification
 - Neural Network for Machine Vision
 - Multi-Layer Perceptron
 - Convolutional Neural Networks
- Pytorch Framework Tutorial (by TAs)




吳致緯




劉致廷



CNN: Local Connectivity




Global connectivity




Local connectivity

- # input units (neurons): 7
- # hidden units: 3
- Number of parameters
 - Global connectivity: $7 \times 3 = 21$
 - Local connectivity: $3 \times 3 = 9$

CNN: Weight Sharing



Without weight sharing




With weight sharing


- # input units (neurons): 7
- # hidden units: 3
- Number of parameters
 - Without weight sharing: 9
 - With weight sharing : 3

Putting them together

- Local connectivity
- Weight sharing
- Handling multiple input channels
- Handling multiple output maps

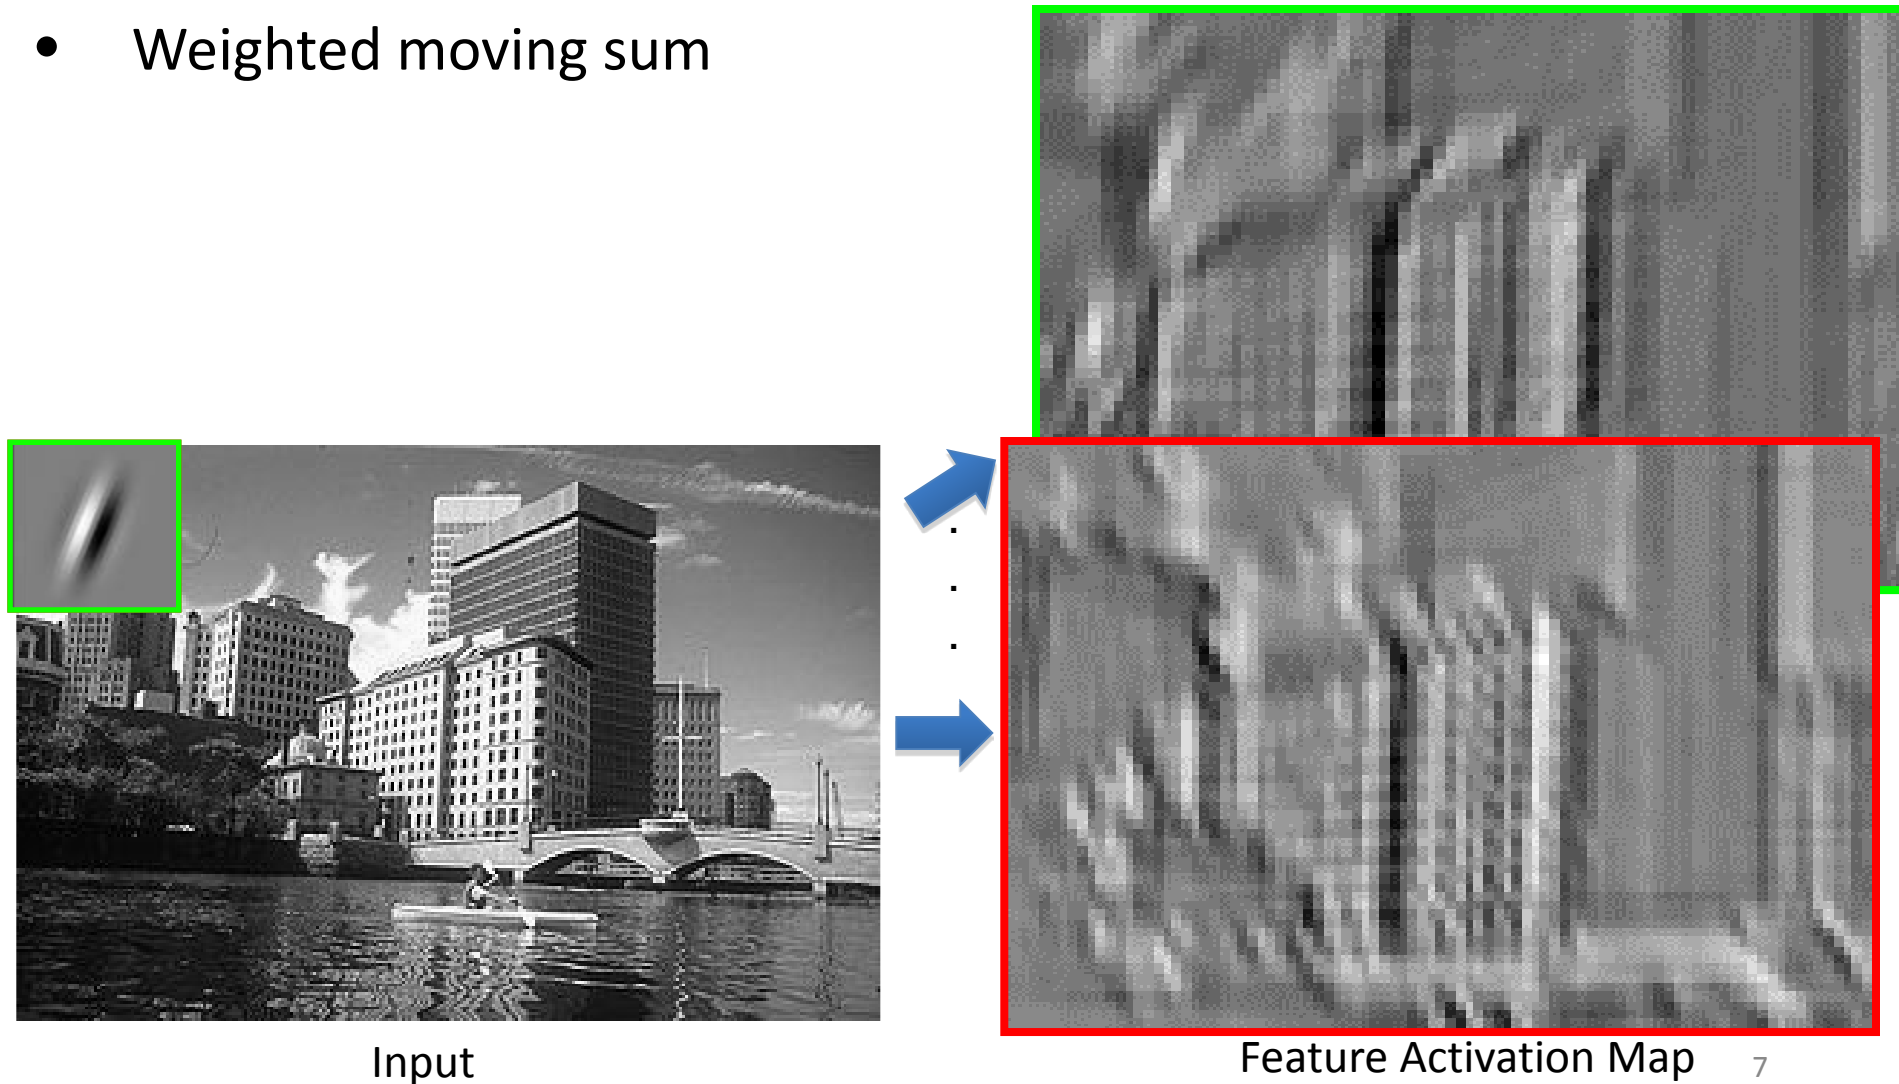


Convolution Layer in CNN



What is a Convolution?

- Weighted moving sum




What is a Convolution?

5	9	2	6	7	9	8
---	---	---	---	---	---	---

Signal

1	0	-1
---	---	----



Filter



Convolution is a local linear operator

Putting them together (cont'd)


- The brain/neuron view of CONV layer



It's just a neuron with local connectivity...

Putting them together (cont'd)

- The brain/neuron view of CONV layer




An activation map is a 28x28 sheet of neuron outputs:

1. Each is connected to a small region in the input
2. All of them share parameters

“5x5 filter” -> “5x5 receptive field for each neuron”

Putting them together (cont'd)

- The brain/neuron view of CONV layer




E.g. with 5 filters,
CONV layer consists of
neurons arranged in a 3D grid
(28x28x5)

There will be 5 different
neurons all looking at the same
region in the input volume

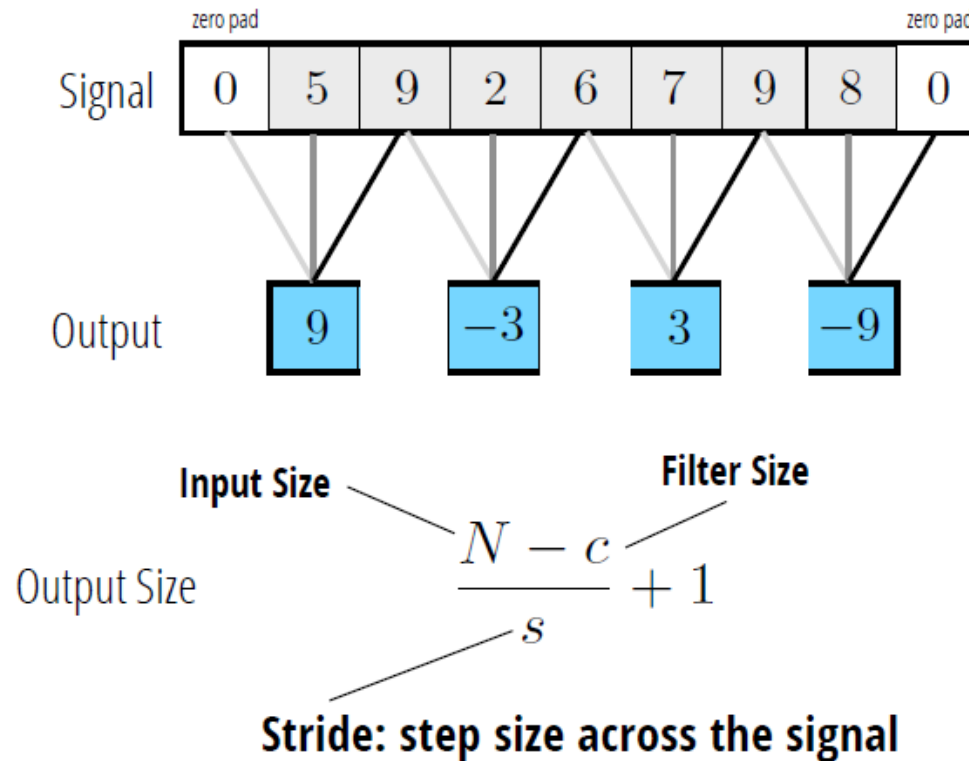
Putting them together (cont'd)

- Image input with 32 x 32 pixels convolved repeatedly with 5 x 5 x 3 filters shrinks volumes spatially (32 -> 28 -> 24 -> ...).




What is a Convolution?

- Stride
 - Step size across signals




Nonlinearity Layer in CNN




Nonlinearity Layer

- E.g., ReLU (Rectified Linear Unit)
 - Pixel by pixel computation of $\max(0, x)$




Nonlinearity Layer

- E.g., ReLU (Rectified Linear Unit)
 - Pixel by pixel computation of $\max(0, x)$

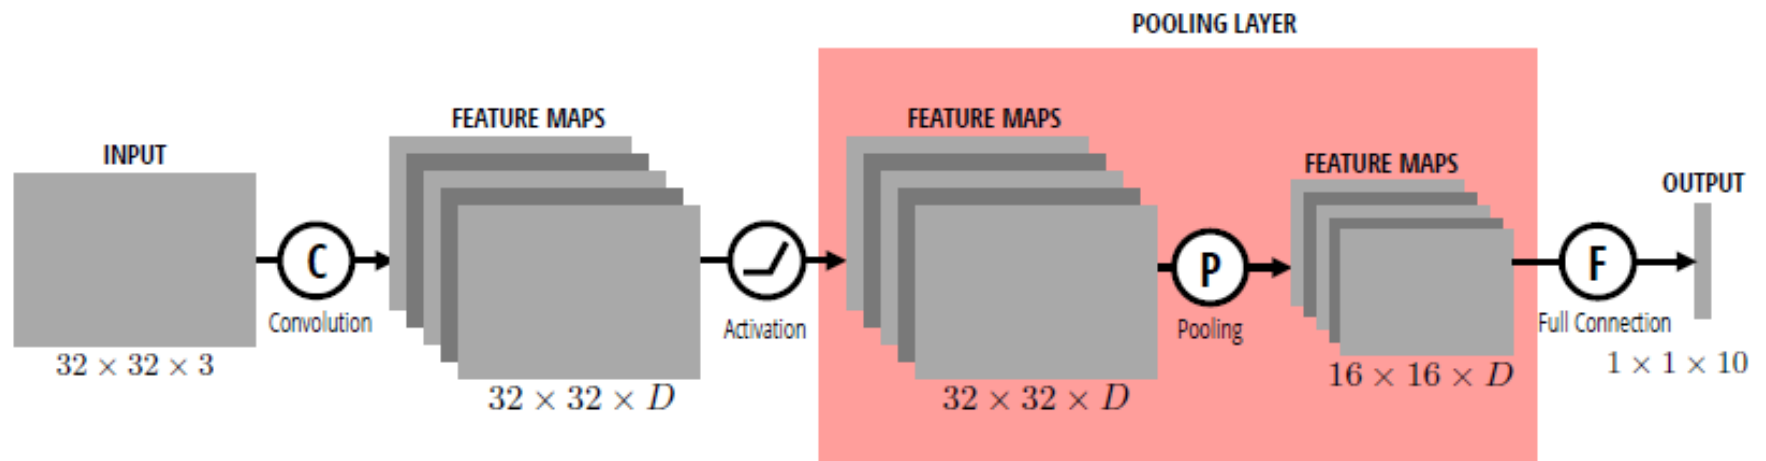


Nonlinearity Layer

- E.g., ReLU (Rectified Linear Unit)
 - Pixel by pixel computation of $\max(0, x)$




Pooling Layer in CNN





Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently
- E.g., Max Pooling




Pooling Layer

- Reduces the spatial size and provides spatial invariance




scale up!

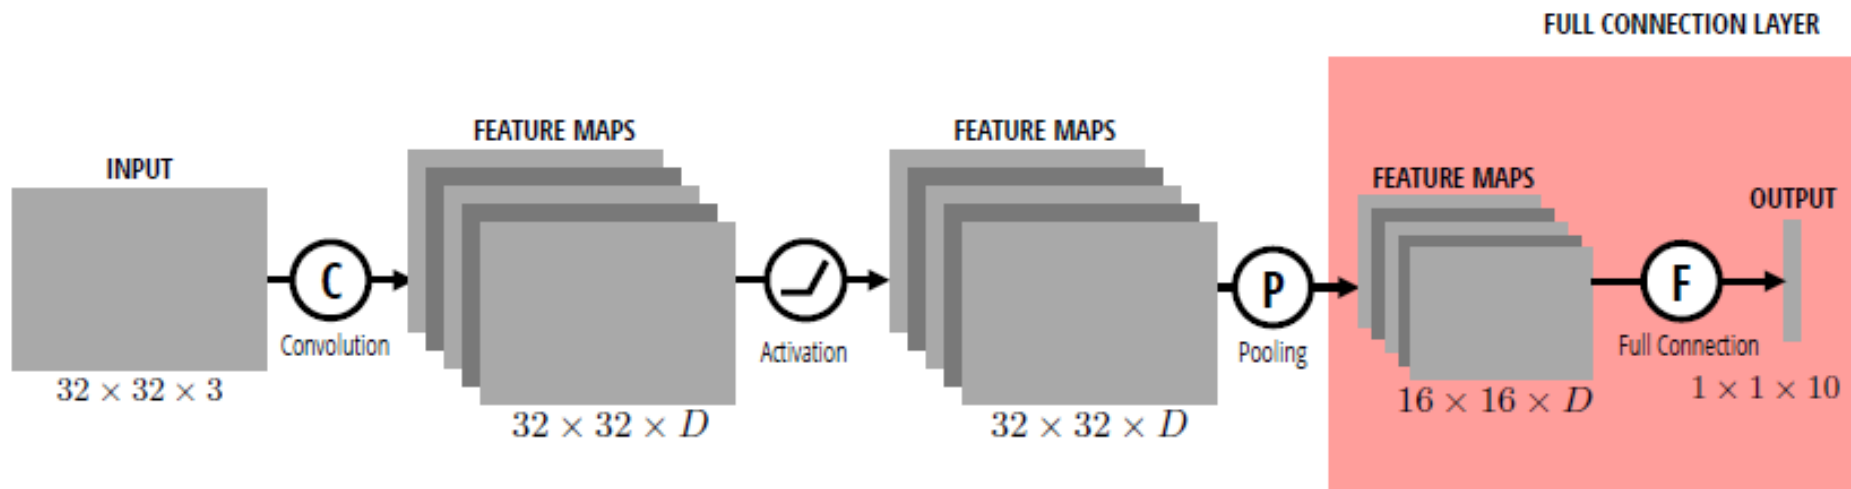
- Example
 - Nonlinearity by ReLU



- Example
 - Max pooling




Fully Connected (FC) Layer in CNN




FC Layer

- Contains neurons that connect to the entire input volume, as in ordinary neural networks



FC Layer

- Contains neurons that connect to the entire input volume, as in ordinary neural networks



CNN






Figure by Andrej Karpathy

LeNet

- Presented by Yann LeCun during the 1990s for reading digits
- Has the elements of modern architectures



AlexNet [Krizhevsky et al., 2012]

- Repopularized CNN by winning the ImageNet Challenge 2012
- 7 hidden layers, 650,000 neurons, 60M parameters
- Error rate of 16% vs. 26% for 2nd place.

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1



[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)



Deep or Not?

- Depth of the network is critical for performance.



AlexNet: 8 Layers with 18.2% top-5 error

Removing Layer 7 reduces 16 million parameters, but only 1.1% drop in performance!


Removing Layer 6 and 7 reduces 50 million parameters, but only 5.7% drop in performance

Removing middle conv layers reduces 1 million parameters, but only 3% drop in performance


Removing feature & conv layers produces a 33% drop in performance

CNN: A Revolution of Depth

AlexNet, 8 layers
(ILSVRC 2012)



VGG, 19 layers
(ILSVRC 2014)




GoogleNet, 22 layers
(ILSVRC 2014)




What is 1x1 Convolution?

- Doesn't 1x1 convolution sound redundant?




What is 1x1 Convolution? (cont'd)

- Doesn't 1x1 convolution sound redundant?
- Simply speaking, it provides...
 - Dimension reduction (?)
 - Nonlinearity




What is 1x1 Convolution? (cont'd)

- ↑ output at i^{th} layer
Example 1
 $\{28 \times 28 \times 192\}$ convolved with 32 $\{5 \times 5 \times 192\}$ kernels into $\{28 \times 28 \times 32\}$
output at $i+1^{th}$ layer
- $(5 \times 5 \times 192)$ muls \times (28×28) pixels \times 32 kernels \sim 120M muls




- Example 2**
 $\{28 \times 28 \times 192\}$ convolved with 16 $\{1 \times 1 \times 192\}$ kernels into $\{28 \times 28 \times 16\}$, followed by convolution with 32 $\{5 \times 5 \times 16\}$ kernels into $\{28 \times 28 \times 32\}$
- $192 \text{ mul} \times (28 \times 28) \text{ pixels} \times 16 \text{ kernels} \sim 2.4\text{M}$
- $(5 \times 5 \times 16) \text{ muls} \times (28 \times 28) \text{ pixels} \times 32 \text{ kernels} \sim 10\text{M}$

• 12.4M vs. 120M




ResNet

- Can we just increase the #layer?





- How can we train very deep network?
 - Residual learning




method	top-5 err. (test)
VGG [41] (ILSVRC'14)	7.32
GoogLeNet [44] (ILSVRC'14)	6.66
VGG [41] (v5)	6.8
PRelu-net [13]	4.94
BN-inception [16]	4.82
ResNet (ILSVRC'15)	3.57

ResNet (cont'd)

- Can we just increase # of layers?



- How to train very deep networks?
 - Residual learning



Non-Bottleneck
(ResNet-18, 34)





Bottleneck
(ResNet-50, 101, 152)

method	top-5 err. (test)
VGG [41] (ILSVRC'14)	7.32
GoogLeNet [44] (ILSVRC'14)	6.66
VGG [41] (v5)	6.8
PReLU-net [13]	4.94
BN-inception [16]	4.82
ResNet (ILSVRC'15)	3.57

Ref: He, Kaiming, et al. "Deep residual learning for image recognition." CVPR, 2016.

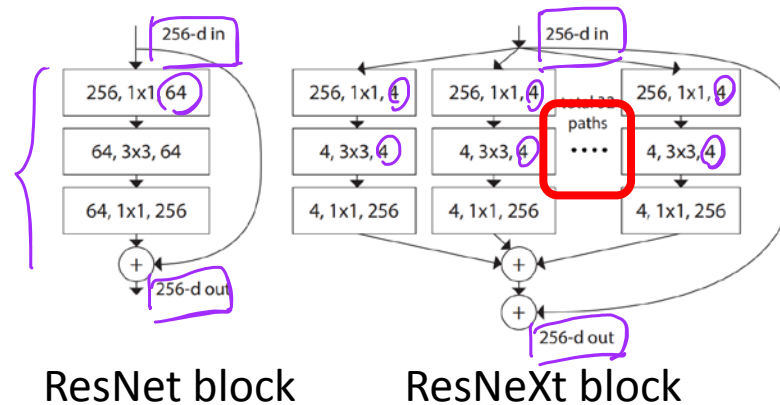
DenseNet

- Shorter connections (like ResNet) help
- Why not just connect them all?




ResNeXT

- Deeper and wider → better...what else?
 - Increase cardinality




	setting	top-1 error (%)
ResNet-50	1 × 64d	23.9
ResNeXt-50	2 × 40d	23.0
ResNeXt-50	4 × 24d	22.6
ResNeXt-50	8 × 14d	22.3
ResNeXt-50	32 × 4d	22.2
ResNet-101	1 × 64d	22.0
ResNeXt-101	2 × 40d	21.7
ResNeXt-101	4 × 24d	21.4
ResNeXt-101	8 × 14d	21.3
ResNeXt-101	32 × 4d	21.2



Squeeze-and-Excitation Net (SENet)

- How to improve acc. without much overhead?
 - Feature recalibration (channel attention)



	original		re-implementation			SENet		
	top-1 err.	top-5 err.	top-1 err.	top-5 err.	GFLOPs	top-1 err.	top-5 err.	GFLOPs
ResNet-50 [13]	24.7	7.8	24.80	7.48	3.86	23.29 _(1.51)	6.62 _(0.86)	3.87
ResNet-101 [13]	23.6	7.1	23.17	6.52	7.58	22.38 _(0.79)	6.07 _(0.45)	7.60
ResNet-152 [13]	23.0	6.7	22.42	6.34	11.30	21.57 _(0.85)	5.73 _(0.61)	11.32
ResNeXt-50 [19]	22.2	-	22.11	5.90	4.24	21.10 _(1.01)	5.49 _(0.41)	4.25
ResNeXt-101 [19]	21.2	5.6	21.18	5.57	7.99	20.70 _(0.48)	5.01 _(0.56)	8.00
VGG-16 [11]	-	-	27.02	8.81	15.47	25.22 _(1.80)	7.70 _(1.11)	15.48
BN-Inception [6]	25.2	7.82	25.38	7.89	2.03	24.23 _(1.15)	7.14 _(0.75)	2.04
Inception-ResNet-v2 [21]	19.9 [†]	4.9 [†]	20.37	5.21	11.75	19.80 _(0.57)	4.79 _(0.42)	11.76

Remarks

- CNN:
 - Reduce the number of parameters
 - Reduce the memory requirements
 - Make computation independent of the size of the image
- Neuroscience provides strong inspiration on the NN design, but little guidance on how to train CNNs.
- Few structures discussed: convolution, nonlinearity, pooling

Training Convolutional Neural Networks

- Backpropagation + stochastic gradient descent with momentum
 - [Neural Networks: Tricks of the Trade](#)
- Dropout
- Data augmentation
- Batch normalization

An Illustrative Example

$$f(x, y) = xy, \quad \frac{\partial f}{\partial x} = y, \frac{\partial f}{\partial y} = x$$

$$\text{Example: } x = 4, y = -3 \Rightarrow f(x, y) = -12$$

Partial derivatives

$$\frac{\partial f}{\partial x} = -3, \quad \frac{\partial f}{\partial y} = 4$$

Gradient

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

$$f(x, y, z) = (x + y)z = qz$$

$$q = x + y$$

$$\frac{\partial q}{\partial x} = 1, \quad \frac{\partial q}{\partial y} = 1$$

$$f = qz$$

$$\frac{\partial f}{\partial q} = z, \quad \frac{\partial f}{\partial z} = q$$

Goal: compute the gradient

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} \right]$$

$$f(x, y, z) = (x + y)z = qz$$

$$q = x + y$$

$$\frac{\partial q}{\partial x} = 1, \quad \frac{\partial q}{\partial y} = 1$$

$$f = qz$$

$$\frac{\partial f}{\partial q} = z, \quad \frac{\partial f}{\partial z} = q$$

Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

set some inputs

x = -2; y = 5; z = -4

perform the forward pass

q = x + y # q becomes 3

*f = q * z # f becomes -12*

perform the backward pass (backpropagation) in reverse order:

*# first backprop through f = q * z*


dfd z = q # df/dz = q, so gradient on z becomes 3

dfd q = z # df/dq = z, so gradient on q becomes -4


now backprop through q = x + y

*dfd x = 1.0 * dfdq # dq/dx = 1. And the multiplication here is the chain rule!*

*dfd y = 1.0 * dfdq # dq/dy = 1*




Backpropagation (recursive chain rule)




Can be computed during forward pass

The gate receives this during backprop

Dropout



(a) Standard Neural Net




(b) After applying dropout.


Intuition: successful conspiracies

- 50 people planning a conspiracy
- Strategy A: plan a big conspiracy involving 50 people
 - Likely to fail. 50 people need to play their parts correctly.
- Strategy B: plan 10 conspiracies each involving 5 people
 - Likely to succeed!



Dropout




(a) Standard Neural Net



(b) After applying dropout.



(a) At training time



(b) At test time

Main Idea: approximately combining exponentially many different neural network architectures efficiently


Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
SVM on Fisher Vectors of Dense SIFT and Color Statistics	-	-	27.3
Avg of classifiers over FVs of SIFT, LBP, GIST and CSIFT	-	-	26.2
Conv Net + dropout (Krizhevsky et al., 2012)	40.7	18.2	-
Avg of 5 Conv Nets + dropout (Krizhevsky et al., 2012)	38.1	16.4	16.4

Table 6: Results on the ILSVRC-2012 validation/test set.

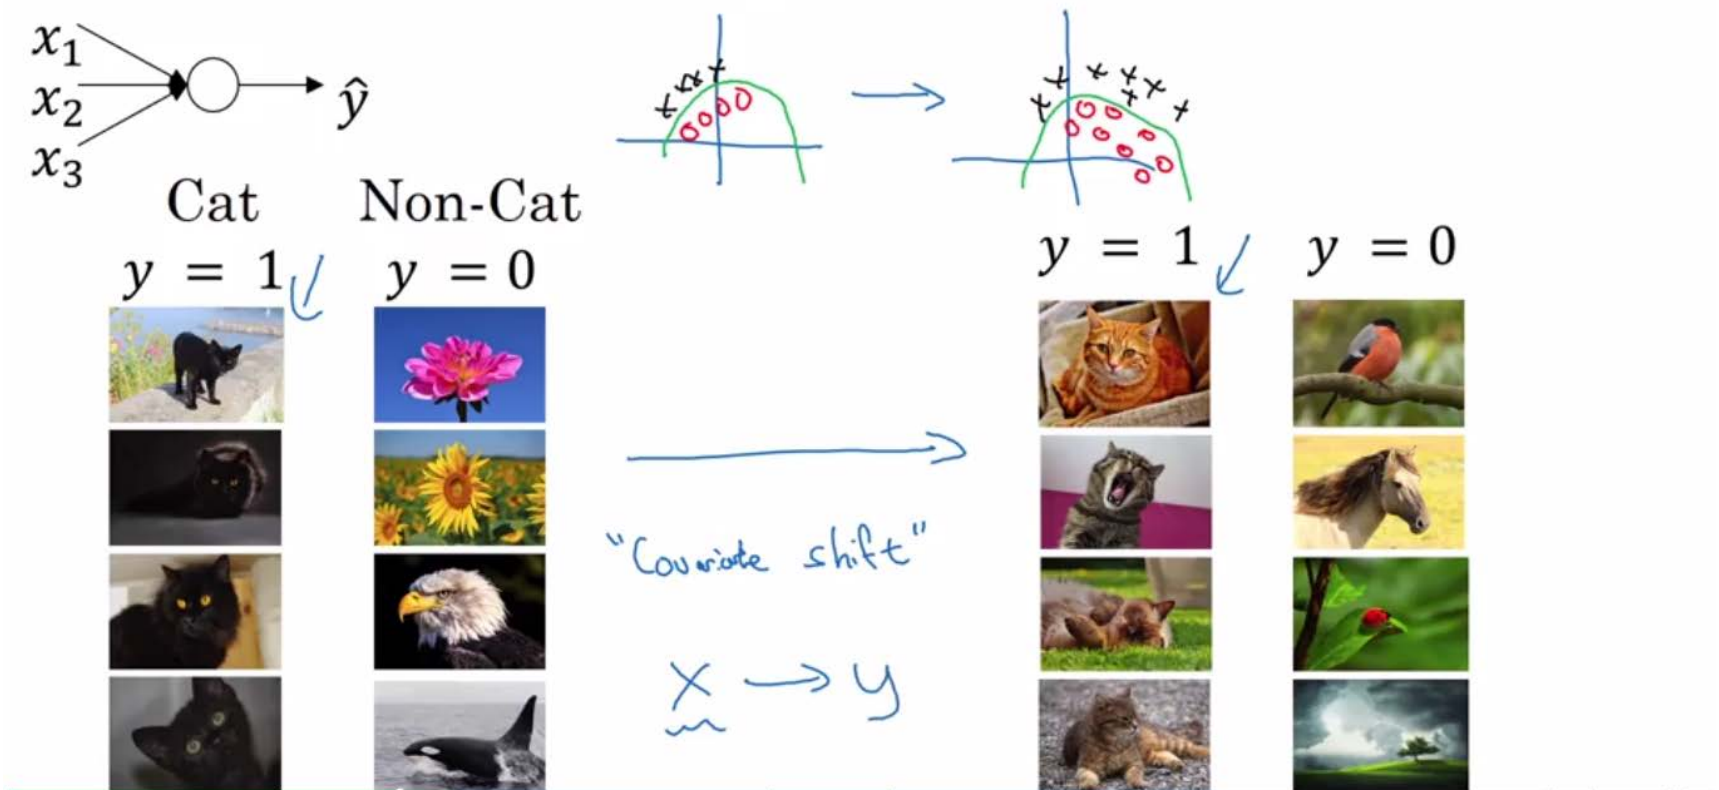
Dropout: A simple way to prevent neural networks from overfitting [[Srivastava JMLR 2014](#)]

Data Augmentation (Jittering)

- Create *virtual* training samples
 - Horizontal flip
 - Random crop
 - Color casting
 - Geometric distortion



Batch Normalization



Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β


Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$


$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$


$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$



(a)



(b) Without BN



(c) With BN



What Will We Cover Next Week?

- **Pytorch Framework Tutorial** (for those who are not familiar with Pytorch)
 - Introduction to Pytorch
 - Installation guide
 - Basic concept of computation graph and back propagation
 - Basic: Module Class
 - How to build complex model with pytorch built-in classes.
 - Basic: DataSet & DataLoader Class
 - How to load data efficiently with pytorch built-in classes.
 - Hands on example : Image Classification Task **(bring your own laptop!)**
 - Advance :
 - Finetuning with pretrained model.
 - Data augmentation
 - Training with multiple GPU
 - Exporting models to other platforms.
- **HW #1 is due 3/23 Sat 3AM & no late submission!!**